



# *Using* **Visual SoftICE™**

Release 3.2

**COMPUWARE.** 

Technical support is available from our Technical Support Hotline or via our FrontLine Support Web site.

Technical Support Hotline:  
1-800-538-7822

FrontLine Support Web Site:  
<http://frontline.compuware.com>

This document and the product referenced in it are subject to the following legends:

Access is limited to authorized users. Use of this product is subject to the terms and conditions of the user's License Agreement with Compuware Corporation.

© 2004 Compuware Corporation. All rights reserved. Unpublished - rights reserved under the Copyright Laws of the United States.

#### U.S. GOVERNMENT RIGHTS

Use, duplication, or disclosure by the U.S. Government is subject to restrictions as set forth in Compuware Corporation license agreement and as provided in DFARS 227.7202-1(a) and 227.7202-3(a) (1995), DFARS 252.227-7013(c)(1)(ii)(OCT 1988), FAR 12.212(a) (1995), FAR 52.227-19, or FAR 52.227-14 (ALT III), as applicable. Compuware Corporation.

This product contains confidential information and trade secrets of Compuware Corporation. Use, disclosure, or reproduction is prohibited without the prior express written permission of Compuware Corporation.

DriverStudio, DriverWorkbench, and SoftICE are trademarks or registered trademarks of Compuware Corporation.

Acrobat® Reader copyright © 1987-2003 Adobe Systems Incorporated. All rights reserved. Adobe, Acrobat, and Acrobat Reader are trademarks of Adobe Systems Incorporated.

All other company or product names are trademarks of their respective owners.

US Patent Nos.: Not Applicable.

November 16, 2004



# Table of Contents

## Preface

Purpose of This Manual .....	ix
Audience .....	ix
Accessibility .....	x
Organization of This Manual .....	x
Typographical Conventions .....	xii
How to Use This Manual .....	xii
Other Useful Documentation .....	xii
Customer Assistance .....	xiii
For Non-Technical Issues .....	xiii
For Technical Issues .....	xiii

---

## Chapter 1

### Choosing Your SoftICE Version

SoftICE or Visual SoftICE? .....	1
Single Machine Debugging: SoftICE .....	2
Dual Machine Debugging: Visual SoftICE .....	3
Which One Should I Use? .....	4
Visual SoftICE Overview .....	5
About Visual SoftICE .....	5

---

## Chapter 2

### Getting Started with Visual SoftICE

Introduction and Overview .....	7
What Is Visual SoftICE? .....	8

Target Installation and Configuration .....	9
Installation .....	9
Transport Selection and Configuration .....	9
Security Considerations .....	11
Master Setup .....	14
Workspaces and Layout .....	15
DriverWorkbench Search Paths .....	19
Visual SoftICE-Specific Settings .....	20
Target Discovery and Connection .....	23
Serial Connection .....	23
Network IP Discovery and Connection .....	23
Connecting to Targets Running in Virtual Machines .....	26
Opening a Crash Dump File .....	26
Remote Target Control .....	27
Missing Images .....	27
Retrieving Exports .....	28
Missing Executable .....	28
Non-Running Target Application .....	29
Stopping an Application on the Target .....	29
Manipulating Drivers on the Target .....	30
General Debugging .....	30
Controlling the Target .....	31
Breakpoints .....	34
Viewing and Editing Data .....	35
Contexts and the Call Stack .....	38
Built-In Commands .....	39

---

## Chapter 3

### Visual SoftICE Target Transports

Visual SoftICE Target Transport Overview .....	41
Serial Target Transport .....	42
Requirements and Characteristics .....	43
Dedicated PCI/CardBus Ethernet Network Interface Cards (NICs) .....	43
Requirements and Characteristics .....	44
Universal PCI/CardBus Ethernet NICs .....	44
Requirements and Characteristics .....	45
OHCI/UHCI USB Host Controller and USB NIC .....	45
Requirements and Characteristics .....	46
1394 (Firewire) .....	46
Requirements and Characteristics .....	47

---

## Chapter 4

### Overview of the Visual SoftICE User Interface

The Visual SoftICE User Interface Overview .....	49
Workspaces .....	50
Browsing Available Targets .....	51
Connecting to Targets .....	53
Visual SoftICE Icons .....	54
About the Status Bar .....	61
About the Context Bar .....	62
About the Debug Toolbar .....	63
Context Menus .....	63
User Interface Preferences .....	64
Global Settings .....	65
Per-Workspace Settings .....	66
Keyboard Settings .....	69
Toolbars & Status Bar Settings .....	71
Fonts & Colors Settings .....	73
Other User Interface Attributes and Features .....	74
Workspace Save and Load .....	74
Special Command Page Features .....	76
Cut, Copy, and Paste .....	79
Drag and Drop .....	79
Saving Contents to a File .....	80
Print and Print Preview .....	81
Script Execution .....	82

---

## Chapter 5

### The Visual SoftICE User Interface Pages

Pads and Pages .....	85
Pads .....	86
Pages .....	87
Visual SoftICE Page Overview .....	88
The Breakpoint Page .....	92
Concepts and Associated Commands .....	92
Page Features .....	93
The Command Page .....	97
Concepts and Associated Commands .....	97
Page Features .....	100
The Debug Message Page .....	103
Page Features .....	104

The Device-Driver Page .....	105
Concepts and Associated Commands .....	106
Page Features .....	113
The Disassembly Page .....	121
Associated Commands .....	122
Page Features .....	122
The Event Page .....	127
Associated Commands .....	128
Page Features .....	129
The Locals Page .....	132
Concepts and Associated Commands .....	133
Page Features .....	133
The Memory Page .....	135
Concepts and Associated Commands .....	136
Page Features .....	138
The Process List Page .....	145
Concepts and Associated Commands .....	146
Page Features .....	148
The Registers Page .....	151
Concepts and Associated Commands .....	151
Page Features .....	153
The Source Page .....	156
Associated Commands .....	158
Page Features .....	160
The Stack Page .....	167
Concepts and Associated Commands .....	167
Page Features .....	169
The Text Scratch Page .....	171
Concepts and Associated Commands .....	172
Page Features .....	172
The Watch Page .....	174
Concepts and Associated Commands .....	175
Page Features .....	175

---

## Chapter 6

### Visual SoftICE Symbol Management

Visual SoftICE Symbol Management .....	179
Where to Put the Symbols .....	179
What Symbols are Supported .....	180

Images — Why you need access to them . . . . .	180
Local Copies of Images . . . . .	180
Tables – Active Table, Loading, and Unloading Commands . . . . .	181
Automatic, On-Demand Loading and Unloading . . . . .	183
Pre-Loading/Persistent Loading . . . . .	184
Integrated MS Symbol Server Access . . . . .	184
Using Exports . . . . .	185
Symbols — Getting Setup in Visual SoftICE . . . . .	185
Setting Up General Paths . . . . .	186
Setting Up Visual SoftICE Paths . . . . .	186
Settings Notes . . . . .	186
Per-Workspace Settings Notes . . . . .	187
Toolbars & Status Bar Settings Notes . . . . .	187
The Symbol Tables Utility . . . . .	188
Filtering . . . . .	189
Table Action Buttons . . . . .	190

---

## Chapter 7

### Using Breakpoints

Introduction . . . . .	191
Types of Breakpoints Supported by Visual SoftICE . . . . .	192
Breakpoint Options . . . . .	193
Execution Breakpoints . . . . .	193
Memory Breakpoints . . . . .	194
Interrupt Breakpoints . . . . .	195
I/O Breakpoints . . . . .	195
Image Load Breakpoints . . . . .	196
Window Message Breakpoints . . . . .	197
Understanding Breakpoint Contexts . . . . .	198
Image-Relative Breakpoints . . . . .	198
Fixed Address Breakpoints . . . . .	198
Virtual Breakpoints . . . . .	198
Setting a Breakpoint Action . . . . .	199
Conditional Breakpoints . . . . .	199
Conditional Breakpoint Count Functions . . . . .	200
Using Local Variables in Conditional Expressions . . . . .	203
Referencing the Stack in Conditional Breakpoints . . . . .	204
Performance . . . . .	206
Duplicate Breakpoints . . . . .	206
Elapsed Time . . . . .	206
Breakpoint Statistics . . . . .	207
Referring to Breakpoints in Expressions . . . . .	207

Manipulating Breakpoints .....	207
Using Embedded Breakpoints .....	208

---

## Chapter 8

### Using Expressions

Expression Values .....	209
Supported Operators .....	210
Operator Precedence .....	211
Forming Expressions .....	211
Numbers .....	212
Registers .....	212
Symbols .....	212
Built-in Casts and Functions .....	213
Expression Evaluator Type System .....	216
Symbol Type .....	216
Address Type .....	217
Evaluating Symbols .....	217
Pointer Arithmetic with Symbols .....	218
Array Symbols In Expressions .....	218

---

## Chapter 9

### Exploring Windows

Overview .....	219
Resources for Advanced Debugging .....	220
Inside the Windows Kernel .....	225
Managing the Intel Architecture .....	226
Windows System Memory Map (on 32-bit Intel Architecture) .....	230
Win32 Subsystem .....	237
Inside CSRSS .....	237
USER and GDI Objects .....	238
32-bit x86 Windows Process Address Space .....	242
Heap API .....	243

---

## Appendix A

### Troubleshooting Visual SoftICE

Troubleshooting .....	255
-----------------------	-----

---

## Appendix B

### Kernel Debugger Extensions

Debugger Extension Overview and VSI Support .....	259
Controlling Debugger Extension DLLs .....	259
Using Debugger Extension Commands .....	260
<b>Glossary</b> .....	263
<b>Index</b> .....	265



# Preface



- ◆ Purpose of This Manual
- ◆ Audience
- ◆ Accessibility
- ◆ Organization of This Manual
- ◆ Typographical Conventions
- ◆ How to Use This Manual
- ◆ Other Useful Documentation
- ◆ Customer Assistance

## Purpose of This Manual

**Note:** Unless stated otherwise, this document will use “Windows” to refer to the Windows, Windows 2000, Windows XP, and Windows Server 2003 operating systems. Characteristics described in this manual apply to all supported operating systems.

Visual SoftICE is an advanced, all-purpose debugger that can debug virtually any type of code including applications, device drivers, EXEs, DLLs, and OCXs.

## Audience

This manual is intended for programmers who want to use Visual SoftICE to debug code for Windows platforms running on either 32- or 64-bit processors.

## Accessibility

Prompted by federal legislation introduced in 1998 and Section 508 of the U.S. Rehabilitation Act enacted in 2001, Compuware launched an accessibility initiative to make its products accessible to all users, including people with disabilities. This initiative addresses the special needs of users with sight, hearing, cognitive, or mobility impairments.

Section 508 requires that all electronic and information technology developed, procured, maintained, or used by the U.S. Federal government be accessible to individuals with disabilities. To that end, the World Wide Web Consortium (W3C) Web Accessibility Initiative (WAI) has created a workable standard for online content.

Compuware supports this initiative by committing to make its applications and online help documentation comply with these standards. For more information, refer to:

- ◆ W3C Web Accessibility Initiative (WAI) at [www.W3.org/WAI](http://www.W3.org/WAI)
- ◆ Section 508 Standards at [www.section508.gov](http://www.section508.gov)
- ◆ Microsoft Accessibility Technology for Everyone at [www.microsoft.com/enable/](http://www.microsoft.com/enable/)

## Organization of This Manual

The *Using Visual SoftICE* manual is organized as follows:

- ◆ Chapter 1, “Choosing Your Visual SoftICE Version”  
Helps you decide which version of Visual SoftICE, Visual SoftICE or Classic Visual SoftICE, is best suited for a given debugging scenario. It also gives an overview of the Visual SoftICE product, and highlights the many new and useful features.
- ◆ Chapter 2, “Getting Started with Visual SoftICE”  
Provides information to help you quickly get started using Visual SoftICE.
- ◆ Chapter 3, “Visual SoftICE Target Transports”  
Provides an overview of the many available transports, and compares the benefits of each.

- ◆ Chapter 4, “Overview of the Visual SoftICE User Interface”  
Explains the GUI for Visual SoftICE, including the meaning of various icons and available menu options. This chapter also discusses how to use the Visual SoftICE configuration settings to customize your environment, pre-load symbols and exports, connect to a target, modify keyboard mappings, create macro-definitions, and set troubleshooting options.
- ◆ Chapter 5, “The Visual SoftICE User Interface Pages”  
Explains the GUI pages available for Visual SoftICE, including the functionality and related commands for each.
- ◆ Chapter 6, “Visual SoftICE Symbol Management”  
Explains how to manage symbols in Visual SoftICE, including dynamically and persistently loaded symbols, MS Symbol Server setup, and path configuration.
- ◆ Chapter 7, “Using Breakpoints”  
Explains how to set breakpoints on program execution, on memory location reads and writes, on interrupts, and on reads and writes to the I/O ports.
- ◆ Chapter 8, “Using Expressions”  
Explains how to form expressions to evaluate breakpoints. This concept differs between Visual SoftICE versions, so pay particular attention to this chapter if you are not familiar with Visual SoftICE.
- ◆ Chapter 9, “Exploring Windows”  
Provides a quick overview of the Windows operating systems.
- ◆ Appendix A, “Troubleshooting Visual SoftICE”  
Explains how to solve typical problems you might encounter.
- ◆ Appendix B, “Kernel Debugger Extensions”  
Explains the available Kernel Debugger Extensions for use with Visual SoftICE, and how to work with them in the Visual SoftICE environment.
- ◆ Glossary
- ◆ Index

## Typographical Conventions

The following conventions are used consistently throughout this manual to identify certain types of information:

Convention	Description
Enter	Indicates that you should type text, then press <b>RETURN</b> or click <b>OK</b> .
italics	Indicates variable information. For example: <i>library-name</i> .
monospaced text	Used within instructions and code examples to indicate characters you type on your keyboard. Also indicates directory names, and file names.
small caps	Indicates a user-interface element, such as a button or menu.
UPPERCASE	Indicates key words and acronyms.

## How to Use This Manual

The following table suggests the best starting point for using this manual based on your level of experience debugging applications.

Experience	Suggested Starting Point
No experience using debuggers.	Read Chapter 2, Chapter 8, and then Chapters 3 through 7.
Experience with other debuggers or another release of SoftICE.	Read Chapter 1 through Chapter 6.

## Other Useful Documentation

In addition to this manual, Compuware provides the following documentation for Visual SoftICE:

- ◆ **Visual SoftICE Command Reference**  
Describes all the Visual SoftICE commands in alphabetical order. Each description provides the appropriate syntax and output for the command as well as examples that highlight how to use it.
- ◆ **Visual SoftICE online help**  
Visual SoftICE provides context-sensitive and topic-oriented HTML help as well as command line help for Visual SoftICE commands in the Command page.

- ◆ Online documentation  
Both the *Using Visual SoftICE* manual and the *Visual SoftICE Command Reference* are available as PDF. To access the PDF version of these books, start Acrobat Reader and open the *Using Visual SoftICE* or the *Visual SoftICE Command Reference* PDF files.

## Customer Assistance

### **For Non-Technical Issues**

Customer Service is available to answer any questions you might have regarding upgrades, serial numbers, and other order fulfillment needs. Customer Service is available from 8:30am to 5:30pm EST, Monday through Friday.

Call:

- ◆ In the U.S. and Canada: 1-888-283-9896
- ◆ International: +1-603-578-8103

### **For Technical Issues**

Technical Support can assist you with all your technical problems, from installation to troubleshooting. Before contacting Technical Support, please read the relevant sections of the product documentation and the release notes.

You can contact Technical Support by:

- ◆ **E-Mail:** Include your serial number and send as many details as possible to:  
<mailto:nashua.support@compuware.com>
- ◆ **World Wide Web:** Submit issues and access additional support services at:  
<http://frontline.compuware.com/nashua/>
- ◆ **Fax:** Include your serial number and send as many details as possible to:  
1-603-578-8401

- ◆ **Telephone:** Telephone support is available as a paid Priority Support Service from 8:30am to 5:30pm EST, Monday through Friday. Have product version and serial number ready.
  - ◊ In the U.S. and Canada, call: 1-888-686-3427
  - ◊ International customers, call: +1-603-578-8100

**Note:** Technical Support handles installation and setup issues free of charge.

When contacting Technical Support, please have the following information available:

- ◆ Product/service pack name and version.
- ◆ Product serial number.
- ◆ Your system configuration: operating system, network configuration, amount of RAM, environment variables, and paths.
- ◆ The details of the problem: settings, error messages, stack dumps, and the contents of any diagnostic windows.
- ◆ The details of how to reproduce the problem (if the problem is repeatable).
- ◆ The name and version of your compiler and linker and the options you used in compiling and linking.

# Chapter 1

## Choosing Your SoftICE Version



- ◆ SoftICE or Visual SoftICE?
- ◆ Single Machine Debugging: SoftICE
- ◆ Dual Machine Debugging: Visual SoftICE
- ◆ Which One Should I Use?
- ◆ Visual SoftICE Overview

### SoftICE or Visual SoftICE?

DriverStudio and SoftICE Driver Suite include two unique debuggers: SoftICE, the single-machine debugger, and Visual SoftICE, a new GUI-based dual-machine debugger. Depending on the debugging task you are facing, it may or may not be obvious which debugger you should use. This section will help you decide which tool best fits your needs.

In some situations, your choice will be simple: some processor architectures and operating systems are only supported by one of the two debuggers. [Table 1-1](#) shows the platforms supported by SoftICE and Visual SoftICE.

[Table 1-1.](#) Supported Platforms

Processor	Operating System	SoftICE	Visual SoftICE
Intel x86 and compatibles	MS-DOS, Windows 3.0/3.1/3.11, Windows 9x	Yes	No
Intel x86 and compatibles	Windows NT 3.x, Windows NT 4.0	Yes	No
Intel x86 and compatibles	Windows 2000, Windows XP, Advanced Server, Windows Server 2003	Yes	Yes

**Table 1-1.** Supported Platforms (Continued)

Processor	Operating System	SoftICE	Visual SoftICE
Intel Itanium1 and Itanium2 (IA64)	Windows XP 64bit Ed., Windows Server 2003 IA64 64-bit Ed.	No	Yes
x64 (64-bit Extended Systems)	Windows XP 64bit Ed., Windows Server 2003 Extended 64-bit Ed.	No	Yes

If you are debugging on DOS or the Windows 9x family, SoftICE is your only choice. If you are working on a 64-bit architecture, only Visual SoftICE will do. If your target is Windows and the x86 or compatible architecture, either debugger will work. In that case, read on for an overview of the differences between these two tools.

## Single Machine Debugging: SoftICE

SoftICE is a single-machine debugger, meaning simply that all of its code runs on the same machine as the code being debugged. When running, SoftICE has two basic states: popped up, where the SoftICE window is displayed, and popped down, where SoftICE is invisible and the machine runs as normal. When SoftICE is popped up, all processes on the machine are stopped, the operating system does not run, and SoftICE commands are available to the user. SoftICE can pop up in response to user input (the CTRL-D hotkey), breakpoints, exceptions, or system crashes. SoftICE is popped down by issuing one of the go or exit commands, at which point the SoftICE screen is erased and all processes in the system resume operation.

The fact that SoftICE halts the operating system when it is popped up means that it must operate without making use of any of the OS services. This has a number of consequences. For one, the SoftICE user interface does not resemble that of a normal Windows application. Although SoftICE supports keyboard and mouse input, it does not use Windows fonts, nor does its interface contain the enhancements common to Windows applications. In addition, SoftICE cannot assume that it is safe to perform disk access whenever it is popped up, so loading or saving symbol information and SoftICE data is done through companion applications, such as Symbol Loader (Loader32.exe).

Another consequence of the SoftICE single machine architecture is that the interface is extremely fast. All the data in the machine is directly

accessible to the debugger, so even tasks involving large amounts of memory access are completed with no noticeable delay.

Because symbols and source code must be loaded ahead of time, SoftICE uses a packaged format for symbols called NMS files. Symbols, translated from the DBG or PDB files output by the linker, can be combined with all or some of the source files used to build the module, and loaded into SoftICE all at once using Symbol Loader or its command-line equivalent, NMSYM. In addition, the new Microsoft Symbol Servers can be accessed using Symbol Retriever utility, which is also capable of translating symbols into NMS files and loading them into SoftICE. These tools make the necessary management of symbols for SoftICE as simple as possible.

SoftICE supports a subset of the available KD Extensions defined by Microsoft. Because the operating system is stopped when the debugger is popped up, SoftICE does not support all the available KD Extensions, since it is not able to make system calls.

There are certain situations where debugging on a single machine is impractical. For instance, if your project is a display driver that is not yet working properly, SoftICE may not be able to display its output. SoftICE does include support for remote debugging, which can be used in many of these situations to redirect the SoftICE input and output over a serial or IP networking link. The remote application in this case is SIRemote, which simply acts as a dumb terminal for SoftICE. The operation of the debugger is not otherwise changed by running remotely.

## Dual Machine Debugging: Visual SoftICE

Visual SoftICE, on the other hand, is a dual-machine debugger. The user interface and nearly all of the interpretive code runs on the “master” machine; the code to be debugged runs alongside a small core of debugging functions on the “target” machine. Master and target machines are connected via a transport, which can be a serial cable, IP network interface device, or IEEE 1394 connection.

Because the master machine is never stopped by the debugger, the Visual SoftICE user interface is free to take advantage of all of the usual Windows UI devices. The Visual SoftICE user interface will be instantly familiar to anyone who has used sophisticated Windows programs before; in addition, the command set has been duplicated (with a few exceptions) from the original SoftICE, so SoftICE users should find much that is familiar about Visual SoftICE as well.

Visual SoftICE is also able to load symbol information on-the-fly at any time – including retrieving symbols from a Symbol Server site – so this

task is generally handled automatically by the debugger. This frees the user from the necessity of manually specifying symbol files to be loaded by the debugger, although that option is still available.

Visual SoftICE supports loading and examining crashdump and minidump files directly, a feature not found in SoftICE (the DriverStudio DriverWorkbench Application also supports this).

Visual SoftICE also provides complete support for the Microsoft KD Extensions, including those that will not run on SoftICE for architectural reasons.

## Which One Should I Use?

If your project falls into the wide overlap between SoftICE and Visual SoftICE, you are probably still wondering which debugger is best for you. Obviously, there is not always a single right answer to this question, but in the remainder of this section we will try to cover some of the scenarios where one debugger might be favored over the other. We are down to guidelines here, though; devotees of either debugger will be quick to point out that their favorite still has advantages, even in cases where the other might appear to be the better choice. We encourage you to try them both, and consider them two similar but distinct tools in your debugging toolbox.

- ◆ If you prefer a full-featured Windows GUI, you will probably want to use Visual SoftICE. The SoftICE interface is fast and powerful, but it has no GUI, and it takes some getting used to.
- ◆ If you are debugging a crashdump file, try Visual SoftICE. You will be able to use many of the debugging commands you are already familiar with, and Visual SoftICE can reveal more information than the crashdump functionality within DriverWorkbench.
- ◆ If you need complete KD Extensions support, use Visual SoftICE. SoftICE provides a limited subset of KD Extensions, but not the whole set.
- ◆ If you are debugging a network driver, and you are concerned that the Visual SoftICE IP transport layer might affect the results, use SoftICE. Conversely, if you are debugging a video driver's mode initialization, a Direct3D or streaming app or driver, or input device driver, try Visual SoftICE.
- ◆ If you want direct access to BoundsChecker events from within the debugger, use SoftICE.

- ◆ If you do not have access to a second machine, or you are traveling and debugging code on a laptop, use SoftICE.
- ◆ If you need the ability to package source code together with symbolic debugging information in NMS files, use SoftICE. Both debuggers are capable of loading source code separately from symbol files, of course.

And if you are still confused about which debugger to use, skim through the documentation for both of them. Chances are that something you see there will point you in the right direction.

## Visual SoftICE Overview

Visual SoftICE is an advanced, all-purpose debugger that can debug virtually any type of code. This includes, but is not limited to, interrupt routines, processor level changes, and I/O. Visual SoftICE combines the power of a hardware debugger with the simplicity of a symbolic debugger. It provides hardware-like breakpoints and sticky breakpoints that follow the memory as the operating system discards, reloads, and swaps pages.

Visual SoftICE displays your source code as you debug, and lets you access your local and global data through symbolic names. Unlike conventional debuggers, which are restricted to application space, Visual SoftICE has complete system access and can trace difficult problems between the system and application layers.

Visual SoftICE has a 32-bit Master that connects to either 32-bit or 64-bit targets over a variety of available transports.

### About Visual SoftICE

Some of the major benefits Visual SoftICE provides include the following:

- ◆ Source level debugging of 32-bit (Win32) and 64-bit applications, and Windows device drivers (both kernel and user mode).
- ◆ Support for x86, x64 Extended Systems, and Itanium (1 and 2) platforms.
- ◆ Supports PAE mode on x86 platforms with 4GB or less of physical memory. This is to support Microsoft's Data Execution Prevention (DEP) initiative, available starting with XP SP2 and W2K3 SP1.
- ◆ Debugging virtually any code, including interrupt routines and the Windows kernels.

- ◆ Setting real-time breakpoints on memory reads/writes, port reads/writes, and interrupts.
- ◆ Setting breakpoints on Windows messages.
- ◆ Setting conditional breakpoints and breakpoint actions.
- ◆ Displaying elapsed time to the breakpoint trigger using the processor clock counters.
- ◆ Displaying internal Windows information, such as:
  - ◊ Complete thread and process information
  - ◊ Virtual memory map of a process
  - ◊ Kernel-mode entry points
  - ◊ Windows NT object directory
  - ◊ Complete driver object and device object information
  - ◊ Heaps
  - ◊ Structured Exception Handling (SEH) frames
  - ◊ DLL exports
- ◆ Using the WHAT command to identify a name or an expression, if it evaluates to a known type.
- ◆ Supporting the MMX, SSE, SSE2, SSE3, x86 instruction set extensions, plus other data decoding such as 3DNow
- ◆ Creating user-defined macros.
- ◆ On-demand symbol handling.
- ◆ Full KD extension support, including Analyze.
- ◆ Remote machine control, including Device Manager, Driver Installation and Removal, and Configuration.
- ◆ Visual SoftICE provides a robust and customizable GUI with dockable pages for debugging applications across all platforms. The Visual SoftICE user interface is designed to be functional and flexible.

# Chapter 2

## Getting Started with Visual SoftICE



- ◆ Introduction and Overview
- ◆ Target Installation and Configuration
- ◆ Master Setup
- ◆ Target Discovery and Connection
- ◆ Remote Target Control
- ◆ General Debugging

### Introduction and Overview

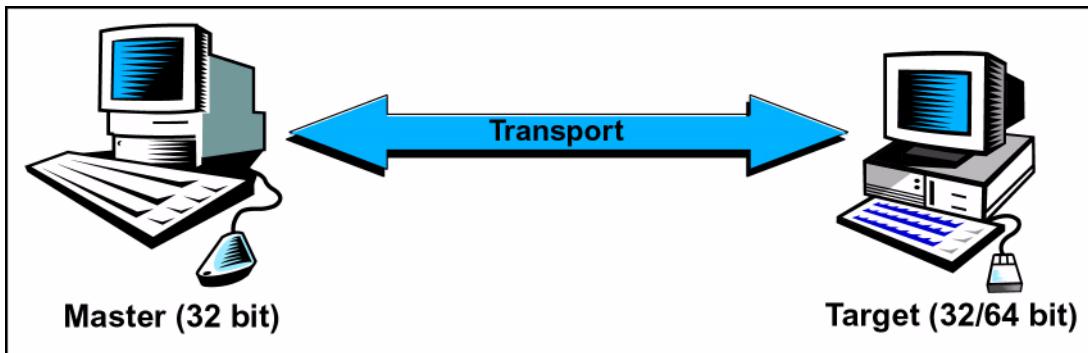
This chapter will introduce you to Visual SoftICE and some of its features, give some advice on getting started quickly, and mention some fundamental debugging steps commonly experienced in the use of the tool. In this chapter we will cover:

- ◆ Installing and starting the target component on a machine
- ◆ Setting up the master for a debugging session
- ◆ Finding and connecting to the target machine
- ◆ Remotely preparing the target machine for a debugging session
- ◆ Starting and stopping the target
- ◆ Automatic loading of source and symbols files on the master
- ◆ Breakpoint usage tips
- ◆ Viewing stack, memory, and register information on the target
- ◆ Exploring the target system (hardware and software) with built-in commands

After completing this chapter you should be well on your way to using and exploring the tool for your specific ends. Please note that other chapters provide more detail on the items introduced here, and you are encouraged to follow the references herein.

## What Is Visual SoftICE?

Visual SoftICE is a dual machine debugger that controls a remote target.



The master resides on a 32bit Windows machine, which can take full advantage of all the OS and system features of a standard application. This allows the operation of things like automatic symbol loading/unloading, script execution, copy and paste to other applications, and a host of common user interface actions. It also allows a single master to connect to a range of different target types (platform, operating system, and data source) without having to reinstall or modify the master. (More about the differences between a single-machine debugger and a dual-machine debugger can be found in [Chapter 1, “Choosing Your SoftICE Version”](#).)

Visual SoftICE can also open operating system crash dump files for exploration, as if these were full fledged targets, allowing common debugging tasks against the available snapshot data.

Usually the master is installed on a user's development machine, where a remote machine will be used for debugging, testing, and exploration. The rest of this chapter will assume this to be the case.

# Target Installation and Configuration

The following sections detail installation of the target software, and configuration of the transports that allow the master and target to communicate.

## Installation

DriverStudio allows for **target-only** installation of the Visual SoftICE software (or as part of a complete install of the package). When installing Visual SoftICE targets, there is support for x86, IA64, and x64 (64-bit Extended Systems) hardware running various Windows operating systems. Each hardware platform, whether it is running a 32-bit or 64-bit operating system, requires a different set of physical bits on the target.

Once selected, the physical installation process is straight forward, but is then followed immediately by transport selection and configuration. This can be bypassed during installation, but is the next important key to success in getting started.

## Transport Selection and Configuration

Visual SoftICE offers a range of transports to communicate between the master and target, and setup can be confusing. What must take place is selection and configuration of a transport type via the DriverStudio Settings application (DSCONFIG) on the target.

Visual SoftICE supports the following transport connection types between the master and a live target:

- ◆ Serial
- ◆ Dedicated PCI Ethernet (Network Interface Card (NIC) or PCMCIA card)
- ◆ Universal PCI Ethernet (NIC or PCMCIA card)
- ◆ OHCI/UHCI USB Host Controller and USB NIC
- ◆ 1394 (Firewire)

Some key points to understanding this process:

- ◆ The selection of a transport type **must** be done on the target itself, and *can not* be done remotely from the master machine.
- ◆ Transport selection normally requires the target to be rebooted to ensure proper OS setup.
- ◆ The transport selection process installs and removes drivers on the target machine. The drivers provided are not signed (nor can they be), so you may see OS warnings and popup messages. You should always allow the drivers to be installed.

- ◆ You must have administrator privileges to load drivers, so this entire process is disabled if you are not running the DSConfig application from a proper OS account.
- ◆ The settings application inspects the available hardware on the target and only displays the transport choices for that available hardware.

---

**Caution:** It is very important to note that *all* the transports will completely take over the hardware they target. Selecting the serial transport means the OS will not be able to use the serial port for general access once the selection is installed; it is now dedicated entirely to Visual SoftICE. This is true for PCI NIC cards and PCMCIA cards as well.

---

**Caution:** This same point is even more important to understand for USB and 1394. In these cases the host controller is taken over, rendering any other devices served by that controller unable to operate. Make sure you do not accidentally disable your keyboard and mouse!

---

**Tip:** For specific step-by-step help configuring a particular transport, please take advantage of the online help (available after master or target installation). Specific configuration information is found in several topics entitled "How to Configure a [transport-name]", under the section entitled "About the VSI Target Transports". This help also includes configuration tips, and troubleshooting directions.

Each transport has its own advantages and disadvantages, as well as performance characteristics. For more detail, and comparison of supported transports, see [Chapter 3, "Visual SoftICE Target Transports"](#).

## Starting the Debugger

Once you have successfully installed and configured the debugger and transport on your target, the debugger may already be running. If it is not active, you can get it started from the DriverStudio taskbar icon on the target itself, or from the master by connecting to the target and using the STARTDEBUGGER command.

## Configuring the Start Mode

On the main page of the DriverStudio Settings application (Control Panel) you can configure the start mode for the Visual SoftICE Remote Debugging Core.

- ◆ **Boot** – Starts Visual SoftICE before Windows loads. This mode is most useful for debugging device drivers.
- ◆ **System** – Starts Visual SoftICE with Windows. This mode is ideal for application developers.
- ◆ **Automatic** – Visual SoftICE is started with Windows, but you will not be able to debug core device drivers.
- ◆ **Manual** – You must start Visual SoftICE manually after each boot.
- ◆ **Disabled** – Disables all Visual SoftICE components and filters.

In addition, two checkboxes are also provided:

- ◆ **Stop on Boot**

Selecting this option causes the target debugger to stop after it initializes. This is useful if you want to connect as early as possible. For example, if you are trying to debug a boot time driver.

- ◆ **Enable Auto-Copy**

When this option is selected, Visual SoftICE allows you to configure an Auto-Copy Script, which executes when the target boots and requests file updates. Configure the Visual SoftICE master to use an Auto-Copy Script via the **Per-Workspace Settings** tab in the **Preferences** dialog.

---



**Your target is now ready to be remotely controlled and debugged!**

---

## **Security Considerations**

Visual SoftICE does not modify the operating system files of the target, nor leave it configured on uninstall in any less secure setup than its original configuration. However, while a debugger is installed on a system, that system is vulnerable and exposed. Visual SoftICE in particular is designed to give the remote user total access to the target machine.

**Note:** We strongly recommend against installing a debugger (of any kind) on business production machines exposed to external access.

Each debugging situation is different, so Visual SoftICE does offer some access control while you are debugging. These security features limit access to a target machine, and limit what Visual SoftICE masters can do to the target remotely.

## **Understanding Network Discovery and Connections**

Network discovery by DriverStudio applications (such as Visual SoftICE) is achieved via a standard network broadcast query. This query is sent out within only the most limited portion of a network, or subnet. The broadcast will not travel across gateways or switches, and is intended by design to be limited to the local network loop.

However, connections are different for different parts of DriverStudio. Many components of DriverStudio use an operating system RPC mechanism to connect to a remote machine. This follows all the standard OS, domain, connection, and ACL security parameters configured on that remote machine and within the OS domain.

Visual SoftICE does *not* make a connection through OS managed subsystems, and therefore does not participate in domain or ACL limitations. Firewalls (software or hardware) will affect connections.

#### Limiting Access

On transport configurations exposed to external network traffic (IP and 1394 transport types), the DriverStudio Settings application allows for a password to be assigned to the target.

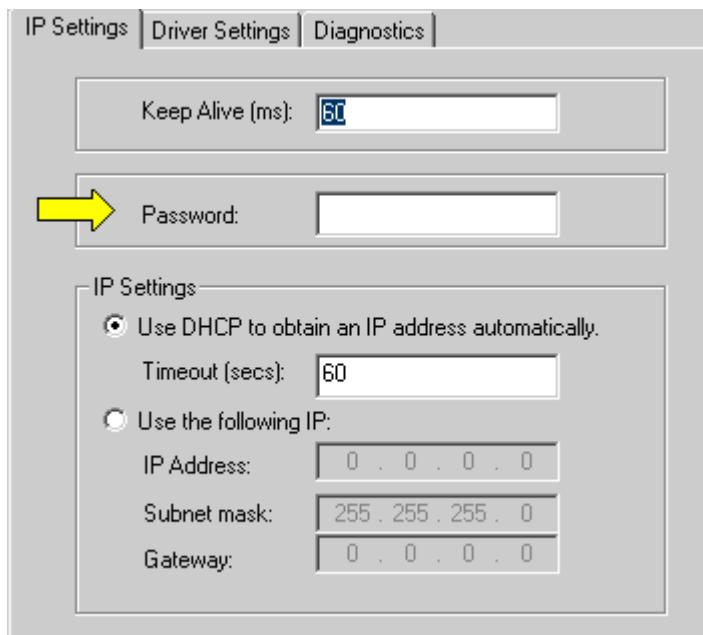


Figure 2-1. IP Settings Tab

The target will continue to be visible to network enumeration, but a Visual SoftICE connection will be disallowed if the password is incorrect. While the password is encrypted for transmission, this should still be considered a very minimal level of security, and is intended only for convenience.

#### Limiting Remote Functionality

Most of the remote control actions that a Visual SoftICE master can take against a running machine can be controlled at the target, from the **Enhanced Debugging** section of the DriverStudio Settings application.

**Note:** These settings can only be modified at the target.

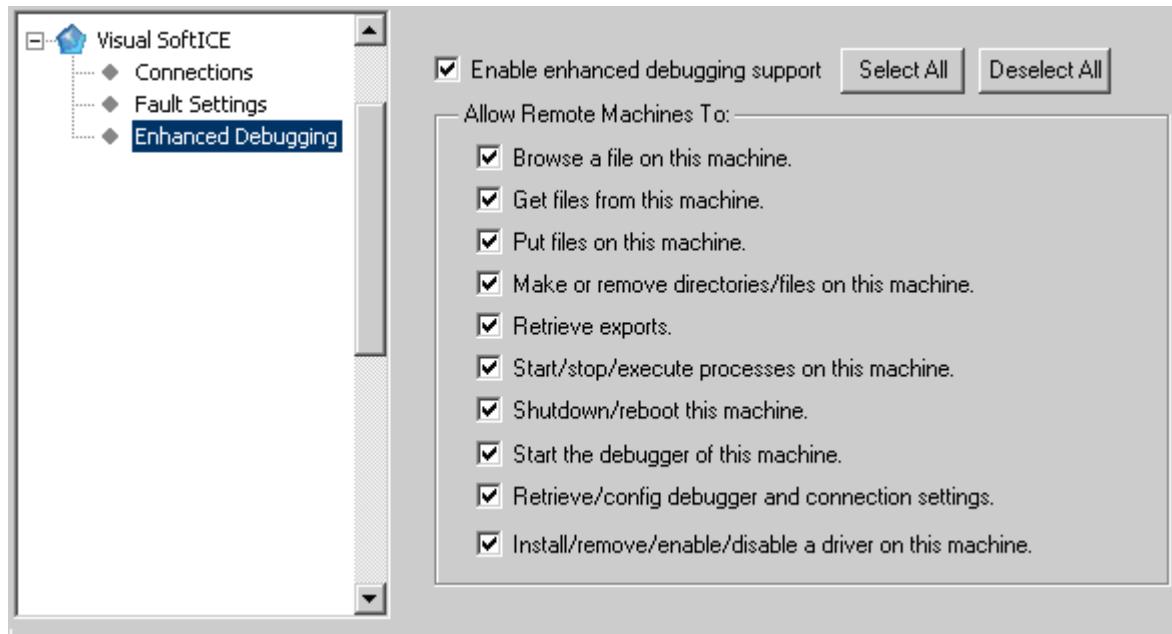


Figure 2-2. Enhanced Debugging Options

- ◆ **Browse a file on this machine.**  
Controls whether file directory searches are allowed on the machine. This directly impacts the FS, TDIR, and TVOL commands.
- ◆ **Get files from this machine.**  
Controls whether files may be retrieved from this machine back to the master. This directly impacts the FGET command.
- ◆ **Put files on this machine.**  
Controls whether files may be sent from the master to this machine. This directly impacts the FPUT, DEVMGR, and SVCMGR commands, Auto-Copy script execution, as well as significant functionality in the Device-Driver page.
- ◆ **Make or remove directories/files on this machine.**  
Controls whether files or directories on this machine may be created or removed, as well as whether files can be renamed, or moved. This indirectly impacts the FPUT command and Auto-Copy script execution. This directly impacts the TATTRIB, TMKDIR, TMOVE, TRENAME, TRMDIR, and TRMFILE commands.

- ◆ **Retrieve exports.**  
Controls whether exports for images in this system can be collected and communicated to the master. This indirectly impacts some symbol functionality when the master can not retrieve other symbol information locally, and as a last attempt tries to retrieve image export information from the target. This directly impacts the GETEXP command.
  - ◆ **Start/Stop/Execute processes on this machine.**  
Controls whether processes can be killed or started on this machine. This directly impacts the EXEC and KILL commands, and indirectly impacts the SVCMGR command.
  - ◆ **Shutdown/Reboot this machine.**  
Controls whether the master can do an orderly OS shutdown or reboot of this machine. This directly impacts the SHUTDOWN and REBOOT commands.
- Note:** This does not impact the HBOOT command, which is not restricted. HBOOT is hardware oriented, and forces a hard reset of the machine.
- ◆ **Start the debugger on this machine.**  
Controls whether the master can start the debugger on this machine. This directly impacts the STARTDEBUGGER command.
  - Note:** You can not shutdown the debugger, once started.

  - ◆ **Retrieve/Config debugger and connection settings.**  
Controls whether the master can see the debugger settings on this machine, and modify some attributes. This directly impacts the TCONFIG command.
  - ◆ **Install/Remove/Enable/Disable a driver on this machine.**  
Controls whether drivers and kernel services may be installed, removed, enabled, or disabled on this machine by the master. This directly impacts the DEVMGR and SVCMGR commands, as well as significant functionality in the Device-Driver page.

## Master Setup

To get up and running with the Visual SoftICE master quickly, there are a few steps you should complete first. You would normally do most of these only once. You can change and override any of these settings in the future.

Regardless of the target type, connection, or debugging tasks, following these steps will get the master prepared to be as functional as possible.

- 1** Customize the workspace.  
Get the workspace display layout set up the way you want by configuring what pages will be visible, as well as rearranging their location, orientation, and size. Also set up the toolbars with any applicable buttons, including custom tools.
- 2** Configure global search paths (Image, KD Extension, Symbols, and Source).
- 3** Configure workspace Load and Save behavior.
- 4** Configure any Visual SoftICE-specific settings.
  - ◊ Search paths (Export, Script)
  - ◊ Disassembly and Source preferences
  - ◊ Breakpoint Save and Restore behavior
  - ◊ Scripts

## **Workspaces and Layout**

DriverWorkbench uses the workspace concept to save and restore your preferred working environment between sessions. This includes the layout of the screen (pads and pages), whether the frame was maximized or not, keyboard definitions and input rules you have defined, your printer preferences, toolbar locations and configurations, and other specific configuration information you have customized.

### **Understanding Workspaces**

Each page instance can store specific configuration information in the workspace. Thus if you have two Debug Message pages, each with different string filters; these filters will be saved to the workspace individually, and restored with the workspace. Almost all multiple-instance pages support some page instance behavior.

The debugger stores configuration information in the workspace as well, including overrides to paths, workspace event script names, event log configuration settings, preferred display sizes for various windows you choose to display, and more.

To edit configuration settings stored in the workspace, you need to access the Per-Workspace Settings in the Preferences dialog.

- 1** Select **Preferences** from the **File** menu.
- 2** Click the **Per-Workspace Settings** tab.

You can use workspaces in many different ways, and these settings allow you to customize the workspace for any purpose. For example:

- ◆ Many users create and save a workspace configuration for different modes of work they do. This may include application debugging, kernel debugging, and/or the development phase versus the debugging and optimization phases.
- ◆ Many users create one workspace per target, allowing them specific views for the types of problems they are debugging on different hardware and operating systems.
- ◆ Many users create specific workspaces for their various development and debugging masters. This may include layout of pages on multi-monitor setups in one case, and single-monitor setups in another.

## Customizing Your Workspace Layout

If you launch Visual SoftICE from the **Start** menu, DriverWorkbench will appear with a default workspace layout. At this point you want to customize the layout of pads and pages until you get something you like. For full detail on the GUI features of Visual SoftICE running inside DriverWorkbench, see [Chapter 4, “Overview of the Visual SoftICE User Interface”](#).

For now, the following sections provide a short overview of pads and pages to get you started quickly. The basic thing to understand is that the pad is a dockable container of pages. You design your workspace layout by arranging pads in docked and floating locations, and then filling them with your preference of pages.

### Pads

Pads have names, can either be docked to the frame or floating, and each have a system menu allowing you to minimize, maximize, and close the window (closing all contained pages). Screen layout is composed of docked and floating pads (and all their contents) and toolbars.

- ◆ You can arrange Pages within the Pad by dragging and dropping them.
- ◆ You can drag and drop Pages between pads.
- ◆ Since pages must exist within a pad, dropping a page on the frame itself will automatically create a pad for it.
- ◆ You can manipulate pads by selecting **Pad** from the **View** menu. This includes creating, deleting, renaming, and walking through the existing pads.

**Note:** When pads are first created, they are always set to floating style. You can switch to a dockable pad by right-clicking on the title bar and selecting **Dockable** from the pop-up menu.

- ◆ Several keystrokes are also useful for shifting display and input focus of pads. [Table 2-1](#) displays the pad control keystrokes.

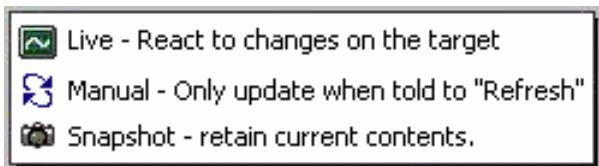
[Table 2-1.](#) Pad Control Keystrokes

Key Sequence	Action
CTRL+F4	Close the current pad.
ALT+F6	Switch to the next pad.
CTRL+F6	Display the next pad.
Shift+CTRL+F6	Display the previous pad.

## Pages

Each page is dedicated to a certain type of functionality, or to a specific type of data, and must exist within a pad. Pages may be single-instanced (only one is ever allowed at a time), or multiple-instanced (any number are allowed). Pages have names, and one state called the Page Mode. Pages may have specific preferences, and provide control over their colors and fonts.

Each page in the DriverWorkbench user interface supports one or more of the following 3 modes:



**Table 2-2.** Available Page Modes

Mode	Description
Live Mode	Events from the target cause the page to refresh automatically.
Manual Mode	You decide to refresh the page whenever you want by selecting Refresh from the pop-up menu.
Snapshot Mode	The page is protecting its current contents from overwrite. You must elect to destroy the data. Changes to the data and actions against the target are not allowed in this mode.

Several keystrokes are also useful for shifting display and input focus of pages. **Table 2-3** displays the page control keystrokes.

**Table 2-3.** Page Control Keystrokes

Key Sequence	Action
CTRL+Tab	Display the next page.
Shift+Ctrl+Tab	Display the previous page.

## Workspace Save and Load Behavior

While you are configuring your workspace layout, it is a good time to set your preferences for how DriverWorkbench will handle startup and shutdown. Complete the following steps to access the Workspace Save and Load preferences.

- 1 Select **Preferences** from the **File** menu.
- 2 Click the **Global Settings** tab.
- 3 Expand the **General** list element.
- 4 Select **Workspace Save/Load** from the list.

Use this tab to configure all of the available options according to your personal preference. As you highlight or hover your mouse pointer over the name of each option, the Settings utility displays a description of the option in the text area below the list.

## **DriverWorkbench Search Paths**

Configuring search paths can make your life significantly easier. Visual SoftICE will use all of the global path information to search the master for image, symbol, source, and KD extensions.

**Note:** Path settings come in two types: a path list (an ordered list of paths to search, semicolon separated), and a single path entry. All path types for Visual SoftICE accept the ellipses (...) syntax at the end of a directory name to indicate Visual SoftICE should search all sub-directories below the one specified.

Complete the following steps to configure the global search paths.

- 1** Select **Preferences** from the **File** menu.
- 2** Click the **Global Settings** tab.
- 3** Expand the **General** list element.
- 4** Select **Paths** from the list.

The Settings utility displays a list of search path entries and MS SymbolServer options the system can use. As you highlight or hover your mouse pointer over the name of each option, the Settings utility displays a description of the option in the text area below the list.

- ◆ **Image Search Paths**  
Where you will put local copies of the executables you will debug on the target.
- ◆ **KD Extension Paths**  
Where to find KD extensions when no absolute path is specified. This should usually point to the top of the directory structure that Microsoft creates on install. If this is the case, Visual SoftICE will attempt to load the proper KD extension for the specific version of the OS on the target.
- ◆ **MS Symbol Server Path**  
Where the MS Symbol Server should be found, and where it should store things its retrieves.
- ◆ **Source Search Paths**  
Where your source code lives.
- ◆ **User Symbol Search Paths**  
Where your symbols (PDB files) live.

**Note:** Carefully read [Chapter 6, “Visual SoftICE Symbol Management”](#) once you become comfortable with the product and are ready to begin debugging.

## **Visual SoftICE-Specific Settings**

It is important to take a moment now to configure a few preferences specific to Visual SoftICE behavior. There are both global and workspace-specific settings you need to configure.

**Note:** In some cases the workspace preference tab appears to have duplicate settings to what is under the global preference tab. This is by design. These are offered as overrides or pre-pending settings to the global values of the same name. There is a checkbox next to these items and if you hover the mouse pointer over the checkbox, a detailed explanation and directions will appear in the text area below.

### **Global Search Paths**

Complete the following steps to access the Visual SoftICE-specific global path settings.

- 1** Select **Preferences** from the **File** menu.
- 2** Click the **Global Settings** tab.
- 3** Expand the **Visual SoftICE** list element.
- 4** Select **Paths** from the list.

There are two settings you need to configure on this tab:

- ◆ **Export Path**  
Where you want the top of the retrieved export directory structure to grow.
- ◆ **Script Search Paths**  
Where to find scripts when no absolute path is specified.

### **Global Disassembly and Source Preferences**

Complete the following steps to access the Visual SoftICE-specific global Disassembly and Source settings.

- 1** Select **Preferences** from the **File** menu.
- 2** Click the **Global Settings** tab.
- 3** Expand the **Visual SoftICE** list element.
- 4** Select **Source/Disassembly** from the list.

There are many settings available on this tab. Some are prefixed with page names that affect just those types of pages, while others impact disassembly in general (throughout the product), and some affect source and file handling.

It is not important to understand every option at this time, but you should be aware that these controls are here to further refine the operation of Visual SoftICE once you become more familiar with it.

For now, leave all of the options set to their default values and take a careful look at the AutoFocusOpen setting.

◆ **AutoFocusOpen**

This parameter controls what happens when a target event occurs (or when you enter commands) that would normally open a new page with data in it. This parameter allows you to control whether Visual SoftICE opens a new page when none exists, or brings an existing page to the top of the pad it is on. This parameter also controls how Visual SoftICE will control the redirection of direct input focus upon the occurrence of target events.

The AutoFocusOpen parameter has the following options available:

◇ **Off**

Do nothing on events. Do not create or open any pages. Do not bring any existing pages to the top of their pads. Do not change the input focus. Commands you enter that should create a new page will still create those pages.

◇ **Disassembly Only**

If a Disassembly page does not exist, create one. If one already exists, bring it to the top of its pad. Make the Disassembly page the current input focus.

◇ **Disassembly Only - No Focus**

If a Disassembly page does not exist, create one. If one already exists, bring it to the top of its pad. Do not change the current input focus.

◇ **Source or Disassembly**

If Visual SoftICE can deduce a source file from symbols, then create or bring forward a Source page, otherwise create or bring forward the Disassembly page. Make the page the current input focus.

◇ **Source or Disassembly - No Focus**

If Visual SoftICE can deduce a source file from symbols, then create or bring forward a Source page, otherwise create or bring forward the Disassembly page. Do not change current input focus.

## Global Breakpoint Save and Restore Behavior

Complete the following steps to access the Visual SoftICE-specific global Breakpoint Save and Restore behavior.

- 1 Select **Preferences** from the **File** menu.
- 2 Click the **Global Settings** tab.
- 3 Expand the **Visual SoftICE** list element.
- 4 Select **Breakpoints** from the list.

This tab provides several settings related to saving and restoring breakpoints for each target you connect to. It is important that you know these settings exist, but you can leave them set to their default values for now.

## Other Global Settings

There are numerous other global settings specific to Visual SoftICE , including those controlling a startup script and other page-specific details. Feel free to explore these at your convenience.

## Per-Workspace Scripts

Complete the following steps to access the Visual SoftICE-specific per-workspace Script settings.

- 1 Select **Preferences** from the **File** menu.
- 2 Click the **Per-Workspace Settings** tab.
- 3 Expand the **Visual SoftICE** list element.
- 4 Select **Scripts** from the list.

You will notice there are a number of event-oriented script settings on this tab. When the applicable events occur, Visual SoftICE will run the script files you have defined. It is important that you know these settings exist, but you can leave them set to their default values for now.

Phew! That is more than enough for now. There are many more configuration settings, but what is important at this stage is that you know that they exist, and how to access them.



---

**Your master is now ready to start debugging. It is time to connect to your target!**

---

## Target Discovery and Connection

It is now time to find and connect to the target that you have up and running. If you configured an IP target transport, you can discover it a number of ways (otherwise follow the section below on connecting via the serial port).

### *Serial Connection*

You can connect from the **Debug** menu or via the command line to a serial target.

- ◆ Select **Connect** from the **Debug** menu.  
Here any master COM ports are enumerated for you to connect to.  
Select the COM port that you have connected to the target machine.
- ◆ On the command line, use the CONNECT command and specify the COM port to establish a connection, e.g.

```
CONNECT COM1
```

### *Network IP Discovery and Connection*

For network-based targets (those configured with a USB, PCI, CARDBUS NIC, or an IP/1394 connection) we can discover and connect to the target either by name or IP address. This can be achieved from the toolbar, or via the command line.

### **Using the Toolbar**



- 1 Click the target browse button on the **Target Select** toolbar to bring up the **Target Selector** dialog.  
Listed on this dialog are all the targets that are currently visible on the network. You may find it easier if you filter the target list by type using the filter combo-box. The filter allows you to view only those target-types your are interested in, such as only the Visual SoftICE targets.
- 2 Double-click on the target of interest, or select the target and click **Make Active** to establish a connection.

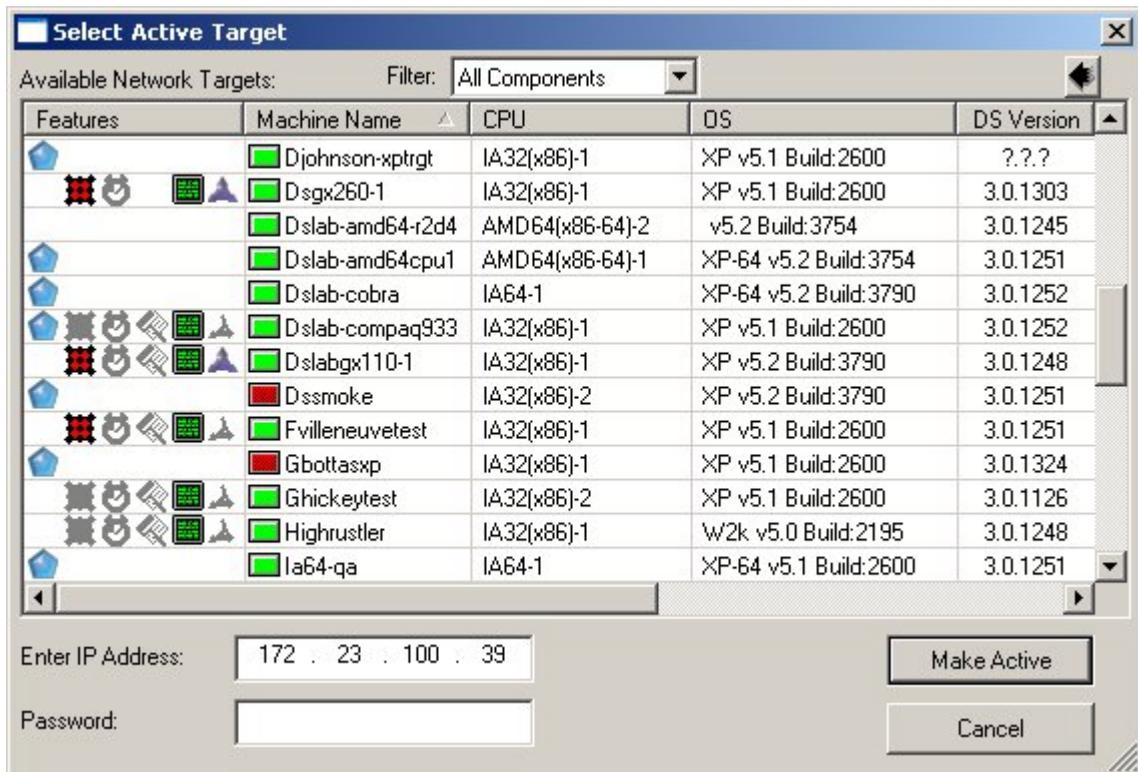


Figure 2-3. Target Selector Dialog

## Using the Command Line

- 1 Enter the NETFIND command to browse the network for a list of Visual SoftICE targets.

This command only shows Visual SoftICE targets, and not all the other DriverStudio target types.

- 2 Use the CONNECT command to establish a connection by either name or IP Address, for example:

```
CONNECT DSLAB-AMDSMP
```

Ip Address	Name	Service			Os Version
		Active	Cpu		
172.023.099.086	DSLAB-CPQ1	yes	IA32(x86)-1	Windows 2000 v5.0 Build:2	
172.023.099.155	DJOHNSON-XPTRGT	yes	IA32(x86)-1	Windows XP v5.1 Build:260	
172.023.099.209	DSLAB-ITAN1-1	yes	IA64-1	64bit (IA64) Windows 2003	
172.023.100.013	DSLAB-MADISON	yes	IA64-4	64bit (IA64) Windows 2003	
172.023.100.180	DSLAB-AMDSMP	yes	AMD64-2	64bit (AMD64/EM64T) Windo	
172.023.101.010	DSSMOKE	no	IA32(x86)-2	Windows XP v5.2 Build:379	
172.023.101.054	GBOTTASXP	yes	IA32(x86)-1	Windows XP v5.1 Build:260	
172.023.101.196	DSLAB-CPQ933	yes	IA32(x86)-1	Windows XP v5.1 Build:260	
172.023.102.119	X32A2600	yes	IA32(x86)-1	Windows XP v5.1 Build:260	
172.023.102.172	DSLAB-CPQ2	yes	IA32(x86)-1	Windows 2000 v5.0 Build:2	

Figure 2-4. Netfind Command Results



---

**Your master is now connected to your target!**

---

**Note:** Upon connection you will notice that the system gets very busy for a time. If you watch the status bar in the lower left hand field, you can see the connection phase events occurring. This includes process and image collection retrieval, and any required symbol searches (locally and via Microsoft Symbol server).

At this point, you can skip to the section “[Remote Target Control](#)” on page 27, or read on for information on connecting to targets in some special circumstances.

## **Connecting to Targets Running in Virtual Machines**

Visual SoftICE can connect to a target running in a VMWare virtual machine. The virtual machine can be located on the same physical computer as the master, or on a different computer connected to the same network. The target within the virtual machine can be configured two different ways:

- ◆ **Using a built-in NIC**

In this case, connecting to the target is just like connecting to a remote IP connected target. The master sees no difference. This is true whether the master itself is running within a VMWare virtual machine or natively in the OS.

- ◆ **Using the serial port and a named pipe**

In this case, connecting to the target is done via a named pipe, and the master must reside natively in the OS. Both the CONNECT command and the **Debug** menu support connections via named pipe.

See *Using Visual SoftICE with VMWare* in the online help for specific instructions on configuring and connecting to VMWare hosted targets.

## **Opening a Crash Dump File**

Visual SoftICE can open a crash dump file as if it were a live target, offering full debugging functionality against this data source. Since a crash dump file is a snapshot in time of a serious OS problem on a specific hardware platform, the target is considered read-only, and no live actions can be taken against it (such as setting breakpoints, stepping, etc...). The available data may also be limited, since there are various types of crash dumps, containing everything from full memory dumps to single-application, and single-thread information.

You can open a crash dump file from the **File** menu or the command line.

### **Using the File Menu**

- 1 Select the **Open** from the **File** menu.
- 2 Navigate to the dump file and open it.

## Using the Command Line

Use the OPEN command to establish a connection to the dump file. You must specify the path and filename of the location on the local machine, e.g.

```
OPEN \\dumps\recentproblem.dmp
```

## Remote Target Control

We will now prepare the remote machine for our debugging session, remotely. This section assumes that you are connected to a target machine, and the target is currently running.

**Note:** If the target is stopped, you can start it with the GO command.

Here are a few common development scenarios and quick steps available for remote control:

### *Missing Images*

In this scenario, the Master machine does not have copies of the images you need, which are on the target.

You want to retrieve image files from the remote file system, and store them in one of the image search paths we defined previously. The FGET command allows you to pull files from the remote target machine to the master file system.

**Usage:** FGET *remote-filename* *local-filename*

For example, assume we had a master side image search path that included C:\TargetImages, and we wanted to get a copy of the main kernel image from the target. The following command would accomplish this:

```
FGET %SYSTEMROOT%\System32\ntoskrnl.exe c:\TargetImages
```

**Note:** Environment variables can be used just as if you were typing into a command shell on the remote system.

**Note:** The local filename need only be the directory, and the remote filename will be used automatically.

## **Retrieving Exports**

In this scenario, you want to retrieve exports from an image on the target.

You are interested in the exports of a given image on the target, and do not have a copy of the image on the master (and may not have symbols). The EXP command will list the exports available on an image, but only if the image (and/or symbols) themselves are locally available on the master. For many reasons, you might not want to retrieve the entire image with FGET (as in the previous scenario), or you may not have the symbols. Use the GETEXP command to retrieve exports from the remote machine.

**Usage:** GETEXP [-s] *image-name*

-s indicates that subdirectories should be searched

If no path is specified, the system directory (.exe) or system\drivers (.sys) directory is used. Exports in the local cache are automatically loaded for an image when symbols are not found.

For example, assume you want all the exports for all the image files in a given directory on the target. You could issue the following command to retrieve them:

```
GETEXP -s %TEMP%\ProjectLibrary\*.*
```

## **Missing Executable**

In this scenario, the Executable to debug is not on the target yet (or is not up to date).

This is a very common situation in development, and the FPUT command is used to place any file from the master filesystem onto the remote target filesystem.

**Usage:** FPUT *local-filespec* *remote-filespec*

**For example:**

```
FPUT c:\myproject\newdriver.sys %SYSTEMROOT%\System32\Drivers
```

## **Non-Running Target Application**

In this scenario, the application on the target is not running.

This is another very common situation in development. An application on the remote target can be started with the EXEC command.

**Usage:** EXEC *application-spec*

**For example:**

```
EXEC %SYSTEMROOT%\notepad.exe
```

## **Stopping an Application on the Target**

In this scenario, the application is running, but you want to stop it.

The command to stop a running process on the remote target is KILL. This command takes a process identifier (PID) as its argument.

**Usage:** KILL *pid*

Use the PROCESS command to get the PID of the application, e.g.

```
proc
```

Name	KPEB	Pid
<hr/>		
System	0xfcda2020	0x8
SMSS.EXE	0xfcbeaa20	0x8c
CSRSS.EXE	0xfcbbc2d60	0xa4
WINLOGON.EXE	0xfc9e9020	0xa0
SERVICES.EXE	0xfc9da020	0xd4
LSASS.EXE	0xfc9d7980	0xe0
svchost.exe	0xff513a00	0x194
spoolsv.exe	0xff50bd60	0x1ac
AClient.EXE	0xff4ea5e0	0x1e4
DSRSvc.exe	0xff4e0540	0x200
svchost.exe	0xff4de020	0x20c
regsvc.exe	0xff4da940	0x230
mstask.exe	0xff4d64a0	0x248
WinMgmt.exe	0xff4c0d60	0x290
svchost.exe	0xff4b1780	0x29c
explorer.exe	0xff44b680	0x3cc
ACIntUsr.EXE	0xff424cc0	0x440
DStrayApp.exe	0xff423760	0x448
wuauctl.exe	0xff3ee660	0x3c8
notepad.exe	0xff413be0	0x474
Idle	0x8046fd60	0x0

Then use the KILL command with the appropriate PID to stop the process, e.g.

```
KILL 474
```

```
Process killed. Pid: 474
```

Alternatively, if you have the Process page open, you can right-click on a process, and select **Kill** from the pop-up menu to stop it.

## Manipulating Drivers on the Target

In this scenario, a Driver executable on the target needs to be installed, removed, or have its state changed.

Drivers and kernel services usually have more complex installation and removal needs than just file copying, and have special cases for enabling and disabling them. Therefore, Visual SoftICE has special support for these image types with the DevMgr and SvcMgr commands.

- ◆ The DevMgr (DM) command allows manipulation of the remote target device manager. Use this command for installation, enabling and disabling, state query, and removal of device drivers that are managed by the device manager.
- ◆ The SvcMgr (SM) command allows manipulation of the remote target service management system. Use this command for installation, removal, starting, and stopping of system services.

The Device-Driver page is designed specifically to aid in the remote management and exploration of these systems, and can make the development and debugging of remote drivers and services much easier.

**Note:** There are numerous remote target control commands, including those to manipulate directories, change file attributes, and search filesystems. You are encouraged to read the *Visual SoftICE Command Reference* thoroughly for details.

## General Debugging

Okay! We have a master fully configured, a target installed and configured, a connection to that target, and we have prepared the target with the right things to debug. Now it is finally time to start using the debugger. There are many books that are specifically designed to teach general debugging approaches, including tips and tricks, so this section will concentrate on introducing how to achieve standard operations within the Visual SoftICE debugger.

## Controlling the Target

Visual SoftICE is a system debugger that allows a great many things to be achieved against a running target. However, to really see the internals of the system, we must stop the target machine. This is akin to hardware debuggers which halt all processing on a system, yet retain the state of the machine at the moment in time when it was halted. Normally, you want to stop the machine at an interesting place, say in the code of your driver or application, and this is achieved via breakpoints. However, it is possible to stop the machine at any time and start peeking around at the components that make up the system.

### Flow Control

#### Stop



Stopping the target can be achieved by issuing the STOP command, by clicking the stop toolbar button, or by using a keyboard definition (the default for STOP/GO toggling is Ctrl+D).

- ◆ When you first issue a stop to a target, you may notice lots of activity in the bottom left-hand status bar field, which is called the *VSI System Activity Message* field. It is useful to watch this, especially on long processes such as symbol retrieval or large data retrieval from the target.

**Note:** Whether this field is displayed, its location, and size are all configurable as a system-wide attribute (not stored in workspace). You can configure the status bar by right-clicking on it.

- ◆ When the system stops, it automatically attempts to load any symbols for the images of the current context, which is where the OS was when it was stopped. You can list the loaded symbol tables by issuing the TABLE command. (For complete details on symbol management, see [Chapter 6, “Visual SoftICE Symbol Management”](#)).
- ◆ Depending on what pages you have open you may see numerous other reactions to the stop event. Most, if not all, pages react to the stopping of a target. Depending on your AutoFocusOpen setting, you may even see a new page open, or a page on a pad come to the top.

**Tip:** A context is defined as a process, thread, and call stack. More on this later...

 <b>Go</b>	<p>Resuming target execution can be achieved by issuing the GO command, by clicking the go toolbar button, or by using a keyboard definition (the default for GO is F5. The default for STOP/GO toggling is Ctrl+D).</p> <ul style="list-style-type: none"> <li>◆ Depending on what pages you have open you may see numerous other reactions to the resume execution event. Many pages will disable functionality they expose, and may indicate their contents are no longer valid. Some change their contents to only a subset of information which remains valid.</li> </ul>
<b>Automatic Stop Conditions</b>	<p>There are situations and configurations where the target will stop automatically, without a user specifically issuing a stop command, or hitting a breakpoint.</p>
<b>Stepping</b>	<ul style="list-style-type: none"> <li>◆ <b>Faults</b> The target will stop on ring3 (application) faults you configure it to catch. This behavior can be configured on the target through the Fault Setting utility (available from the DriverStudio Settings application), or can be remotely set from the command line via the FAULTS, I1HERE, and I3HERE commands. The settings configured directly on the target (via the Fault Settings utility) become the default each time a master connects. The settings configured remotely are only good for the duration of that session.</li> <li>◆ <b>Bug Check</b> When the OS encounters a condition that disallows further operation, the system will halt in a special way, called a bug check (also commonly referred to as a system crash, a kernel error, a stop error, blue screen, or BSOD). When Visual SoftICE is connected to a target that bug checks, it will stop before the OS fully processes this event (before the blue screen).</li> <li>◆ <b>GLOBALBREAK</b> Visual SoftICE can be configured to stop on every load in the system. Unlike standard breakpoints, this behavior is global to everything in the system (drivers, services, applications), and occurs whenever the OS loads an image for execution. This behavior can be controlled via the SET GLOBALBREAK command.</li> </ul>

#### ◆ **Source Page Stepping**

The source page loads a file for display, and matches this to an image and the current instruction pointer. It can also display source code plus per line disassembly (called mixed mode), if you so choose. In the source page, you can choose to:



- ◊ step-over (F10) the current source line.
- ◊ step-into (F11) the current source line.
- ◊ step-out-of (Shift+F11) the current source line.

If in mixed mode, you can choose instruction stepping as well. Use the options under the **Debug** group of the pop-up menu for the page, the function keys as defined in the Preference settings, or the stepping buttons on the main toolbar, to perform your stepping functions. If you are source stepping and attempt a step-into when the related source file is available, Visual SoftICE opens the file and automatically highlights the line containing current IP. If no source file is available, the step-into action behaves the same as step-over.

#### ◆ **Disassembly Page Stepping**

In the disassembly page you can choose to:



- ◊ step-over (F10) the current instruction or function call.
- ◊ step-into (F11) a function call.
- ◊ step-out-of (Shift+F11) the current function.

Use the options under the **Debug** group of the pop-up menu for the page, the function keys as defined in the Preference settings, or the stepping buttons on the main toolbar, to perform your stepping functions.

#### ◆ **Command Line Stepping**

From the command line, you can instruction step using the T command. This allows tracing one or more instructions, and even setting the start address. To source step, issue the P command. The P command can step until a return instruction is encountered. This can be dangerous though. The function context in the OS can be lost when a routine jumps to functionality that is not a call, and it never returns. You can still stop the machine when this occurs.

## Boot Control

The following commands give you control over booting on the target machine.

#### ◆ **SHUTDOWN**

Orderly shutdown of the machine, as if you selected it from the **Start** menu.

- ◆ **REBOOT**  
Orderly reboot of the machine, as if you selected it from the **Start** menu.
- ◆ **HBOOT**  
Hard reset of the machine, as if you hit the reset button on the motherboard.

## *Breakpoints*

Breakpoints are the key to any debugging session, and one of the strengths of the SoftICE family of debuggers. From within most pages of the Visual SoftICE interface, you can set and modify breakpoints with a mouse click or from a right-click context menu. You can also set any breakpoint type from the command line. Visual SoftICE has a large range of breakpoint types and this concept is covered in detail in [Chapter 7, “Using Breakpoints”](#), which you are encouraged to read.

First, we will explore a scenario providing a quick and easy way to have automatic symbol loading work for you.

### **Setting Breakpoints before an Image Loads**

Many times it is necessary or desirable to be able to set a breakpoint on an image not yet in memory, so that you will hit it very quickly when it loads.

The classic mechanism to achieve this is to preload the symbols into memory, set the breakpoint based on a symbol in the preloaded table, and then start the process or driver on the target to hit the breakpoint. This method is possible with Visual SoftICE by using the ADDSYM command to preload symbols, the BPX command to set the breakpoint, and then the GO command to instruct the target to continue booting (and load your driver), e.g.

```
ADDSYM \mysymbols\driver\mydriver.pdb
BPX mydriver!DriverEntry
GO
```

However, this approach does have a downside for developers in the “build->test->debug->rebuild” cycle on a remote machine. It is very easy to update the image on the target, and forget to load the right version of symbols on the master (or vice-versa). The ADDSYM command complicates things further, in that it overrides ALL automatic behavior, and stays loaded in the master until it is unloaded manually via the DELSYM command.

An easier approach is to use BPLOAD. This command causes a break to occur when any image matching the name provided is loaded by the operating system. For the Visual SoftICE user, this makes symbol management simple, in that the proper symbols for this image will automatically be loaded at the break, handling all the standard image matching and versioning for you. BPLOAD generally works for most situations, as it stops the image well before DriverEntry, “main” or another first entry point. Then you need only set the next breakpoint you want, and tell the system to resume execution, e.g.

```
BUPLOAD mydriver.sys  
BPX DriverEntry  
GO
```

### Per-Target Breakpoint Save and Restore

As covered above in the master setup section, you can control the behavior of saving and restoring breakpoints on a per-target basis each time you disconnect and reconnect to a target. This can be very helpful when debugging a project over several days; it gives you the ability to start up the next day right where you left off.

## ***Viewing and Editing Data***

Visual SoftICE has numerous data displays and commands that allow you to view information in the format you desire. Data from the target exists in one of three places: standard memory (including heaps), registers, or the stack. Function locals can be stored in any of these places.

### Memory

Modern operating systems use a virtual memory mechanism to provide a large range of addressable memory to each executable in the system. Each executable has perceived access to the entire memory of the machine without regard for any other executable running concurrently. The OS maps virtual memory to real (physical) memory for each executable, imposing access rules, and trying to optimize access to this map for the best execution performance of the system.

Visual SoftICE provides access to virtual or physical memory via a dedicated memory page, or via the command line using the following commands:

- ◆ **D** — Display memory (virtual or physical) in various scalar and floating point formats.
- ◆ **C** — Compare memory.
- ◆ **E** — Edit memory (virtual or physical) in various scalar and floating point formats.
- ◆ **F** — Fill memory.
- ◆ **M** — Move data.
- ◆ **PEEK** — Read physical memory.
- ◆ **POKE** — Write physical memory.
- ◆ **S** — Search memory.

**Note:** Refer to the *Visual SoftICE Command Reference* for details on using these commands.

The memory page provides a dual-display editable view of a block of memory on the target. You can format each view in a range of scalar or character types. Additionally, the memory page can format blocks of memory into structure or class shapes (think of this as a cast operation). Once in this shape, memory can be walked as through a simple list (array). Any structure pointers within a structure can be followed to the new address, and that new memory block at that address will also be formatted to a new datatype (cast). This makes exploration of known structures within a stopped machine much easier.

See “[The Memory Page](#)” on page 135 for more details.

Additionally, you can discover more about how the OS is managing virtual and physical memory using the following commands:

- ◆ **PAGE** — Display page table information.
- ◆ **PHYS** — Display all virtual addresses that correspond to a physical address.

## Registers

The registers of the target machine are available for display whenever the machine is stopped. Some processor architectures have a few registers, while others have hundreds. To help clarify and find the registers you are interested in, Visual SoftICE has a concept called register groups, which are collections of registers accessible by name. The available groups for a given target can be viewed by issuing the RG command, or from the Registers page context menu.

Visual SoftICE provides access to all registers via the command line, using the following commands:

- ◆ **MSR** — Enumerate\Read\Write\Discover Model Specific Register(s).
- ◆ **R** — Display or change register or groups of register values.
- ◆ **RG** — Display available register groups.

The register page provides formatted display and editing of register values and can be customized to user preferences. See “[The Registers Page](#)” on page 151 for more details.

## Locals

Locals are the variables of a given function, available while execution is active within the function code block. The storage for locals, configuration of local descriptor information, and data differs according to the processors and operating systems. Most modern compilers try to optimize code to use registers for local storage whenever possible. Success of this approach varies widely based on situation and platform.

Visual SoftICE provides a means to display locals by issuing the LOCALS command. Additionally a dedicated locals page is provided to display and edit local variables. It should be noted that local variables can be known by the system during function prolog and epilog, yet be unavailable for display or editing. In these cases, the locals may be listed, but their current values may be displayed as an error message.

See “[The Locals Page](#)” on page 132 for more details.

## **Contexts and the Call Stack**

The word “Context” is used throughout the documentation of the product and user interface, referring to a collection of three pieces of data: a process, one of its threads, and a call stack frame.

A process can have any number of threads, but must have at least one. A stopped machine will have a call stack associated with each thread. This stack is made up of frames, which describe the function, location within that function, and other function attributes (such as locals, registers, etc...) available at the point in time the machine was stopped. A call stack is an ordered collection of stack frames (one function calling another through the thread execution). Usually the topmost frame (frame zero) is the default for a thread.

When the term “Stopped Context” is used, it refers to the process and thread the current CPU was in when the machine was stopped. On machines with multiple CPUs, each CPU has its own stopped context information.

When the term “Current Context” or “UI Context” is used, it means whatever process, thread, and stack frame you have currently selected. When a target is first stopped, the “Current Context” is equal to the “Stopped Context”. You can change from the stopped context to other threads, frames, or processes, using either the command line or the context bar.

Most pages display data relative to the current context. This includes commands in the command page, unless specifically overridden. Some pages (e.g. Stack and Locals) allow a specific context to be selected, and to *not* follow the current context as it changes. These pages provide their own context selection bar when in this mode.

**Note:** These modes are not on by default.

The stack of the current context can be displayed from the command line by issuing the STACK command. You can also use this command to show stack information for threads other than the current context.

Visual SoftICE also offers a stack page displaying the complete call stack of the current context, which supports a number of actions and additional functionality.

See “[The Stack Page](#)” on page 167 for more details.

## Built-In Commands

Visual SoftICE contains a wealth of commands for exploration of the system. This includes hardware, operating system structures, objects and constructs, system state, and shared resources. One of the real differences between the SoftICE family and other debuggers is the amount of built-in information support available, both for hardware and the operating system.

The most important command in the product is the HELP command (or H), which will list all commands by category and provide descriptive and input help on a particular command. You can even use the HELP command with wildcards to narrow the list of commands to only those that match. There is no better way to quickly find out what commands are available.

The most important operating system commands to get started with are PROCESS, DRIVER, IMAGE/MOD, THREAD, and STACK. These commands show you the top level entities in the system.

For driver and kernel level development, Visual SoftICE contains commands for finding and decoding common kernel objects and links (APC, ARBITER, DEVICE, DEVNODE, DPC, FOBJ, SYMLINK, and TIMER), interrupts (IDT/IT, INTOBJ, and IRP), and kernel internal structures (KEVENT, KMUTEX, KJOBCT, KSEM, and OBJDIR).

For those interested in hardware state and exploration, Visual SoftICE provides the following built-in commands (CPU, GDT, IDT/IT, PCI, DEVNODE, DEVICE, PAGE, RG, R, and MSR).

For those interested in data stream and packet decoding, Visual SoftICE provides the following built-in symbol commands (IRB, SRB, URB, and PACKET).

There are over 250 command forms built into Visual SoftICE natively, as well as support for many commands for those familiar with KD (in a separate dialect). You are encouraged to explore the online and printed *Visual SoftICE Command Reference* in addition to the command line help available via the HELP or H command.



**Good luck, and happy debugging!**

---



# Chapter 3

## Visual SoftICE Target Transports



- ◆ Visual SoftICE Target Transport Overview
- ◆ Serial Target Transport
- ◆ Dedicated PCI/CardBus Ethernet Network Interface Cards (NICs)
- ◆ Universal PCI/CardBus Ethernet NICs
- ◆ OHCI/UHCI USB Host Controller and USB NIC
- ◆ Virtual NIC Driver (optional)
- ◆ 1394 (Firewire)

### Visual SoftICE Target Transport Overview

Visual SoftICE offers several debug transport choices, all of which have their own advantages and disadvantages. Some transports might not be available to you depending on your hardware configuration, or a particular problem you are trying to debug. The only time the master does not use a distinct transport is when it opens a crash dump file.

Visual SoftICE supports the following transport connections between the master and a live target:

- ◆ Serial
- ◆ Dedicated PCI or CardBus Ethernet Network Interface Card (NIC)
- ◆ Universal PCI or CardBus Ethernet NIC
- ◆ OHCI/UHCI USB Host Controller and USB NIC
- ◆ Virtual NIC Driver (optional)
- ◆ 1394 (Firewire)

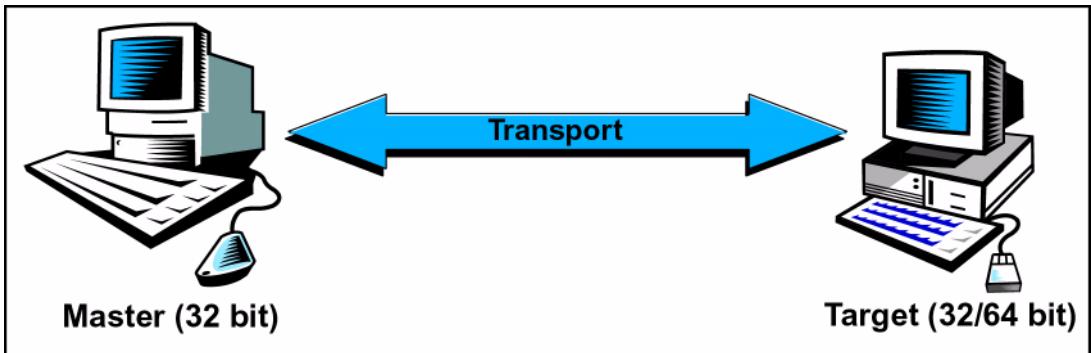


Figure 3-1. Transport Overview

For more information and tips on configuring your transports, refer to the Transport Configuration help.

## Serial Target Transport

Standard serial cable into dedicated serial port or PCI card serial port on the target. Debugging using the serial interface is quite slow, but probably the easiest to install and configure. Legacy free PC's, however, may not have serial hardware and may require a different transport or an add-in PCI card. If serial is your required transport and you do not have serial hardware built in, you can install the SIIG CyberSerial PCI card.

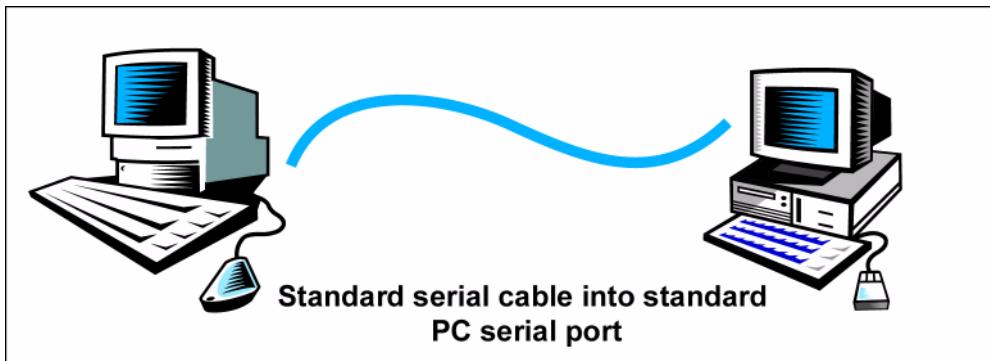


Figure 3-2. Serial Transport

## **Requirements and Characteristics**

**Target:** UART 8250/16450/16550 or SIIG CyberSerial PCI card

**Master:** Any serial adapter

**Connection Cable:** Null modem cable

**Speed:** ~ 10 KB/sec

**Driver:** SISERIAL.SYS

## Dedicated PCI/CardBus Ethernet Network Interface Cards (NICs)

This transport is probably the most flexible choice you have. You can connect the target NIC into your LAN and have access to your machine from anywhere on your network. In addition, you can use a crossover cable to connect the master and target NICs together, thus completely restricting access to your target machine. The primary drawback of this debug transport affects you only under the circumstances where you have only one NIC installed in your target machine. After installing our driver for the single NIC, you will not be able to use it for normal windows networking (unless you configure a Virtual NIC (VNIC)).

**Note:** Please be advised that the VNIC (shared mode) is not recommended, and greatly impacts performance of the target networking.

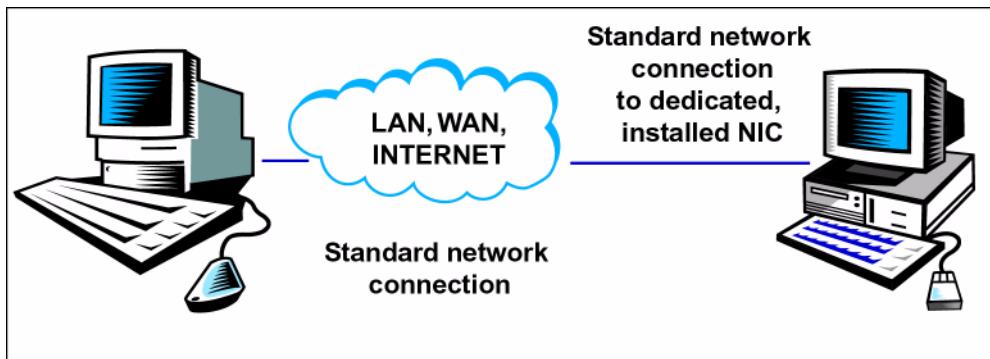


Figure 3-3. Dedicated PCI and CardBus Ethernet NIC

## **Requirements and Characteristics**

**Target PCI NIC:** One of the following supported PCI NIC cards:

- ◊ Based on RTL8029 chip
- ◊ Based on RTL8139 chip
- ◊ Intel EtherExpress E100x series
- ◊ 3Com 3C90X series
- ◊ Based on AMD PCNET chip
- ◊ Based on Lite-On chip

**Target CardBus NIC:** One of the following supported CardBus NIC cards:

- ◊ Based on RTL8139 chip
- ◊ Intel EtherExpress E100x series

**Master:** Any Ethernet NIC

**Connection Cable:** Standard Ethernet cable or Crossover Ethernet cable

**Speed:** ~ 1-2 MB/sec

**Driver:** One of the following supported drivers:

- ◊ SI8029.SYS
- ◊ SI8139.SYS
- ◊ SIE100.SYS
- ◊ SI3C90X.SYS
- ◊ SIPCNET.SYS
- ◊ SILITE.SYS
- ◊ SIVNIC.SYS (optional)

## **Universal PCI/CardBus Ethernet NICs**

This transport may be right for you if you have a PCI or CardBus NIC that is not directly supported by our dedicated drivers (refer to the “Dedicated PCI/CardBus Ethernet Network Interface Cards (NICs)” on page 43). The major disadvantages to using this transport configuration are:

- ◆ A larger memory footprint on the target than you get with the dedicated NIC drivers.
- ◆ This driver loads later than a dedicated driver, because it depends on the NDIS subsystem. For situations where early-stopping of the target is critical, you should use a dedicated NIC driver.

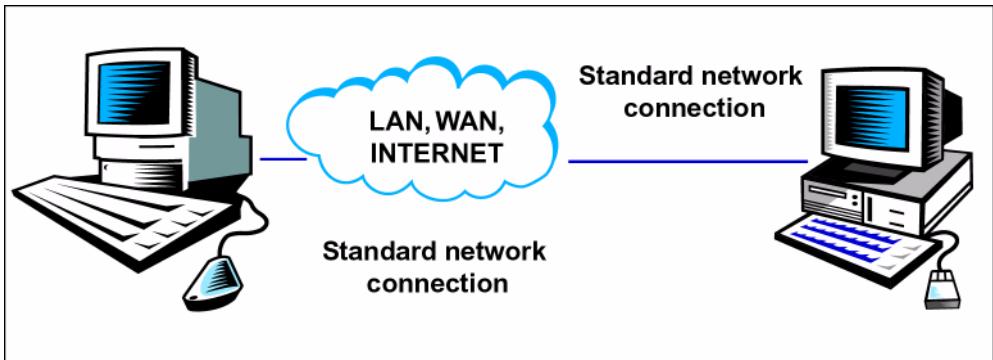


Figure 3-4. Universal PCI and CardBus Ethernet NIC and Optional Virtual NIC

### **Requirements and Characteristics**

**Target:** Most PCI Ethernet NICs. Any of the supported CardBus NICs (listed in “[Dedicated PCI/CardBus Ethernet Network Interface Cards \(NICs\)](#)” on page 43).

**Master:** Any Ethernet NIC

**Connection Cable:** Standard Ethernet cable or Crossover Ethernet cable

**Speed:** ~ 1-2 MB/sec

**Drivers:** One of the following supported drivers:

- ◊ SIDN.SYS
- ◊ SINIC.SYS
- ◊ SIVNIC.SYS (optional)

### **OHCI/UHCI USB Host Controller and USB NIC**

This transport may be the only one available to you in some situations; for example, if your target is a legacy free laptop without a cardbus slot. In this case you do not have a serial port or any NICs, thus using an external USB NIC is your only option. The disadvantage of this transport solution is that we need to install our own driver, not only for the USB NIC, but also for the USB controller. This means that you cannot plug in any other USB devices under this host controller, even if it has more than one port. A solution to this limitation would be to install an additional USB host controller, for devices like a USB keyboard or mouse, if needed.

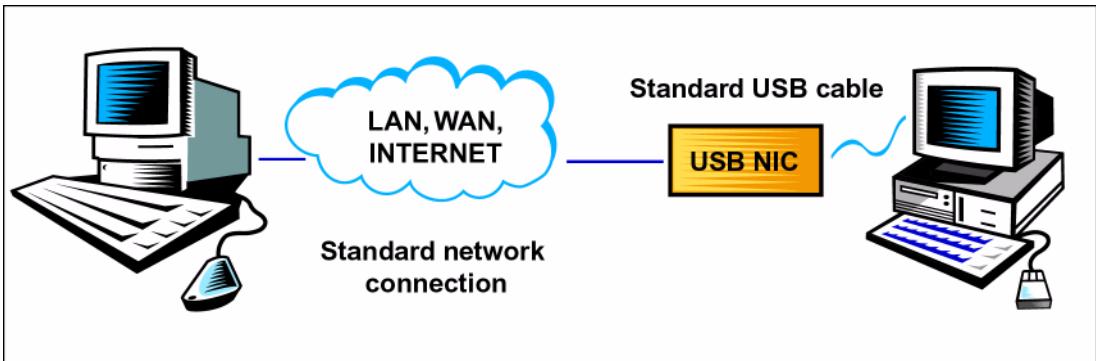


Figure 3-5. OHCI/UHCI USB Host Controller and USB NIC

### **Requirements and Characteristics**

**Target:** USB 1.1 OHCI\UHCI host controller and Admtek based USB Ethernet NIC

**Master:** Any Ethernet NIC

**Connection Cable:** Standard Ethernet cable or Crossover Ethernet cable

**Speed:** ~ 100-200 KB/sec

**Drivers:** SIUSB.SYS

## 1394 (Firewire)

We use IP over 1394 as a protocol between the master and a target, thus the target machine appears to be on the local IP network. This transport solution is not shared, and once you install our driver for the 1394 host controller you cannot use it for additional 1394 devices.

**Note:** 1394 Targets are not visible to any other machine except the one that is physically connected.

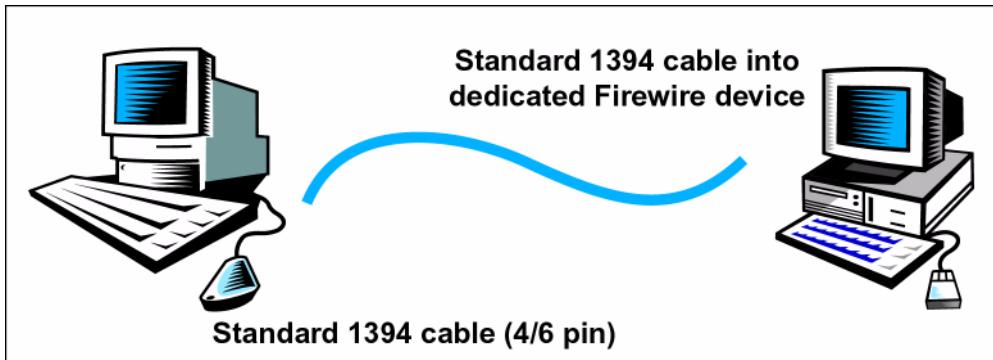


Figure 3-6. 1394 (Firewire)

### ***Requirements and Characteristics***

**Target:** OHCI 1394 host controller

**Master:** Any 1394 host controller and Windows XP or later for OS.

**Connection Cable:** standard 1394 cable

**Speed:** ~ 2 MB/sec

**Driver:** SI1394.SYS



# Chapter 4

## Overview of the Visual SoftICE User Interface



- ◆ The Visual SoftICE User Interface Overview
- ◆ User Interface Preferences
- ◆ Other User Interface Attributes and Features

### The Visual SoftICE User Interface Overview

Various user interface pages support a number of ease of use features, including the ability to:

- ◆ Set and follow a main context for all pages, and overriding page-specific contexts for Locals and Stack pages
- ◆ Use standard Cut, Copy, and Paste operations
- ◆ Drag and Drop pages and items (Single element, and element collections between pages)
- ◆ Save the contents of a page to a file
- ◆ Print and Print Preview the contents of a page
- ◆ Change page modes (snapshot, live, and manual)
- ◆ Use special Command page features, such as:
  - ◇ The C++ line comment symbol to append comment text to input (or provide comment lines in scripts)
  - ◇ Redirection of command output to other pages for display handling (including creation of new pages, on demand)
  - ◇ Automatic redirection of appropriate commands from the command page to other existing pages
- ◆ Create custom keyboard definitions
- ◆ Load a custom workspace
- ◆ Create and save a custom workspace
- ◆ Create and save script file paths
- ◆ Discern Target State and Breakpoint Types via special icons

## Workspaces

DriverWorkbench uses the workspace concept to save and restore your preferred working environment between sessions. This includes the layout of the screen (pads and pages), whether the frame was maximized or not, keyboard definitions and input rules you have defined, your printer preferences, toolbar locations and configurations, and page, plugin, and debugger specific configuration information you have customized.

The customization of your workspace preferences deserves further explanation:

- ◆ Each page instance can store specific configuration information in the workspace.
- ◆ Thus if you have two Debug Message pages, each with different string filters, these filters will be saved to the workspace individually, and restored with the workspace.
- ◆ Almost all multiple-instance pages support this behavior.

The debugger stores configuration information in the workspace as well, including overrides to paths, workspace event script names, event log configuration settings, preferred display sizes for various windows you choose to display, and more. To edit configuration settings stored in the workspace, launch the preferences dialog from the **File** menu (**File->Preferences->Per-Workspace Settings**).

## Keyboard Definitions

DriverWorkbench allows you to create custom keyboard definitions, or edit existing keyboard definitions, and save them in your workspace file. DriverStudio components that work within the DriverWorkbench environment provide default keyboard definitions on a workspace level and sometimes on a page-component level. You cannot edit default definitions, however you can override them by adding your own custom keyboard definitions for the same keystroke. The following definitions are global to the DriverWorkbench frame and you cannot override them:

Table 4-1. Global Keyboard Definitions

Keystroke	Definition
CTRL-TAB	Next page on a pad.
SHIFT-CTRL-TAB	Previous page on a pad.
CTRL-F4	Close the current pad.

For more details about keyboard settings, refer to the online help.

## Workspace Uses

Workspaces can be used in many different ways, for example:

- ◆ Many of our users create and save a workspace configuration for different modes of work they do, including application debugging, kernel debugging, and/or the development phase versus the debugging and optimization phases.
- ◆ Many users create one workspace per target, allowing them specific views for the types of problems they are debugging on different hardware and operating systems.
- ◆ Many users create specific workspaces for their various development and debugging master machines, including layout of windows on multi-monitor setups in one case, and single monitor setups in another.

For instructions on loading and saving a custom workspace, refer to “[Workspace Save and Load](#)” on page 74.

In addition to loading and saving custom workspaces manually, you can configure several automatic workspace loading and saving behaviors in DriverStudio. For more information on configuring the automatic workspace saving and loading behaviors, refer to “[Workspace Save and Load](#)” on page 74.

## Browsing Available Targets

You can browse the available targets by clicking the Target Browser toolbar icon, or by selecting **Browse** from the Target menu.

### Toolbar Icon



Figure 4-1. Target Browser Button

## Target Menu



Figure 4-2. Target Menu

## Target Browser Utility

Once you open the Target Browser Utility, select **Visual SoftICE** using the **Network Target Filter** drop-down list to display only Visual SoftICE targets. You can also find a target by entering its IP Address in the field at the bottom or the utility. Select a target from the list by clicking on it and then clicking the **Make Active** button.

You can refresh the target list at any time by clicking the **Refresh Target List** button.

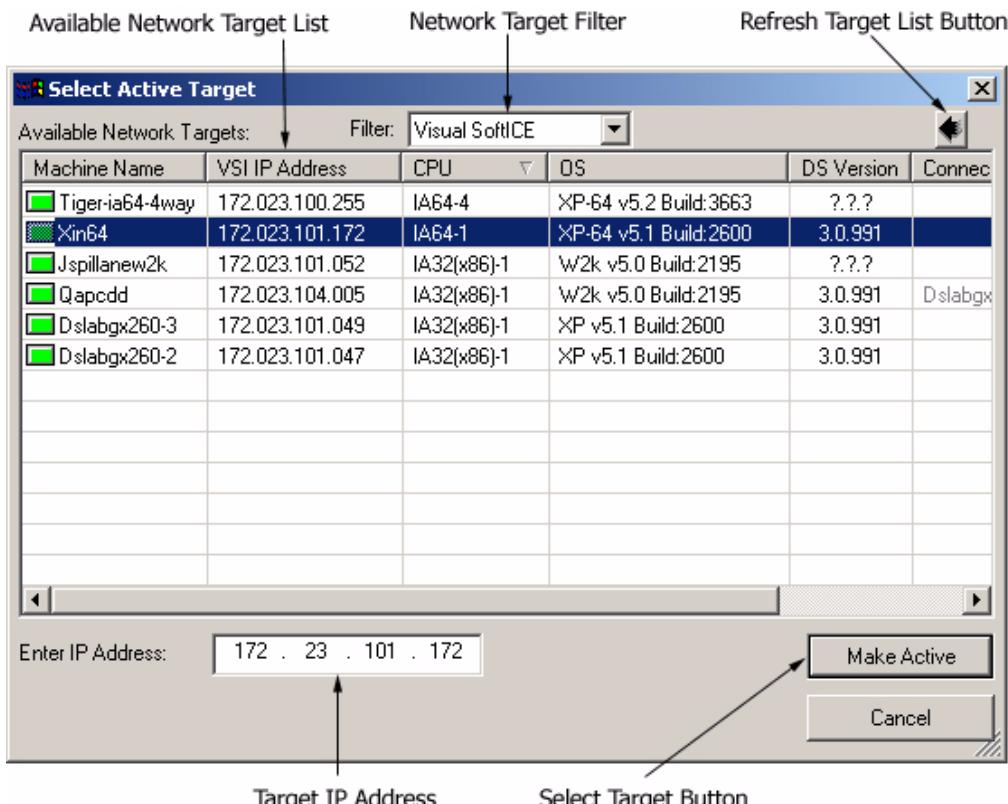


Figure 4-3. Active Target List

## Connecting to Targets

The Debug menu is a Visual SoftICE-specific menu that provides you with the functionality to connect, disconnect, and start the debugger on targets. The Debug menu is dynamic, in that:

- ◆ If you are currently connected to a target, Disconnect becomes available on the menu. Also refer to the DISCONNECT command in the *Visual SoftICE Command Reference*.

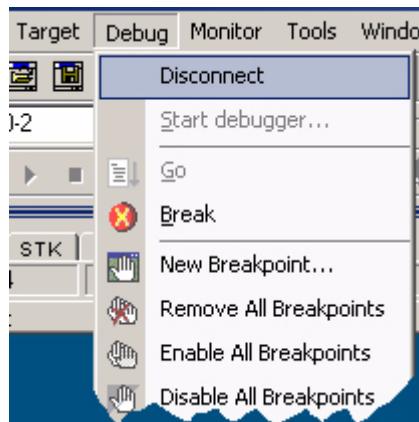


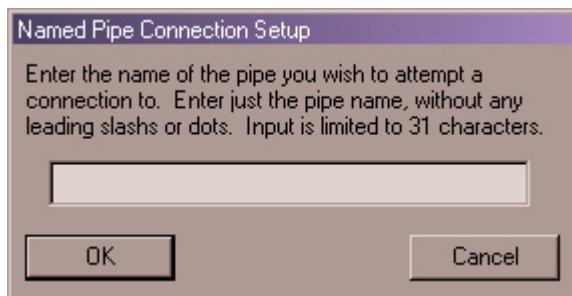
Figure 4-4. Active Disconnect Menu Option

- ◆ If you have previously connected to a target, or selected a target by highlighting it in the Target Browser utility, that target appears as a valid option in the **Connect** menu, along with any available COM ports on the system, and the option to connect via Named Pipe.



Figure 4-5. Connect Menu Options

- ◆ Visual SoftICE also provides an option for you to connect via named pipe within a virtual OS product that allows inter-session communication between virtual OS sessions. In such scenarios, Visual SoftICE can run both the master and target on a single machine. To connect via named pipe, select **Named Pipe** from the **Connect** sub-menu, and Visual SoftICE opens the **Named Pipe Connection** dialog.



**Figure 4-6.** Named Pipe Connection Dialog

**Note:** See also the CONNECT and WCONNECT commands in the *Visual SoftICE Command Reference*.

## Visual SoftICE Icons

Visual SoftICE makes use of many different icons within the GUI to access the many pages, and to identify such things as breakpoints (type and state), target state, and current context (target IP and UI context). The following sections will help you identify the various icons and their meanings.

### Page Access

The following icons are used to access Visual SoftICE pages.

**Table 4-1.** Page Access Icons

Icon	Description
	This icon accesses <a href="#">The Breakpoint Page</a> .
	This icon accesses <a href="#">The Command Page</a> .
	This icon accesses <a href="#">The Debug Message Page</a> .

**Table 4-1.** Page Access Icons (Continued)

Icon	Description
	This icon accesses <a href="#">The Device-Driver Page</a> .
	This icon accesses <a href="#">The Disassembly Page</a> .
	This icon accesses <a href="#">The Event Page</a> .
	This icon accesses <a href="#">The Locals Page</a> .
	This icon accesses <a href="#">The Memory Page</a> .
	This icon accesses <a href="#">The Process List Page</a> .
	This icon accesses <a href="#">The Registers Page</a> .
	This icon accesses <a href="#">The Source Page</a> .
	This icon accesses <a href="#">The Stack Page</a> .
	This icon accesses <a href="#">The Text Scratch Page</a> .
	This icon accesses <a href="#">The Watch Page</a> .

## Breakpoints

The following icons are used within various pages to indicate the different breakpoint types and states.

Table 4-1. Breakpoint Icons

Icon	Description
	This icon represents a Fixed Address breakpoint that has been hit.
	This icon represents a Fixed Address breakpoint that has not been hit.
	This icon represents an Image Load breakpoint that has been hit.
	This icon represents an Image Load breakpoint that has not been hit.
	This icon represents an Image Relative breakpoint that has been hit.
	This icon represents an Image Relative breakpoint that has not been hit.
	This icon represents an I/O breakpoint that has been hit.
	This icon represents an I/O breakpoint that has not been hit.
	This icon represents a disabled breakpoint.

## Target State

The following icons are used to indicate the state of the target.

Table 4-2. Target State Icons

Icon	Description
	This icon indicates the target is in a sleep (power management) mode.
	This icon indicates the target is running.
	This icon indicates the target is stopped.
	This icon indicates the target is stopped due to hitting a breakpoint.
	This icon indicates the target is stopped due to a fault.
	This icon indicates the target is instruction or source stepping.
	This icon indicates the target is stopped due to a BugCheck.
	This icon indicates the target is in an unknown state.

## Current Context

The following icons are used within various pages to identify such things as the target IP and UI context.

Table 4-3. Current Context Icons

Icon	Description
	This icon indicates the line you are on when you do a Goto in the Source or Disassembly page.

**Table 4-3.** Current Context Icons (Continued)

Icon	Description
	This icon indicates the Current UI Context (Process, Thread, or Stack Frame depending on the display it appears in).
	This icon indicates the Current UI Context when a Snapshot of the information was taken.
	This icon indicates the Current Target Instruction Pointer (IP).
	This icon indicates the Current Target IP when a Snapshot of the information was taken.
	This icon indicates the Current Target IP location, and that the next step will move it downward (forward or increasing) into memory.
	This icon indicates the Current Target IP location, and that the next step will move it downward into memory, at the time when a Snapshot of the information was taken.
	This icon indicates the Current Target IP location, and that the next step will move it upward (backward or decreasing) into memory.
	This icon indicates the Current Target IP location, and that the next step will move it upward into memory, at the time when a Snapshot of the information was taken.
	This icon indicates the line where the Current IP will be upon executing the next instruction. Displayed when the Current IP and the next IP locations are available in the same screen (Source or Disassembly pages).

## Symbol Table State

The following icons are used to indicate the state of the symbol table.

**Table 4-4.** Symbol State Icons

Icon	Description
	This icon indicates that on-demand symbol loading is enabled (See SET SYMTABLEAUTOLOAD in the <i>Visual SoftICE Command Reference</i> ).
	This icon indicates that on-demand symbol loading is disabled (See SET SYMTABLEAUTOLOAD in the <i>Visual SoftICE Command Reference</i> ).
	This icon indicates that searching for symbols via symbol servers is enabled (See SET SYMSRVSEARCH in the <i>Visual SoftICE Command Reference</i> ).
	This icon indicates that searching for symbols via symbol servers is disabled (See SET SYMSRVSEARCH in the <i>Visual SoftICE Command Reference</i> ).
	This icon indicates that all tables loaded without warnings (See TABLE in the <i>Visual SoftICE Command Reference</i> ).
	This icon indicates that at least one table loaded with a warning condition (See TABLE in the <i>Visual SoftICE Command Reference</i> ).
	This icon indicates that at least one table was not loaded (See TABLE in the <i>Visual SoftICE Command Reference</i> ).

## Source Page Symbol State Icons

The following icons are used to indicate the state of the symbols for a source file loaded into a source page. They only have the following meanings when in the status bar of a source page.

**Table 4-5.** Source Page Symbol State Icons

Icon	Description
	This icon indicates that symbols are loaded for the indicated image, and that the source file is part of that image and has line information available.

**Table 4-5.** Source Page Symbol State Icons (Continued)

Icon	Description
	This icon indicates that symbols are loaded for the indicated image, and that the source file is part of that image, but there is no line information available.
	This icon indicates that the source file is not part of any image in the current/active process.

## Cursor Types

The following icons are used to indicate various states of the cursor.

**Table 4-6.** Cursor Icons

Icon	Description
	This icon indicates you have selected a single item and are dragging it from a page.
	This icon indicates you have selected multiple items and are dragging them from a page.
	This icon indicates the cursor is over a window that does not allow dropping of the data you are dragging.
	This icon indicates the cursor is over a window that allows dropping the single item being dragged into it.
	This icon indicates the cursor is over a window that allows dropping the multiple items being dragged into it.
	This icon indicates the cursor is over the breakpoint column, on the Source or Disassembly page.
	This icon indicates the cursor is over the bookmark column.

## About the Status Bar

Visual SoftICE provides main status bar fields, an optional one of which displays the status of the symbol engine. The status bar provides the following information:

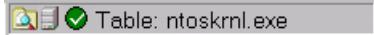
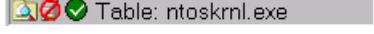
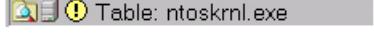
- ◆ Whether on-demand loading of symbols is enabled or disabled.
- ◆ Whether using the symbol server is enabled or disabled.
- ◆ Whether any symbol table failed to load, or loaded with a warning, or all symbol tables loaded successfully.
- ◆ The name of the current default symbol table

The Source page also has a page-specific symbol status bar.

## Status Bar Examples

The following examples show various states of the symbol engine, as indicated by the status bar:

Table 4-7. Visual SoftICE Status Bar Examples

Status Bar	Meaning
 Table: ntoskrnl.exe	Auto-load is enabled, the symbol server is enabled, and all tables loaded successfully.
 Table: ntoskrnl.exe	Auto-load is disabled, the symbol server is enabled, and all tables loaded successfully.
 Table: ntoskrnl.exe	Auto-load is enabled, the symbol server is disabled, and all tables loaded successfully.
 Table: ntoskrnl.exe	Auto-load is disabled, the symbol server is disabled, and at least one table failed to load.
 Table: ntoskrnl.exe	Auto-load is enabled, the symbol server is enabled, and at least one table loaded with a condition.

## Symbol Settings

You can access various symbol settings by right-clicking on the status bar, or via the Debug menu. These settings allow you to:

- ◆ Enable or Disable Automatic Table Loading  
This is equivalent to using the SET SYMTABLEAUTOLOAD command.
- ◆ Enable or Disable Symbol Server Searching  
This is equivalent to using the SET SYMSRVSEARCH command.
- ◆ Manage Symbol Tables  
Allows you to open the Symbol Tables Utility.

## Source Page Status Bar

The source page has its own page-specific status bar that uses icons to display the state of symbols for the source file loaded into that page. Icons indicating successful loading of symbols are followed by the image name.

## About the Context Bar

Visual SoftICE provides a main context bar which displays the current context the debugger is following, and allows you to change the current context using its drop-down list. The Locals page and Stack page also each have page-specific context bars available. The page-specific context overrides the main context. If you activate a page-specific context bar, you change the context that page is following. Both the main context bar and any page-specific context bars are comprised of three controls and a button.

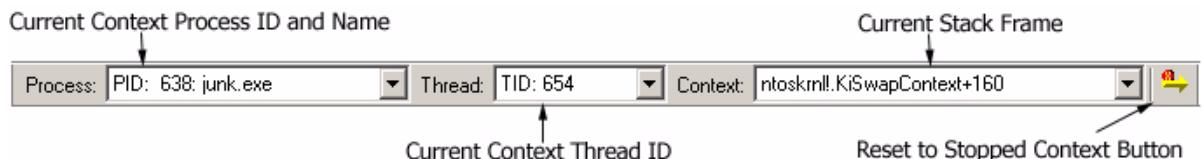


Figure 4-1. Visual SoftICE Context Bar

## Context Bar Controls

The main context bar is comprised of controls allowing you to track the current context right down to the stack frame level. The page-specific context bars are comprised of a sub-set of two of the controls, depending on which page the context bar belongs to. The controls and their functions are as follows:

Table 4-8. Context Bar Controls

Control	Description
Process	Displays the <b>Process ID and Name</b> . You can use this control to select a new process from the drop-down list of available processes. This is equivalent to changing the current process from the Process List page or via the ADDR command.
Thread	Displays the <b>Thread ID</b> . You can use this control to select a new thread from the drop-down list of available threads, if the current process is multi-threaded. This control is the most convenient way to change threads within a multi-threaded process.

Table 4-8. Context Bar Controls (Continued)

Control	Description
Context	Displays the current <b>Stack Frame</b> , including any symbolic representation of the frame instruction pointer. You can use this control to select a new stack frame from the drop-down list of available stack frames. This control is the most convenient way to change the current stack frame.
Reset to Stopped Context Button	Switches you back to the context that was current when you stopped the target. This button is located on the main context bar, and becomes active when you stop the target, and switch contexts. Clicking the button does not restart the target.

## About the Debug Toolbar

This dockable toolbar supports the most common debugger actions, including stop and go, stepping, breakpoints, access to fault, thread, and module lists, and symbol and source operations.



Figure 4-2. Debug Toolbar

## Context Menus

Most Visual SoftICE pages have a context or pop-up menu providing various action items associated with it. The action items for any page will vary depending on the page type, contents, and focus within that page. The intent of the page-specific context menu items is to give you faster and easier access to those actions that you would most likely be executing within that page.

For example, if you have focus on a specific address in the Stack page, you may have action items to view the source or disassembled code at that address, set a breakpoint at that address, or view memory at that address.

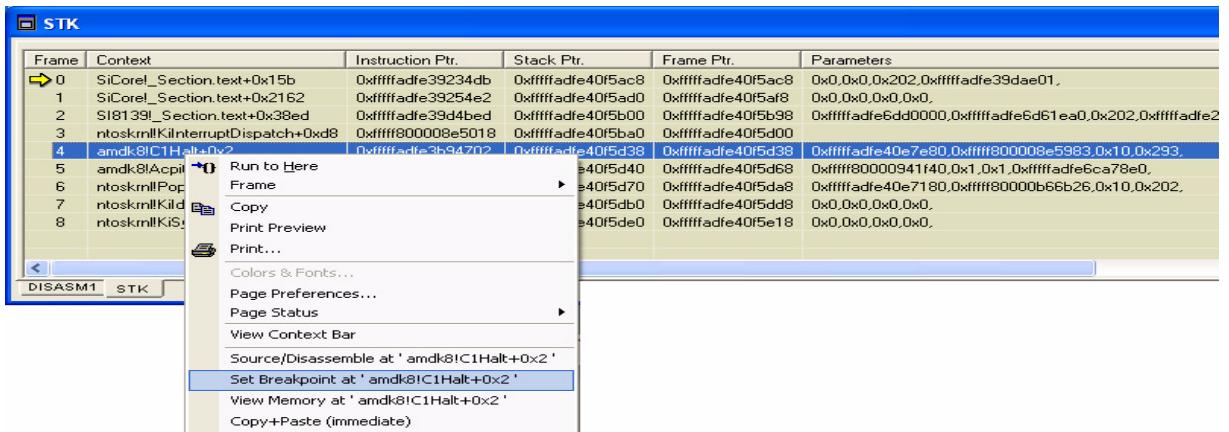


Figure 4-3. Page and Content-specific Context Menu

## User Interface Preferences

You can set global, per-workspace, toolbar, status bar, font, color, and keyboard preferences using the Preferences dialog. Preferences are set on a global basis for all of DriverStudio and its components, as well as on a per-workspace basis for specific components. This dialog provides the following groups of preferences you can set.

- ◆ Global Settings
- ◆ Per-Workspace Settings
- ◆ Keyboard
- ◆ Toolbars & Status Bar
- ◆ Fonts & Colors

The Per-Workspace settings are provided for users who wish to use different workspaces for different debugging conditions, targets, or target types. Available per-workspace settings fall into two categories: Paths and Scripts. The Per-Workspace paths values prepend or replace the global path values of the same name.

## Global Settings

Global Settings are global to the executable, and do not change with different workspaces.

### General Global Settings

When you select an element, the settings for that element are listed along with any value they may currently have. If you click on a setting, a description of that setting, any ranges (if applicable), and other rules on data entry (if applicable), are displayed in the Setting Explanation window below the list. The types of behavior controlled by these settings ranges from path definitions, global page properties for specific pages, and workspace saving and loading behavior.

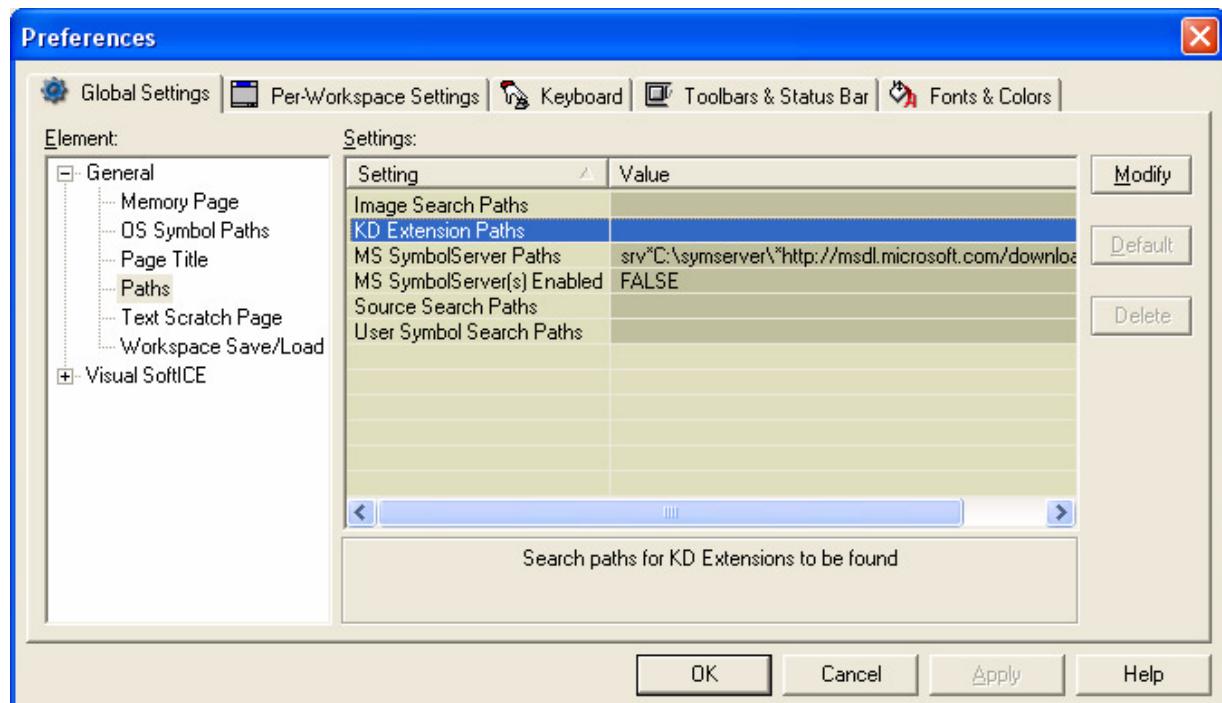


Figure 4-4. General Global Settings

## Visual SoftICE Global Settings

Visual SoftICE global settings act on the Visual SoftICE component alone, and do not change with different workspaces. When you select an element, the settings for that element are listed along with any value they may currently have. If you click on a setting, a description of that setting, any ranges (if applicable), and other rules on data entry (if applicable) are displayed in the Setting Explanation window below the list.

Most of the settings deal with page properties, with the exception of Event Handling, Paths, and Scripts.

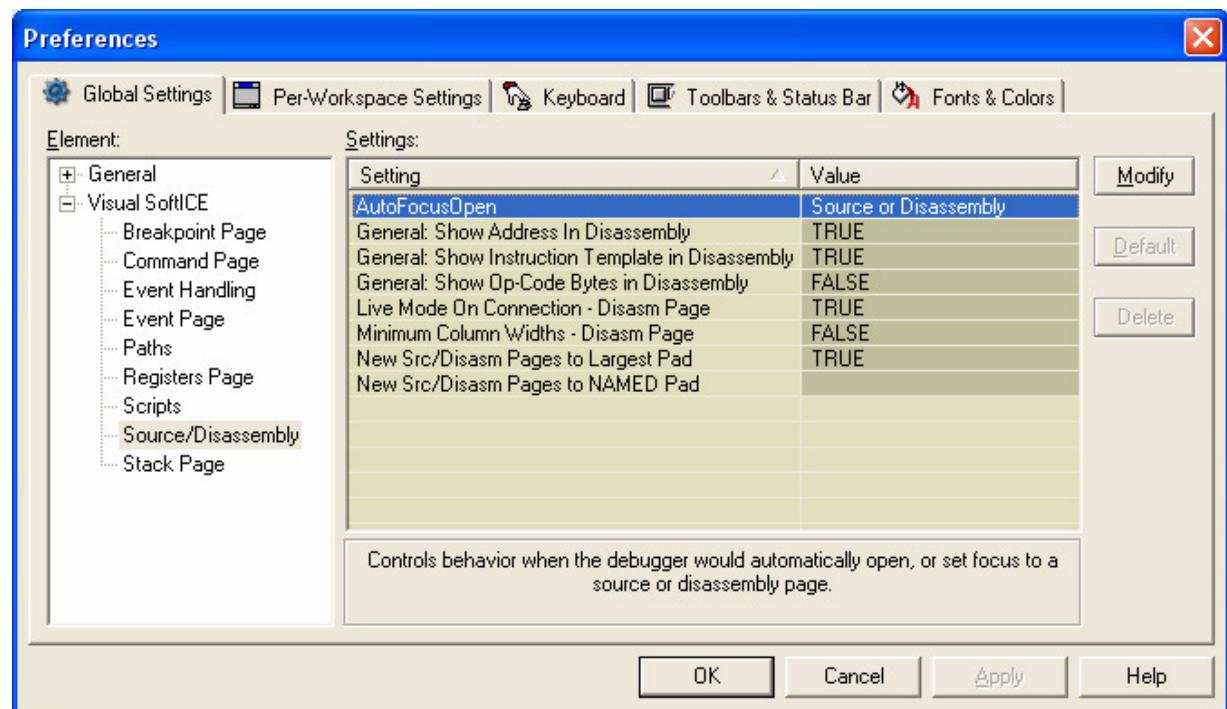


Figure 4-5. Visual SoftICE Global Settings

## Per-Workspace Settings

Per-Workspace settings are provided for users who wish to use different workspaces for different debugging conditions, targets, or target types. The properties you set here are specific to the current workspace. Available per-workspace settings fall into two categories: Paths and Scripts.

## Paths

The Paths element in the Per-Workspace Settings tab allows you to define specific paths to prepend or override the global path settings. These paths, along with the global settings they prepend or replace, define the Visual SoftICE configuration. The current configuration can be seen by issuing the SET command.

If you click on a path type, a description of that path, any ranges (if applicable), and other rules on data entry (if applicable), are displayed in the Setting Explanation window below the list.

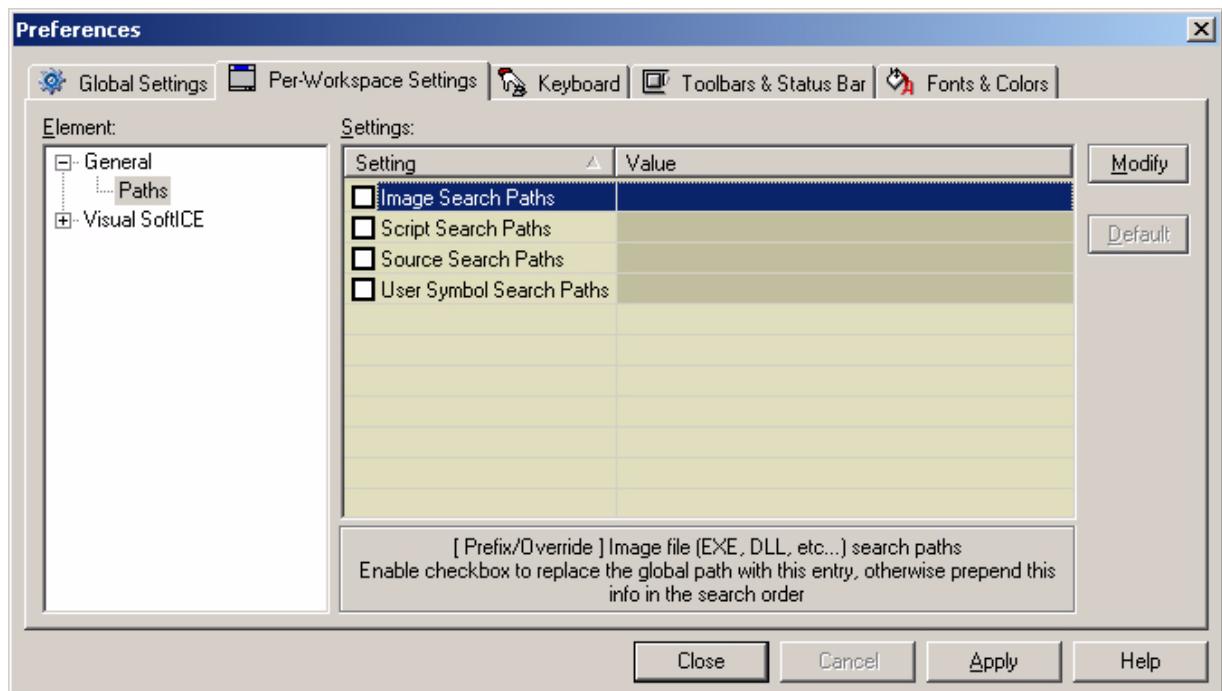


Figure 4-6. General Per-Workspace Path Preferences

To edit the value of a path, click on its **Value** field. Visual SoftICE opens the **Path List Edit** utility.

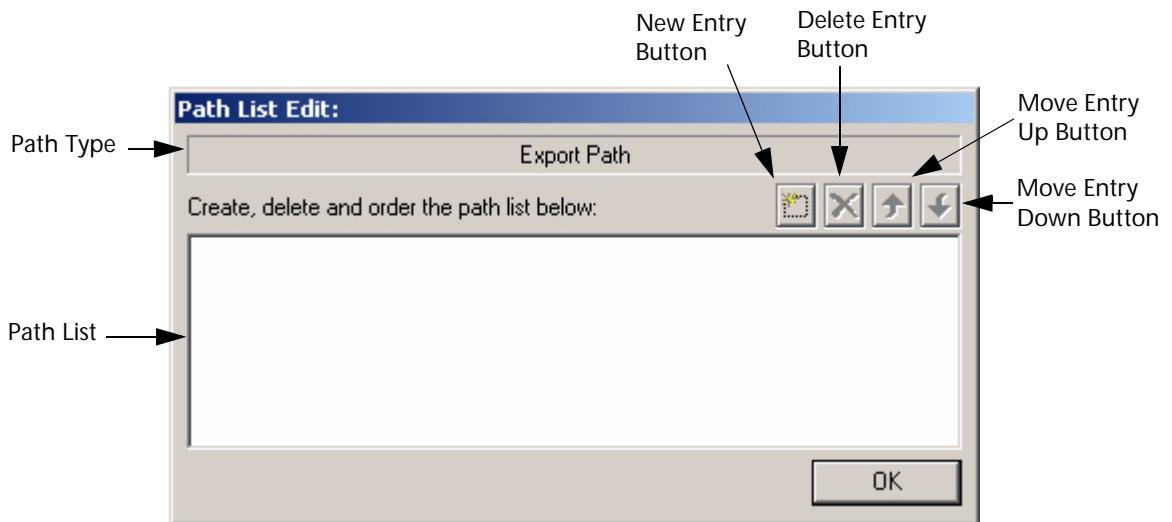


Figure 4-7. Path List Edit Utility

Using the Path List Edit utility, you can add new path definitions, delete obsolete path definitions, and move list items up and down to change the order in which they are searched.

## Scripts

The Scripts element in the Per-Workspace Settings tab allows you to associate scripts with certain events in Visual SoftICE. If you click on a script type, a description of that event association and any other rules on data entry (if applicable) are displayed in the Setting Explanation window below the list.

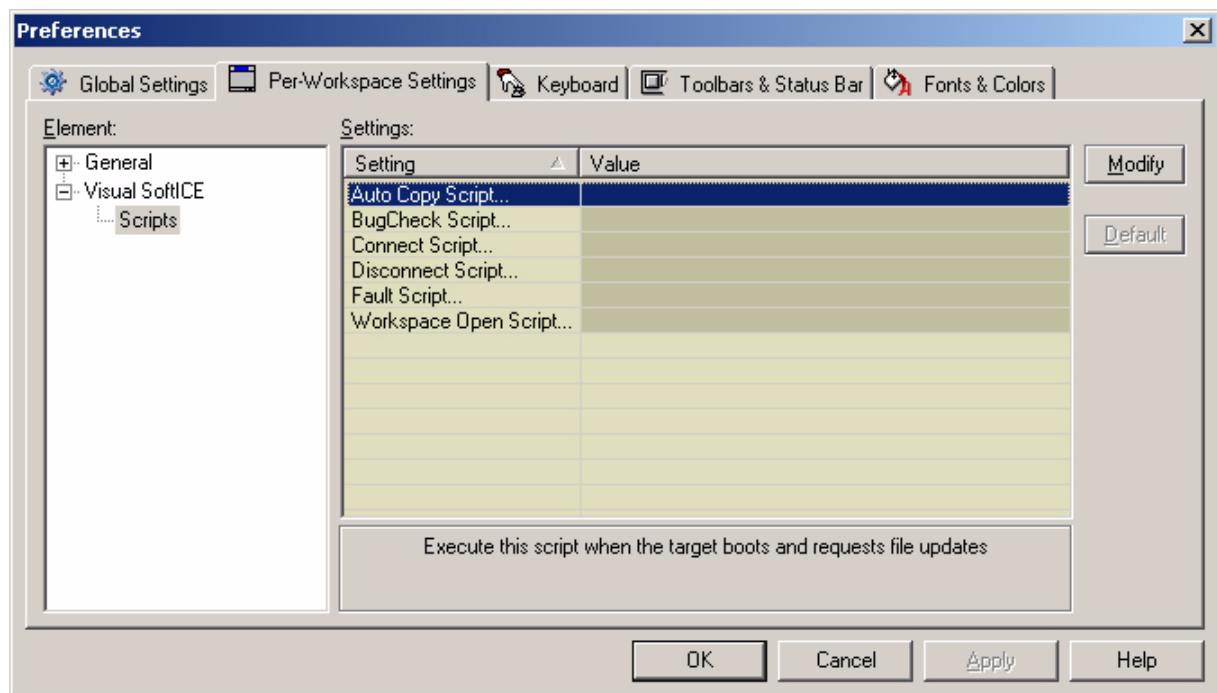


Figure 4-8. Visual SoftICE Per-Workspace Script Preferences

To edit the value of a script path, click on its **Value** field. Visual SoftICE opens the value field for editing, allowing you to enter the script file location and name. It also provides a browse button in the right margin, which you can use to browse your file system and select the script file. For more information on Script File Execution, refer to “Script Execution” on page 82.

## Keyboard Settings

DriverStudio allows you to create custom keyboard definitions, or edit existing keyboard definitions, and save them in your workspace file. DriverStudio components that work within the DriverWorkbench environment provide default keyboard definitions on a workspace level and sometimes on a page component level. You cannot edit default definitions; however, you can override them by adding your own custom keyboard definitions for the same keystroke. Table 4-9 on page 70 provides the definitions that are global to the DriverWorkbench environment, which you cannot override.

Table 4-9. Global Keyboard Definitions

Keystroke	
CTRL-TAB	Next page on a pad.
SHIFT-CTRL-TAB	Previous page on a pad.
CTRL-F4	Close the current pad.

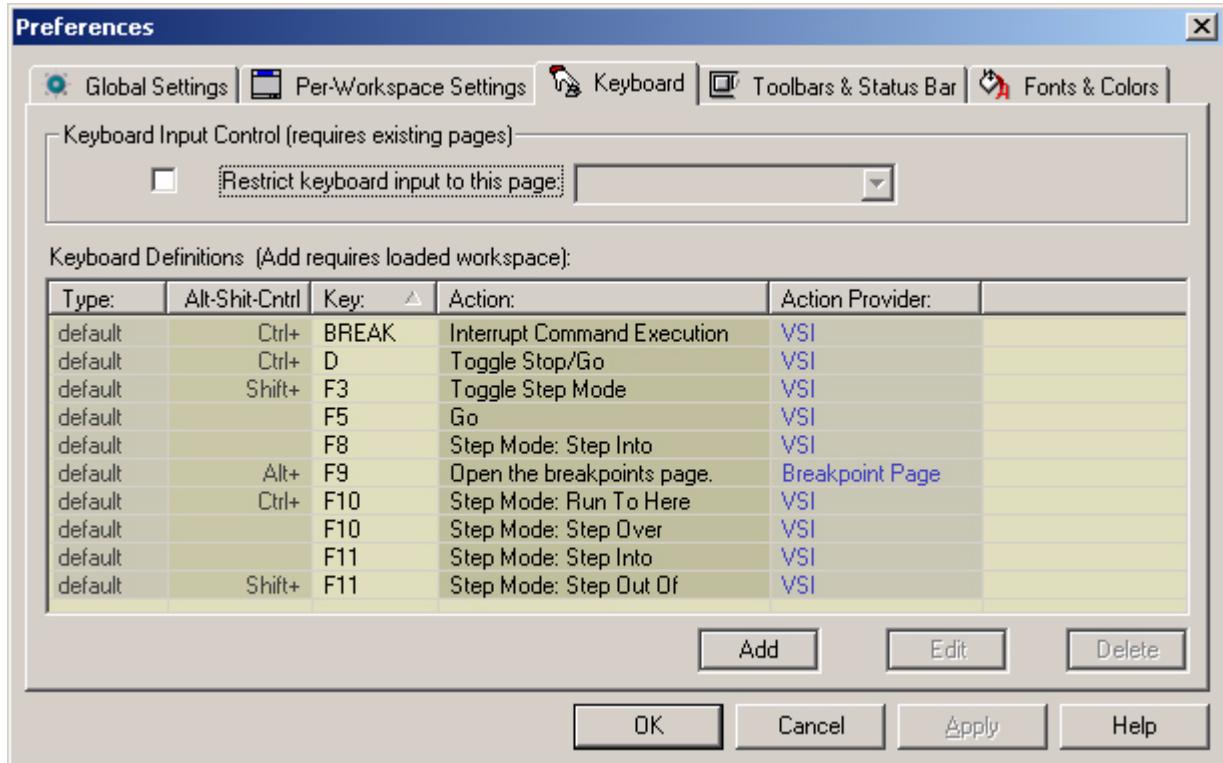


Figure 4-9. Keyboard Preferences Tab

Keyboard definitions are global, and will function regardless of what page or control has focus.

## Toolbars & Status Bar Settings

DriverStudio has specific page-access icons that appear in the Pages toolbar, and a Status Bar that is displayed in the frame. You can also add icons to the the Pages toolbar representing Custom tools you have previously added (via the Tools menu). The Toolbars & Status Bar tab is composed of a Toolbars list displaying available toolbar types, and a right-side section for customizing the status bar.

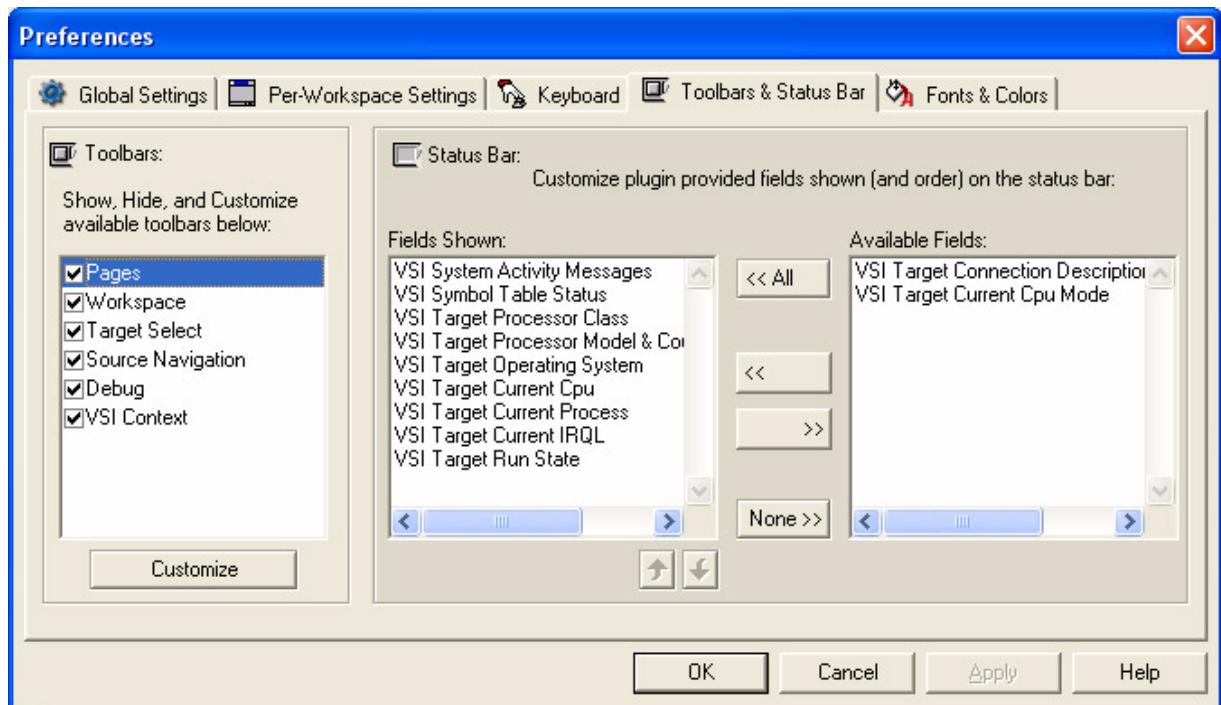


Figure 4-10. Toolbars & Status Bar Preferences Tab

### Customizing the Status Bar

DriverStudio allows you to customize the information displayed on the status bar. To customize the information displayed by the status bar, select or deselect field types by moving them into the appropriate column. You can arrange the fields in the order you would like them to appear in the status bar by using the **Up** and **Down** buttons for the **Current Field Shown** column.

## Customizing Toolbars

You can configure which of the toolbars or page-access icons you want to display, and you can add further options on certain toolbars. Use the **Customize** button to open the **Customize Toolbar** utility.

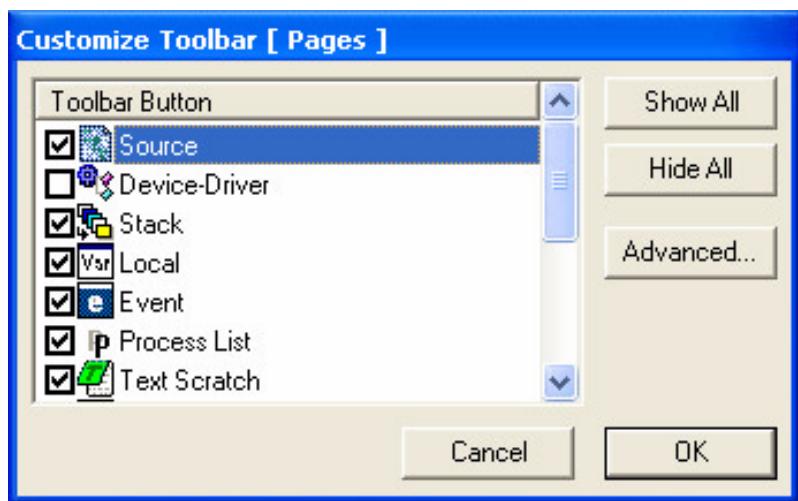


Figure 4-11. Customize Toolbar [Pages] Utility

Select or deselect Toolbar Buttons using the check-boxes to the left of each. Toggle display of all Toolbar Buttons using the **Show All** and **Hide All** buttons.

Certain toolbars allow you to further customize them via the Advanced button. For example, the Pages toolbar allows you to add any toolbar button or available tools to its user-defined area (to the right of the system-defined area).

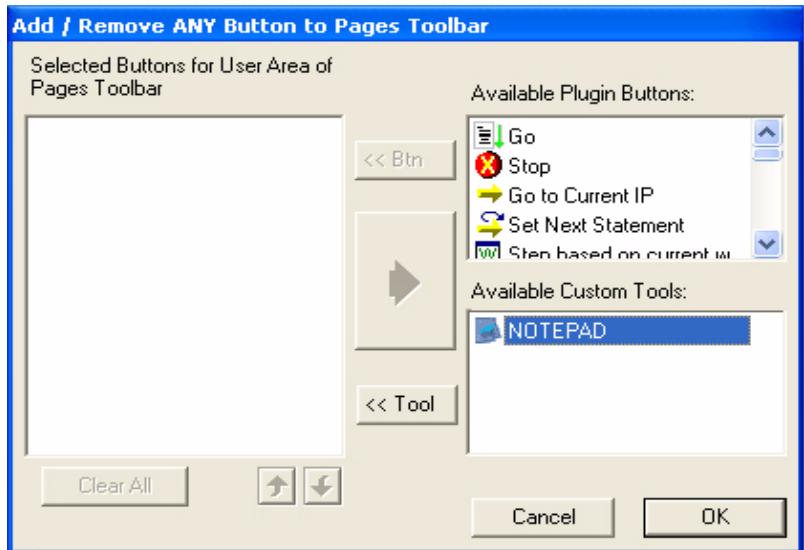


Figure 4-12. Add/Remove Any Button Utility

Double-click on list entries, or use the left and right arrow buttons, to select the ones you want to display on the toolbar. You can arrange the order in which they appear on the toolbar by using the up and down arrow buttons to shift the order in the **Pages Toolbar** list.

### *Fonts & Colors Settings*

Many pages allows you to configure the properties of text that appears in them. DriverStudio provides a Fonts & Colors tab in the Preferences dialog where you can configure various aspects of different pages. The Fonts & Colors tab is composed of a left-side Element column displaying a list of available pages and their settings, and a series of controls (font, size, color, and emphasis). Each page determines the settings available for it, and governs what you can do with each setting.

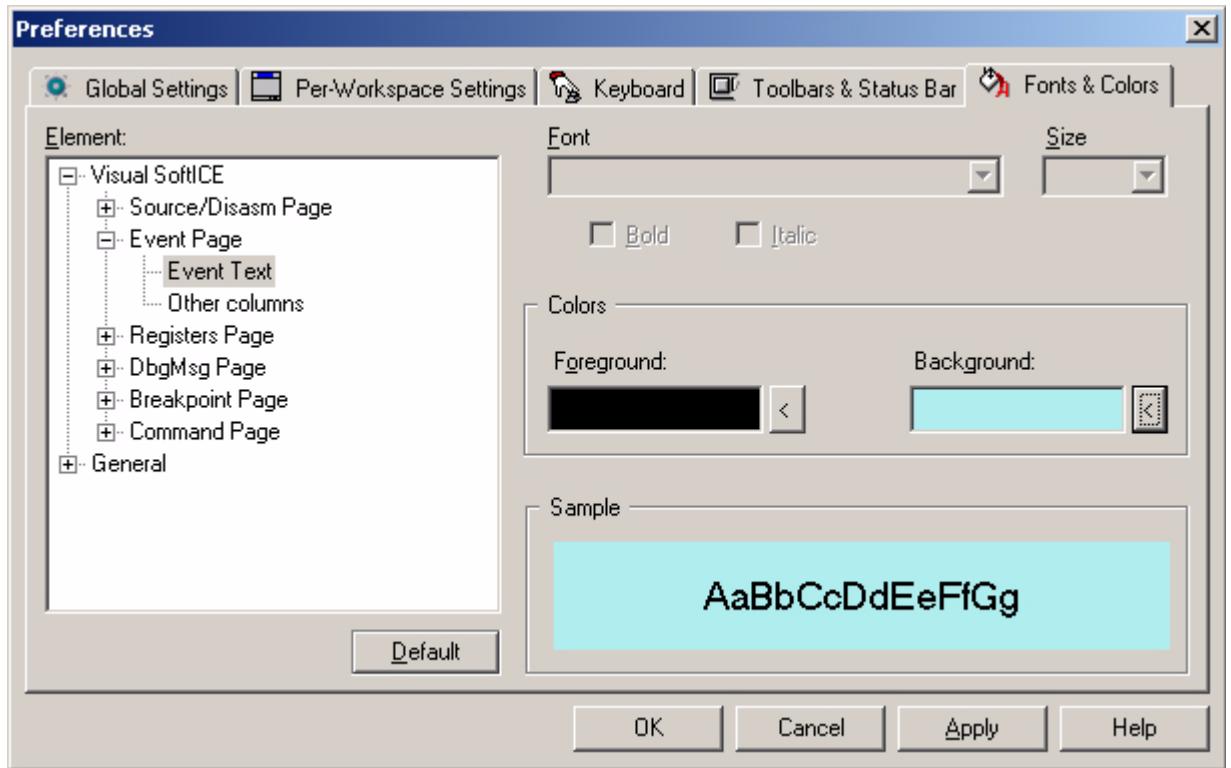


Figure 4-13. Fonts & Colors Preferences Tab

## Other User Interface Attributes and Features

The following sections describe other attributes and features of the Visual SoftICE user interface, and how to use them to customize your workspace.

### ***Workspace Save and Load***

In addition to loading and saving custom workspaces manually, you can configure several automatic workspace loading and saving behaviors in DriverStudio. To configure workspace behavior, access the **Workspace Save/Load** element on the Global Settings tab of the Preferences dialog.

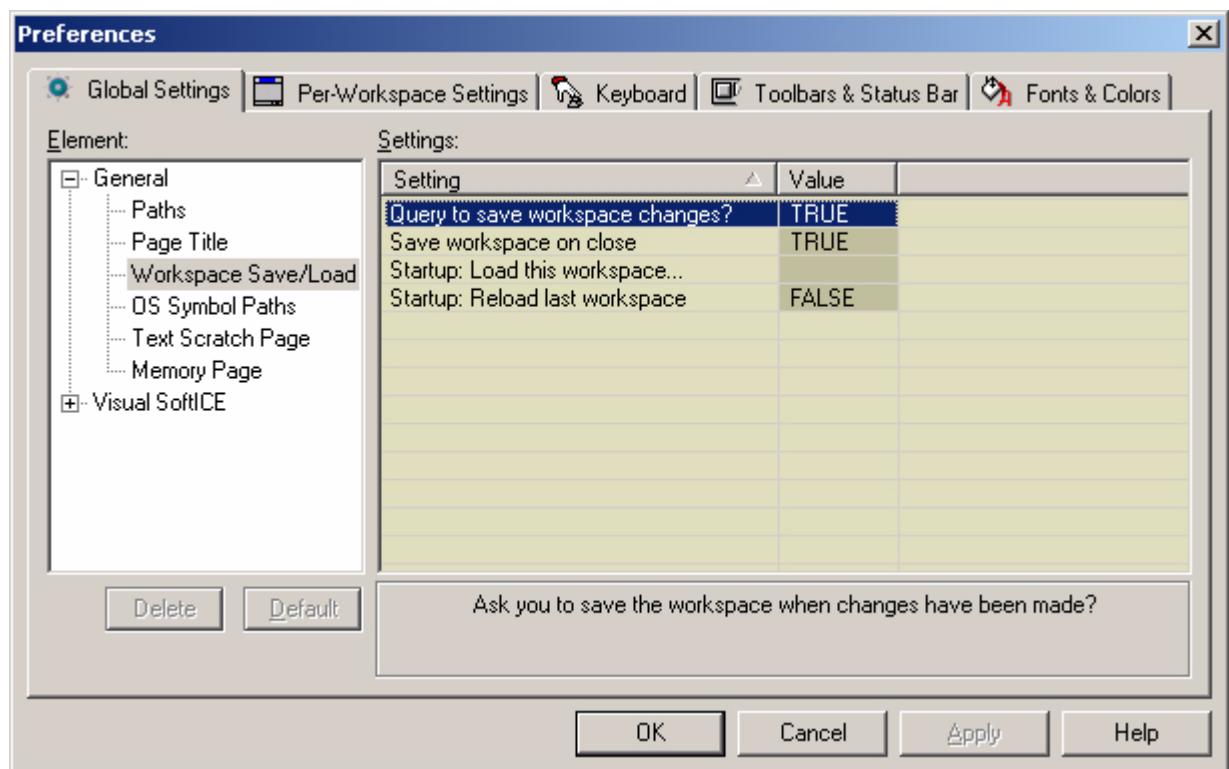


Figure 4-14. Workspace Save/Load Preferences

You can configure the following workspace properties:

Table 4-10. Workspace Save/Load Properties

Setting	Description
Query to save workspace changes?	This option causes DriverStudio to ask you if you want to save the workspace, as a reminder, if you have made changes.
Save workspace on close	This option automatically saves the workspace when you close it, or when you exit DriverStudio.
Startup: Load this workspace	This option allows you to designate a workspace to automatically load on startup.
Startup: Reload last workspace	This option allows you to configure DriverStudio to automatically load the last workspace you had open on startup.

To change a True or False value for a workspace property, click in the **Value** field to toggle it. To specify a workspace for the **Startup: Load this workspace** property, click in the **Value** field to open it for edit, and enter the workspace path and name.

You can also use the browse button to browse your file system and select a saved workspace. If you click on a setting, a description of that setting, any ranges (if applicable), and other rules on data entry (if applicable) are displayed in the Setting Explanation window below the list.

## Special Command Page Features

### Command Comments

The input dialects support a line comment style (as opposed to block style comments). Each dialect can have its own syntax; but, to date, all supported command dialects use the C++ form:

```
// This entire line is a comment - useful in scripts
proc // Command displays all currently running processes
```

### Command Syntax Output Redirection

You can redirect the output of a given command to another page by using the dialect's redirection syntax. To date, all the supported command dialects use the following form:

```
Cmd-text /> page-type
```

- ◆ The Cmd-text is any normal command that is recognized in the command dialect (e.g. “proc”).
- ◆ The redirection syntax “/>” indicates that the remainder of the input stream is a target page type. This form has been chosen so as not to conflict with operands that the expression evaluator acts on.
- ◆ The page-type can be either of the following:
  - ◊ **The name of an existing page.** If the page does not exist, the redirection will fail. If the page exists, it is passed the output for display within its domain. If it cannot handle the data, that page will ignore it, and nothing is displayed.
  - ◊ **The “type name” of a known page plugin.** This is a short mnemonic of the page type, such as CMD for the command page, or DBG for the debug output page, (refer to [Table 4-11](#)). In this case, a page of “type-name” will be created, and the output of the command passed to it. This page will be created in the current PAD, and will get the text of the command in the Cmd-text as its title.

**Table 4-11.** Page Plugin Type Names

Page	Name
Command Page	CMD
Debug Message Page	DBG
Device Page	DVC
Event Page	EVT
Text Scratch Page	TXT
Locals Page	LOC
Memory (Data) Page	MEM
Process List Page	PROC
Register Page	REG
Disassembly Page	DISASM
Source Page	SRC
Stack Page	STK
Watch Page	WAT

**Examples** The following example sends the output of the proc command to a page named text1:

```
SI>proc /> text1
```

The following example sends the output of the proc command to a new text page:

```
SI>proc /> TXT
```

## Automatic Output Redirection

The command page supports a mode whereby commands can have output results automatically handled by an appropriate page other than the command page itself. Toggling the “Auto Command Redirection” button on the command page itself enables this.

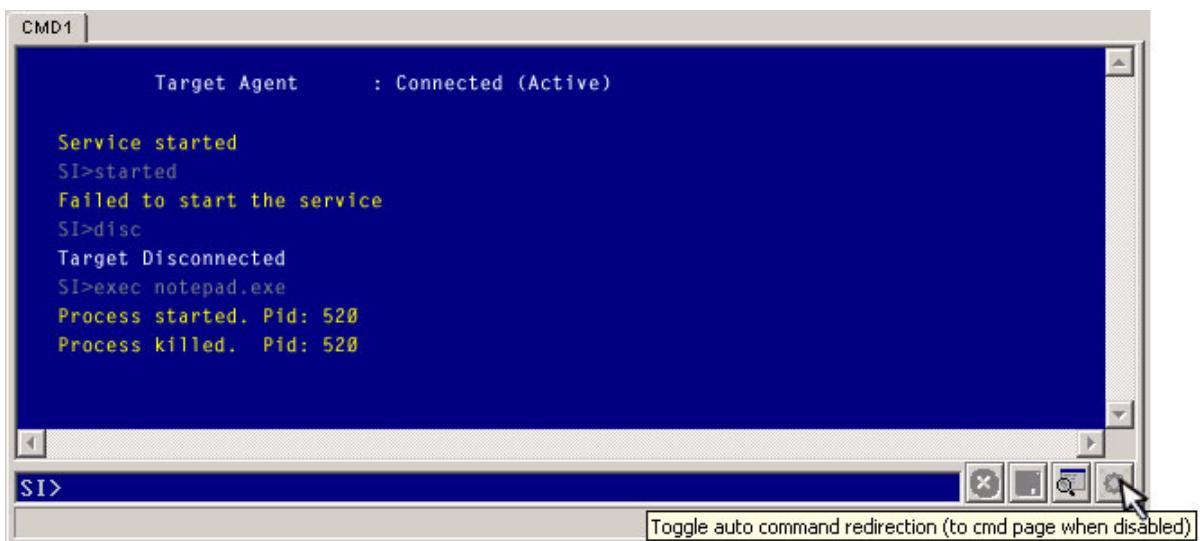


Figure 4-15. Auto Command Redirection Button

When active, after a command is successfully processed, the system looks for registered page handlers for the command. If found, the command output is automatically passed to the first instance of a page found that handles the command.

For example, if your workspace contains a command page and a registers page, you have auto command redirection enabled, and you type the following:

```
r general
```

The register page will change to reflect the current contents of the general register group. If the auto command redirection switch had not been enabled, the output of the command would have appeared within the command page.

## Cut, Copy, and Paste

All the controls in all the pages that allow highlight selection support the standard copy operation.

Many input controls that support highlight selection support the standard cut operation.

Most, if not all, input controls that take text support the standard paste operation.

All the user interface controls use the operating systems clipboard buffer.

## Drag and Drop

There are 2 types of drag and drop supported in the user interface:

- ◆ Pages between and within pads
- ◆ Data items from one page to another

### Pad Level Page Drag and Drop

All Pads allow any number of pages to be placed within them.

You may rearrange the order of the pages within a pad, by simply dragging and dropping the page to a new place in the horizontal order. You may also drag a page out of one pad and into another pad.

If you press the Esc key while dragging a page, the action is cancelled immediately. If you press and hold the Ctrl key while dragging a page, the action is temporarily cancelled: Releasing the mouse button cancels the drag and drop, whereas releasing the Ctrl key continues the drag and drop.

## Data Item Drag and Drop

Many pages of the user interface support dragging a selection from that page to another page's input area. There are 2 different cases: dragging a single element from one page to another, and dragging multiple elements (a collection) between pages. For a table that correlates the graphical representation of each cursor type with an explanation of its meaning, refer to "Visual SoftICE Icons" on page 54.

## Saving Contents to a File

Many pages of the user interface support saving their contents to file. Saving is normally available via right-clicking on a page and selecting **Save Output To File** from the pop-up menu, or by clicking a button within the page itself (as shown below):

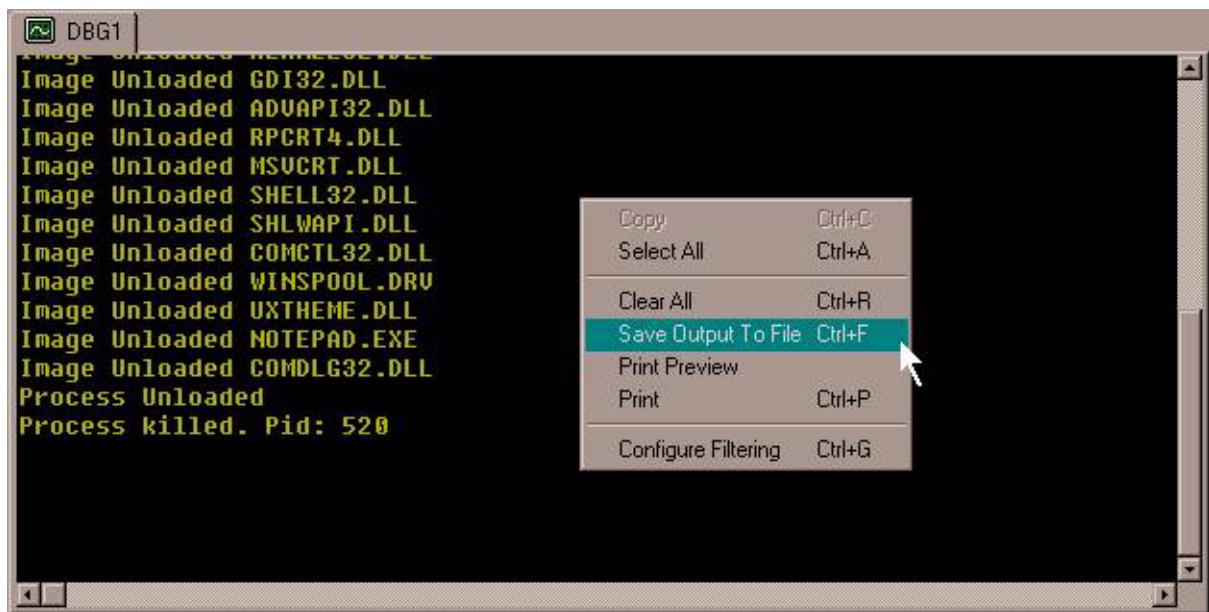
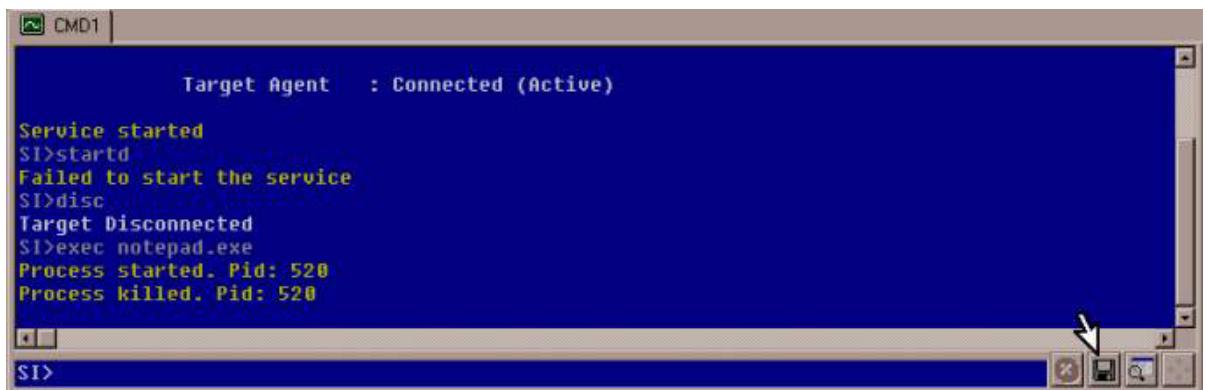


Figure 4-16. Saving Output Via the Pop-up Menu



```
CMD1  
Target Agent : Connected (Active)  
Service started  
SI>startd  
Failed to start the service  
SI>disc  
Target Disconnected  
SI>exec notepad.exe  
Process started. Pid: 520  
Process killed. Pid: 520  
SI>
```

Figure 4-17. Saving Output Via the Save Button

### Print and Print Preview

Most of the pages of the user interface support printing their contents. This is normally available from the main **File** menu.

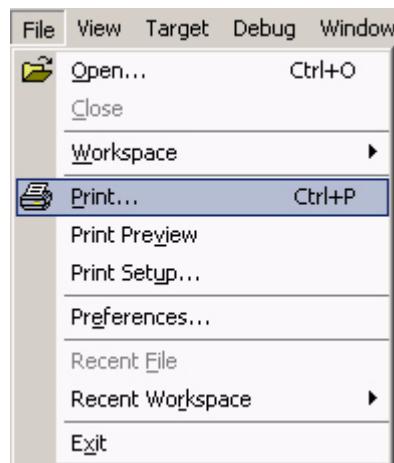


Figure 4-18. Printing from the File Menu

Printing may also be available via right-clicking on a page and selecting **Print** from the pop-up menu.

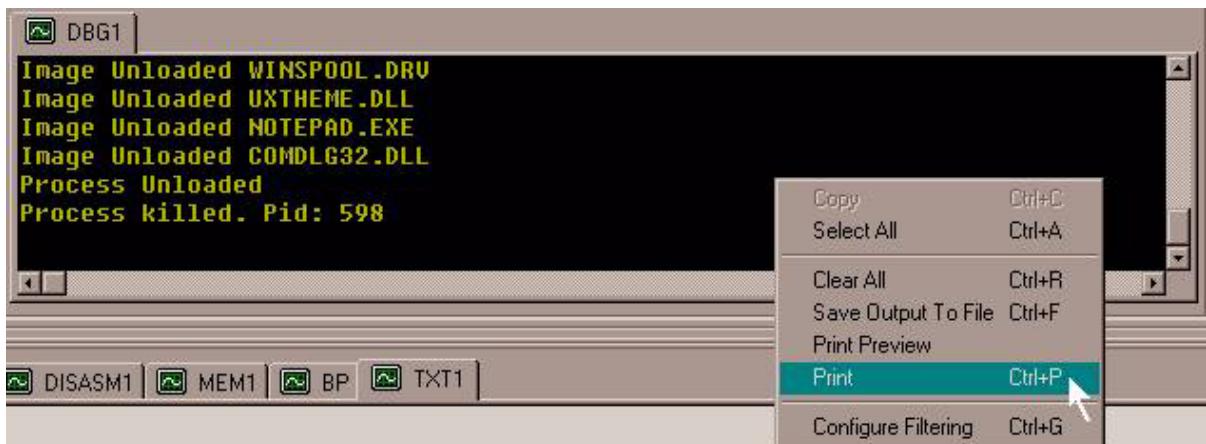


Figure 4-19. Printing from the Pop-up Menu

## Script Execution

You can configure Visual SoftICE to use scripts you have written upon the triggering of certain events, such as opening a workspace or connecting to a target. You configure scripts via the **Scripts** item in the **Per-Workspace Settings** tab under **Preferences**.

To configure Visual SoftICE to use a script, you need to define the path and script file name in the field corresponding to the trigger event for the script, and activate the event by checking its check-box. You can define scripts for the trigger events listed in [Table 4-12](#).

Table 4-12. Script Trigger Events

Script	Trigger Event
Auto Copy Script	Execute a script when the target boots and requests file updates.
BugCheck Script	Execute a script upon encountering a BugCheck event (blue screen).
Connect Script	Execute a script upon connecting to a target.
Disconnect Script	Execute a script upon disconnecting from a target.
Fault Script	Execute a script upon encountering a Fault.
Workspace Open Script	Execute a script upon opening a workspace.

**Note:** Visual SoftICE does not save actual scripts, but saves the paths to the scripts, so you must make sure the scripts exist in the defined location for them to succeed. If you move your workspace to another machine, be certain you also move the script files referenced by the workspace.

## Writing Scripts for VSI

Visual SoftICE script files are any ASCII text files containing at least one Visual SoftICE command per line, or separated by semi-colons on the same line.

You can use any basic text editor to write Visual SoftICE script files, as long as you save them as plain text (.txt). There is no specific directory where you are required to save Visual SoftICE script files, but in order to use them you must configure VSI with their location.



# Chapter 5

## The Visual SoftICE User Interface Pages



- ◆ Pads and Pages
- ◆ The Breakpoint Page
- ◆ The Command Page
- ◆ The Debug Message Page
- ◆ The Device-Driver Page
- ◆ The Disassembly Page
- ◆ The Event Page
- ◆ The Locals Page
- ◆ The Memory Page
- ◆ The Process List Page
- ◆ The Registers Page
- ◆ The Source Page
- ◆ The Stack Page
- ◆ The Text Scratch Page
- ◆ The Watch Page

### Pads and Pages

Many applications follow the Microsoft Multiple Document Interface (MDI) standard, which makes perfect sense for editors and manipulation programs that work against multiple instances of the same document types. A debugger, on the other hand, requires very few instances of a lot of different window types, where most of the window types are not really documents in the classic sense (editable, file-backed, display surfaces).

Instead of forcing the product to live within existing standards, we chose to embrace a metaphor more like a clipboard or pad of paper. The idea was to have pads that contain pages of interest, and be able to have an unlimited number of pads and pages.

The pad is a dockable container of pages. You design your workspace layout by arranging pads in docked and floating locations, and then filling them with your preference of pages. Since the pad is acting like an MDI application (a frame to the pages), the DriverWorkbench system can be thought of as an MDI application inside an MDI application.

## Pads

Pads are containers for any number of pages. Pads have names, can either be docked to the frame or floating, and each have a system menu allowing you to minimize, maximize, and close the window (closing all contained pages). Workspace screen layout is composed of docked and floating pads (and all their contents) and toolbars.

- ◆ Pages may be arranged within the Pad by dragging and dropping them.
- ◆ Pages may be dragged and dropped between pads.
- ◆ Since pages must exist within a pad, dropping a page on the frame itself will automatically create a pad for it.

Of special note:

- ◆ When pads are first created, they are always set to floating style. You can switch to a dockable pad by bringing up the pad's system menu (right-click the title bar) and selecting **Dockable**.
- ◆ You can manipulate pads via the **View->Pad** menu, including creating, deleting, renaming, and walking through the existing pads.
- ◆ The following keystrokes are also useful for shifting display and input focus of pads:

**Table 5-1.** Keystrokes for Pad Control

Keystroke	Definition
CTRL+F4	Close the current pad.
ALT+F6	Switch to the next pad.
CTRL+F6	Display the next pad.
Shift+CTRL+F6	Display the previous pad.

## Pages

Pages are the real workhorses of the product. Each page is dedicated to a certain type of functionality, or to a specific type of data, and must exist within a pad.

Pages may be single-instanced (only one is ever allowed at a time), or multiple-instanced (any number are allowed). Pages have names, and one state called the Page Mode. Most pages support formatted print and print preview functionality. Pages may provide specific preferences, and control over their colors and fonts.

### Page Modes

Each page in the DriverWorkbench user interface supports one or more of the following three page modes.

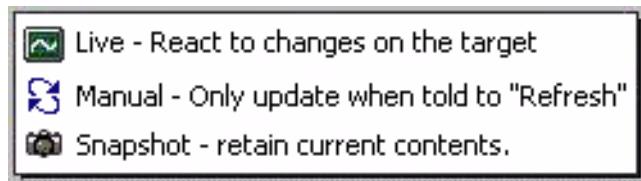


Figure 5-1. Page Modes

**Note:** Page mode icons are only displayed when the font size is Normal.

#### Live Mode

Events from the target cause the page to refresh automatically.

#### Manual Mode

You decide to refresh the page whenever you want by selecting Refresh from the pop-up menu.

#### Snapshot Mode

The page is protecting its current contents from overwrite. You must elect to destroy the data. Changes to the data, and actions against the target are not allowed in this mode.

For more details about page modes, refer to the online help.

Of special note:

- ◆ The display of a page name within a containing pad is controlled from the **File->Preferences->General->Page Title** area. You can select the font and the position of the page names in a pad (top, bottom, left, or right).
- ◆ You can manipulate pages via the **View->Page** menu, including creating, deleting, renaming, and the walking of existing pages within a pad.

- ◆ The following keystrokes are also useful for shifting display and input focus of pages:

**Table 5-2.** Keystrokes for Page Control

Keystroke	Definition
CTRL+Tab	Display the next page.
Shift+CTRL+Tab	Display the previous page.

### Automatic Changes in Mode State

Each page that supports Snapshot mode will automatically switch to this mode when a target disconnection is detected. This is done to indicate that the connection to a target is gone, and that changes to the data contained in a page are no longer valid, or supported.

### Live Mode On Connection

Many pages provide a feature called “Live Mode On Connection” (in the preferences > settings dialog). This feature controls how a pre-existing page behaves when a connection to a target is established. If this feature is active, the page will switch to Live Mode when the connection is detected, and any contents it previously held will be abandoned. This feature is on by default, for pages that support it.

**Note:** If you wish the interface to always preserve the contents of a page, turn this feature off. Doing so will ensure that when the page changes mode to Snapshot, no automatic event will cause that data to be lost. You can always manually change the state back to Manual or Live, and you will get a confirmation dialog box about abandoning the page’s data, if it has any.

## Visual SoftICE Page Overview

Visual SoftICE provides the following dedicated page types:

### Breakpoint (BP)

The Breakpoint page is a single-instance display of all breakpoints currently active, and a log of breakpoint events since the opening of the page. You can create, edit, disable, enable, and/or delete breakpoints using this page. See “The Breakpoint Page” on page 92.

### Command (CMD)

SoftICE users should be right at home in this page, as it provides command execution and result display in a console like surface. The Command page is a multiple-instance container supporting multiple command syntax's (dialects), and currently ships with SoftICE and KD command sets.

There are literally hundreds of commands available, and help is directly available for all commands through three mechanisms - command line help (typing HELP keyword at the prompt), input help (command syntax is updated in the display as you type), and *Visual SoftICE Command Reference* help (online HTML help available via a button and/or F1).

Every function Visual SoftICE offers is available using a command, so learning the power of the available features is well worth the effort.

See “The Command Page” on page 97.

## Debug Message (DBG)

The Debug Message page is a multiple-instance container for the display of text debug messages from the target. It will show only the messages received since the page was opened, and can be configured to only show strings that match a filter. See “The Debug Message Page” on page 103.

## Device-Driver (DVC)

The Device-Driver page is a single-instance tool for exploration and remote system manipulation. It provides a view of the target machine's operating system device nodes, and a view of active loaded drivers and kernel services. The DevNode view can be considered a remote Device Manager, showing available data in a tree that identifies hierarchical relationships. The Device-Driver page offers the ability to remotely install, update, remove, enable, or disable drivers and kernel services, quickly and easily from your master machine. See “The Device-Driver Page” on page 105.

## Disassembly (DISASM)

The Disassembly page is a multiple-instance container for display of disassembled target memory, from live machines or crash-dump files, in the appropriate form for the hardware, OS, and process type (IA32, IA64, x64). The Disassembly page provides helpful data about jump direction, locations, current instruction pointer, and evaluation of highlighted expressions. It supports actions like setting a breakpoint, moving the instruction pointer, and running to an address, as well as standard editor functionality like inserting bookmarks, and searching text. See “The Disassembly Page” on page 121.

## Event (EVT)

The Event page is a single-instance display of the connected target session event log. The display provides a 1 line summary of each event (details are available) in the order they occurred. Deciding which events the page displays is user-configurable. The Event page can also open pre-existing logs from old sessions for review. Unlike the Debug Message page, the Event page shows all the events since the beginning of a session.

The contents of the event log is automatically cleared on disconnection from the target, however you can save the output in text format, and can configure the page to remind you to do so. See “The Event Page” on page 127.

## Locals (LOC)

The Locals page is a single-instance display of the current frame's local variables. It allows selection of other available frames (contexts) and variable editing. Datatypes understood from symbol files can be fully explored via structure expansion. See “The Locals Page” on page 132.

## Memory (MEM)

The Memory page is a multiple-instance container displaying memory from the target. It can display memory in a number of scalar formats, as well as in any type defined by available symbols. Memory can be viewed from live targets as well as crashdump files, and can be accessed via virtual or physical address. The Memory page has two panes for dual-display of the same address in different formats (optional), and supports byte memory searches of its contents. See “The Memory Page” on page 135.

## Process List (PROC)

The Process List page is a single-instance display of processes running on the connected target. It shows the process that is the currently active debugger context (see the ADDR command), and, when stopped, it shows the process the current instruction pointer resides in. You can explore process-specific information from this page, as well as kill desired processes. See “The Process List Page” on page 145.

## Registers (REG)

The Registers page is a multiple-instance container supporting the display and editing of target processor registers for a specific CPU. What registers to show, as well as the data formatting and CPU source, is fully configurable. See “The Registers Page” on page 151.

## Source (SRC)

The Source page is a read-only multiple-instance container for a source file. It supports integration with the symbol engine and debugger to disassemble, manipulate breakpoints, step through code, run to lines, and jump to addresses. It supports some standard editor functionality, such as inserting bookmarks and searching text. See “The Source Page” on page 156.

## Stack (STK)

The Stack page is a single-instance display of the data on the stack for a selected context (process, thread, and frame). The page can display data for the current, or any user-selected, context. It also provides a means to explore frame-specific registers and locals. This page also supports breakpoint manipulation for frame addresses. See “The Stack Page” on page 167.

## Text Scratch (TXT)

The Text Scratch page is a multiple-instance container used to store and manipulate simple text. It can be used to capture redirected command output, or to hold other “copy and paste” or “drag and drop” information. It is not a source editor, but does allow you to load the contents of a file into it. See “The Text Scratch Page” on page 171.

## Watch (WAT)

The Watch page is a single-instance display of the current value of user input expressions. It allows you to select from available contexts, and edit both expressions and results. Datatypes understood from symbol files can be fully explored via structure expansion. See “The Watch Page” on page 174.

# The Breakpoint Page

The Breakpoint page displays statistics on the breakpoints, breakpoint history, and indicates if any of the breakpoints are currently active.

The screenshot shows the 'Breakpoint Details' page in Visual SoftICE. At the top left is a 'Mode Indicator' showing 'BP'. Below it is a table with columns: ID, Type, Size, Msg(s), Logged, Address, Condition, and Trigger Cmds. The table lists five breakpoints:

ID	Type	Size	Msg(s)	Logged	Address	Condition	Trigger Cmds
0	BPX (EXECUTE-INSTR)	NA		<input type="checkbox"/>	userobjs_32.exe!CTestHandles::DisplayValue ...		
1	BPM/R (READ)	000...		<input type="checkbox"/>	810d5540		
2	BUPLOAD (Image Load)	NA		<input type="checkbox"/>	userobjs_32.exe		
3	BPIO (READ-WRITE)	000...		<input type="checkbox"/>	00000027		
4	BMSG (Wnd Msg Proc)	NA	WM_PAINT	<input checked="" type="checkbox"/>	userobjs_32.exe!CTestHandles::CTestHandle...		
5	BPM/R (EXECUTE)	000...		<input checked="" type="checkbox"/>	00000000		

Annotations on the left side of the table identify breakpoint types: Active (green circle), Fixed (yellow circle), I/O (blue circle), Image (green circle), Relative (orange circle), and Disabled (grey circle). Arrows point from these labels to their corresponding columns in the table. To the right of the table is a vertical 'Breakpoint List' and below it is a 'Breakpoint History' window containing two entries:

```
Breakpoint [ 2 ] hit. CPU : 0 Address: 00407f80 Elapsed Time: 4 min 17.870 sec
Breakpoint [ 0 ] hit. CPU : 0 Address: 00402193 Elapsed Time: 24.689 millisec
```

Figure 5-2. Visual SoftICE Breakpoint Page

## Concepts and Associated Commands

### Breakpoint Types

Visual SoftICE supports the following types of breakpoints:

- ◆ Code Execution
- ◆ Interrupt
- ◆ Memory Location Access
- ◆ I/O
- ◆ Memory Range Access
- ◆ Image Load
- ◆ HWND and OS Application Message
- ◆ Deferred or Virtual

Each breakpoint type and state is represented graphically at the place you set it by a Breakpoint Icon. Use the New Breakpoint utility to set any of the breakpoints.

### Display Breakpoint Information

There are two commands you can issue at the command line to display breakpoint information.

- ◆ You can display Breakpoint information by issuing the BL command.
- ◆ You can display the Breakpoint page by issuing the WB command.

## Conditions

All breakpoint types, except the Image Load type, support a conditional IF statement as an optional trigger. If the condition evaluates to true when the breakpoint is encountered, it triggers the breakpoint (in the manner you specified) by stopping, logging the breakpoint without stopping, and/or executing a triggered command list.

## Triggered Command Lists

All breakpoints support optional execution of a list of commands when the breakpoint is triggered. You can enter a single command, or list of commands, that Visual SoftICE will trigger upon hitting a successful breakpoint.

## Global Break

You can configure Visual SoftICE to break upon loading image files by using the SET GLOBALBREAK command.

## *Page Features*

### Breakpoint History

You can access the Breakpoint History dialog by clicking the **Set History** button. The input history will list all breakpoints that you previously created in this session, and previous breakpoints you had set for the currently connected target. From the history list you can select any number of breakpoints and attempt to restore them.

You can configure Visual SoftICE to save and restore any breakpoints you have set upon target disconnection and connection. The preferences controlling breakpoint behavior are configured under Global Settings for Visual SoftICE.

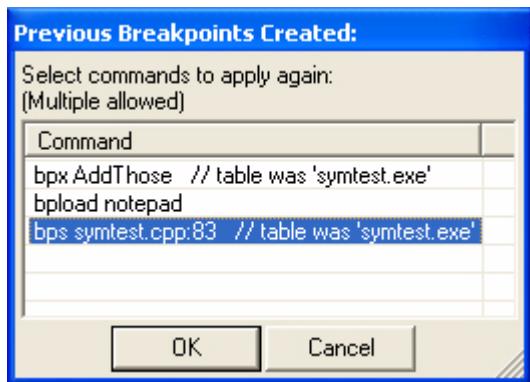


Figure 5-3. Breakpoint History Dialog

## Breakpoint Event Log

The Breakpoint page allows you to display breakpoint related event messages as they are generated on the target. The breakpoint history window also supports the following functionality:

- ◆ You can save all of the current breakpoint history to a specified text file by right-clicking on the window and selecting **Save Output To File** from the pop-up menu.
- ◆ You can clear the breakpoint history by right-clicking on the window and selecting **Clear All** from the pop-up menu.

## New / Edit Breakpoint

The Breakpoint page allows you to create a new breakpoint by right-clicking on the breakpoint list and selecting **New Breakpoint** from the pop-up menu. Visual SoftICE opens the New Breakpoint utility, which provides an easy interface to create breakpoints.

You can also edit an existing breakpoint with the New Breakpoint utility by double-clicking on a breakpoint in the breakpoint list.

## Enable, Disable, Delete Breakpoint

The Breakpoint page allows you to enable, disable, or delete a breakpoint by right-clicking on that breakpoint and selecting the desired operation (**Enable**, **Disable**, or **Delete**) from the pop-up menu. For any breakpoint that supports disabling, you can also double-click the breakpoint entry icon to toggle it between enabled and disabled states. Likewise, highlighting the breakpoint entry icon and pressing <Delete> will remove that breakpoint.

You can also enable, disable, or delete a breakpoint by issuing the following commands at the command line:

- ◆ You can enable a breakpoint by issuing the BE command with the appropriate breakpoint index.
- ◆ You can disable a breakpoint by issuing the BD command with the appropriate breakpoint index.
- ◆ You can delete a breakpoint by issuing the BC command with the appropriate breakpoint index.

**Note:** To get the appropriate breakpoint index, issue the BL command.

## Display Breakpoint Statistics

The Breakpoint page allows you to display the statistics for a breakpoint by right-clicking on the address and selecting Display Statistics. The Breakpoint Statistics dialog displays the following information:



Figure 5-4. Breakpoint Statistics Dialog

- ◆ You can reset the breakpoint log count by clicking Reset.
- ◆ You can view breakpoint statistics from the command line by issuing the BSTAT command with the appropriate breakpoint index.
- ◆ You can optionally display statistics columns by right-clicking on the column headings and selecting **Customize** from the pop-up menu. Customizations to the columns displayed are saved in the workspace.

**Note:** To get the appropriate breakpoint index, issue the BL command.

## Toggle Breakpoint Logging

The Breakpoint page allows you to toggle whether Visual SoftICE is only logging a breakpoint, or if it will stop the target. To toggle logging, right-click on a breakpoint and select **Toggle Logging** from the pop-up menu.

## Open Only One Breakpoint Page

You can only have one Breakpoint page open.

## Copy and Drag

The page only supports the Copy and Drag functions to retrieve its data.

- ◆ You can select any text in the breakpoint history window and copy it to the clipboard.
- ◆ You can drag any text from the breakpoint history window to any page that accepts dropped items.
- ◆ You can select any cell and copy its text to the clipboard.
- ◆ You can drag any cell to any page that accepts dropped items.

**Note:** Select a cell by placing the mouse cursor on the cell and then right-clicking on the element you wish to copy.

## Customize the User Interface

There are multiple attributes of the page that you can customize. They are divided into two categories: per-page and application wide settings. Per-page attributes are always remembered by the workspace when you save it.

### Per-Page Settings

- ◆ You can sort the Breakpoint list by any of the displayed columns by clicking on the column heading.
- ◆ You can reorganize and resize the individual columns displayed.

### Application Wide Settings

- ◆ You can use the **Live Mode On Connection** setting to control whether or not the page switches automatically to Live mode when a connection to a new target is established.

## Print

The page supports printing, and print-previewing of its contents.

## The Command Page

The Command page is used to execute Visual SoftICE and KD commands. It is composed of an input command line, an output window, and several control buttons. There is no limit to the number of lines that can be retained in the output pane.

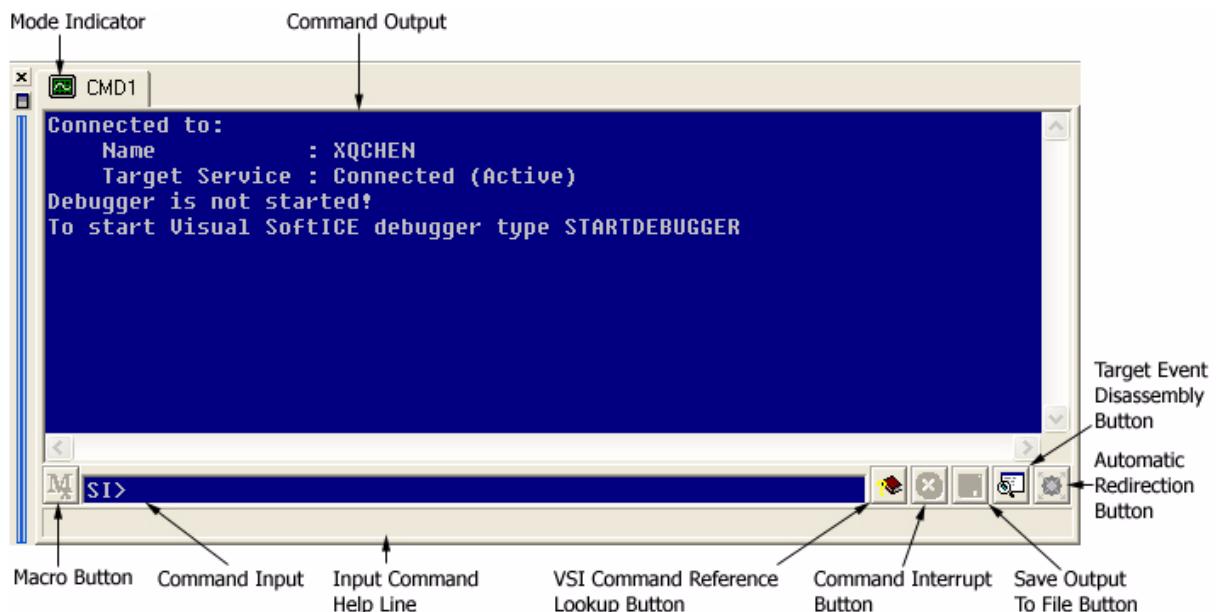


Figure 5-1. Visual SoftICE Command Page

## Concepts and Associated Commands

### Redirecting Output



The command page supports command redirection of its output either automatically or explicitly to a particular page.

- ◆ **Automatic Redirection** - allows you to toggle the automatic redirection of the command output to other pages, whether they are Command pages or other page types.
- ◆ **Explicit Redirection** - allows you to redirect the output of a particular command to the specified page or page type.

## Message Output

The Command page supports displaying event messages from the target. If the Message Level setting is set to **on** or **verbose**, then messages from the target will be displayed in the output window. For more information on controlling the display of messages in Visual SoftICE, refer to the SET MSGLEVEL command.

## Logging

The Command page supports the echoing of the input and/or output of the page to a file. The output window of the page where you execute the command will echo to the specified file in the manner you specify. For more information on controlling the logging of the input and/or output of a Command page, refer to the LOG command.

## Scripts and Macros

The Command page supports execution of scripts, and macro definition, execution, and deletion.

- ◆ You can echo script commands to the current console by using the SET SCRIPTECHO command. You can set this command on a per-page basis.
- ◆ You can set file system search path for scripts by using the SET SCRIPTPATH command. You can set this command on a per-workspace basis.
- ◆ You can control whether scripts automatically stop execution when an error occurs by using the SET SCRIPTSTOPONERROR command. You can set this command on a per-page basis.

Refer to the MACRO command in the *Visual SoftICE Command Reference* for more details.

## Output Formatting

You can configure the Command page output format by modifying the following attributes:

- ◆ Modify the radix of data for input and output by using the SET RADIX command.
- ◆ Modify the uppercase hex disassembly output by using the SET UPPERCASE command.
- ◆ Modify the way Visual SoftICE displays 64-bit addresses by using the SET ADDRESSFORMAT command.

- ◆ Modify the mnemonics used for register names by using the SET REGNAME command.
- ◆ Enable or disable the formatting of FP registers by using the SET FLOATREGFORMAT command.
- ◆ Modify the format of the PACKET command by using the SET PACKETFORMAT command.

## Customization Settings

You can use the Command page to configure global system settings. To view the current settings use the SET command without any parameters. You can configure the following global system settings:

- ◆ Automatically stop the target when a command is issued by using the SET STOPONCMD command.
- ◆ Configure the target to stop on embedded INT 1 instructions by using the I1HERE command.
- ◆ Configure the target to stop on embedded INT 3 instructions by using the I3HERE command.
- ◆ Configure the cache size by using the SET CACHE command.
- ◆ Configure the input command dialect by using the SET DIALECT command.
- ◆ Configure all fault trapping by using the FAULTS command.
- ◆ Configure the way the symbol engine and target attempt to match symbolic data by using the SET IMAGEMATCH command.
- ◆ Control thread-specific stepping by using the SET THREADP command.
- ◆ Configure the warning and confirmation level by using the SET WARNLEVEL command.
- ◆ Configure paths used by Visual SoftICE by using the SET EXEPATH, SET EXPORTPATH, SET KDEXTPATH, SET SYMPATH, SET SRC PATH, and SET SYMSRVPATH commands.

## MACRO Command

Define Visual SoftICE macros using the MACRO command and you can use the **Macro** button provided by this page to quickly and easily access them.

### Open Any Number of Command Pages

There are no restrictions on the number of command pages that you can open; however, Visual SoftICE executes commands from all the pages in serial fashion. For more information on Visual SoftICE commands, refer to the *Visual SoftICE Command Reference*.

### Automatic Command Completion

You can configure the Command page to automatically complete recognized commands as you enter them at the Input line. This occurs after you have entered a partial but unique command name and press the space bar. If Visual SoftICE recognizes the command you are entering, and you have Automatic Command Completion enabled, it will complete the command for you. To enable Automatic Command Completion, access the Command page **Preferences** and set **Automatic Command Completion** to **True**.

### VSI Command Reference Lookup



Look up a specific command, or go to the Table of Contents, in the online *Visual SoftICE Command Reference* via this button. To look up a specific command, enter that command name and click the command lookup button (or press F1). To open the *Visual SoftICE Command Reference* at the TOC, leave the input line blank and click the command lookup button (or press F1).

You can also configure the Command Page to interpret the HELP command as equivalent to F1 and use the *Visual SoftICE Command Reference* instead of displaying the brief command line help. To enable online lookup for the HELP command, access the Command page **Preferences** and set **Help command using Online lookup** to **True**.

### Command Interruption



The command page supports interruption of commands that have been entered on the command line. Some Visual SoftICE commands may take a long time to execute, and are designed to be interruptible. Not all commands support interruption. If execution can be interrupted, then the Command Interrupt button becomes active. Complete one of the following procedures to interrupt execution of a command.

- ◆ Click the **Command Interrupt** button

- ◆ Press <Ctrl> <Break> (or your custom key-sequence for “interrupt command execution”)



## Event Disassembly Output

The Command page supports automatic output of disassembly when certain events are received from the target. The Target Event Disassembly button toggles this disassembly off and on. If you enable target event disassembly, then events on the target that should show disassembly will display it in the output window. If you disable target event disassembly, then no disassembly will appear in the Command page.



## Saving and Clearing Output

You can save all of the output on the Command page to a specific text file by clicking on the **Save Output** button, or right-clicking on the page and selecting **Save Output To File** from the pop-up menu.

Since Visual SoftICE saves the entire contents of the Output window, you may wish to clear the Output window before executing the command whose results you wish to save. You can clear all of the output on this page by issuing the CLS command, or right-clicking on the page and selecting **Clear All** from the pop-up menu.

**Note:** If you select an existing file, Visual SoftICE replaces the contents of the file with contents of the Output window. It does not append to the selected file.



## Access Defined Macros via the Macro Button

The Command page provides a **Macro** button to give you fast and easy access to any macros you have defined. The button is enabled once you have created macros using the MACRO command. Clicking the **Macro** button accesses a list of the available macros. If you have any input at the command prompt when you click the button, Visual SoftICE attempts a partial match against the list of available macros. When you select a macro from the list, the macro is placed in the input line for execution, replacing any text that was previously entered there.

## Custom Keyboard Definition Provider

The Command page is a provider for Custom Keyboard Definitions. You can designate any single keystroke to pass a string (any valid VSI command under 40 characters in length) to a command page, if one is open. If an open command page does not exist, the keystroke will have no effect.

## Finding Text

Using the context menu **Find** option, or via the **CTRL+F** key sequence, you can search for a text string value within the output pane.

- ◆ You can search by trying to match the whole word only, or by finding a partial string match.
- ◆ You can search by trying to match the case of the specified search string.

## Go To Line

Using the context menu **Go To Line** option, or via the **CTRL+G** key sequence, you can move to a specific line number in the Command page output.

## Cut, Copy, Paste, Drag, and Drop

The output window only supports Copy to retrieve its data. The input field supports Cut, Copy, Paste, Drag, and Drop.

- ◆ You can select any text within the output window and copy it to the clipboard.
- ◆ You can drag a selection from the output window (or other parts of the VSI UI) and drop it on the input field. This is useful when entering values for commands on the input field.
- ◆ You can highlight any text within the output window and select **Copy+Paste (immediate)** from the pop-up menu to copy that text to the input line.
- ◆ You can use **Select All** from the pop-up menu to highlight all the lines on the output window before copying them to the clipboard.

## Customize the User Interface

There are multiple attributes of the page that you can customize. They are divided into two categories: per-page and application wide settings. Per-page attributes are always remembered by the workspace when you save it.

### Per-Page Settings

- ◆ You can set the state of automatic command redirection.
- ◆ You can set event disassembly options.
- ◆ You can display or hide line numbers.
- ◆ You can display or hide bookmarks.

## Application Wide Settings

- ◆ You can use the **Colors & Fonts** tab to modify the colors and fonts used for Normal text, Highlighted text, Title text, Dim text, Target Notifications, and Internal Notifications.
- ◆ You can use the **Command Completion** setting to enable or disable automatic command completion.
- ◆ You can use the **HELP Command Using Online Lookup** setting to enable or disable *Visual SoftICE Command Reference* lookup via the HELP command.

## Print

The page supports printing, and print-previewing of its contents.

## The Debug Message Page

The Debug Message page is a read-only receiver of information from the target and from other sources. The page is composed of a single window, which receives and displays text messages, and a pop-up utility controlling the filtering of the messages that are displayed in the window.

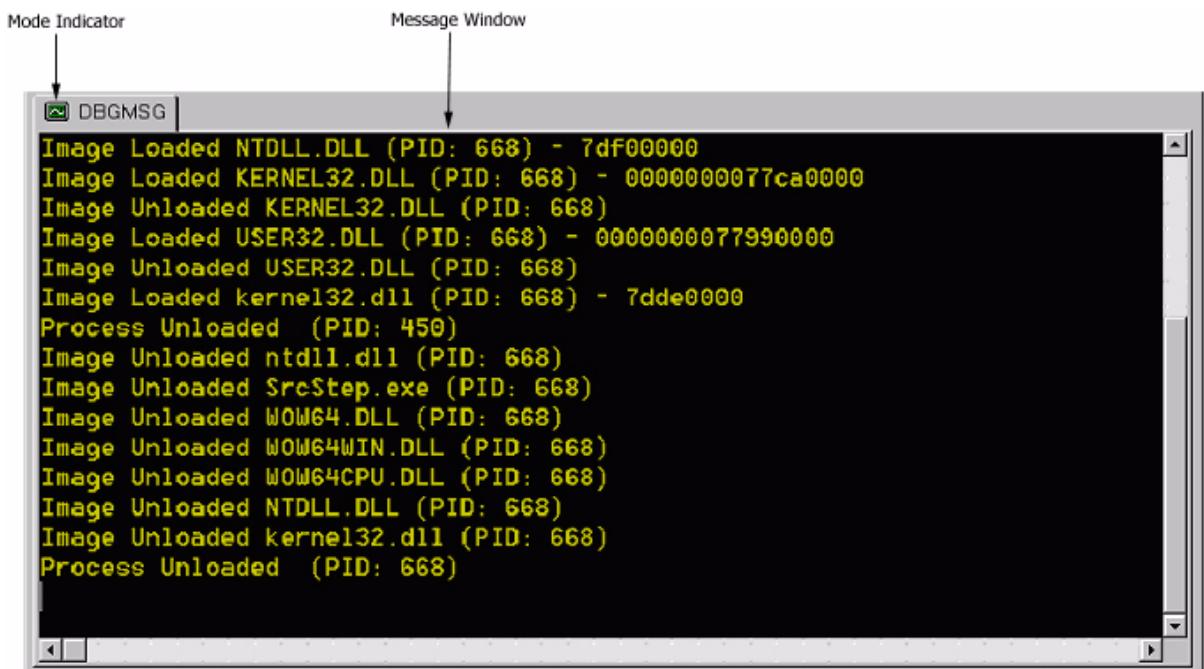


Figure 5-2. Visual SoftICE Debug Message Page

### Filtering Messages

The Debug Message page supports the filtering of messages output to the page. Message filtering is controlled by the Debug Message Filter utility, which is available by right clicking on the page and selecting **Configure Filtering** from the pop-up menu.



Figure 5-3. Debug Message Filter Utility

You can filter messages based upon a substring match of the message content. The filtering is limited to file system level matching criteria. That is, the '?' and '\*' operators are the only supported matching wildcards. For example, to display the debug messages that your driver sends, prefix your DBGPRINT statements with a tag like MYDRIVER:, then set the regular expression message matching to "MYDRIVER: \*".

### Open Any Number of Debug Message Pages

There are no restrictions on the number of Debug Message pages you can have open, so you can build a workspace with multiple views containing just the debug messages you are interested in, where you want them displayed. You may find this useful when debugging problems with a process on the target, as it allows non-intrusive observation of the behavior of the process.

### Copy

The page only supports copy command to retrieve its data.

- ◆ You can select any text within the page and copy it to the clipboard.
- ◆ You can use **Select All** from the pop-up menu to highlight all the lines on the page before copying them to the clipboard.

## Customize the User Interface

There are multiple attributes of the page that you can customize. They are divided into two categories: per-page and application wide settings. Per-page attributes are always remembered in the workspace when you save it.

### Per-Page Settings

- ◆ You can filter the messages that are output to a particular page. For more information, refer to “[Filtering Messages](#)” on page 104.

### Application Wide Settings

- ◆ You can use the **Fonts & Colors** tab to select the font used for all the text.

## Print

The page supports printing, and print-previewing of its contents.

## The Device-Driver Page

The Device-Driver page is an exploration and remote system manipulation tool. It provides a view of the target machine's operating system device nodes, or a view of active loaded drivers and kernel services. The DevNode view can be considered a remote Device Manager, where a best attempt to identify each entry by PNP or PCI information is done. The DevNode view shows the results in a tree that identifies the hierarchical relationship.

**Note:** You can access the same kind of Device Node view natively in XP via the Device Manager (Start>My Computer>Manage>Device Manager), by selecting View Devices By Connection and then Show Hidden Devices.

For remote manipulation, the Device-Driver page offers seven large buttons at the top, four of which act on the active selection in the current view.



Figure 5-4. Device-Driver Page Button Bar

- ◆ **New** — Install a driver, or kernel service on the target machine, from the master. Allows you to start with an INF, SYS file, or any executable, and do a service or driver install to it, remotely.
- ◆ **Update** — Update an existing, running driver or service on the target machine from the master. Update supports running a local script before the update, tracks local source directories, and works on your selection in the DevNode or Driver display.
- ◆ **Un-Install** — Remove an existing, running driver or service on the target machine. Un-Install supports running a local script before the removal, and works on your selection in the DevNode or Driver display.
- ◆ **Enable/Disable** — Enable or disable a running driver or service on the target machine. For some device driver types, this is equivalent to an uninstall. These functions work on your selection in the DevNode or Driver display.
- ◆ **Logging** — Manage the SetupAPI logging level on the target machine.
- ◆ **Refresh** — Refresh the Device-Driver page contents. While driver loads, unloads, and other actions you take on this page are tracked, it is possible that the target state may not be fully reflected. This button forces a content refresh of the target state.

## *Concepts and Associated Commands*

### **DevNode Exploration**

The left side pane of the page contains a tree that shows the relationships between a DevNode and its siblings, parent, and children. Seeing this relationship is especially helpful for understanding devices that exist on a particular bus or host controller. If a DevNode instance path can be identified as a standard PNP device, or a known PCI device (see note on PCIDEV.CSV file below), the entry will get a descriptive name, and potentially an icon for certain hardware classes. You can search the tree display to find a partial name match by right-clicking within the tree and selecting **Find**. The Find functionality for the tree searches downwards through the branches, either from the current position, or from the tree root.

The right side pane of the page contains a summary of information about the selected DevNode, including flags, and current state (the state as of the last event that refreshed the page, or since the last manual refresh). State will only be meaningful and valid for exploration when the target is stopped. However, getting a snapshot from a running machine can be also be useful.

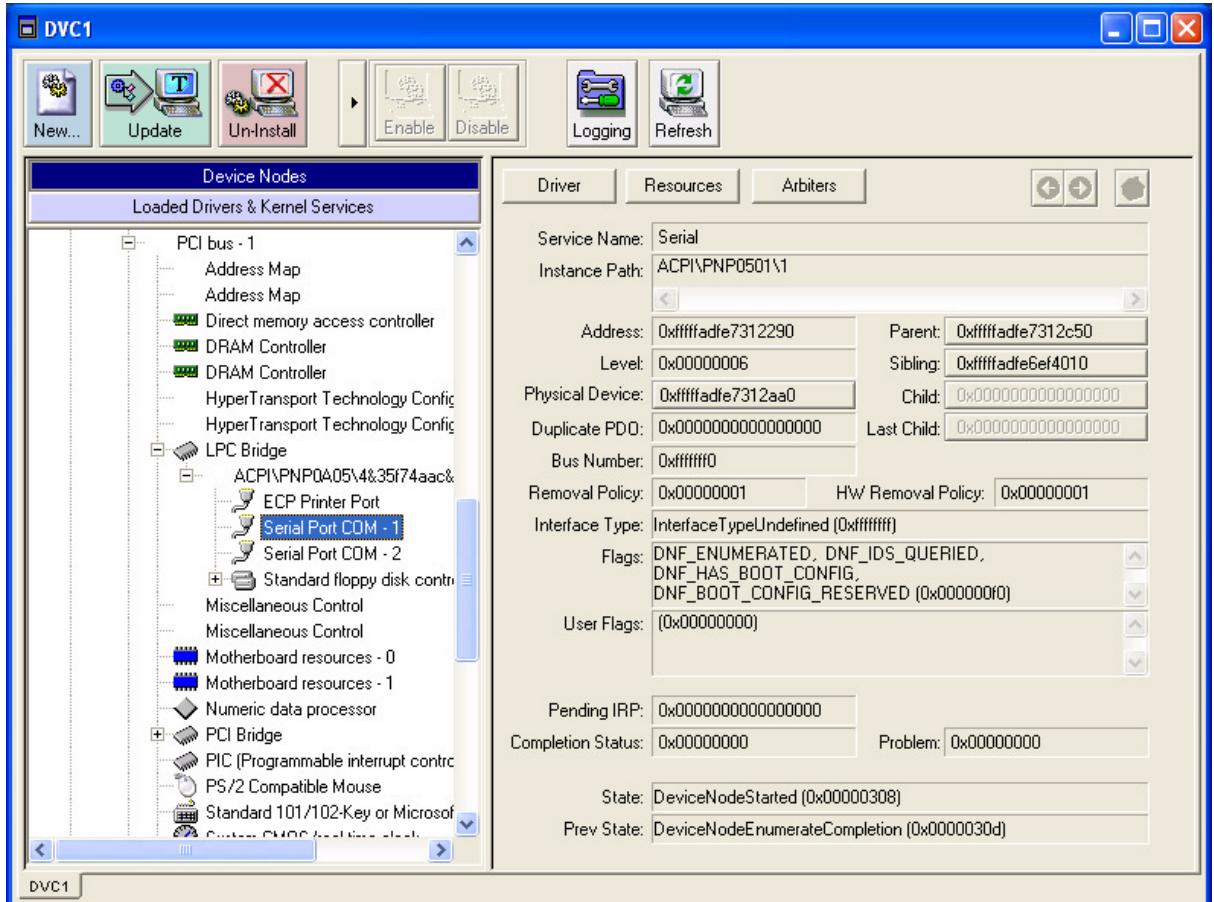


Figure 5-5. Device-Driver Node View

- Driver Button** This button attempts to find the appropriate driver for the current DevNode. If found, it will change the entire view to the Driver/Kernel Service display. There are certain DevNodes which do not have any driver at all. If this is the case, the button will beep, and then disable itself. There are other instances where a DevNode is provided by the OS in a built-in fashion. In this case, the Driver button will often lead to the parent device driver, but be forewarned that any actions on these base drivers can have drastic affects to the target PC (up to and including forcing a re-install of the operating system). Appropriate warnings are displayed by the action buttons, when the situation can be detected.
- Resources Button** This button accesses a dialog displaying a list of resources (raw or translated) associated with the DevNode. Note that this button is only enabled if the DevNode is not marked as DNF\_NO\_RESOURCE\_REQUIRED.

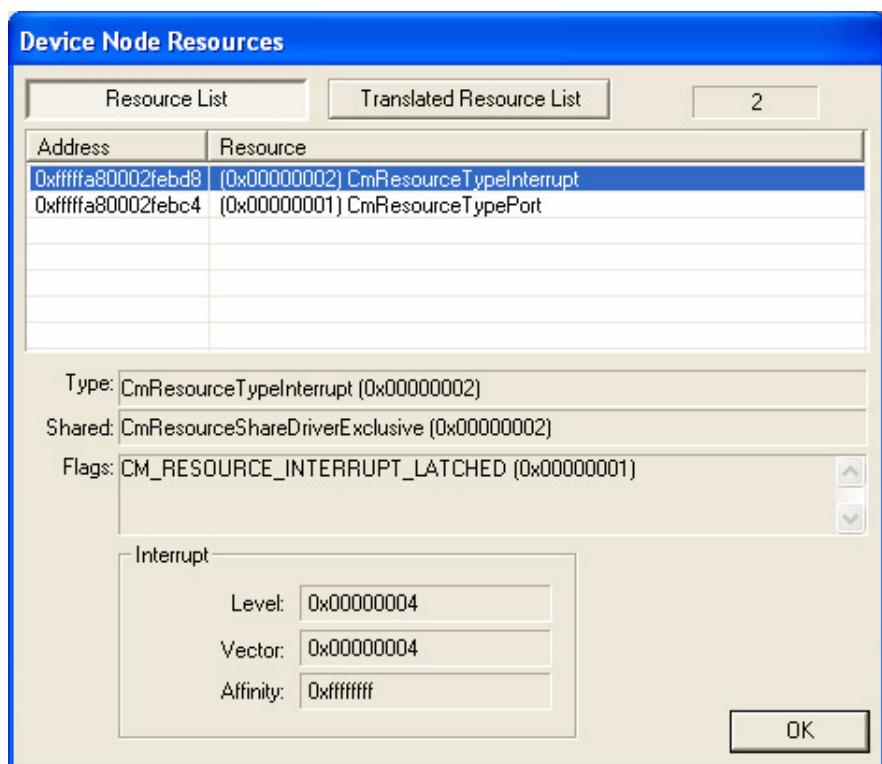


Figure 5-6. DevNode Resources View

Selecting an entry from the resource list displays specific information in the fields provided.

### Arbiters Button

This button accesses a dialog showing a list of all current arbiters associated with the DevNode. Note that this list can be, and often is, empty.

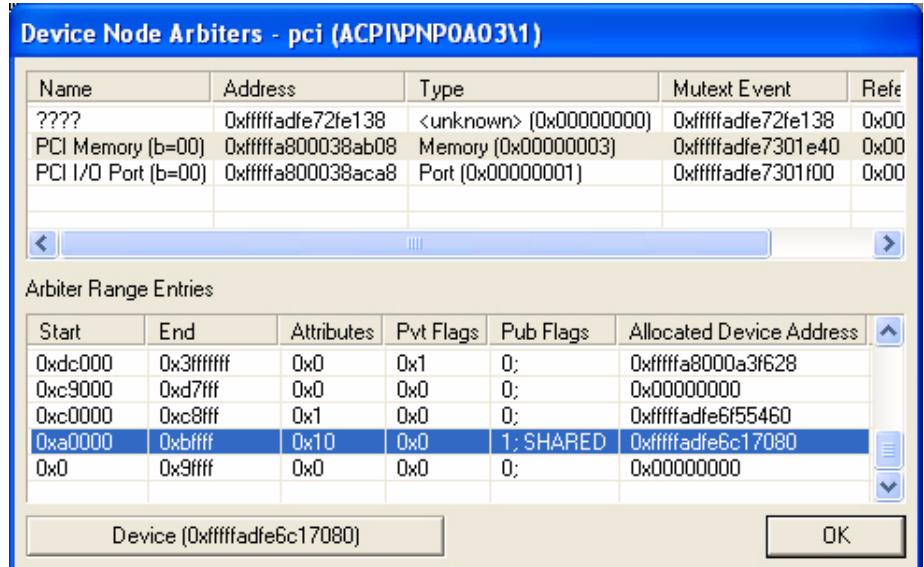


Figure 5-7. DevNode Arbiters View

Selecting a particular entry from the arbiter list will display a list of arbiter range entries. Selecting an entry from the range list with a valid Allocated Device address will enable the **Device** button. Clicking the **Device** button accesses the operating system device object information (See “[Device Object Display](#)” on page 112.).

## Driver/Kernel Service Exploration

The left side pane of the page contains an alphabetically sorted list of system drivers, kernel services, and any shared kernel images. Shared kernel images (DLL's or libraries) will normally only be in this list when the device page was open during a boot cycle. This list is adjusted as drivers are loaded and unloaded on the target.

**Note:** Watching which images are loaded into kernel drivers and when they get loaded can be helpful when exploring system behavior.

The right side pane of the page contains a user-selectable display of information about the selected image. The combo box in the upper left hand corner can be changed to any one of the following views: Summary, Code, Entries, or Devices.

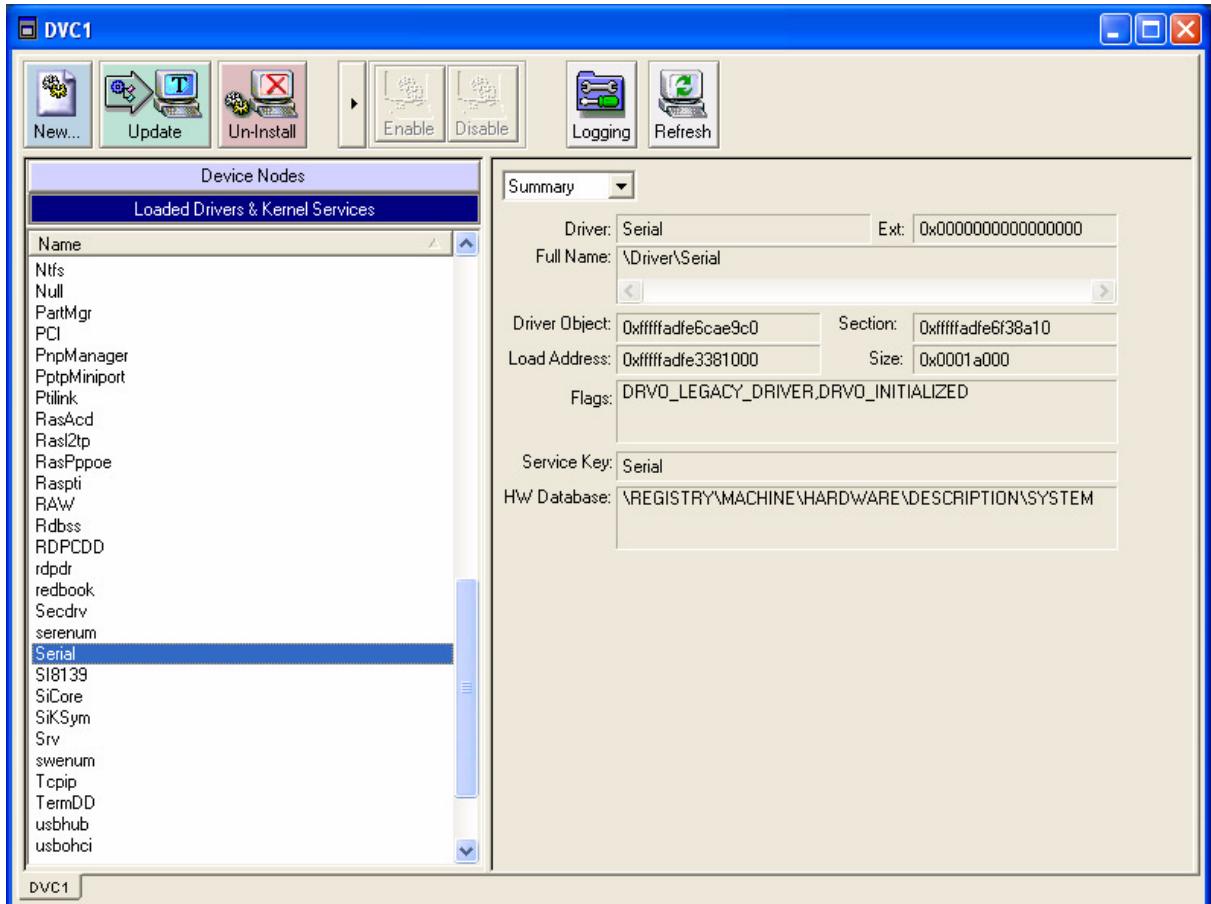


Figure 5-8. Device-Driver Page Driver View

#### Summary View

This view (shown in the main image above) contains general data about the kernel image. All summary fields that contain valid data support copy and paste, and potentially debugger actions, via the context menu (right-click in the field of interest).

#### Code View

This view displays the major IRP codes supported by the driver and the handler routine. The code and location columns support copy and paste. The location column additionally supports debugger actions, via the context menu (right-click on the entry of interest).

Code	Location
IRP_MJ_CREATE	serial!SerialCreateOpen
IRP_MJ_CREATE_NAMED_PIPE	ntoskrnl!IoPlnvalidDeviceRequest
IRP_MJ_CLOSE	serial!SerialClose
IRP_MJ_READ	serial!SerialRead
IRP_MJ_WRITE	serial!SerialWrite
IRP_MJ_QUERY_INFORMATION	serial!SerialQueryInformationFile
IRP_MJ_SET_INFORMATION	serial!SerialSetInformationFile
IRP_MJ_QUERY_EA	ntoskrnl!IoPlnvalidDeviceRequest
IRP_MJ_SET_EA	ntoskrnl!IoPlnvalidDeviceRequest
IRP_MJ_FLUSH_BUFFERS	serial!SerialFlush
IRP_MJ_QUERY_VOLUME_INFORMATION	ntoskrnl!IoPlnvalidDeviceRequest
IRP_MJ_SET_VOLUME_INFORMATION	ntoskrnl!IoPlnvalidDeviceRequest
IRP_MJ_DIRECTORY_CONTROL	ntoskrnl!IoPlnvalidDeviceRequest
IRP_MJ_SYSTEM_CONTROL	serial!SerialSystemControlDispatch
IRP_MJ_DEVICE_CONTROL	serial!SerialIoControl
IRP_MJ_INTERNAL_DEVICE_CONTROL	serial!SerialInternalIoControl
IRP_MJ_SHUTDOWN	ntoskrnl!IoPlnvalidDeviceRequest
IRP_MJ_LOCK_CONTROL	ntoskrnl!IoPlnvalidDeviceRequest
IRP_MJ_CLEANUP	serial!SerialCleanup
IRP_MJ_CREATE_MAILSLOT	ntoskrnl!IoPlnvalidDeviceRequest
IRP_MJ_QUERY_SECURITY	ntoskrnl!IoPlnvalidDeviceRequest
IRP_MJ_SET_SECURITY	ntoskrnl!IoPlnvalidDeviceRequest
IRP_MJ_POWER	serial!SerialPowerDispatch
IRP_MJ_SYSTEM_CONTROL	serial!SerialSystemControlDispatch
IRP_MJ_DEVICE_CHANGE	ntoskrnl!IoPlnvalidDeviceRequest
IRP_MJ_QUERY_QUOTA	ntoskrnl!IoPlnvalidDeviceRequest
IRP_MJ_SET_QUOTA	ntoskrnl!IoPlnvalidDeviceRequest
IRP_MJ_PNP	serial!SerialPnpDispatch

Figure 5-9. Code View

#### Entries View

This view displays a set of common entry points in the driver. The location column supports copy and paste, and potentially debugger actions, via the context menu (right-click in the entry of interest).

EntryPoint	Location
DriverInit	serial!GsDriverEntry
DriverStartIo	0x0000000000000000
DriverUnload	serial!SerialUnload
AddDevice	serial!SerialAddDevice

Figure 5-10. Entries View

## Devices View

This view displays a current enumeration of associated device objects for the driver. Most device information is available on a given line, but you can also select a row and click the **Device Details** button to access the **Device Object Display**. If the selected device object contains a DevNode link, the **Device Node** button is enabled, and switches the page to DevNode mode for that entry. (See “[DevNode Exploration](#)” on page 106). Applicable columns are enabled for copy and paste, or debugger actions, via the context menu (right-click on the row and column of interest ).

Name	Link Target	Device Object	Device Object Extension	Device Extension	Type
Serial0	COM1	0xfffffadfe6eba2f0	0xfffffadfe6ebae80	0xfffffadfe6eba440	FILE
Serial1	COM2	0xfffffadfe6eb5080	0xfffffadfe6eb5c10	0xfffffadfe6eb51d0	FILE

Figure 5-11. Devices View

## Device Object Display

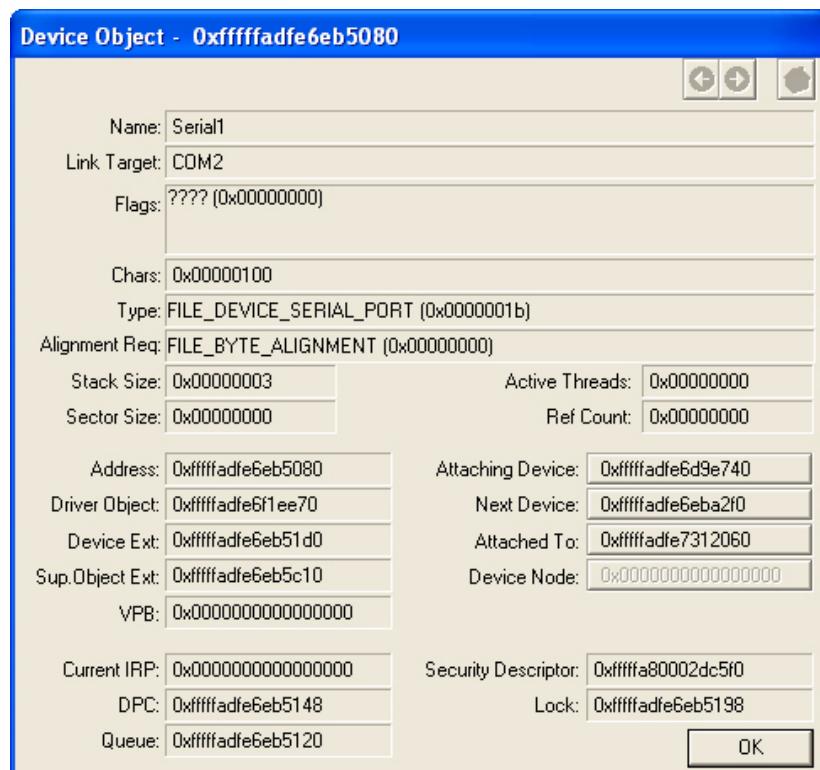


Figure 5-12. Device Object View

From several places in the Device-Driver page, the device object dialog is available. This dialog displays a summary, including flags and current state (the state as of the last event that refreshed the system, or since the last manual refresh). State will only be meaningful and valid for exploration when the target is stopped. However, getting a snapshot from a running machine can be also be useful.

### Summary Fields

All summary fields containing valid data support copy and paste, and potentially support debugger actions, via the context menu (right-click in the field of interest).

Some fields (e.g. Attaching Device, Next Device, Attached To, Device Node) contain a link to (address of) another Device Object (or DevNode) with a given relation. Clicking on these buttons will take you to that object.

### Link Buttons



If you follow a link to another Device Object, note that the link and home buttons at the top of the summary become active. These operate much like a web browser, allowing you to return back or move forward within the context of views you have already visited. Use the **Home** button to take you to the Device Object you started with.

**Note:** Clicking the Device Node button will dismiss the Device Object dialog before taking you to the device node of interest.

## System Information Commands

Visual SoftICE supports exploration of the target from the command line for much of the features in the Device-Driver page. If you prefer the command line interface, you should examine the DEVNODE, DRIVER, IMAGE, DEVICE, and ARBITER commands. For more information on all the available commands and what they do, refer to the System Information Commands section of the *Visual SoftICE Command Reference*.

## Page Features

### Open Only One Device-Driver Page

You can only have one Device-Driver page open.



### New (Remote Installation)

The **New** button opens a dialog that allows you start from an INF, a SYS, or other executable file, and put together the appropriate files and settings to accomplish a remote installation on the target.

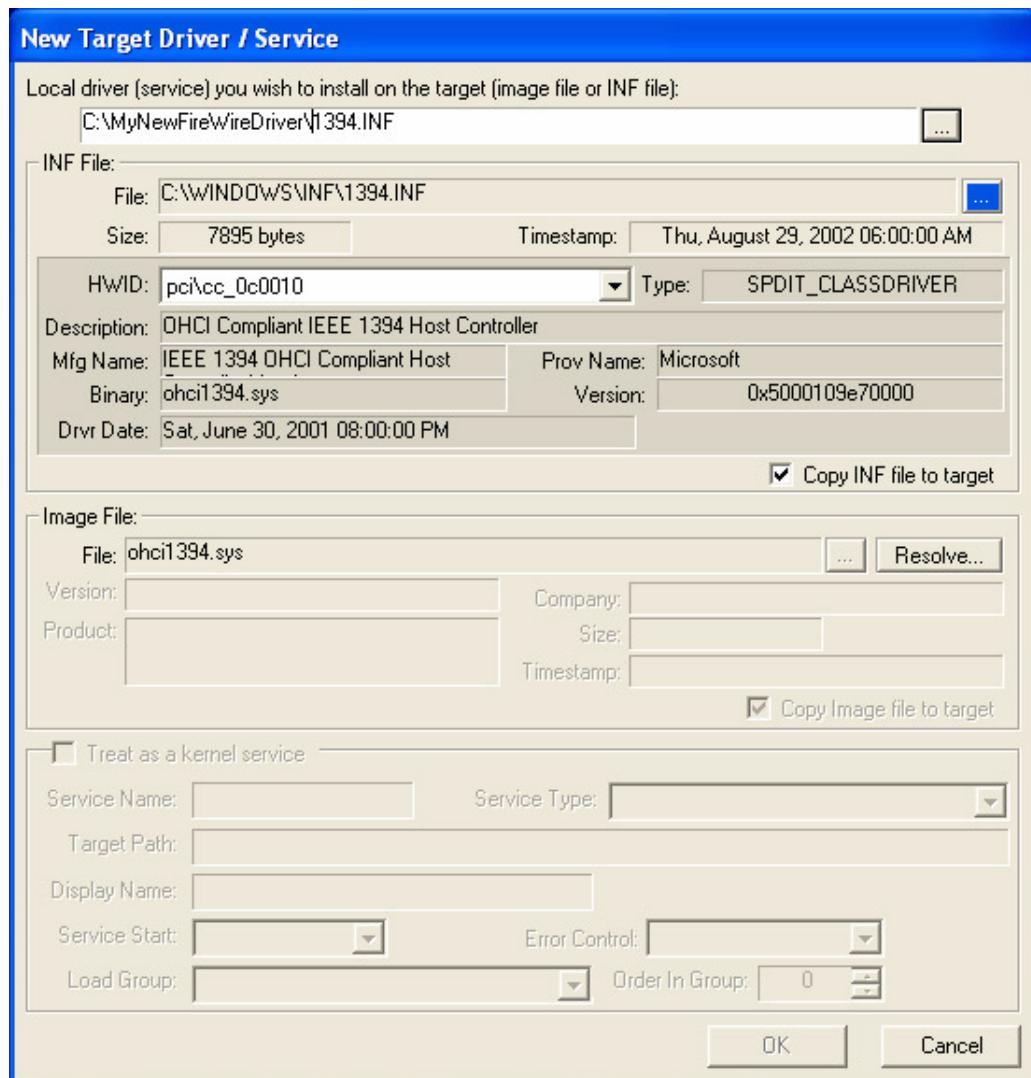


Figure 5-13. Remote Installation Dialog

The dialog is configured into 3 main sections (INF, IMAGE, and SERVICE), with a general file path entry field at the very top. Enter a path and filename in this top field and press <Enter> to start the process (or use the browse button (...) to select a file).

#### INF Files

If the entered filename is an INF, the dialog will show file information about it, parse it, and show available hardware ID's in the INF combo box, as well as displaying the appropriate service image filename in the IMAGE section of the dialog.

To continue with the installation process:

- 1 Validate the date and time on the INF itself.  
If the INF file you see is not the correct one, you can use the main input line, or use the browse button (...)within the INF section directly.
- 2 Select whether to copy the INF file to the target using the check-box in the INF section.
- 3 Select an appropriate hardware ID from the combo box.  
Changing this can change the service executable name in the IMAGE section.
- 4 Validate the HWID information.
- 5 Click the **Resolve** button to find a local instance of the named image file.  
This will work if the executable is in the same directory as the INF file, or you have set up the **Local Driver Search Path** preference for the Device-Driver page. If a matching instance is found, the image information section is filled in. If multiple instances are found, you are prompted to select one of the set. If none are found, the browse button (...) is enabled and you must find the file manually. Once selected, the image information section will be filled in.
- 6 Validate the image date, time, and other data.
- 7 Select whether to copy the image file to the target using the check-box in the IMAGE section.  
At this point the **OK** button should be enabled, as well as the **Treat As Kernel Service** check-box.
  - ◊ If this is a standard driver, click **OK** to install it on the target. If there are mismatches between what the INF file says and the image selected, you will receive appropriate warnings.
  - ◊ If this image should be a kernel service, check the **Treat As Kernel Service** check-box, fill in the appropriate information (defaults are generated), and click **OK**.

#### Executables

If the entered filename is an EXECUTABLE, the dialog will fill in the IMAGE section and enable the **Treat As Kernel Service** check-box.

To continue with the installation process:

- 1 Validate the image date, time, and other data.
- 2 Select whether to copy the image file to the target using the check-box in the IMAGE section.

It is presumed at this point that there is no INF file associated with this image, and you want to install this executable as a kernel service.

- 3 Check the **Treat As Kernel Service** check-box, fill in the appropriate information (defaults are generated), and click **OK**.



## Update (Remote Image and INF File Update)

The **Update** button starts a process to update the selected device driver on the target. This can be done from either the DevNode or Driver/Kernel Service view. From either view, the process is the same, but the source of information has a different starting point. An **Action Results** view displays results, and the update process follows this flow:

- ◆ The HWID of the selection is retrieved.
- ◆ The driver's load address is found.
- ◆ The driver is found in the active image list and the executable filename extracted (if not cached from a previous update).

**Note:** Update only works on active images on the target (Use FPUT for inactive images).

- ◆ An attempt to locate a local image filename is done. This uses the **Local Driver Search Path** preference, and if successful, operation continues. If no match is found, you are prompted to find one manually. If multiple matches are found in the search path, you are prompted to select one.

**Note:** An image file must be successfully found to continue.

- ◆ The local filename selected is cached for this device for future updates.
- ◆ An attempt to locate a local INF file is done. This uses the **Local INF Search Path** preference, and if successful, operation continues. If no match is found, you are prompted to find one manually. If multiple matches are found in the search path, you are prompted to select one.

**Note:** An INF file need not be found to continue.

- ◆ If so configured, a Pre-Update Script is run (configured under the per-page settings). This might be used to prepare a new version of any of these files (such as building, copying, or managing INF data).
- ◆ The active driver on the target is disabled.
- ◆ The new image and INF file (if matched) are pushed to the target.
- ◆ The driver is re-enabled on the target.

The results of the Update process remain in the **Action Results** view until you dismiss it, or directly switch to another main view (DevNode or Driver/Kernel Service).



## Un-Install (Remote Removal)

The **Un-Install** button starts a process to remove the selected device driver on the target. This can be done from either the DevNode or Driver/Kernel Service view. From either view, the process is the same, but the source of information has a different starting point. An **Action Results** view displays results, and the update process follows this flow:

- ◆ The HWID of the selection is retrieved.
- ◆ Warnings are presented if the device-driver is marked by certain flags, is a bus driver, or has a series of other devices dependent on it.
- ◆ If the driver is marked as built-in, no removal is supported.
- ◆ If so configured, a Pre-Removal Script is run (configured under the per-page settings)
- ◆ If a reboot is indicated for the removal, you are given control over whether this will actually occur.
- ◆ The removal action is attempted.

The results of the Un-Install process remain in the **Action Results** view until you dismiss it, or directly switch to another main view (DevNode or Driver/Kernel Service).



## Get State/Enable/Disable

These buttons are grouped together, the large arrow key to the left being **Get State** and other named buttons either being disabled (meaning state has not yet been retrieved), or reflecting the known state of the selected DevNode or driver. The Device-Driver page does not automatically retrieve the state of every selection since this operation is slow. However, once state is retrieved for a given selection, that information is cached.

- ◆ Click the arrow key to get, or refresh, the state of the selected DevNode or driver.
- ◆ Click the **Enable** button to enable a disabled DevNode or driver.
- ◆ Click the **Disable** button to disable an active DevNode or driver.

Any errors that might occur are reported in a pop-up dialog or in the **Action Results** view.



## Logging

The **Logging** button opens a dialog that allows you to configure the logging attributes of the SetupAPI used when doing an installation, enable, disable, or removal through Visual SoftICE.

## SetupAPI Logging Level

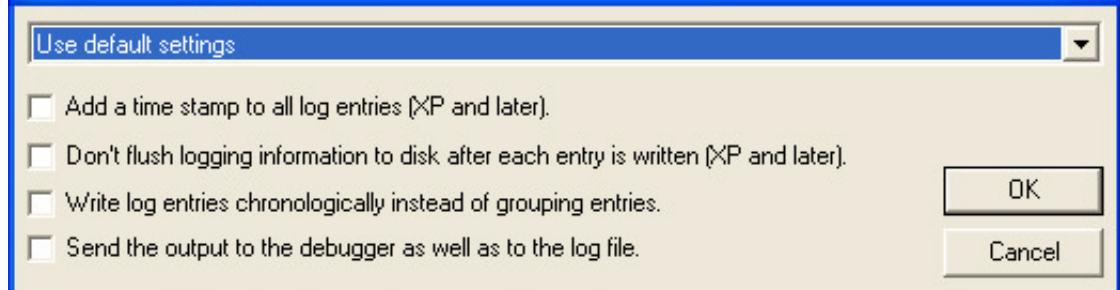


Figure 5-14. SetupAPI Logging Level Dialog

**Note:** See the Windows DDK documentation for details on the Windows Setup API.



### Refresh

The **Refresh** button forces the data in DevNode and Driver list areas of the Device-Driver page to be updated. This is normally not necessary. However, installation, removal, and enabling actions on certain drivers can make the target unstable, and may cause Visual SoftICE to not capture important state events. If you believe the display does not correctly reflect the expected state of the target, use this button to refresh it.

### Copy, Paste, Drag, and Drop

The Device-Driver page supports many ways of retrieving its data.

- ◆ You can copy most any field to the clipboard.
- ◆ You can drag certain fields to other pages and drop the contents on them.

### Customize the User Interface

There are multiple attributes of the page that you can customize. They are divided into two categories: per-workspace and application wide settings. Per-workspace attributes are specific to a particular workspace and not global to the application.

#### Per-Workspace Settings

- ◆ You can use the **Local Driver Search Path** setting to control where the page searches for drivers and other executables for actions such as Update and New. If the path ends in ellipses (...), then all sub-directories are also searched.

## Application Wide Settings

- ◆ You can use the **Local INF File Search Path** setting to control where the page searches for INF files for actions such as Update and New. If the path ends in ellipses (...), then all sub-directories are also searched.
- ◆ You can use the **Target Working Directory** setting to control the location on the target system for files to be pushed for installations. The installation process itself will place INF and executables in the appropriate locations. The default is the target system's %TEMP% environment variable.
- ◆ You can use the **Pre-Remove Script** setting to specify the name of a script to run on the master before removing a selected device-driver on the target.
- ◆ You can use the **Pre-Update Script** setting to specify the name of a script to run on the master before updating a selected device-driver on the target.
- ◆ You can use the **Live Mode On Connection** setting to control whether or not the page switches automatically to Live mode when a connection to a new target is established.

The Device-Driver page uses a separate file named PCIDEV.CSV that is located in the same directory as the main executable (DS.EXE). This file contains matching information for thousands of PCI devices, and can be extended and maintained by hand. It is an ASCII data file composed of a C++ comment section which gives directions and attribution, followed by a quote surrounded, comma delimited table generated from the web site <http://pcidatabase.com>.

As shipped this file contains directions for editing and maintaining it, encouragement to share information through the web site, and appropriate acknowledgement of the source, which is repeated here for the record.

**Table 5-3.** Source Acknowledgement

---

```
//-----  
//Comma Delimited PCI Device Table  
//-----  
//Created automatically from the web using the following URL:  
//http://pcidatabase.com/  
//Software to create and maintain the PCICODE List written by:  
//Jim Boemler (jboemler@halcyon.com)  
//  
//Too many people have contributed to this list to acknowledge them all, but  
//a few have provided the majority of the input and deserve special mention:  
// Frederic Potter, who maintains a list for Linux.  
// Chris Aston at Madge Networks.  
// Thomas Dippon of Hewlett-Packard GmbH.  
// Jurgen ("Josh") Thelen  
// William H. Avery III at Altitech  
// Sergei Shtylyov of Brain-dead Software in Russia  
//  
// NOTE that the 0xFFFF of 0xFF entries at the end of some tables below are  
// not properly list terminators, but are actually the printable definitions  
// of values that are legitimately found on the PCI bus. The size  
// definitions should be used for loop control when the table is searched.  
//-----  
//  
// You may extend this file at will, however you are encouraged to enter new  
// information into the web site at PCIDATABASE.COM in order that the entire  
// industry may benefit.  
//  
//-----
```

---

## The Disassembly Page

The Disassembly page is a read-only container to disassemble target memory. By selecting different modes, the page can be used as a live display, for viewing disassembly code on a specified target memory address, or for capturing a snapshot of a chunk of disassembled lines. Use the Disassembly page to:

- ◆ Display a chunk of disassembled lines, including current Instruction Pointer (IP), or at a specific target address.
- ◆ Set, remove, enable, and disable breakpoints on target addresses.
- ◆ Set and remove bookmarks on target addresses.
- ◆ Step through instructions.
- ◆ Perform **Go To** and **Run To** by target address or function name, current IP, and bookmarks.
- ◆ Search for strings inside the current view.
- ◆ Print out the page.

You can open a disassembly page by clicking the Disassembly page icon on the toolbar. Visual SoftICE can also automatically open a disassembly page while stepping or receiving WC or T commands if it cannot find a matching source file for the current IP. The automatic behavior of this page is controlled by **AutoFocusOpen** setting under **Source/Disasm Page** on the Global Settings tab of the Preferences dialog.

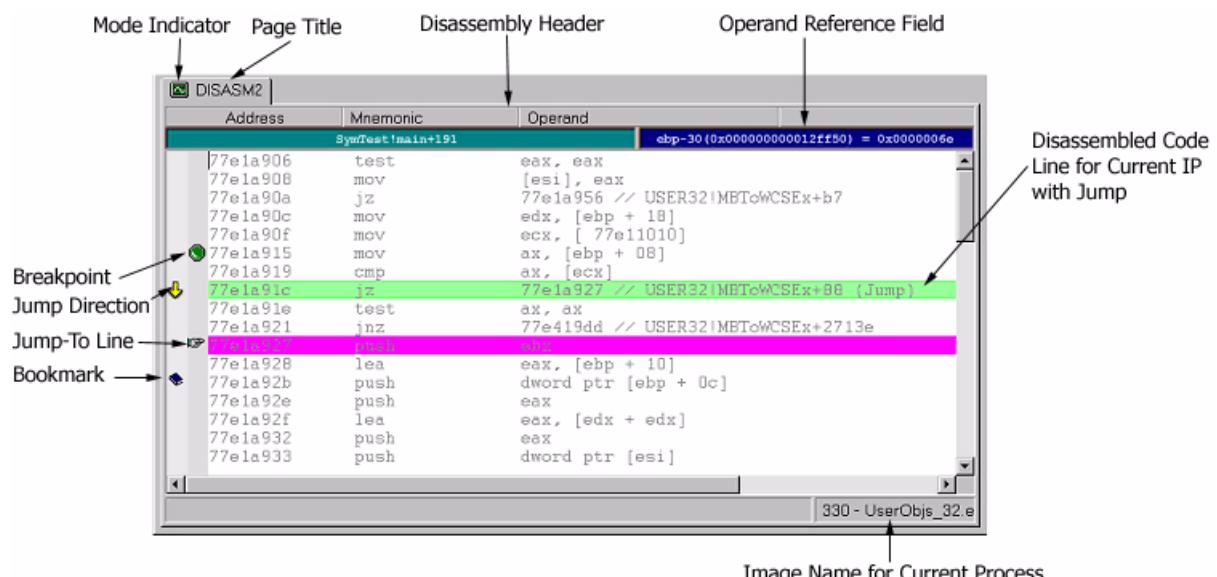


Figure 5-15. Visual SoftICE Disassembly Page

## **Associated Commands**

### **T**

The T command steps one instruction when the target is stopped. Visual SoftICE searches through the symbols to locate a file containing the current instruction pointer (IP). If it finds the file, and a source page for that file is already available, Visual SoftICE updates the existing source page. If it finds the file and no source page for that file is open, Visual SoftICE creates a new source page.

If Visual SoftICE does not find a source file containing the current IP, and a live disassembly page for the current IP is available, it updates that page. If no live disassembly page is available, Visual SoftICE opens a new disassembly page.

### **WC**

When you type WC [address], Visual SoftICE searches through the symbols to locate the source file containing the specified address. If Visual SoftICE finds the source file, and a source page for the file is already open, it brings the opened page forward. If the source file is not already open, Visual SoftICE creates a new source page for the file and brings it forward.

If the address is not contained in any source file, Visual SoftICE creates a disassembly page and brings it forward.

## **Page Features**

### **Open Any Number of Disassembly Pages**

There are no restrictions on the number of Disassembly pages you can open.

### **Mode Selection**

Page modes have some specific affects to the Disassembly page:

- ◆ **Live Mode** - If the target is connected and stopped, it automatically updates the display to include the disassembly line for current IP upon target events. Instruction stepping can be performed.
- ◆ **Manual Mode** - The display is not updated upon target events. You can choose to display a chunk of disassembly lines for a specific target address. Instruction stepping is not allowed.

- ◆ **Snapshot Mode** - The page is only used to view the current chunk of disassembly lines. It will not update the contents upon target events. Instruction stepping is not allowed.

For more detailed information on Page Modes and how they work in Visual SoftICE, refer to “[Page Modes](#)” on page 87.

## Breakpoints

When the master is connected to a target, you can set breakpoints on the disassembly page at any viewable target address. You can set breakpoints by locating the cursor on the line and using one of the following methods:

- ◆ Selecting **Set Breakpoint** under the Breakpoint group on the pop-up menu for the page
- ◆ Pressing the F9 function key
- ◆ Clicking the Breakpoint button on the main toolbar
- ◆ Right-clicking in the second column to the left of the code to access the Breakpoint pop-up menu.

Breakpoints are represented at the location you set them by a Breakpoint Icon that is specific to the type and state of the breakpoint. Once you set a breakpoint, you can disable or remove it via one of the following methods:

- ◆ Selecting the applicable options from the breakpoint pop-up menu
- ◆ Using the function keys
- ◆ Using the buttons on the main toolbar

Any breakpoint changes made elsewhere in the system are synchronously updated on the disassembly page.

## Operand Reference Field

If an operand to an instruction is a reference, then its value and its dereferenced value (what it points to) are displayed in the Operand Reference field. The contents of this field are updated as the instruction pointer moves through the list of disassembled instructions.

## Instruction Stepping

When the master is connected to a target that is stopped at an address, the disassembly line of the current IP is highlighted. You can choose to step-over, step-into, or step-out of the current instruction. Use the options under the **Debug** group of the pop-up menu for the page, the function keys as defined in the Preference settings, or the stepping buttons on the main toolbar, to perform your stepping functions.

If the current instruction has a jump-to address to branch to, the yellow arrow on the left of the highlighted line changes its direction to up or down, depending on the branch address location, indicating that the next instruction will be at a specific address. To view the jump-to line, you can click the arrow, press **<Ctrl><J>**, or select **Go To Jump** from the pop-up menu for the page. Visual SoftICE highlights the jump-to line when it displays it.

You can always view the current IP line by pressing **<Ctrl><I>**, or by selecting **Go To Current IP** from the pop-up menu for the page.

To run-to and stop at a specified line, you can set the cursor to the line where you want to stop and press **<Ctrl><R>**, or select **Run To Here** from the pop-up menu for the page.

## Bookmarks

You can set bookmarks at any viewable disassembly line, but only if the target is connected. Bookmarks in the disassembly page are represented by a target address. To set a bookmark:

- ◆ Right-click on the line and select **Set Bookmark** from the pop-up menu for the page
- ◆ Right-click in first column to the left of the code and select **Set Bookmark** from the pop-up menu

To remove bookmarks:

- ◆ You must have focus on a line with a bookmark on it. You can then remove the bookmark for the line, or all bookmarks in the file, by right-clicking and selecting the applicable option from the page's pop-up menu.
- ◆ You can locate the cursor in the first column to the left of the code and right-click to access the Bookmark pop-up menu.

You can view the bookmark list in the **Go To** dialog. From that dialog you can view or run-to the specified bookmark via the **Go To** and **Run To** buttons.

Visual SoftICE automatically removes all bookmarks when the target is disconnected.

## Go To

You can use the Go To dialog to:

- ◆ View or run-to a disassembly line using its target address or function name
- ◆ View a line containing the current IP
- ◆ View or run-to a selected bookmarked line

## Go To Current IP

You can right-click on the page and select **Go To Current IP** from the pop-up menu to view the line containing the current IP.

## Go To Jump

When you are on a jump statement, you can right-click on the page and select **Go To Jump** from the pop-up menu to move to the line where the jump will go.

## Run To Here

You can right-click on the page and select **Run To Here** from the pop-up menu, and Visual SoftICE will run from the current IP to the line that currently contains the cursor.

## Set Next Statement

The Set Next Statement command moves the IP to the next statement. You can execute the Set Next Statement command via one of the following methods:

- ◆ Right-clicking to access the pop-up menu for the page
- ◆ Pressing <Ctrl><Shift><F10>
- ◆ Selecting the command from the Source Page drop-down menu
- ◆ Clicking the **Set Next Statement** toolbar button

If you attempt to execute the Set Next Statement command on a specified line with an address outside of the function scope of the current IP, and the warning level is not set to off, then Visual SoftICE displays a warning message asking you to confirm that you want to set the IP to the new address.

## Find

You can use the Find dialog to search for a string in the current disassembly page.

## Copy

The page only supports copy command to retrieve its data.

- ◆ You can select any text within the page and copy it to the clipboard.
- ◆ You can use **Select All** from the pop-up menu to highlight all the lines on the page before copying them to the clipboard.

## AutoFocus

You can use the AutoFocusOpen setting to customize the behavior of the disassembly page when a target event occurs. If you select **Source or Disassembly** or **Source or Disassembly - No Focus**, then when Visual SoftICE can locate a source file for the symbols, it creates a source page. If Visual SoftICE cannot locate a source file, it creates a disassembly page.

For more information on available settings and how they behave, refer to AutoFocusOpen Settings in the online help.

## Customize the User Interface

There are multiple attributes of the page that you can customize. They are divided into two categories: per-page and application wide settings. Per-page attributes are always remembered in the workspace when you save it.

### Per-Page Settings

### Application Wide Settings

There are no per-page settings for the Disassembly page.

- ◆ You can use the **AutoFocusOpen** settings to control whether Visual SoftICE opens a new page when none exists (or brings an existing page to the top of the pad it is on), and places input focus on that page.
- ◆ You can use the **New Src/Disasm Pages to Largest Pad** setting to configure Visual SoftICE to place new Disassembly pages on the largest pad.
- ◆ You can use the **New Src/Disasm Pages to NAMED Pad** setting to configure Visual SoftICE to place new Disassembly pages on a specific named pad.

- ◆ You can use the **Live Mode On Connection** setting to configure Visual SoftICE to put Disassembly pages into live mode upon connection to a live target.
- ◆ You can use the **Minimum Column Widths** setting to configure Visual SoftICE to restrict the minimum column width on the Disassembly page to the actual field length.
- ◆ You can configure Visual SoftICE to display or hide the machine language (op-code) bytes column by right clicking on the Disassembly page, selecting **Disassembly**, and toggling **Show Op-Code Bytes** in the sub-menu. You can also toggle this setting under the **Source/Disassembly** element in the Global Settings tab of the Preferences dialog.
- ◆ You can configure Visual SoftICE to display or hide the address column by right clicking on the Disassembly page, selecting **Disassembly**, and toggling **Show Addresses** in the sub-menu. You can also toggle this setting under the **Source/Disassembly** element in the Global Settings tab of the Preferences dialog.
- ◆ You can configure Visual SoftICE to display or hide the instruction template field (on IA-64 only) by right clicking on the Disassembly page, selecting **Disassembly**, and toggling **Show Instruction Template** in the sub-menu. You can also toggle this setting under the **Source/Disassembly** element in the Global Settings tab of the Preferences dialog.

### Print

The page supports printing, and print-previewing of its contents.

## The Event Page

The Event page is a read-only display of information from the target. The page is composed of a single window, which displays events, and a pop-up utility controlling the filtering of the events that are displayed in the window.

You may have only one Event page per target. If the connection to the target is lost, the collection of events is cleared from the page.

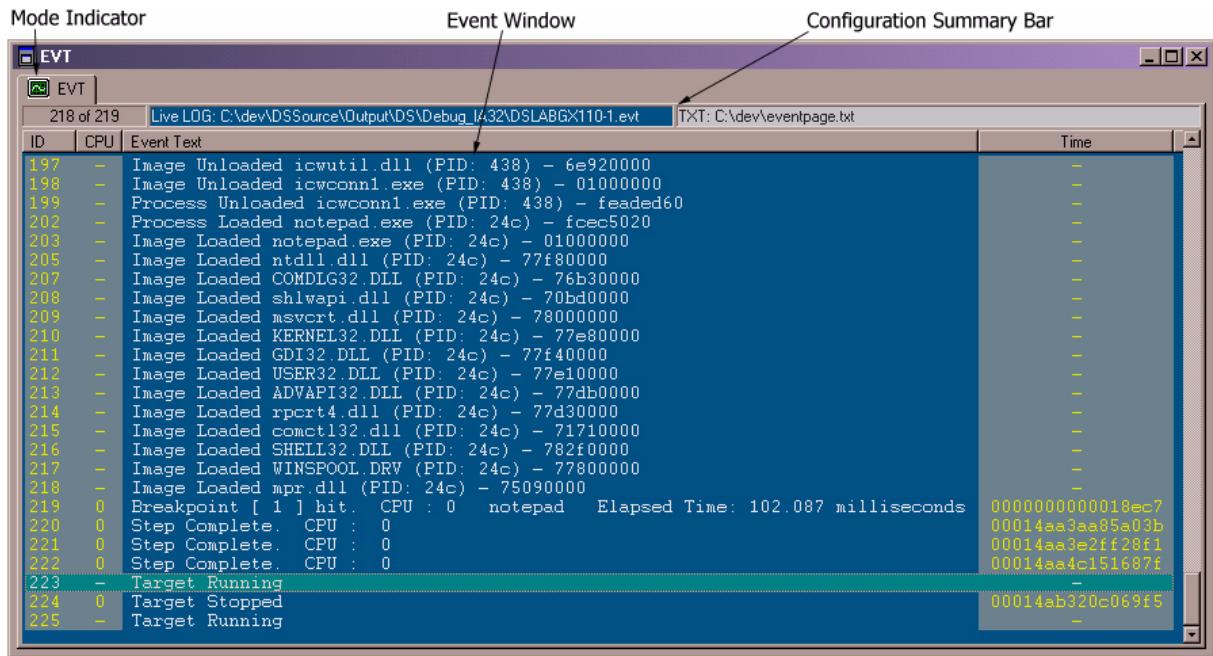


Figure 5-16. Visual SoftICE Event Page

## Associated Commands

### SET EVENTLOGBACKUPS

Use the SET EVENTLOGBACKUPS command to specify how many event log files to maintain on a per-target basis. An event log is a binary file containing all the events that occurred on the target, for the duration of the debugging session. Normally these files are named `TARGETNAME.EVT` (e.g. a target named `MyUnitTestMachine` would end up having a file named `MyUnitTestMachine.EVT` created).

The default value of `EVENTLOGBACKUPS` is zero, which indicates that any pre-existing event log file of the same name, will be overwritten. Otherwise event log files are given a datetime stamp appended to the filename, up to the number of files you specify, before overwriting the oldest.

## SET EVENTLOGPATH

Use the SET EVENTLOGPATH command to specify the directory where you want to create and maintain any event log files for the target. The default value of EVENTLOGPATH is set to the operating system directory where the master is running.

## Page Features

### Configuring Filtering

The Event page supports the filtering of events output to the page. Event filtering is controlled by the Target Event Display Filter utility, which is available by right clicking on the page and selecting **Configure Filtering** from the pop-up menu.

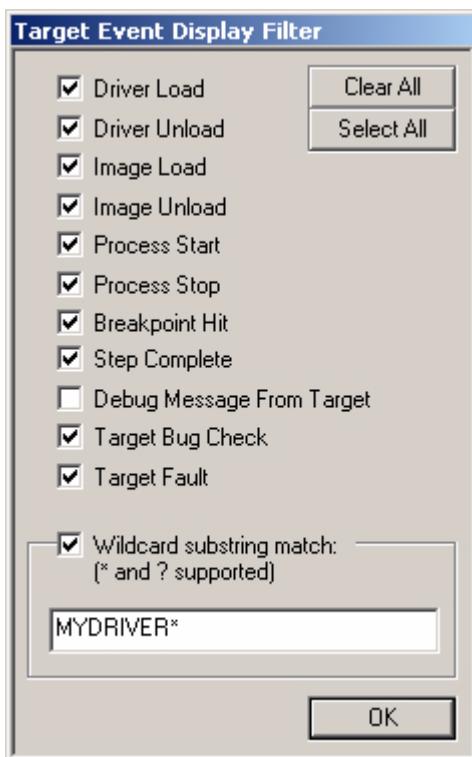


Figure 5-17. Target Event Display Filter Utility

- ◆ You can enable or disable specific Target Notification events, such as Driver Load, Driver Unload, or Debug messages.

- ◆ You can filter events based upon a substring match of the event content. The filtering is limited to file system level matching criteria. That is, the '?' and '\*' operators are the only supported matching wildcards. For example, to display the debug messages that your driver sends, prefix your DBGPRINT statements with a tag like MYDRIVER:, then set the regular expression event matching to "MYDRIVER: \*".

**Note:** The Event page will always attempt to apply any filtering you implement to the entire event collection. With extremely large event collections this process can take some time, so if the collection is very large and the target is running, the Event page will halt the target while it synchronizes the event collection.

## Viewing Event Details

The Event page allows you to access details for a selected event listed in the page. To open the event details dialog, double-click on an event or right-click on the event and select **Details** from the pop-up menu. Not every event will have further details available.

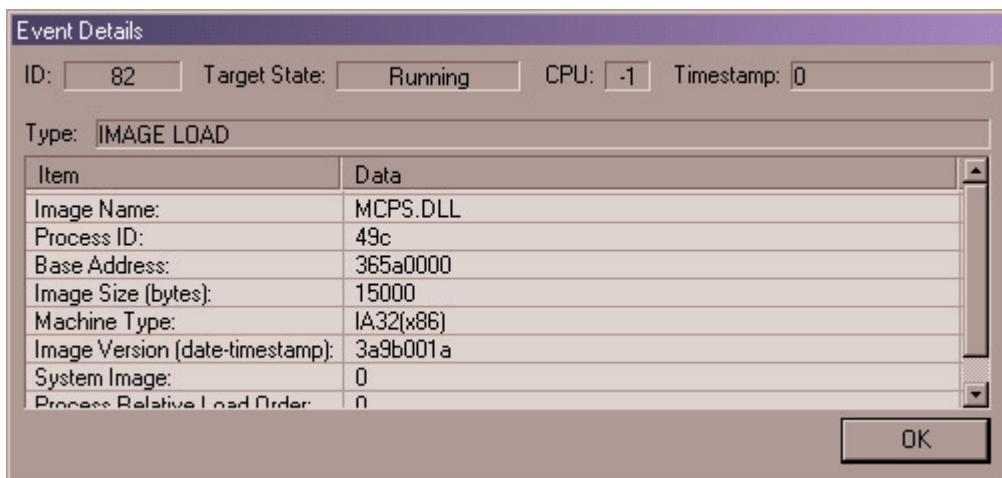


Figure 5-18. Event Page Details Dialog

## Open Only One Event Page

You can only have one Event page open.

## Open Old Log Files for Display

You can open old log files for display by right-clicking on the page and selecting **Open Old Event Log** from the pop-up menu.

## Configuration Summary Bar

The Configuration Summary bar displays the event counter, Log file name, and Text auto-save filename. To display or hide the Configuration Summary bar, right-click on the page and toggle the **View Config Summary** option in the pop-up menu. Access the Page Preferences to set a new auto-save text file name by clicking on the **Auto Save Text Log Filename** field.

## Copy

The page only supports copy command to retrieve its data.

- ◆ You can select any text within the page and copy it to the clipboard.
- ◆ You can use **Select All** from the pop-up menu to highlight all the lines on the page before copying them to the clipboard.

## Customize the User Interface

There are multiple attributes of the page that you can customize. They are divided into two categories: per-page and application wide settings. Per-page attributes are always remembered in the workspace when you save it.

### Per-Page Settings

- ◆ You can filter the events that are output to the page. For more information, refer to “[Configuring Filtering](#)” on page 129.
- ◆ You can display or hide the Configuration Summary bar at the top of the Event page by right-clicking on the page and toggling the **View Config Summary** option.
- ◆ You can show or hide the ID, CPU, and Time columns by right-clicking on the column heading and selecting Customize. Use the Customize dialog that appears to toggle columns as displayed or hidden.

### Application Wide Settings

- ◆ You can use the **Fonts & Colors** tab to select the background and foreground color for event text and all other columns.
- ◆ You can configure an **Auto Save Text Log Filename** from the Global Settings tab by right-clicking on the page and selecting **Page Preferences** from the pop-up menu. Alternatively, if you have the Configuration Summary bar displayed you can click on the **Auto Save Text Log Filename** field, and VSI automatically opens the Page Preferences to allow you to modify that field.

- ◆ You can configure the page to **Auto Save Text Log on Disconnect** from the Global Settings tab by right-clicking on the page and selecting **Page Preferences** from the pop-up menu.
- ◆ You can use **Edit/Copy Text Field** from the Global Settings tab to enable or disable the ability to access a text-field edit control. When enabled, clicking on a field, or double-clicking on the text field, opens the edit control where you can select all or a portion of the string. When disabled, clicking on an entry does nothing, and double-clicking accesses the Details dialog. This option is disabled by default.
- ◆ You can configure the page to **Warn on Disconnect Data Loss** from the Global Settings tab by right-clicking on the page and selecting **Page Preferences** from the pop-up menu.
- ◆ You can configure the **Text Log File Write Mode** from the Global Settings tab by right-clicking on the page and selecting **Page Preferences** from the pop-up menu.

## Print

The page supports printing, and print-previewing of its contents.

## The Locals Page

The Locals page is a read-write container for displaying and editing the local variables for the current context. The Locals page can display the variable name, address, type, size, location, and value. The Locals page also has a page-specific context bar, allowing you to change threads or stack frames, which overrides the main context bar.

Mode Indicator   Thread ID   Stack Frame

**LOC**

Thread: TID: 1b4   Context: SymTest!main+db0

Locals Information

Name	Type	Value
_argc	int	0x00000001
+ _argv	char**	0x000006fbff661210
+ hHandle	struct HINSTANCE__*	0x0000000000000000
c	int	0x000a2fb2
+ classA	class A	{...}
+ IntArray[10]	int	0x000006fbff8fd8
+ szText[256]	char	0x000006fbff8cf8 "Yes, we are here to test the local page"
+ pszText	char*	0x000006fbff660ef0
b	int	0x00000014
+ enumArray[10]	enum MYENUM	0x000006fbff8fc0
+ pClassB	class B*	0x000006fbff661160
+ myStruct	struct _MY_STRUCT	{...}
iMyInt	int	0x0000000d
dwMYDword	float	99.000000
+ szMtArray[27]	char	0x000006fbff8fc98 "ABCDEFHIJKLMNOPQRSTUVWXYZ!"
+ BitField	struct _BIT_FIELD	{...}
ubit1	unsigned long	0x00000000
ubit2	unsigned long	0x00000000
ubit3	unsigned long	0x00000001
ubit4	unsigned long	0x00000000
i	int	0x0000001a
+ pns	struct tag_NewStructor*	0x000006fbff661040
a	int	0x000a2fb2
+ ppns	struct tag_PointerNewStructor*	0x000006fbff6610a0
enumData	enum MYENUM	ENUM_THREE (3)
+ ns	struct tag_NewStructor	{...}
m_intA	int	0x000a2fb2
m_intB	int	0x00000014

Figure 5-19. Visual SoftICE Locals Page

## Concepts and Associated Commands

### WL

The WL command opens the Locals page. If a Locals page is already open, WL activates that page.

## Page Features

### Live Mode Variable Tracking

When the Locals page is in live mode, Visual SoftICE displays the variable values in red when they change. You can expand the variable if it contains child-level components.

## Open Only One Locals Page

You can only have one Locals page open.

## Toggle Values Displayed Between Decimal and Hexadecimal

You can toggle displaying values in hexadecimal or decimal. To toggle the display format for the values, right-click and select **Hexadecimal** from the pop-up menu. When it is checked, Visual SoftICE displays values in hexadecimal. When it is unchecked, Visual SoftICE displays non-address values in decimal.

## Edit Variable Values

Visual SoftICE allows you to edit the value of a variable displayed in this page. To modify the value of a variable, double-click the value column, and enter the new value in the field. Press <Enter> or click any place in the Local page to save the value, or press <Esc> to cancel editing.

## Select a Context

Visual SoftICE allows you to select a different stack frame, or context, whose variables you wish to display. To switch contexts for the current thread, click the Context combo-box to access the drop-down list of available stack frames for the current thread, and select the stack frame you want to change to. Visual SoftICE displays the variables for the selected context.

## Select a Thread

Visual SoftICE allows you to select a different thread whose variables you wish to display. To select a different thread, click the Thread combo-box to access the drop-down list of available threads for the current process, and select the thread you want to change to. Visual SoftICE automatically updates the stack frames for the selected thread, and displays the latest variables for the first stack frame.

## Copy and Drag

The page supports copying and dragging of data to other Visual SoftICE pages. How Visual SoftICE handles the variables once you drop them depends on the page you drop them into.

## Customize the User Interface

There are multiple attributes of the page that you can customize. They are divided into two categories: per-page and application wide settings. Per-page attributes are always remembered by the workspace when you save it.

### Per-Page Settings

- ◆ You can configure the page to display any of six available columns of data (Name, Address, Type, Size, Location, and Value). Name, Type, and Value are displayed by default. To customize which columns are displayed, right-click on a column heading to access the pop-up menu.
- ◆ You can resize the individual columns displayed. Double-clicking on any column heading will auto-resize the columns to fit the content.

### Application Wide Settings

- ◆ You can toggle the display format for variable values between decimal and hexadecimal.
- ◆ You can hide or display the page-specific context bar. To toggle display of the context bar on or off, right-click on the page and select **View Context Bar** from the pop-up menu.

## Print

The page supports printing, and print-previewing of its contents.

## The Memory Page

The memory page is a container for displaying and editing blocks of memory on a target machine, or from file sources. This is a shared resource for many DriverStudio components, including BoundsChecker for Drivers, and Visual SoftICE. Memory loaded into this page can be viewed in scalar modes (BYTE, WORD, DWORD...), string, and as structures and classes. Access to physical and virtual memory, display of symbols, formatting to non-scalar types, editing, and data drag and drop to other pages is all driven by the source of the memory loaded into the page (for example, BoundsChecker for Drivers does not support formatting memory into types known to the symbol engine, while Visual SoftICE does). You can retrieve memory by physical or virtual address.

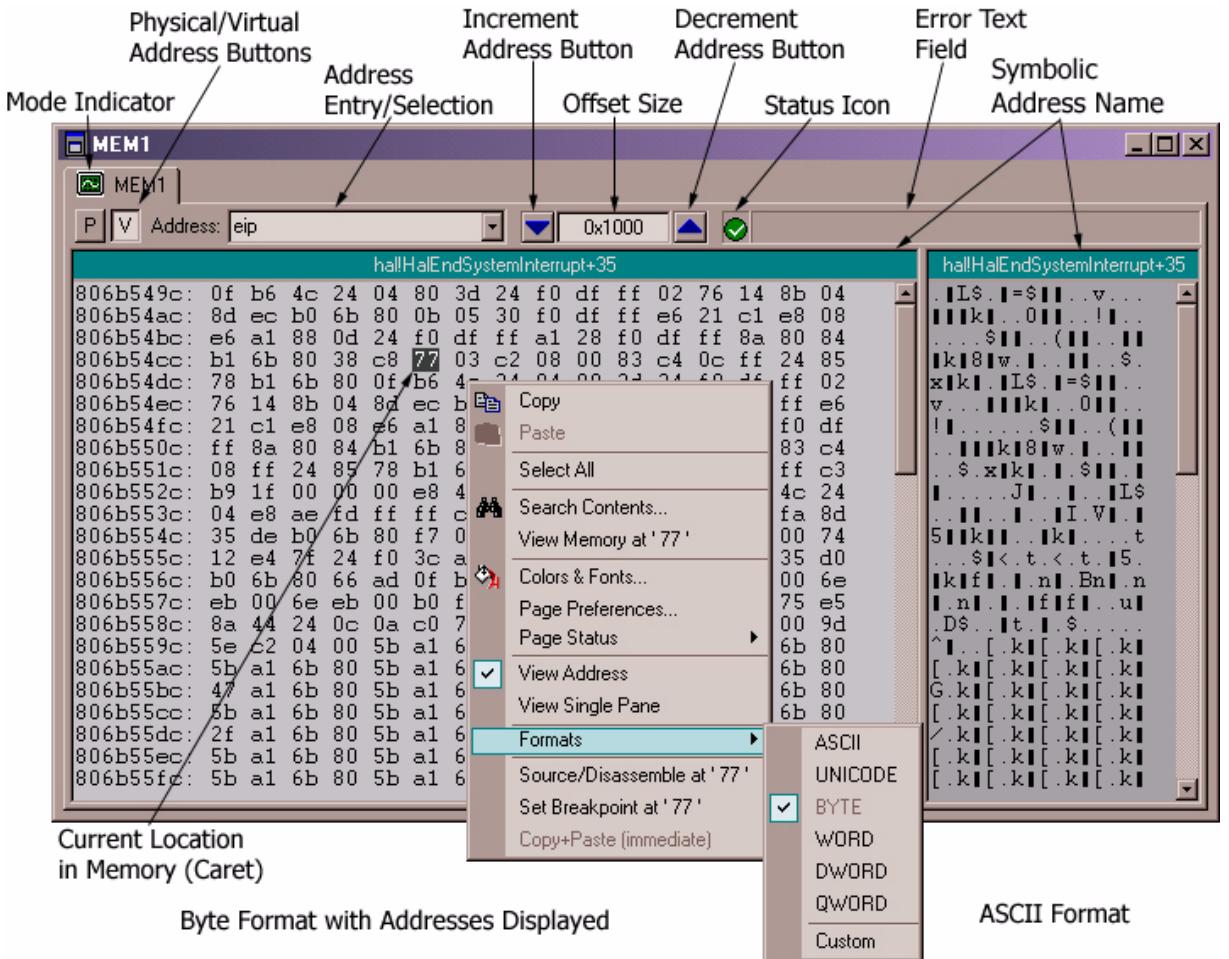


Figure 5-20. Visual SoftICE Memory Page

## ***Concepts and Associated Commands***

## Address Types

You can access memory on the target machine via its physical location (the location on the memory bus where the memory unit of interest is located), or via virtual address per process (as implemented by the OS). Access to memory on the target can fail if one of the following situations is true:

- ◆ The virtual address requested is paged out
  - ◆ The virtual address is unmapped (invalid) in the current process
  - ◆ The memory is inaccessible based on processor state

Virtual memory access is the most common form of memory access, and is connected with static and dynamic allocations for a given process or the kernel itself.

Physical memory access is useful for looking at I/O registers, video memory, or other regions physically mapped (memory mapped) by the processor. This means the memory of these devices appears to be part of the standard address space of the processor.

You cannot access unmapped memory on the target, or memory held by a specific device on some other bus, in the same manner. Refer to the PCI command for exploring device-managed resources off that bus.

Visual SoftICE can read memory while the target is running or stopped. Reading memory while the target is running retrieves a snapshot, and is not guaranteed to be accurate to current values.

## Virtual Memory

You can display virtual memory by:

- ◆ Selecting the virtual address type (**V**) button and then entering an address in the **Address** field and pressing <Enter>.
- ◆ Issuing the D command in Visual SoftICE.

## Physical Memory

You can display physical memory by:

- ◆ Selecting the physical address type (**P**) button and then entering an address in the **Address** field and pressing <Enter>.
- ◆ Issuing the PHYS command in Visual SoftICE (to display the virtual memory addresses that map to a physical address).
- ◆ Issuing the D command with the -p switch in Visual SoftICE.
- ◆ Issuing the PEEK command in Visual SoftICE.

## Symbols

All addresses may have a symbolic representation. Visual SoftICE automatically updates and displays the symbolic representation on the current address title bar as you move the caret through various memory locations. You can disable symbolic lookup (for performance improvements) for all memory pages via the global setting (Symbolic Address Lookup) under the Memory page preferences (File > Preferences > Global Settings > Memory Page).

### **Open Any Number of Memory Pages**

There are no restrictions on the number of Memory pages you can open. You can build a workspace with multiple views containing just the memory blocks you are interested in, displayed the way you want, where you want.

### **View Memory in the Format You Want**

Memory can be interpreted for display in a number of ways, including built-in scalar formats (BYTE, WORD, DWORD, QWORD, ASCII, and UNICODE), and custom types known by the symbol engine (type, structure, or class).

The default view of memory is 2 panes, BYTE view on the left with addresses, ASCII view on the right with no addresses. The Memory page remembers all scalar formatting preferences, the last custom format name you selected (per pane), and last address viewed (on a per-page basis) in the workspace when you save it. The last custom format name and the last address viewed are displayed in the pane pop-menu and address combo box respectively at the time DriverStudio loads the workspace.

You can customize the page to display memory in 2 separate panes or just one via the pop-up menu. If you display memory in 2 panes, the Memory page remembers the splitter location (on a per-page basis) in the workspace when you save it.

You can select a custom format off the pop-up menu, per pane. You should consider this formatting a cast operation on the block of memory. The format must be a known type in the symbol engine (only available in Visual SoftICE).

The selection dialog allows you to search using wildcards, and/or limit the types displayed based on your preferences. Additionally, if you know the format name, you can enter it directly and press <Enter>. If it is a known type in the symbol engine, Visual SoftICE will immediately apply it.

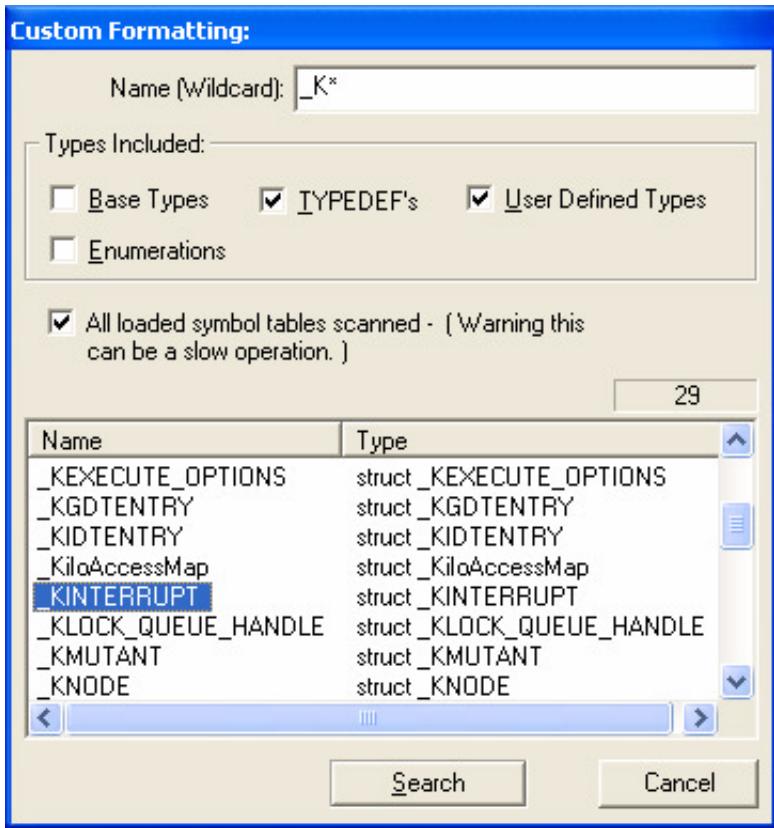
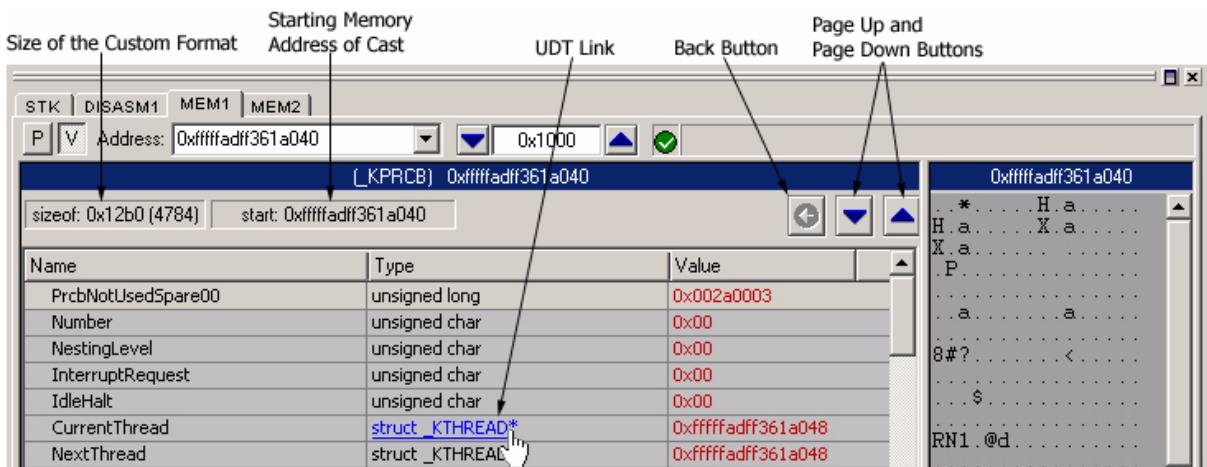


Figure 5-21. Custom Formatting Dialog

When a custom format is shown, the name of the format, size of the format, and two additional buttons (**PageUp**, **PageDn**) are displayed to walk through memory.



**Figure 5-22.** Custom Memory Formatting

Please note this is a cast operation on existing memory. Therefore these buttons may be useful for walking through simple arrays of fixed size types, structures, or classes, but do not walk lists (linked lists, pointers, or other collection mechanisms). See “Jumping via Active UDT Links”.

If you save the workspace with a pane displaying a custom format (type, structure, or class), Visual SoftICE remembers the format name, but does not apply it at the time it loads the workspace. This is because the format may be unknown in the current context (not in the loaded or active symbol tables). Once the workspace is loaded, you may reassign the formatting directly via the pop-up menu.

### Jumping via Active UDT Links

When viewing memory for a User Defined Type (UDT) structure, pointers to other structures are displayed as active links. When you click on a UDT link, Visual SoftICE will attempt to read memory at that address. If Visual SoftICE is able to read memory, it adds the new address to the address list, takes you to the new address, and then attempts to format it in the new UDT typedef.

Clicking the Back button will step you back to the previous address and previous typedef.

Links only appear on pointers of UDT type where the address is non-zero.

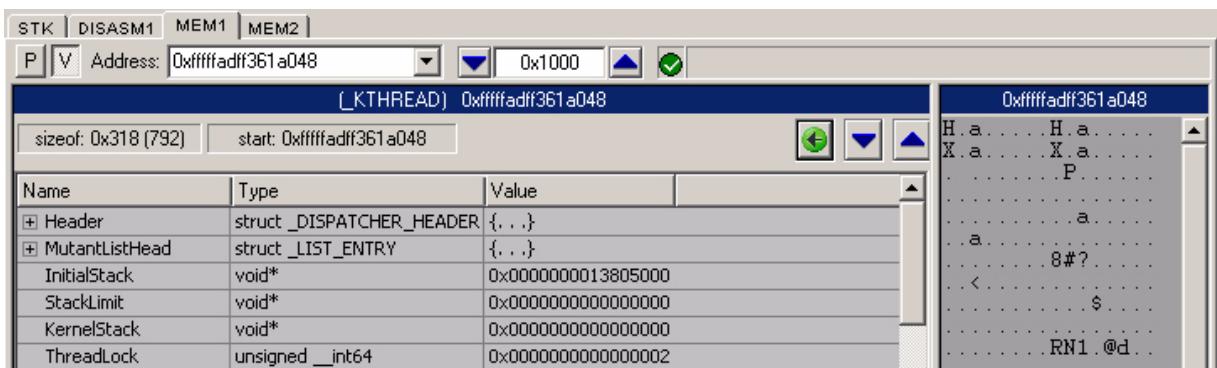


Figure 5-23. Following an Active UDT Link

## Search Page Contents

This page provides a utility to conduct byte pattern searching. The Search utility for the Memory page allows you to search the page's contents forward and backward, but restricts the searching to a 32-byte pattern length.

This utility behaves like the S command in Visual SoftICE, however it is restricted to the contents of the memory page, and can search forwards or backwards.



Figure 5-24. Memory Page Search Dialog

## Modify Starting Memory Address by Offset

This page provides a control bar that allows you to increment or decrement the starting memory address by a specified offset size. The control bar consists of three elements:

Table 5-4. Modify Starting Memory Control Bar Elements

Element	Description
Offset Size Field	This is the offset size (in bytes) to jump from the current address. The size entered in the field can be any expression, including expressions such as <code>sizeof(my_struct)</code> . What is displayed after evaluation is the hex value of the input expression. If the expression cannot be properly evaluated, the previous value is restored. The size you designate is saved and restored for each instance of the Memory page in the workspace. The default is 0x1000.
	Increment the starting memory address by the specified offset size. The <i>offset +current address</i> must be greater than zero.
	Decrement the starting memory address by the specified offset size. The <i>current address-offset</i> must be greater than zero.

## Status Icons

The Memory page has its own status bar that uses icons to display the resultant state of any actions taken on that page. Error text appears in the Error Text Field (or not) as appropriate for each state.

Table 5-5. Memory Page Status Icons

Icon	Meaning
	This icon indicates that everything is OK.
	This icon indicates a warning.
	This icon indicates a failure.

## Error Message Display

The Memory page has its own status bar that displays any error messages and warnings generated by actions taken within the page.

## Edit Memory

There are several ways of editing memory data, both in scalar and custom formatting.

- ◆ In scalar modes, you can double-click or select a unit of memory and press <Enter>, or the <Backspace> key, to get an in-place edit control. You can cancel editing by pressing the <Esc> key, or you can apply any changes by pressing <Enter>. If your edit fails to be accepted, the Memory page displays an error message. If accepted, the Memory page updates all panes displaying that memory location, using the modified color to highlight the change.
  - ◇ In non-string formats, this edit control will be constrained to the unit of memory (BYTE, WORD, DWORD, QWORD), and allows hexadecimal input only.
  - ◇ In string formats (ASCII, UNICODE), the edit control will be constrained to a full line of memory (as displayed) and allows string input. The current character at the start of editing will be highlighted within the line. On acceptance of your input, the string will be converted to an appropriate byte stream (ASCII or UNICODE) before being sent to the target.
- ◆ In custom format mode, you can directly edit the members of the type, structure, or class by clicking on the member value cell. The Memory page opens an in-place edit control. You can cancel editing by pressing the <Esc> key, or you can apply any changes by pressing <Enter>.
- ◆ You can edit virtual memory from the command line in Visual SoftICE by issuing the E command.
- ◆ You can edit physical memory from the command line in Visual SoftICE by issuing the POKE command.

## Copy, Paste, Drag, and Drop

The page supports many ways of retrieving its data.

- ◆ You can copy a single scalar unit, or range of lines, by highlighting them and then selecting **Copy** from the pop-up menu.
- ◆ You can drag a single scalar data value out of the memory page and drop it on any other page that accepts data this way.
- ◆ You can drop a single data value from any other page on the memory page. Visual SoftICE interprets the data value as an address, and places it in the address field (along with keyboard focus). While focus is in the Address field, you can edit the address, or append to it before pressing <Enter> to accept and display that address.

## Customize the User Interface

There are multiple attributes of the page that you can customize. They are divided into two categories: per-page and application wide settings. Per-page attributes are always remembered by the workspace when you save it.

### Per-Page Settings

- ◆ You can customize the page to display 1 or 2 views of memory (panes).
- ◆ You can customize the page to display addresses in each pane.
- ◆ You can customize the page to display a range of scalar or custom formats per pane.

### Application Wide Settings

- ◆ You can use the **Fonts and Colors** tab to select the font used, as well as the background and foreground color, for each pane.
- ◆ You can use the **Fonts & Colors** tab to select the color used to indicate a change in the value of memory between updates of the page.
- ◆ You can enable or disable **Symbolic Address Lookup** of every address location displayed in the page, as you move the caret through memory.
- ◆ You can use the **Default Pane 1 Scalar Format** setting to select the scalar format used for pane 1 when the page is created. The default is Byte format.
- ◆ You can use the **Default Pane 2 Scalar Format** setting to select the scalar format used for pane 2 when the page is created. The default is ASCII format.

- ◆ You can use the **Default to Virtual Memory** setting to configure the page to default to physical or virtual memory addressing when it's created. The default is True, which selects Virtual memory for the page.
- ◆ You can use the **Display Bytes Per Line** setting to set the number of bytes per line the Memory page displays in each pane (the valid range is 16 to 128 bytes, and must be evenly divisible by 2).
- ◆ You can use the **Read Block Size** setting to control the size of memory blocks the page requests from the target when you display memory. Larger memory blocks provide better performance, while smaller memory blocks reduce the chances of hitting paged out conditions (the valid size range is 128 to 8192 bytes).
- ◆ You can use the **Live Mode On Connection** setting to control whether or not the page switches automatically to Live mode when a connection to a new target is established.
- ◆ You can use the **New Memory Pages to Largest Pad** setting to designate new Memory pages to always open on the largest pad.
- ◆ You can use the **New Memory Pages to NAMED Pad** setting to designate new Memory pages to be created on a specific named pad.

## The Process List Page

The Process List page displays a list of the processes running on the target machine. It also uses a triple green arrow to indicate which process is the current process context, and a yellow arrow to indicate the process containing the current IP.

The diagram illustrates the Visual SoftICE Process List Page. At the top left is the 'Mode Indicator' showing 'PROC'. To its right are 'Columns with Process Details' labeled: Name, KPEB, Pid, Threads, Priority, UserTime, KernlTime, and State. A vertical bar on the left contains 'Process Names' (spoolsv..., msdtc.e...). Below it are 'Active Process' (junk.ex...) and 'Current IP' (Idle). Arrows point from the labels to their respective parts of the interface.

Name	KPEB	Pid	Threads	Priority	UserTime	KernlTime	State
spoolsv....	e0000165e9851bf0	3b4	9	8	4	49	Out Of Memory
msdtc.e...	e0000165e9c53260	3e8	15	8	d	34	Out Of Memory
DSRSv....	e0000165e9c102a0	450	7	8	6	26	Out Of Memory
llssrv.ex...	e0000165e9818f90	474	9	8	4	25	Out Of Memory
svchost....	e0000165e9ad8ef0	498	2	8	0	e	Out Of Memory
dfssvc....	e0000165e9c57030	51c	9	8	1	23	Out Of Memory
explorer...	e0000165e9803210	5ec	a	8	1d8	5df	In Memory
DStray....	e0000165e9879450	624	1	8	0	14	Out Of Memory
svchost...	e0000165e9ad8750	64c	d	8	7	2a	In Memory
wmiprvs...	e0000165e9bb8540	7a8	4	8	2d	41	Out Of Memory
wuauctl...	e0000165e9b68ad0	7f4	5	8	2	2b	Out Of Memory
junk.ex...	e0000165e9b019b0	638	1	8	8	9	In Memory
Idle	e000000081aedd00	0	1	0	0	ee4889	In Memory

Figure 5-25. Visual SoftICE Process List Page

## Concepts and Associated Commands

### Current Context

You can use the Process List page to make a process the “current” one. You can always make a process current, and that becomes the current context for Visual SoftICE. Refer to “About the Context Bar” on page 62 for more information on the main context bar, and page-specific contexts.

### Stopped vs. Running State

The page supports displaying only two levels of information for a process: basic and extended. The levels are tied to the target state. When the target is running, Visual SoftICE displays only the basic information for the processes. When the target is stopped, Visual SoftICE displays the extended information for each process.

The following basic information is displayed: **Name**, **KPEB address**, and **PID**.

The following extended information is displayed: **Name**, **KPEB address**, **PID**, **Thread Count**, **Priority Level**, **User Mode Time**, **Kernel Mode Time**, and **Process State**.

## Images

You can use the Process List page to view the images for a selected process. This works the same as the IMAGE command. To view the images for a process, right-click on the process, and select **Images** from the pop-up menu.

Images List for SERVICES.EXE			
Address	Size	Name	FullName
01000000	1b000	services.exe	C:\WINDOWS\system32\services.exe
5f770000	e000	NCDobjAPI.DLL	C:\WINDOWS\system32\NCDobjAPI.DLL
629c0000	8000	LPK.DLL	C:\WINDOWS\system32\LPK.DLL
71aa0000	8000	WS2HELP.dll	C:\WINDOWS\system32\WS2HELP.dll
71ab0000	15000	WS2_32.dll	C:\WINDOWS\system32\WS2_32.dll
71c20000	4000	NETAPI32.dll	C:\WINDOWS\system32\netapi32.dll
72fa0000	5a000	USP10.dll	C:\WINDOWS\system32\USP10.dll
758a0000	f000	EVENTLOG.DLL	C:\WINDOWS\system32\eventlog.dll
758e0000	1c000	umpnpmgr.dll	C:\WINDOWS\system32\umpnpmgr.dll
758e0000	4d000	SCESRV.dll	C:\WINDOWS\system32\SCESRV.dll
75a70000	a3000	USERENV.dll	C:\WINDOWS\system32\USERENV.dll
76360000	f000	WINSTA.dll	C:\WINDOWS\system32\WINSTA.dll
76390000	1a000	IMM32.DLL	C:\WINDOWS\System32\IMM32.DLL
76bf0000	b000	PSAPI.DLL	C:\WINDOWS\System32\PSAPI.DLL
76cc0000	10000	AUTHZ.dll	C:\WINDOWS\System32\AUTHZ.dll
76f50000	8000	WTSAPI32.dll	C:\WINDOWS\System32\wtsapic32.dll
76f90000	10000	Secur32.dll	C:\WINDOWS\System32\secur32.dll
77c10000	53000	msvcr7.dll	C:\WINDOWS\System32\msvcr7.dll
77c70000	40000	GD132.dll	C:\WINDOWS\System32\GD132.dll
77cc0000	75000	RPCRT4.dll	C:\WINDOWS\System32\RPCRT4.dll
77d40000	8d000	USER32.dll	C:\WINDOWS\System32\USER32.dll
77dd0000	8b000	ADVAPI32.dll	C:\WINDOWS\System32\ADVAPI32.dll
77e60000	e5000	KERNEL32.dll	C:\WINDOWS\System32\kernel32.dll
77f50000	a9000	ntdll.dll	C:\WINDOWS\System32\ntdll.dll
804d0000	1e300	ntoskrnl.exe	\WINDOWS\system32\ntoskrnl.exe
806b4000	13280	hal.dll	\WINDOWS\system32\hal.dll
bf800000	1b7580	win32k.sys	\???\C:\WINDOWS\System32\win32k.sys
bfec0000	8ee80	i81xdht5.dll	\SystemRoot\System32\i81xdht5.dll
bf800000	10a80	dxa.svs	\SystemRoot\System32\drivers\dxa.svs

Figure 5-26. Image List

## Threads

You can use the Process List page to view the threads for a selected process. This works the same as the THREAD command. To view the threads for a process, right-click on the process and select **Threads** from the pop-up menu.

**Note:** The target must be stopped to view threads for a process.

Thread List for EXPLORER.EXE									
Tid	KernlTEB	KernlStkBtm	KernlStkTop	StackAddr	Ip Address	UserTEB	Name	Pid	
254	80de5da8	f470f000	f4712000		00000000	7ffd5000	EXPLORER.EXE	12c	
25c	80d81da8	f41fe000	f4201000		00000000	7ffaf000	EXPLORER.EXE	12c	
558	f9b3408	f41ed000	f41fc80		00000000	7ffac000	EXPLORER.EXE	12c	
130	f966020	f4902000	f4906000		00000000	7ffd000	EXPLORER.EXE	12c	
184	f95fc68	f4b99000	f4b9c000		00000000	7ffd000	EXPLORER.EXE	12c	
188	f95f990	f48c1000		f48c5d00	804eff8b (ntoskrnl!_Section.text+1fa0b)	7ffd000	EXPLORER.EXE	12c	
18c	f95f4a8	f473b000	f473e000		00000000	7fdb000	EXPLORER.EXE	12c	
190	f95a020	f4893000	f4896000	f4895cc0	804eff8b (ntoskrnl!_Section.text+1fa0b)	7ffd000	EXPLORER.EXE	12c	
194	f95ada8	f4733000	f4736000		00000000	7ffd9000	EXPLORER.EXE	12c	
1a8	f959b30	f48a3000	f48a6000		00000000	7fd8000	EXPLORER.EXE	12c	
280	80d4370	f4ba9000	f4bac000		00000000	7fd6000	EXPLORER.EXE	12c	
3fc	80d8f298	f42ed000	f42fc80		00000000	7ffae000	EXPLORER.EXE	12c	
400	80d8bc18	f42fe000	f4301000		00000000	7ffad000	EXPLORER.EXE	12c	
4c0	f991020	f46ae000	f46b1c80		00000000	7ffab000	EXPLORER.EXE	12c	

OK

Figure 5-27. Thread List

## Page Features

### Open Only One Process List Page

You can only have one Process List page open.

### View Detailed Process Information

You can use the Process List page to view the known details of a process. The Details view for a given process lists all the known information about that process. You can sort the list of details by Item or by Value by clicking on the appropriate column heading.

Item	Value
KPEB	86653620
Pid	2e0
Name	explorer.exe
Threads	b
Priority	8
UserTime	26
KernelTime	75
CreateTime	1c20cce02a7b630; 06/05/2002 20:17:26.713
ExitTime	0
Parent	70; Unknown
State	1; Out Of Memory
Processor	0
Affinity	1
Quantum	6
WinVer	400
ErrorMode	0
DebugPort	00000000
ExceptionPort	e19fbcaa0
UserPeb	7ffd000
Win32 Process	e1bb5e68
ForkInProgress	00000000
PageDir	067f2000
VadRoot	86d8c908
MRU Vad	
EmptyVad	85f11888
ModifiedPages	2c0
PrivatePages	159
VirtualSize	20fc000
PeakVirtualSize	224f000
LastTrimTime	1c20cce02a7b630
PageFaultCount	b8c
PeakWorkingSetSize	461
WorkingSetSize	152
MinimumWorkingSetSize	32

Figure 5-28. Detailed Process Information

## Rearrange and Customize Columns

You can rearrange and customize the display of columns in any order that best suits your needs. Any modifications you make to the display of the columns will be saved in the workspace. The following columns are available to the Process List page:

- ◆ Name
- ◆ KPEB
- ◆ Pid
- ◆ Threads
- ◆ Priority
- ◆ UserTime
- ◆ KernelTime
- ◆ State
- ◆ DateTimeVersion
- ◆ OS Type (Image Type x86, 64-bit, etc.)
- ◆ Filespec

## Kill a Process

This page provides a way of terminating a process. You can only kill a process if the target is running. The killing of a process may fail if the OS refuses to kill the selected process; however, Visual SoftICE does not limit what processes you can attempt to kill.

## Copy and Drag

The page supports copying and dragging to retrieve its data.

- ◆ You can copy data from any column to the clipboard.
- ◆ You can drag data from any column to another page.

## Print

The page supports printing, and print-previewing of its contents.

## The Registers Page

The Registers page is a container for displaying the names, fields, contents, descriptions, and symbols of the target processor registers, and editing other values.

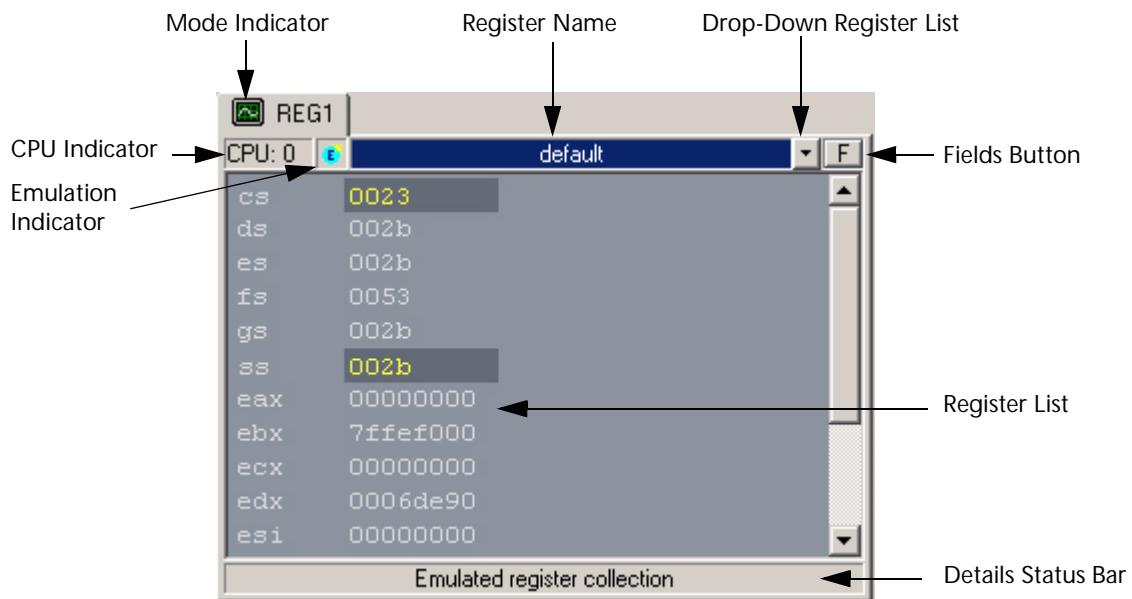


Figure 5-29. Visual SoftICE Registers Page

## Concepts and Associated Commands

### Types

Different processors contain different registers and register types. Registers are always bit containers, but can be interpreted as bytes, integers, a collection of flags, or floating point values meeting various standard representations and precisions.

### Fields

Registers of a given size (for example 32bits wide) very often contain a mixture of information using various subsets of their bits, called fields.

- ◆ You can display fields for any register (or group of registers) in the page by clicking the **Fields** button on the upper right hand corner.
- ◆ You can display register fields from the command line by issuing the R command with the -f switch.

## Descriptions

Registers all have additional description information, primarily drawn from details provided by the processor manufacturer, and in some cases as commonly used by operating systems.

- ◆ You can view descriptions for any register in the details status bar at the bottom of the page (if the register does not currently map to a symbol).
- ◆ You can display register descriptions from the command line by issuing the R command with the -d switch.

## Symbols

Registers can contain data that currently maps to a known symbol in the current context.

- ◆ If a symbol is available for the current register contents, the symbol is displayed in the details status bar at the bottom of the page.
- ◆ You can display symbols for registers from the command line by issuing the R command with the -s switch.

## Allowed Operations

Some registers are read-only, except by the processor itself. Some registers are not available, based on the state of the processor. Modifying some registers can be a dangerous action. As a level of protection, registers are marked in Visual SoftICE as being read-write, or read-only, and can have associated warnings on write operations.

For more information on controlling the display of warnings in Visual SoftICE, refer to the SET WARNLEVEL command.

## Groups

Registers that are associated with one another, or are commonly useful to see together, have been gathered into collections called Register Groups.

- ◆ All processors support the GENERAL and ALL groups.
- ◆ You can display a pop-up menu listing the available register groups by right-clicking on the page (when connected to a stopped target) and selecting Groups.
- ◆ You can display register groups from the command line by issuing the RG command.

- ◆ You can read from groups, or write to groups. However, if any member is read-only, or has warnings associated with it, the write will fail for any read-only registers, and warnings will be displayed for any warning-associated registers.
- ◆ You can read from, or write to, groups from the command line by issuing the R command.

## Page Features

### View the Registers You Want

The page supports a title area, which tells you the group, register, or set of registers being displayed within it.

- ◆ The group, register, or set of registers you have in any page is remembered on a per-page basis in the workspace when you save it.
- ◆ You can click on this area to type in a register name, register group name, or list of registers (space or comma delineated) you want displayed.
- ◆ You can right-click on this area or select the drop-down register list to choose from the available register groups.
- ◆ You can set the displayed registers from the command line by issuing the R command and redirecting its output:  
`R eip, esp, eflags >reg`
- ◆ You can set the displayed registers from the command line automatically when you issue the R command if you have enabled auto-command redirection.

### Open Any Number of Register Pages

There are no restrictions on the number of Register pages, so you can build a workspace with multiple views containing just the registers you are interested in, displayed the way you want, where you want.

## Register Group Emulation Indicator

The Register Group Emulation Indicator is an icon that VSI displays when the contents of the register page is a group that is emulated (via hardware or other means). This can be useful, as both IA64 and x64 platforms provide an x86 emulated register group. Register groups named "general" and "default" exist for all register sets, even emulated collections, so a group name alone could be misleading. This icon normally appears when the debugger stops in a 32-bit image running in a 64-bit operating system.

## View the Registers from a Specific Processor

The page supports a CPU display and selector field. On multi-processor machines you can select the specific processor for which you want to display the registers. Depending on the state of the target, some processors on a multi-processor machine may be inaccessible.

- ◆ You can select the CPU via the CPU indicator on the Registers page.
- ◆ You can view the registers for a specific processor from the command line by issuing the R command with the -c switch.
- ◆ You can control whether the page automatically switches to the stopped CPU on a multi-processor target from the preferences utility.

## View Register Data in the Format You Want

Register data can be interpreted for display in a number of ways. The default for display in the page is integer information in the datalength appropriate for the length of the register itself (or length of the field when displaying fields).

- ◆ All formatting preferences are remembered on a per-page basis in the workspace when you save it.
- ◆ You can force the display of data to always be BYTE, WORD, DWORD, or QWORD using the pop-up menu.
- ◆ You can enable registers that contain floating point data to display that data in a floating point representation using the pop-up menu.
- ◆ You can select which floating point representation Visual SoftICE applies to floating point data from the following representations: Double (64bit real), IEEE 80bit, 82bit, and AMD 3DNow.

## Edit a Register or Set of Registers

This page provides several ways of editing register values, both for individual registers and groups of registers.

- ◆ You can double-click on a single register to open the Modify Register utility.
- ◆ You can right-click on a selected register, or group of registers, and select Modify from the pop-up menu.
- ◆ You can select the Zero option from pop-up menu to set the selected register, or group of registers, to zero.
- ◆ You can write to registers, register groups, and sets of registers from the command line by issuing the R command.

## Copy, Paste, Drag, and Drop

The page supports many ways of retrieving its data.

- ◆ You can copy register names and values to the clipboard for single and multiple registers.
- ◆ You can drag a single register (in which case just the value of the register is placed in the clipboard).
- ◆ You can drag multiple registers (in which case the names and values of all the registers are collected in the clipboard, and can then be used by other pages in raw or formatted form. For example, the Text Scratch page prints the collection of names and values in table form).

## Customize the User Interface

There are multiple attributes of the page that you can customize. They are divided into two categories: per-page and application wide settings. Per-page attributes are always remembered by the workspace when you save it.

### Per-Page Settings

- ◆ You can hide or display the details status bar.
- ◆ You can display the registers and their values in a vertical list only (default) or in a wrapping horizontal collection.

### Application Wide Settings

- ◆ You can use the **FONT & COLORS** tab to select the font used, background and foreground color, and the color used to indicate a change in the value of a register between updates of the page.
- ◆ You can use the **LIVE MODE ON CONNECTION** setting to control whether or not the page switches automatically to Live mode when a connection to a new target is established.

- ◆ You can use the **Always Switch to Stopped CPU** setting to control whether the page will always switch to the stopped CPU on multi-processor machines.

## Print

The page supports printing, and print-previewing of its contents.

## The Source Page

The Source page is a read-only container for displaying source code files. This page tracks the current context on the target, and whether its source file is associated with an image of the current process (when symbol information is available). Once the symbol information for the source file is available, you can:

- ◆ Disassemble the entire source file, or individual source lines that have associated disassembly
- ◆ Set, remove, enable, and disable breakpoints
- ◆ Step through source code or associated disassembly
- ◆ Perform Go To and Run To by line number, target address, function name, current Instruction Pointer (IP), and bookmarks

Whether or not symbol information for the source file is available, you can:

- ◆ Set and remove bookmarks
- ◆ Perform Go To and Run To by line number and bookmarks
- ◆ View or hide line numbers
- ◆ Perform string searches
- ◆ Print out the page

You can open a source file by using the FILE command, or by clicking the Source page icon in the toolbar. You can have any number of source pages open at any given time. Automatic behavior of this page is controlled by **AutoFocusOpen** setting under **Source/Disasm Page** on the Global Settings tab of the Preferences dialog.

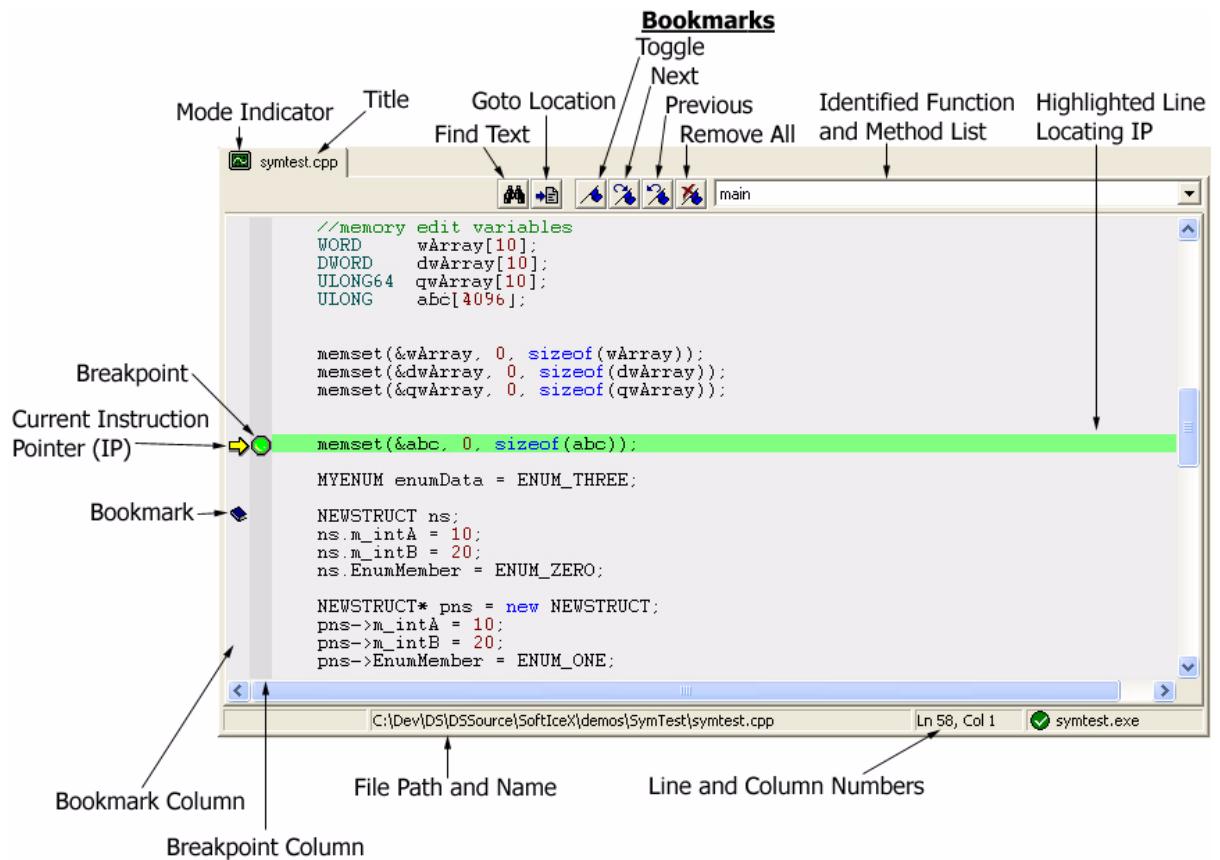


Figure 5-30. Visual SoftICE Source Page

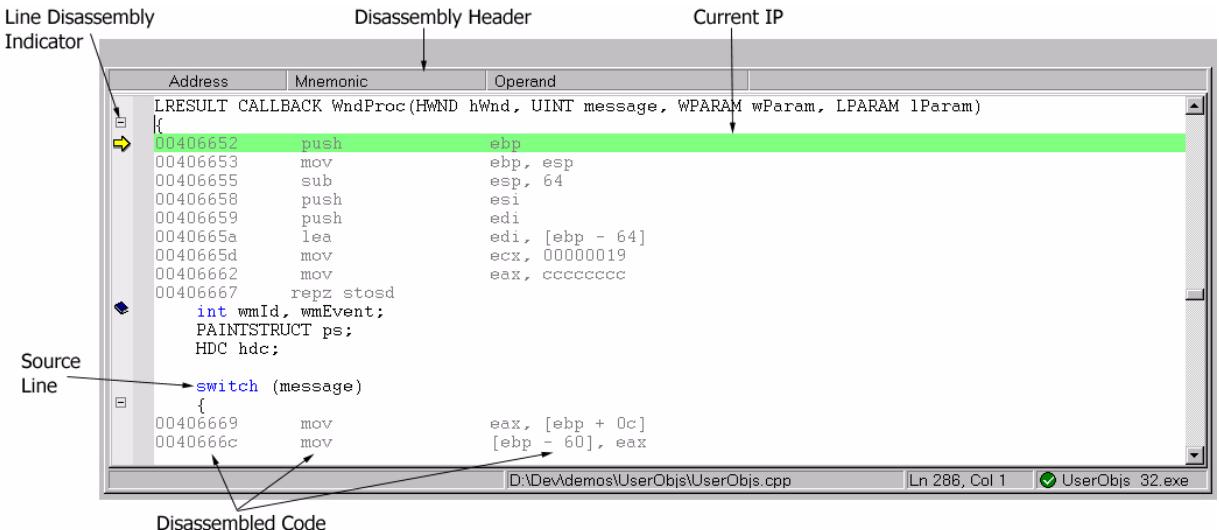


Figure 5-31. Visual SoftICE Source Page with Disassembly

## Associated Commands

### FILE

The FILE command is useful for such actions as setting a breakpoint, where the sequence of steps would be:

- 1 Use FILE to bring the desired file into a Source page
- 2 Use the SS command to locate a string or move the cursor to the specific line
- 3 Enter BPX or press F9 to set the breakpoint

When using the FILE command, consider the following:

- ◆ If you specify file-name, that file becomes the current file and the start of the file displays in a Source page.
- ◆ If you do not specify file-name, the name of the source file for the current IP, if any, displays.
- ◆ If you specify the \* (asterisk), all files in the current symbol table display (also see the FILES command).

- ◆ If you specify an address, Visual SoftICE attempts to retrieve the source file name and line number data. The address must be in the current process. This form of the command is useful in instances where, for example, you have to debug a crash dump file without source, but with symbols.

When you specify a file name in the FILE command, Visual SoftICE switches address contexts if the current symbol table has an associated address context.

## GO

The GO command executes code starting from either a specified starting address or the current IP until the specified break address or another breakpoint is hit.

## LOAD

The LOAD command loads symbols for a specified image-name.

## SS

The SS command searches all the open files for a specified string. If you specify a line number, Visual SoftICE starts the search at that line. If you do not specify a line number, then the search will start at the current top view line. If you do not use the -a option, the search will wrap around within the file (performing the equivalent of Find Next). If you use the -b option, Visual SoftICE will bookmark anything found.

## T

The T command steps one instruction when the target is stopped. Visual SoftICE searches through the symbols to locate a file containing the current instruction pointer (IP). If it finds the file, and a source page for that file is already available, Visual SoftICE updates the existing source page. If it finds the file and no source page for that file is open, Visual SoftICE creates a new source page.

If Visual SoftICE does not find a source file containing the current IP, and a live disassembly page for the current IP is available, it updates that page. If no live disassembly page is available, Visual SoftICE opens a new disassembly page.

## WC

When you type WC [address], Visual SoftICE searches through the symbols to locate the source file containing the specified address. If Visual SoftICE finds the source file, and a source page for the file is already open, it brings the opened page forward. If the source file is not already open, Visual SoftICE creates a new source page for the file and brings it forward.

If the address is not contained in any source file, Visual SoftICE creates a disassembly page and brings it forward.

## Page Features

### Open Only One Source Page Per File

You can only open one Source page per file.

### File or Line Disassembly

When symbol information is available, Visual SoftICE displays the process ID and image name on the status bar at the bottom right of the page. When this is the case, the Disassemble File and Disassemble Line options become enabled under the Disassembly group of the pop-up menu for the page.

- ◆ You can view or hide the disassembly code for the entire source file by toggling the Disassemble File option on or off.
- ◆ You can view or hide a single source line by toggling the Disassemble Line option on or off.

When disassembly code is visible in the source page, the header bar appears at the top of the page. You can choose to view or hide the Address and Op-code fields of the disassembly lines by toggling the **Show Address** and **Show Opcode Bytes** options under the Disassembly group of the pop-up menu for the page. You can also use the Page Preference option to set the default mode for header display.

If the target is an IA64 CPU, Visual SoftICE displays the Predicate/Prefix field on the disassembly header in addition to the other fields.

### Source and Instruction Stepping

When the master is connected to a target that is stopped at the address contained in the source page, Visual SoftICE highlights the line of the current instruction pointer (IP). You can choose to step-over, step-into, or step-out of the current source line.

If the page is in the mixed mode in which disassembly lines are available, you can elect to do source stepping or instruction stepping. Use the options under the Debug group of the pop-up menu for the page, the function keys as defined in the Preference Settings, or the stepping buttons on the main toolbar, to perform your stepping functions.

If you are source stepping and attempt a step-into when the related source or .asm file is available, Visual SoftICE opens the file and automatically highlights the line containing current IP. If no source file is available, the step-into action behaves the same as step-over.

If you are instruction stepping in the mixed mode, Visual SoftICE highlights the disassembly line for the current IP. If the current instruction has a jump-to address to branch to, the yellow arrow on the left of the highlighted line changes its direction to up or down, depending on the branch address location, indicating that the next instruction will be at a specific address. To view the jump-to line, you can click the arrow, press **<Ctrl><J>**, or select **Go To Jump** from the pop-up menu for the page. Visual SoftICE highlights the jump-to line when it displays it.

You can always view the current IP line by pressing **<Ctrl><I>**, or by selecting **Go To Current IP** from the pop-up menu for the page.

To run-to and stop at a specified source or disassembly line, you can set the cursor to the line where you want to stop and press **<Ctrl><R>**, or select **Run To Here** from the pop-up menu for the page.

## Symbol Status Bar

The source page has its own status bar that uses icons to display the state of symbols that match the file. Icons indicating the status of symbols are followed by an image name.

**Table 5-6.** Symbol Status Bar Icons

Icon	Meaning
	This icon indicates that symbols are loaded for the indicated image, and that the source file is part of that image and has line information available.
	This icon indicates that symbols are loaded for the indicated image, and that the source file is part of that image, but there is no line information available.
	This icon indicates that the source file is not part of any image in the current/active process.

## Buttons and Identified Function and Method List

The Source page has buttons and a list of any identified functions and methods. The buttons allow you to manipulate bookmarks, locate strings, and move to specific locations in the source. The list displays any identified functions and methods in the source file, and is active for C and C++ source only.

Table 5-7. Source Page Buttons

Icon	Meaning
	Use this button to find text in the source file.
	Use this button to go to a specific location in the source file.
	Use this button to toggle the bookmark at the current line.
	Use this button to move to the next bookmark.
	Use this button to move to the previous bookmark.
	Use this button to remove all bookmarks.

## Breakpoints

With symbol information available, you can set breakpoints at a source line containing actual code via one of the following methods:

- ◆ Selecting Set Breakpoint under the Breakpoint group on the pop-up menu for the page
- ◆ Pressing the F9 function key
- ◆ Clicking the Breakpoint button on the main toolbar
- ◆ Locating the cursor in the second column to the left of the code, and right-clicking to access the Breakpoint pop-up menu.

- ◆ Double-clicking in the second column to the left of the code to set, enable, or disable breakpoints.

If you attempt to set a breakpoint on a line that does not contain actual code, Visual SoftICE sets the breakpoint at the next line containing code. If a line is in the source/disassembly mixed mode, you can set the breakpoint at any of the visible disassembly lines.

Once you set a breakpoint, you can disable or remove it via one of the following methods:

- ◆ Selecting the applicable options from the breakpoint pop-up menu
- ◆ Using the function keys
- ◆ Using the buttons on the main toolbar
- ◆ Double-clicking in the second column to the left of the code to enable or disable breakpoints

Any breakpoint changes made elsewhere in the system are synchronously updated on the source page.

## View Line Numbers

You can right-click on the Source page and select View Line Numbers from the pop-up menu to display line numbers for every line currently contained in the Source page. The current state of the line numbers (whether or not they are displayed) is preserved and restored in the workspace. The workspace also preserves and restores what the top displayable line was.

## Evaluate a Selection

Positioning the mouse cursor within a selection will cause Visual SoftICE to evaluate that selection and display the results in the hover-text pop-up.

## Bookmarks

In the source page, you can set a bookmark at any line in the source file, but you cannot set bookmarks at the disassembly lines in mixed mode. Bookmarks in the source page are associated with file line numbers. You can view or hide line numbers by toggling the **Show Line Numbers** option on and off from the pop-up menu for the page, or by using **<Ctrl><L>** to toggle the selection.

The following keystrokes control bookmarks:

- ◆ <CTRL><F2> toggles a bookmark
- ◆ <Shift><Ctrl><F2> clears all bookmarks
- ◆ <F2> goes to the next bookmark
- ◆ <Shift><F2> goes to the previous bookmark

To remove bookmarks:

- ◆ You must have focus on a line with a bookmark on it. You can then remove the bookmark for the line, or all bookmarks in the file, by right-clicking and selecting the applicable option from the page's pop-up menu.
- ◆ You can locate the cursor in the first column to the left of the code and right-click to access the Bookmark pop-up menu.
- ◆ You can double-click in the first column to the left of the code to enable or disable a bookmark.

You can view the bookmark list in the Go To dialog. From that dialog you can view or run-to the specified bookmark via the **Go To** and **Run To** buttons. You can save up to 25 bookmarks per Source page instance when you save the workspace. If you exceed 25 bookmarks within an instance of a Source page, the first 25 are saved.

## Go To (CTRL-G)

You can use the Go To dialog to:

- ◆ View or run-to a source line using its target address or function name
- ◆ View or run-to a line using its line number
- ◆ View a line containing the current IP
- ◆ View or run-to a selected bookmarked line

## Go To Current IP

You can right-click on the page and select **Go To Current IP** from the pop-up menu to view the line containing the current IP.

## Go To Jump

When you are on a jump statement, you can right-click on the page and select **Go To Jump** from the pop-up menu to move to the line where the jump will go.

## Run To Here

You can right-click on the page and select **Run To Here** from the pop-up menu, and Visual SoftICE will run from the current IP to the line that currently contains the cursor.

## Set Next Statement

The Set Next Statement command moves the IP to the next statement. You can execute the Set Next Statement command via one of the following methods:

- ◆ Right-clicking to access the pop-up menu for the page
- ◆ Pressing <Ctrl><Shift><F10>
- ◆ Selecting the command from the Source Page drop-down menu
- ◆ Clicking the **Set Next Statement** toolbar button

If you attempt to execute the Set Next Statement command on a specified line with an address outside of the function scope of the current IP, and the warning level is not set to off, then Visual SoftICE displays a warning message asking you to confirm that you want to set the IP to the new address.

## Find (CTRL-F)

You can use the Find dialog to search for a string in the current source page.

## Copy

The page only supports copy command to retrieve its data.

- ◆ You can select any text within the page and copy it to the clipboard.
- ◆ You can use **Select All** from the pop-up menu or <CTRL><A> to highlight all the lines on the page before copying them to the clipboard.

## AutoFocus

You can use the AutoFocusOpen setting to customize the behavior of the source page when a target event occurs. If you select **Source or Disassembly** or **Source or Disassembly - No Focus**, then when Visual SoftICE can locate a source file for the current IP, it creates a source page. If Visual SoftICE cannot locate a source file, it creates a disassembly page.

For more information on available settings and how they behave, refer to AutoFocusOpen Settings in the online help.

## Customize the User Interface

There are multiple attributes of the page that you can customize. They are divided into two categories: per-page and application wide settings. Per-page attributes are always remembered in the workspace when you save it.

### Per-Page Settings

- ◆ Line number display is persisted and restored in the workspace.
- ◆ Up to 25 bookmarks are persisted and restored in the workspace.
- ◆ Display location in the file is persisted and restored in the workspace.

### Application Wide Settings

- ◆ You can use the **AutoFocusOpen** settings to control whether Visual SoftICE opens a new page when none exists (or brings an existing page to the top of the pad it is on), and places input focus on that page.
- ◆ You can use the **New Src/Disasm Pages to Largest Pad** setting to configure Visual SoftICE to place new Source pages on the largest pad.
- ◆ You can use the **New Src/Disasm Pages to NAMED Pad** setting to configure Visual SoftICE to place new Source pages on a specific named pad.
- ◆ You can configure Visual SoftICE to display or hide the machine language (op-code) bytes column by right clicking on the Source page, selecting **Source**, and toggling **Show Op-Code Bytes** in the sub-menu. You can also toggle this setting under the **Source/Disassembly** element in the Global Settings tab of the Preferences dialog.
- ◆ You can configure Visual SoftICE to display or hide the address column by right clicking on the Source page, selecting **Source**, and toggling **Show Addresses** in the sub-menu. You can also toggle this setting under the **Source/Disassembly** element in the Global Settings tab of the Preferences dialog.
- ◆ You can configure Visual SoftICE to display or hide the instruction template field (on IA-64 only) by right clicking on the Source page, selecting **Source**, and toggling **Show Instruction Template** in the sub-menu. You can also toggle this setting under the **Source/Disassembly** element in the Global Settings tab of the Preferences dialog.

## Print

The page supports printing, and print-previewing of its contents.

## The Stack Page

The Stack page is a read-only page that displays the data on the stack for a selected process and thread context. The Stack page also has a page-specific context bar, allowing you to change the process or thread, which overrides the main context bar.

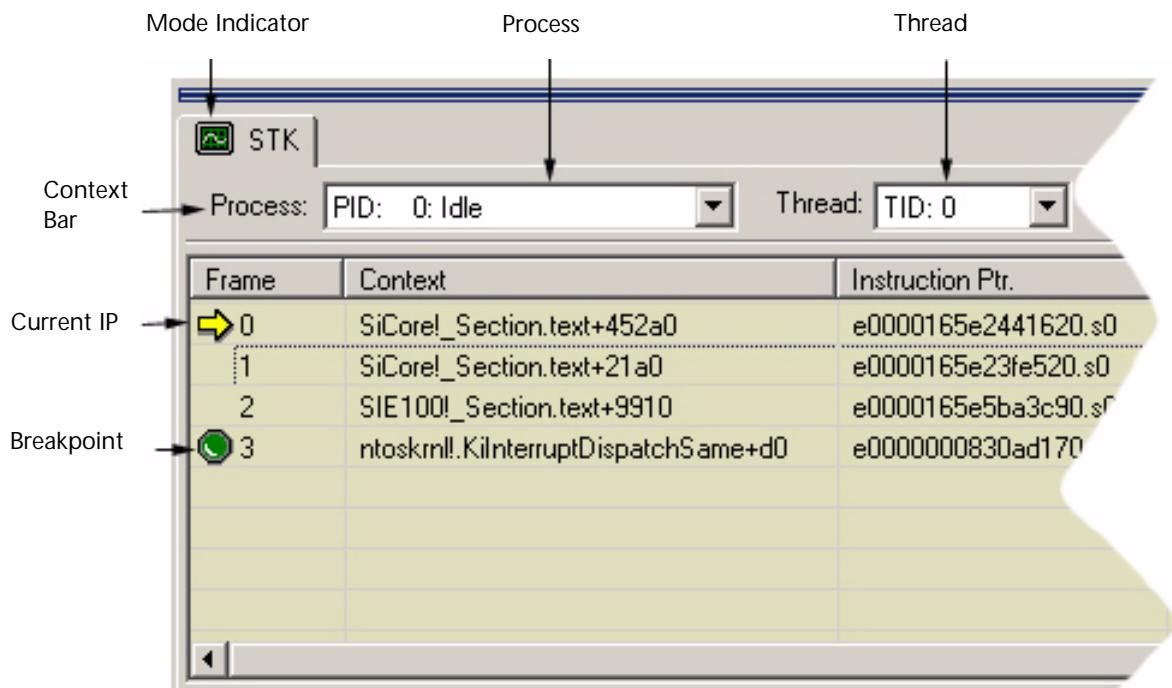


Figure 5-32. Visual SoftICE Stack Page

## Concepts and Associated Commands

### Displaying Stack Information

There are several commands that can display stack information.

- ◆ You can display the stack information from the command line by issuing the STACK command.
- ◆ You can display the stack information for a particular thread from the command line by issuing the STACK command with thread-id.

- ◆ You can display the stack information for a fiber from the command line by issuing the FIBER command with -s switch.
- ◆ You can display the stack page from the command line by issuing the WS command.
- ◆ You can force a refresh of the stack frame by issuing the STACK command with the -r switch.
- ◆ You can display detailed information on the stack frame, including function table (FPO data or Unwind data) and parameters by issuing the STACK command with the -d switch.

## Changing Current Context

The stack page displays information based upon the context that is set by either the main context bar, or the page-specific context bar. There are several ways to change the context that the Stack page is following.

If you change the process for the main context, the debugger and all pages follow that new context, except for any pages for which you have set an overriding page-specific context. You can change the current process via any of the following methods:

- ◆ Select a new process using the **Process** drop-down list on the **Context** bar.
- ◆ Use the ADDR command.
- ◆ Double-click on a process listed within the **Process List** page.
- ◆ Right-click on a process listed within the **Process List** page, and select **Make Current** from the drop-down menu.

## Frames

The stack page displays a list of items called frames. The frames are listed from the current frame to starting frame of the thread when data is available. The stack page will display the following information for each frame: frame number, current context, instruction address, stack address, and frame address.

- ◆ Frame numbers are listed from the current frame to top-most frame of the thread. Each time a new function is called, Visual SoftICE renumerates the frame accordingly.
- ◆ Current context displays the function name plus the byte offset from the start of the function of to the instruction address for that frame. If no symbols are available for the frame's function address, then the field will be blank or show section-offset data. Refer to the SET SYMPATH command for more information on finding symbols.

### Context Bar

The Stack page has a page-specific context bar allowing you to select the context for which the page displays stack frames. This page-specific context bar overrides the main context bar and allows the Stack page to track an independent context.

You can switch to a different context to display its stack frames. Use the **Process** drop-down list to select the process, and the **Thread** drop-down list to select a specific thread for that process (if the process is multi-threaded).

### Source / Disassemble At

The stack page supports displaying a Source or Disassembly page at the instruction address for the selected stack frame. To open either a Source or Disassembly page, right-click on a stack frame and select **Source / Disassemble At**. Visual SoftICE displays a Source page if the symbols and source files can be located for the instruction address. Otherwise, Visual SoftICE displays a Disassembly page starting at the instruction pointer.

- ◆ You can change the path that Visual SoftICE uses to search for symbols by issuing the SET SYMPATH command at the command line.
- ◆ You can change the path that Visual SoftICE uses to search for source files by issuing the SET SRCPATH command at the command line.

### View Memory At

The stack page supports displaying a Memory page for a selected instruction, stack, or frame address. To display a Memory page, right-click on an address and select **View Memory At**.

### Breakpoints

The stack page supports the setting, enabling, disabling, or clearing of breakpoints at the instruction address for the selected stack frame. To complete any of these breakpoint operations from the Stack page, right-click on a frame in the **Frame column**, select **Set Breakpoint**, **Enable Breakpoint**, **Disable Breakpoint**, or **Remove Breakpoint**.

- ◆ You can set a breakpoint for an instruction address of a frame by issuing the BPX command and the address at the command line.

- ◆ You can enable, disable, or clear a breakpoint for an instruction address of a frame from the command line by issuing a BL command to get the breakpoint index, and using the BE, BD, or BC command along with that index.

## View Locals and Registers for a Frame

The stack page supports displaying either a Register or a Locals page for the selected stack frame. To display one of these pages, right-click on a frame and select **Frame** and then the desired page type (**Registers** or **Locals**).

- ◆ You can display the registers for individual frames of a thread from the command line by issuing the THREAD -r [-f frame] TID command.
- ◆ You can display the locals for the current stack frame from the command line by issuing the LOCALS command.

## Stack Status Column

The stack status column provides extra information about the stack walking process. It notifies you about potential problems Visual SoftICE encountered while walking the stack. If Visual SoftICE encountered no problems, and the stack was walked fully, no status messages are displayed. If the stack walk terminates prematurely there will be a status message describing the terminating condition. The most important message is the "Unwind info unavailable" message on IA64.

"Unwind info unavailable" usually means that Visual SoftICE could not read the unwind information from the target. This occurs because the part of the image on the target containing the unwind info is paged out and Visual SoftICE could not access it. A solution for this problem is to have a local copy of all executables (including files in the system32 directory). Add these directories to the exepath so the local copies can be found instead of having to be retrieved from the target.

## Open Only One Stack Page

You can only have one Stack page open.

## Copy

The page only supports the Copy function to retrieve its data. You can select any cell (intersection of a row and column) and copy its text to the clipboard.

**Note:** Select a cell by placing the mouse cursor on the cell and then right-clicking.

## Customize the User Interface

There are multiple attributes of the page that you can customize. They are divided into two categories: per-page and application wide settings. Per-page attributes are always remembered by the workspace when you save it.

### Per-Page Settings

- ◆ You can hide or display the page-specific context bar.
- ◆ You can sort the stack frames by any of the displayed columns.
- ◆ You can reorganize and resize the columns.

### Application Wide Settings

- ◆ You can use the **Live Mode On Connection** setting to control whether or not the page switches automatically to Live mode when a connection to a new target is established.

## Print

The page supports printing, and print-previewing of its contents.

## The Text Scratch Page

The Text Scratch page is used to capture redirected output from various commands that can be executed in the Command page. The Text Scratch page is also a convenient place to type, cut, and paste output to.

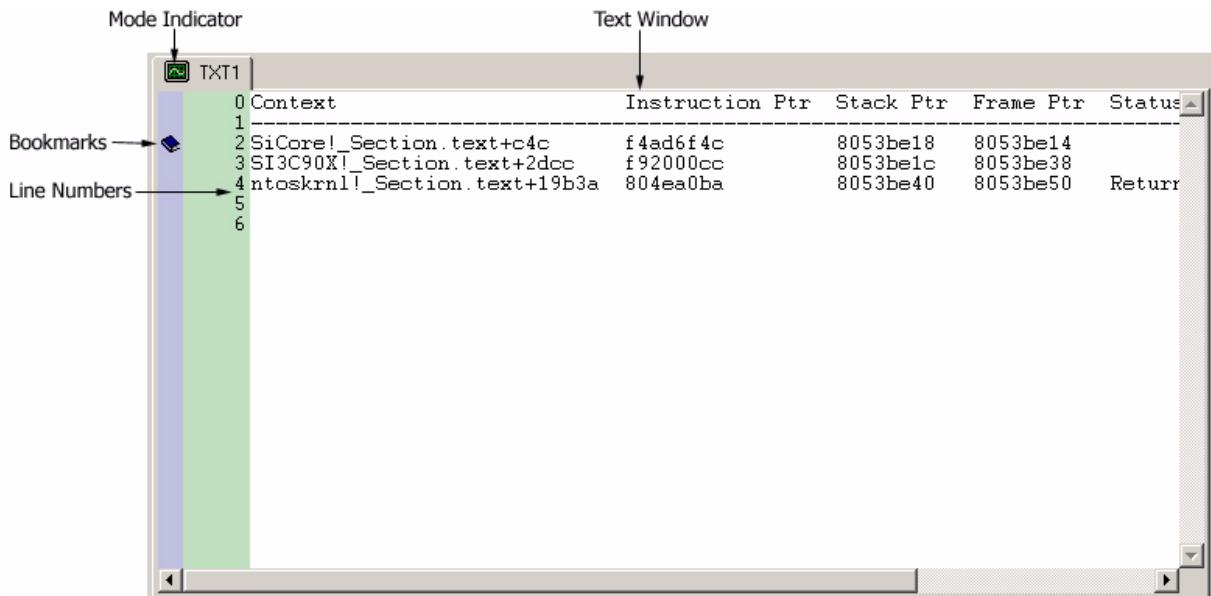


Figure 5-33. Visual SoftICE Text Scratch Page

## Concepts and Associated Commands

### Command Redirection

The page supports command redirection from a command page. This allows output (such as from a kd extension, or a script file) to be captured and saved. Refer to Command Redirection in the main help file for more information.

## Page Features

### Save and Clear Output

You can save all of the output on this page to a specified text file by right-clicking on the page and selecting **Save Output To File** from the pop-up menu. You can clear all of the output on this page by right-clicking on the page and selecting **Clear All** from the pop-up menu.

## Open Any Number of Text Scratch Pages

There are no restrictions on the number of Text Scratch pages you can open. You may find this useful when capturing the output from a specific command. Simply specify /> TXT as the page to redirect the output to. If there is no existing Text Scratch page with that name, a new page is created with the command as its name.

## Find Text (CTRL-F)

The Text Scratch page provides a utility for finding a text string value within the page.

- ◆ You can search by attempting to match the whole word, only or partial string.
- ◆ You can further narrow a search by attempting to match the case of the specified search string.

## Go To a Specific Line (CTRL-G)

You can right-click on the Text Scratch page and select **Go To Line** from the pop-up menu to advance the cursor to any line of text currently contained in the Text Scratch page.

## Copy, Paste, Drag, and Drop

The page supports many ways of retrieving its data.

- ◆ You can select any text within the page and copy it to the clipboard.
- ◆ You can choose **Select All** or press <CTRL><A> to select all the lines on the page and copy them to the clipboard.
- ◆ You can paste any text from the clipboard into the page.
- ◆ You can drag items from other pages and drop them onto the page. For example, you can drag multiple registers from a register page, and Visual SoftICE inserts a table of names and values into the page.

## Append a File

You can right-click on the Text Scratch page and select **File Append** from the pop-up menu to append the contents of any text file to the Text Scratch page contents.

## **View Line Numbers**

You can right-click on the Text Scratch page and select **View Line Numbers** from the pop-up menu to display line numbers for every line of text currently contained in the Text Scratch page.

## **View Bookmarks**

You can right-click on the Text Scratch page and select **View Bookmarks** from the pop-up menu to display the bookmarks that may be set for the lines of text currently contained in the Text Scratch page.

## **Customize the User Interface**

There are multiple attributes of the page that you can customize. They are divided into two categories: per-page and application wide settings. Per-page attributes are always remembered by the workspace when you save it.

### **Per-Page Settings**

### **Application Wide Settings**

There are no per-page settings for this page.

- ◆ You can use the **Save/Restore Contents** in Workspace settings to configure Visual SoftICE to automatically save any text contained in the Text Scratch pages when you save your workspace. Each Text Scratch page saves and restores its own content.
- ◆ You can use the **Fonts & Colors** tab to select the font, foreground color, and background color used for the text on the page.

## **Print**

The page supports printing, and print-previewing of its contents.

## **The Watch Page**

The Watch page is a read-write container for displaying and editing any expression. You can enter any expression into this page and Visual SoftICE returns the resulting value. The Watch page displays the expression, result type, and value.

When you switch the context in the Locals page, the Watch page also switches expression result types or values, remaining synchronized with the Locals page. When this happens, some expressions may not have a value at the current context even though they displayed a value previously.

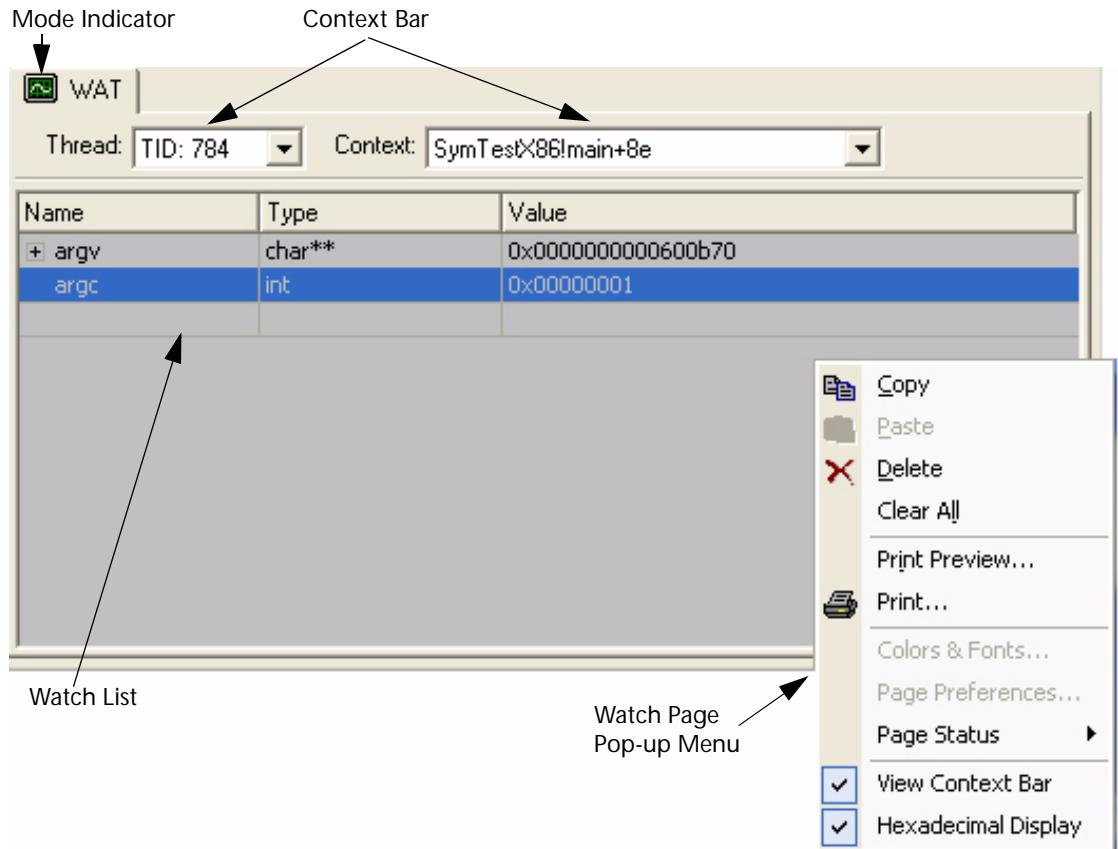


Figure 5-34. Visual SoftICE Watch Page

## *Concepts and Associated Commands*

### **WATCH**

Use the WATCH command to display the results of expressions.

## *Page Features*

### **Live Mode Variable Tracking**

When the Watch page is in live mode, Visual SoftICE displays the variable values in red when they change. You can expand the variable if it contains child-level components.

## Open Only One Watch Page

You can only have one Watch page open.

## The Context Bar

By default the Watch page follows the current system context. You may override this behavior by right-clicking on the page and selecting **View Context Bar** from the pop-up menu. The context bar allows you to view a specific thread and context by selecting from those available in the drop-down lists.

## Toggle Values Displayed Between Decimal and Hexadecimal

You can toggle displaying values in hexadecimal or decimal. To toggle display format for the values, right-click on the Value column and select **Hexadecimal** from the pop-up menu. When it is checked, Visual SoftICE displays values in hexadecimal. When it is unchecked, Visual SoftICE displays non-address values in decimal.

## Edit Expression Values

Visual SoftICE allows you to edit the value of expressions displayed in this page. To modify the value of an expression, double-click the Value column of the expression to highlight the value, and enter the new value in the Edit field. Press <Enter> or click anyplace on the Watch page to save the new value.

## Save Expressions

When you have a workspace file opened, Visual SoftICE will store all the expressions on the Watch page in your workspace file, when saved.

## Add Expressions

Visual SoftICE allows you to add expressions directly into the Watch page, or by using the WATCH command from the Command page. To add a new expression directly into the Watch page, double-click on an empty row to make the **Name** field editable, and enter your new expression name into the field. Press <Enter> or click anyplace on the Watch page to save the new expression.

## Modify Expressions

Visual SoftICE allows you to modify expressions directly on the Watch page. Visual SoftICE re-evaluates the expressions and displays the new result types and values in each column. To modify an expression, double-click the **Name** column of the expression you want to modify, and enter the new expression name in the Name field. Press <Enter> or click anyplace on the Watch page to save the value.

## Delete or Clear Expressions

Visual SoftICE allows you to delete or clear expressions using the Watch page. To delete an expression, select the expression and press <Delete>, or right-click on the expression and select **Delete** from the pop-up menu.

## Copy, Paste, Drag, and Drop

The Watch supports Copy, Paste, Drag, and Drop operations.

## Customize the User Interface

There are multiple attributes of the page that you can customize. They are divided into two categories: per-page and application wide settings. Per-page attributes are always remembered by the workspace when you save it.

### Per-Page Settings

- ◆ You can configure the page to display any of six available columns of data (Name, Address, Type, Size, Location, and Value). Name, Type, and Value are displayed by default. To customize which columns are displayed, right-click on a column heading to access the pop-up menu. Customized column configurations are not preserved in the workspace for the Watch page.
- ◆ You can resize the individual columns displayed. Double-clicking on any column heading will auto-resize the columns to fit the content.
- ◆ You can toggle the display format for variable values between decimal and hexadecimal.

### Application Wide Settings

## Print

The page supports printing, and print-previewing of its contents.



# Chapter 6

## Visual SoftICE Symbol Management



- ◆ Visual SoftICE Symbol Management
- ◆ Images — Why you need access to them
- ◆ Symbols — Getting Setup in Visual SoftICE
- ◆ The Symbol Tables Utility

### Visual SoftICE Symbol Management

Due to the dual-machine nature of Visual SoftICE, symbol management is very different from SoftICE, the single machine system debugger. In some ways, symbol management is simplified, allowing dynamic table loading and unloading, as well as automatic symbol retrieval. In other ways (for those familiar with SoftICE) some preparation is required, in that the system must be properly configured to provide access to data that is normally packaged up in the SoftICE NMS file creation process.

#### *Where to Put the Symbols*

With Visual SoftICE, the master loads and manages all symbols, exports, and image data (PE headers, etc...) locally. This can be on a network drive, or via a company MS Symbol Server, but the clear distinction to understand is that symbols are used on the master side, not on the target.

**Note:** The one exception to this general rule is providing data to the target core debugger about OS structures and locations for boot mode operation. This is handled via a dedicated target side driver for the target platform.

## **What Symbols are Supported**

Visual SoftICE supports PDB files (version 5 or later), and DBG files, for each target platform and OS it supports. This includes support and identification of images that have been optimized (usually kernel components), support for DDK compiler/linker generated symbols, and application generated symbols. Modern compilers and linkers support features like incremental compilation, incremental linking, and “edit and continue”. These useful rapid application development enhancements complicate symbol file information, and sometimes can generate files with missing data, or data that is out of date.

Visual SoftICE also supports user-defined symbols, of address or value type. These are considered global in nature, are always accessible, and always loaded. To learn more about user-defined symbols, refer to the NAME command in the *Visual SoftICE Command Reference*.

## **Images — Why you need access to them**

Symbols are only half the story, and can be considered meta-data about the image that will be loaded into memory by the operating system. The image itself contains information about the type and location of accurate symbol files generated for it, section layout (memory partitioning within the image), and export data (what publicly visible functions there are, and their locations). Depending on the target architecture, and the type of image, this file may also contain UNWIND data (useful stack-walking information on IA64 and x64 platforms) and other useful information.

## **Local Copies of Images**

Keeping local copies of images you are debugging on the target is highly recommended. This allows the master to look things up locally, instead of attempting to retrieve data from the target over a transport. Also, keep in mind that any image or parts of its memory may be paged out on the target. This generally happens just when you need to access the paged out data on a stopped machine.

Additionally, when preloading symbols to set breakpoints in drivers or applications that have not yet been loaded into memory, the section table of the image is required to fix up image relative addresses (especially with symbols that contain OMAP data on x86 platforms).

Keeping a local copy accessible on the master allows access to the necessary data, regardless of the targets current state.

**Note:** Refer to the FGET command in the *Visual SoftICE Command Reference* for retrieving files from the target to the master.

## Tables – Active Table, Loading, and Unloading Commands

Visual SoftICE, like SoftICE, uses the metaphor of a table to manage symbols to images. The TABLE command, without parameters, will display the currently loaded symbol tables, and highlight the current or active table. Entering the TABLE command with the verbose switch (-v) shows that there is a table entry for every image in the active/current process on the target.

### Active Table

Many symbol commands will restrict their search to the active table for performance reasons, unless user input specifically references another table, or indicates a request to search all known tables (refer to the following commands in the *Visual SoftICE Command Reference*: FILE, SYM, TYPES).

The active or current table will automatically change as the user changes the current UI context, either through the context bar, or the ADDR command.

The active table can be directly changed by entering the TABLE command with the name of one of the loaded tables.

### Loading and Unloading

The following commands and actions apply to dynamic tables (those symbol resources that are needed on demand, and that the system can unload if not necessary).

- ◆ You can load a table for any in-memory image by use of the LOAD command. You can optionally specify the exact path to the file (otherwise the MS Symbol Server and/or search paths are used). The file will only be loaded for images that exist, and for which a table has not yet been loaded.
- ◆ You can force a RELOAD of a table (equivalent to an UNLOAD followed by a LOAD command) for any image that currently exists.
- ◆ You can force a table to be unloaded by issuing the UNLOAD command.

The following commands and actions apply to persistent tables (those symbol resources that you have hand loaded, and stay loaded until you unload them by hand).

- ◆ You can persistently load a table by use of the ADDSYM command.
- ◆ You can remove a persistent table by use of the DELSYM command.

## Matching

Tables that are loaded (dynamically, or persistently) have matching criteria to images in the operating system. The match between a given symbol file and an image is achieved differently depending on the image, the type of symbol file loaded, and by user image match settings.

Images have a name, extension, target architecture type (64,32,16bit), signature type (PDB6, PDB7, etc...), timestamp, and in some cases, age and/or GUID data for reference to symbol files. The image data will be retrieved and compared to the symbol file to be loaded, using the IMAGEMATCH setting, to conclude if the symbol file meets the minimum criteria of a match.

Differences in target architecture are not allowed:

- ◆ 32bit symbols that match a 64bit image of the same name will not be loaded.
- ◆ Symbols for a 64bit Itanium of the same name of that for an x64 platform will not be loaded.

The IMAGEMATCH setting (refer to SET IMAGEMATCH in the *Visual SoftICE Command Reference*) can be **BEST** (default), or **EXACT**.

- ◆ When this setting is **BEST**, symbols that are out of date, mismatched on timestamp, GUID, signature, or only partially match the image name and extension, are still allowed. This supports having symbols that are “close” but not exact, which is sometimes good enough for debugging.
- ◆ When **BEST** is not good enough, (or confusing), **EXACT** mode is available, which means that image name, extension, signature, and version (age, GUID, timestamp) must match exactly, or the table will not load. This is useful to ensure you absolutely have the right symbols for a given situation.

Additionally, the TABLE command will display other pertinent information that is discovered about loaded tables for an image; for example, showing the “(PERF)” indicator in the status field, which informs you that the image has been performance optimized since the symbol file was generated, and some lookups may not match exactly (refer to the TABLE command in the *Visual SoftICE Command Reference* for more details).

## **Automatic, On-Demand Loading and Unloading**

### **Automatic Loading**

Visual SoftICE loads only the tables on-demand that are necessary.

- ◆ Tables are loaded for “current” images (images referenced in the current process, and/or referenced in the current stack). As the OS loads images, and those images are part of the current context, the Symbol Engine will automatically load them.
- ◆ As debugger commands are issued that reference other tables, these are loaded on-demand as well.
- ◆ All known images are tracked, and those not currently loaded yet are considered in a “Deferred” state. Once loaded, the table remains loaded until the image itself unloads, or the user forces it to unload (refer to the UNLOAD command in the *Visual SoftICE Command Reference*).
- ◆ Automatic, dynamic loading presumes that the path information has been setup so it can retrieve symbols in local (or accessible) directories, or from any configured MS Symbol Server, and that it may find local exports and images as well.

### **Automatic Unloading**

When the OS unloads an image, if the Symbol Engine has a matching table, that table is dynamically unloaded as well.

### **Controls**

To disable or enable automatic symbol table loading, use the SET SYMTABLEAUTOLOAD on/off command. Automatic loading is enabled by default.

## **Pre-Loading/Persistent Loading**

It is often desirable to set breakpoints in images that have not yet been loaded by the operating system. You can accomplish this in Visual SoftICE by preloading symbols. These symbols are considered persistent. Their match state is updated whenever image collections change, but these tables are loaded into the symbol engine, even if there is no match. Persistent symbols do not participate in automatic load or unload behavior.

### **Commands**

- ◆ To load persistent symbols, use the ADDSYM command.
- ◆ To unload persistent symbols, use the DELSYM command.

## **Integrated MS Symbol Server Access**

Retrieval of operating system symbols using the Microsoft Symbol Server technology is a great boon to productivity when working with a heterogeneous set of target hardware and OS versions. This same technology is reusable by end users to provide their own symbol files (and file types) in private servers.

Visual SoftICE has directly integrated MS Symbol Server access into its automatic, on-demand loading logic. If so configured, when Visual SoftICE does not find the file it is looking for in the local search directories, it will attempt to retrieve the file from a symbol server. The paths to use, symbol server to attempt a connection to, and whether this behavior is enabled or not, is fully configurable.

### **GUI Configuration**

You can configure the MS Symbol Server Access and the server and local directory information from the **Paths** element under **Global Settings** in the Preferences dialog.

- ◆ To set the server and local directory Visual SoftICE will bring files into, configure the **MS Symbol Server Paths** element.
- ◆ To enable symbol server retrieval (disabled by default), configure the **MS Symbol Server(s) Enabled** element.

## Equivalent Commands

You can configure the MS Symbol Server Access and the server and local directory information by issuing some Visual SoftICE commands at the command line in the Command page.

- ◆ To modify or add a symbol server, use the SET SYMSRVPATH command.
- ◆ To enable symbol server retrieval (disabled by default), use the SET SYMSRVSEARCH on/off command.

## Using Exports

It is often useful to have exports in the absence of, or in addition to, full or partial symbolic data. Visual SoftICE supports a means to use exports, when available, in addition to any symbol files.

- ◆ A path to the top of an exports tree is specified (refer to the SET EXPORTPATH command in the *Visual SoftICE Command Reference*), and exports can then be extracted on demand from the connected target (refer to the GETEXP command in the *Visual SoftICE Command Reference*).
- ◆ Exports are stored with as much matching information as is available in the image on the target, in a local directory structure similar to the storage mechanism MS Symbol Servers use for symbols
- ◆ You can specify a wildcard “\*” in the GETEXP command, getting exports for everything in the specified directory. You can also specify the flag (-s) to walk subdirectories in the retrieval. Please note that retrieving exports for every image file on a target can be a very lengthy process (but is possible).

## Symbols — Getting Setup in Visual SoftICE

This section provides a couple steps to get you up and running quickly with symbols and Visual SoftICE.

Once these are complete, you should be able to easily manage the tables loaded into Visual SoftICE via the TABLE, ADDSYM, DELSYM, LOAD, UNLOAD, and RELOAD commands.

Once proper symbols are loaded, you should be able to explore available data via the EXP, FILE, LOCALS, SYM and TYPES commands.

## **Setting Up General Paths**

Configure the following Global General Path Settings:

- ◆ **Image Search** — Where you will put copies of the executables you will debug on the target.
- ◆ **Source Search** — Where your source code lives.
- ◆ **User Symbol Search Path** — Where your symbols live.
- ◆ **MS Symbol Server Paths** — Where the MS Symbol Server should be found, and where it should store things its retrieves.
- ◆ **KD Extension Paths** (optional) — Where to find KD extensions when no absolute path is specified.

## **Setting Up Visual SoftICE Paths**

Configure the following Global Visual SoftICE Path Settings:

- ◆ **Export** — Where you want the top of the retrieved export directory structure to grow.
- ◆ **Script Search Path** (optional) — Where to find scripts when no absolute path is specified.

## **Settings Notes**

### **Path Settings**

Path settings come in two types, a path list (which is an ordered list of paths to search, each one semicolon separated) and a single path entry. All path types for Visual SoftICE accept the ellipses (...) syntax at the end of a directory name to indicate Visual SoftICE should search all directories below the one specified.

Thus to get the source search path to include all your projects and source code, the following example might be useful:

```
SET SRCPATH c:\development\driver_projects\...
```

In this case, Visual SoftICE searches all the files in c:\development\driver\_projects, and any subdirectories underneath that directory, for a matching source file.

## **Per-Workspace Settings Notes**

Workspaces are intended to be containers not only of GUI layout, but also of user preferences. These might be used for particular targets, or types of targets (one for x86, one for IA64, etc...), or even one workspace for debugging, and one for development.

Whatever way they are used, the per-workspace settings provide a means to augment or override global (application wide) settings. The case of interest here is for Path Settings, specifically **Per-Workspace General Path** settings.

Under this section of preferences, the Image, Script, Source, and User Symbol paths are repeated. These entries are provided so that you may *prepend* additional search paths to the global search paths already specified under Global General Path Settings. Optionally (and on a per-path basis), you can replace the global settings of the same name. Visual SoftICE provides a checkbox allowing you to choose whether the path you specify here will prepend or replace the Global General Path setting.

To evaluate the effect these settings have, after modifying them, issue the SET SYMPATH command. The output will allow you to see how Visual SoftICE is using your configuration.

## **Toolbars & Status Bar Settings Notes**

Visual SoftICE provides two optional status bar fields to give you quick updates of the Symbol Engine's current status and actions.

### **VSI System Activity Messages**

This field will be updated when the Symbol Engine takes an action to retrieve, load, or make a given table active. This status bar field is shown in the default workspaces provided with the product.

## VSI Symbol Table Status

This field shows:

- ◆ The current table name, within the set of loaded tables
- ◆ The status of all active tables for the current context
  - ◊ A green icon indicates all tables for the current context are loaded and matched.
  - ◊ A yellow warning icon indicates that some of the tables for the current context may have mismatch or other pertinent warning information.
  - ◊ A red error icon indicates some tables for the current context are missing.
- ◆ The status of automatic table loading
  - ◊ A folder and magnifying glass icon is displayed when enabled, the same icon with a red circle and line over it indicates it is disabled.
- ◆ The status of symbol server retrieval (enabled or disabled)
  - ◊ A server icon is displayed when enabled, the same icon with a red circle and line over it indicates it is disabled

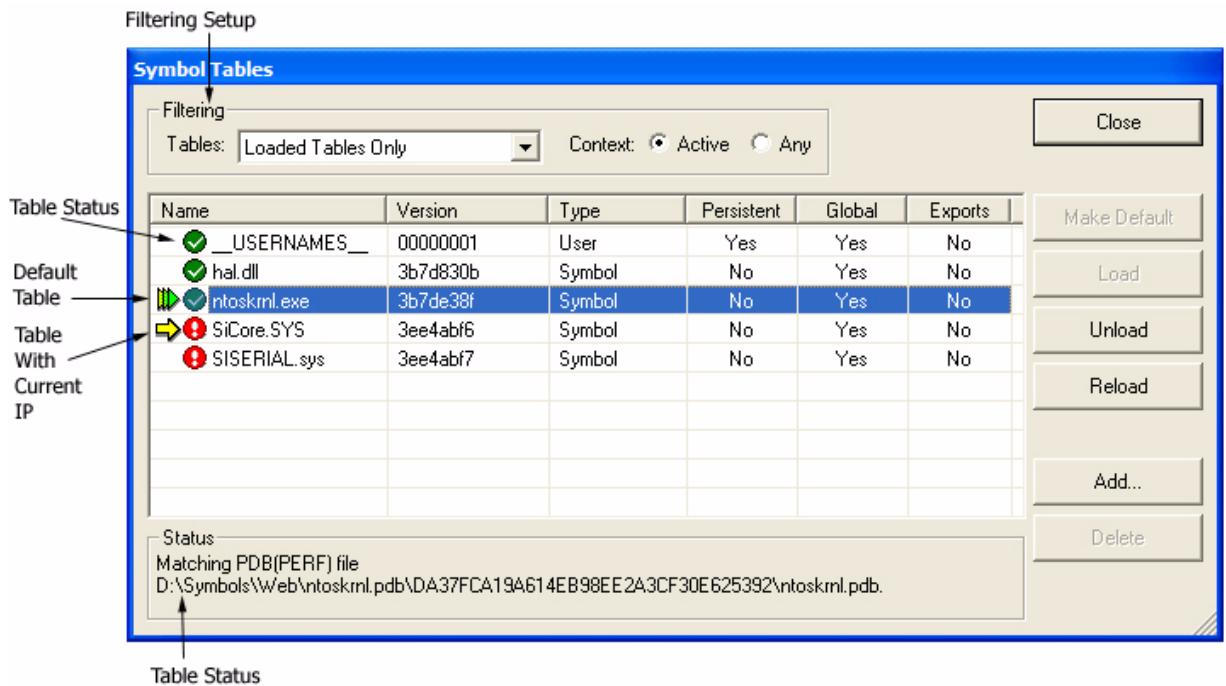
To display these fields on your status bar, click the **Customize** button under the Status Bar section of the **Toolbars & Status Bar** tab in the Preferences dialog.

Visual SoftICE will save your status bar configuration in the active workspace, and restore it any time you open that workspace.

## The Symbol Tables Utility

The Symbol Tables utility allows you to manage the symbol tables with the same available options as the various Visual SoftICE symbol commands. In addition to the basic symbol table information (name, version, type, persistent, and global), the Symbol Tables utility also displays whether exports are loaded for the table. The icon to left of the table name indicates the status of that table. The default table, and table containing the stopped IP, are indicated to left of the status icon. For more detailed information on these icons and their meanings, refer to “Visual SoftICE Icons” on page 54.

For more information on the functionality of the various table management functions available, refer to the symbol commands in the *Visual SoftICE Command Reference*.



**Figure 6-35.** Symbol Tables Utility

The Symbol Tables utility is composed of the following sections and elements:

## Filtering

The Filtering section allows you to select what tables are displayed in tables list. The following selections are available:

**Table 6-8.** Filtering Options

Option	Description
All Tables	Displays the entire list symbol tables that Visual SoftICE is aware of.
User-mode (Ring 3) Tables Only	Displays all the User-mode tables that are loaded into the current process (refer to changing the current process for more information).

**Table 6-8.** Filtering Options (Continued)

Option	Description
System (Ring 0) Tables Only	Displays all the System tables that are loaded into the current process (refer to changing the current process for more information).
Loaded Tables Only	Displays only those tables which are loaded, or that Visual SoftICE currently wants to load that are not loaded. This may include tables that are part of the current thread's stack list. Refer to the LOAD command for more information on dynamic symbol loading.

### **Table Action Buttons**

The table action buttons allow you to make default, load, unload, reload, or delete the currently selected table in the list. Only those actions valid for the select table are enabled. The **Persistent Add** button is always enabled; it allows you add a persistent table to the Visual SoftICE table list.

# Chapter 7

## Using Breakpoints



- ◆ Introduction
- ◆ Types of Breakpoints Supported by Visual SoftICE
- ◆ Understanding Breakpoint Contexts
- ◆ Virtual Breakpoints
- ◆ Setting a Breakpoint Action
- ◆ Conditional Breakpoints
- ◆ Elapsed Time
- ◆ Breakpoint Statistics
- ◆ Referring to Breakpoints in Expressions
- ◆ Manipulating Breakpoints
- ◆ Using Embedded Breakpoints

### Introduction

You can use Visual SoftICE to set breakpoints on program execution, memory location reads and writes, interrupts, and reads and writes to I/O ports. Visual SoftICE assigns a breakpoint index to each breakpoint. You can use this breakpoint index to identify breakpoints when you delete, disable, enable, or reference them.

All Visual SoftICE breakpoints are sticky, which means that Visual SoftICE tracks and maintains a breakpoint until you intentionally clear or disable it using the BC or the BD command. Breakpoints on a target machine are removed when you disconnect, or the connection times out.

The number of breakpoints you can set on memory location (BPMs) and I/O ports (BPIOs) is limited by the number of debug registers the target processor supports (e.g. x86 supports 4).

Where symbol information is available, you can set breakpoints using function names. When in a disassembly, source, or stack page, you can set breakpoints via popup context menus, or at anytime from the breakpoint page, or the debug main menu. A valuable feature is that you can set breakpoints in an image before it is loaded.

## Types of Breakpoints Supported by Visual SoftICE

Visual SoftICE provides a powerful array of breakpoint capabilities that take full advantage of the target architecture, as follows:

- ◆ **Execution Breakpoints:** Visual SoftICE replaces existing instructions with architecture appropriate breaks. You use the BPX or BPS command to set execution breakpoints.
- ◆ **Memory Breakpoints:** Visual SoftICE uses the target processor debug registers to break when a certain byte/word/dword/qword of memory is read, written, or executed. You use the BPM command to set memory breakpoints on scalar data types. You use the BPR command (where supported) to set memory breakpoints on a larger range of memory.
- ◆ **Interrupt Breakpoints:** Visual SoftICE allows setting breakpoints on OS fault handlers or driver provided interrupt service routines (see INTOBJ command). You use the BPINT command to set interrupt breakpoints.
- ◆ **I/O Breakpoints:** Visual SoftICE uses target processor debug registers to watch for an IN or OUT instruction going to a particular port address. You use the BPIO command to set I/O breakpoints.
- ◆ **Image Load Breakpoints:** Visual SoftICE supports two mechanisms to stop on an image load – either when an image matching a name is loaded (using the BLOAD command), or when any image in the system loads (using the SET GLOBALBREAK command).
- ◆ **Window Message Breakpoints:** Visual SoftICE triggers these breakpoints when a particular message or range of messages arrives at a window. This is not a fundamental breakpoint type; it is just a convenient feature built on top of the other breakpoint primitives. You use the BMSG command to set window message breakpoints.

## **Breakpoint Options**

Visual SoftICE supports logging, conditions and actions (triggered) on most breakpoint types.

You can qualify each type of breakpoint with the following options:

- ◆ An optional logged flag can be specified [-l] which indicates that statistics should be updated, but execution continued when the breakpoint is encountered.
- ◆ A conditional expression [IF expression]: The expression must evaluate to non-zero (TRUE) for the breakpoint to trigger. Refer to “Conditional Breakpoints” on page 199.
- ◆ A breakpoint action [DO “command1;command2;...”]: A series of Visual SoftICE commands can automatically execute when the breakpoint triggers. You can use this feature in concert with user-defined macros to automate tasks that would otherwise be tedious. Refer to “Setting a Breakpoint Action” on page 199.

**Note:** For complete information on each breakpoint command, refer to the *Visual SoftICE Command Reference*.

## **Execution Breakpoints**

An execution breakpoint traps executing code such as a function call or language statement. This is the most frequently used type of breakpoint. By replacing an existing instruction with an architecture appropriate break instruction, Visual SoftICE takes control when execution reaches the breakpoint.

Visual SoftICE provides many ways for setting execution breakpoints: using the BPX or BPS command, the breakpoint page, the main debug menu, and off numerous context menus on other pages. The following section describes one way to use commands for setting breakpoints.

### **Using the BPX Command to Set Breakpoints**

Use the BPX command with any of the following parameters to set an execution breakpoint:

```
BPX [-l] address [IF conditional-expression] [DO  
"command1;command2;..."]
```

IF expression

Refer to “Conditional Breakpoints” on page 199.

DO “command1;command2;...”

Refer to “Setting a Breakpoint Action” on page 199.

To set a breakpoint on your application's WinMain function, use this command:

```
BPX WinMain
```

## Memory Breakpoints

A memory breakpoint uses the debug registers found on the target processor to monitor access to a certain memory location. This type of breakpoint is extremely useful for finding out when and where a program variable is modified, and for setting an execution breakpoint in read-only memory. You are limited by the number of processor debug registers to how many breakpoints may be set at one time.

Use the BPM command to set memory breakpoints on scalar data type sized memory blocks:

```
BPM[size] [-l] address [ R | W | RW | X ] [IF conditional-expression][DO "command1;command2;..." ]
```

size	B (byte) W (word) D (dword) Q (qword).
BPM and BPMB	Set a byte-size breakpoint.
R, W, and RW	Break on reads, writes, or both.
X	Breaks on execution; this is more powerful than a BPX-style breakpoint because memory does not need to be modified, enabling such options as setting breakpoints in ROM or setting breakpoints on addresses that are not present.
IF expression	Refer to "Conditional Breakpoints" on page 199.
DO "command1;command2;..."	Refer to "Setting a Breakpoint Action" on page 199.

The following example sets a memory breakpoint to trigger when a value of 5 is written to the Dword (4-byte) variable MyGlobalVariable.

```
BPMD MyGlobalVariable W IF MyGlobalVariable==5
```

If the target location of a BPM breakpoint is frequently accessed, performance can be degraded regardless of whether the conditional expression evaluates to FALSE.

Use the BPR command to set a memory breakpoint on a larger range of memory.

**Note:** The BPR command is currently available only on the IA64 processor.

```
BPR [-l] start-address end-address [R|W|RW|X] [IF conditional-expression] [DO "command1;command2;..." ]
```

```
BPR [-l] start-address L length [R|W|RW|X] [IF conditional-expression] [DO "command1;command2;..." ]
```

The following example sets a memory breakpoint to trigger when a value of 5 is written to the Dword (4-byte) variable MyGlobalVariable.

```
BPMD MyGlobalVariable W IF MyGlobalVariable==5
```

If the target location of a BPM breakpoint is frequently accessed, performance can be degraded regardless of whether the conditional expression evaluates to FALSE.

## Interrupt Breakpoints

Sets an execution breakpoint on an interrupt handler, provided by either the OS or Driver.

Use the BPINT command to set interrupt breakpoints:

```
BPINT [-l] interrupt-number [service-address] [IF conditional-expression] [DO "command1;command2;..." ]
```

interrupt-number	Architecture supported number.
service-address	Often an interrupt is shared between different handler services. This optional parameter allows for setting the breakpoint on the exact service. If excluded, the breakpoint will be set on all service handlers associated with the interrupt.
IF expression	Refer to "Conditional Breakpoints" on page 199.
DO "command1;command2;..."	Refer to "Setting a Breakpoint Action" on page 199.

Visual SoftICE will stop either on the OS fault handler, or the entry point to the driver provided Interrupt Service Routine. You can list all OS interrupts and their handlers by using the IT/IDT command.

## I/O Breakpoints

An I/O breakpoint monitors reads and writes to a port address. The breakpoint traps when an IN or OUT instruction accesses the port.

Use the BPIO command to set I/O breakpoints:

```
BPIO [-l] port [R|W|RW] [IF conditional-expression] [IF expression] [DO "command1;command2;..." ]
```

R, W, and RW	Break on reads (IN instructions), writes (OUT instructions), or both, respectively.
IF expression	Refer to "Conditional Breakpoints" on page 199.
DO "command1;command2;..."	Refer to "Setting a Breakpoint Action" on page 199.

When an I/O breakpoint triggers and Visual SoftICE stops, the current instruction is the instruction following the IN or OUT that caused the breakpoint to trigger. Unlike BPM breakpoints, there is no size specification; any access to the port-number, whether byte, word, dword, or larger triggers the breakpoint. Any I/O that spans the I/O breakpoint will also trigger the breakpoint. For example, if you set an I/O breakpoint on port 2FF, a word I/O to port 2FE would trigger the breakpoint.

Use the following command to set a breakpoint to trigger when a value is read from port 3FEH with the upper 2 bits set:

```
BPIO 3FE R IF (AL & C0)== C0
```

The condition is evaluated after the instruction completes. The value will be in AL, AX, or EAX because all port I/O, except for the string I/O instructions (which are rarely used), use the EAX register.

I/O breakpoints appear as blue throughout the GUI.

## ***Image Load Breakpoints***

An image load breakpoint is a stop request when a given image loads in the operation system and the module or dll contains an entry point that gets executed. This can be filtered to when a named image is loaded, or when ANY image loads, with the following commands:

<code>BUPLOAD [-once] image-name [DO "command1;command2;..." ]</code>	
<code>-once</code>	One-shot breakpoint. Once this breakpoint is hit, it is automatically deleted.
<code>image-name</code>	Name to match for the image or module. This can be a partial match, but does not support wildcards.
<code>DO "command1;command2;..."</code>	Refer to “Setting a Breakpoint Action” on page 199.

```
SET GLOBALBREAK [off/load]
```

This is a general setting for the debugger, and if set to LOAD, the target will stop on every image loaded into the system.

## Window Message Breakpoints

Use a window message breakpoint to trap a certain message or range of messages delivered to a window procedure. Although you could implement an equivalent breakpoint yourself using BPX with a conditional expression, the following BMSG command is easier to use:

```
BMSG [-l] window-handle [begin-msg [end-msg]] [IF conditional-expression] [DO "command1;command2;..."]
```

window-handle	Value returned when the window was created; you can use the HWND command to get a list of windows with their handles.
begin-message	Single Windows message or the lower message number in a range of Windows messages. If you do not specify a range with an end-message, then only the begin-message will cause a break.
end-message	For both begin-message and end-message, the message numbers can be specified either in hexadecimal or by using the actual ASCII names of the messages, for example, WM_QUIT.
IF expression	Higher message number in a range of Windows messages.
DO "command1;command2;..."	Refer to "Conditional Breakpoints" on page 199. Refer to "Setting a Breakpoint Action" on page 199.

When specifying a message or a message range, you can use the symbolic name, for example, WM\_NCPAINT. Use the WMSG command to get a list of the window messages that Visual SoftICE understands. If no message or message range is specified, any message will trigger the breakpoint.

To set a window message breakpoint for the window handle 1001E, use the following command:

```
BMSG 1001E WM_NCPAINT
```

Visual SoftICE is smart enough to take into account the address context of the process that owns the window, so it does not matter what address context you are in when you use BMSG.

## Understanding Breakpoint Contexts

A breakpoint context consists of the address context in which the breakpoint was set, and in what image the breakpoint is in, if any. The concept of breakpoint context applies to all breakpoints, and there are two fundamental types of breakpoints supported by Visual SoftICE, Fixed Address, and Image-Relative.

### ***Image-Relative Breakpoints***

Image-Relative breakpoints are the most common, and most useful type — where the address is understood to be in the context of a given image (even if that image has not loaded, or its memory paged out). Any image-relative breakpoint will trigger in any instance of the image that loads, regardless of its address or process (e.g. a breakpoint set in an OS image like `KERNEL32.DLL` breaks in every process context that has the image loaded, regardless of what context the breakpoint was initially set in). Image relative breakpoints require accurate symbols to be available in order to set them. These types of breakpoints appear as green throughout the GUI.

### ***Fixed Address Breakpoints***

Fixed Address breakpoints are uncommon, and not as useful. They are set at an address in a particular process context. They stay at that address, and are not shared unless their placement accidentally coincides with another OS mapping. Fixed addresses are not recommended, as they can be system fatal if used incorrectly. These types of breakpoints appear as yellow or brown throughout the GUI. Their appearance generally means missing or mismatched symbols are in use.

## Virtual Breakpoints

In Visual SoftICE, you can set breakpoints in images before they load, and it is not necessary for a page to be present in physical memory for a breakpoint to be set. In such cases, the breakpoint is virtual; it will be automatically planted when the image loads or the page becomes present. Virtual breakpoints can only be set when symbols are available.

## Setting a Breakpoint Action

You can set a breakpoint to execute a series of Visual SoftICE commands, including user-defined macros, after the breakpoint is triggered. You define these breakpoint actions with the DO option, which is available with every breakpoint type:

```
DO "command1;command2;..."
```

The body of a breakpoint action definition is a sequence of Visual SoftICE commands, or other macros, separated by semicolons. You need not terminate the final command with a semicolon.

Breakpoint actions are closely related to macros. Breakpoint actions are essentially unnamed macros that do not accept command-line arguments. Breakpoint actions, like macros, can call upon macros. In fact, a prime use of macros is to simplify the creation of complex breakpoint actions.

If a breakpoint is marked as log-only, the action will not be executed.

The following examples illustrate the basic use of breakpoint actions:

```
BPX EIP DO "dd eax"  
BPX EIP DO "data 1;dd eax"  
BPMB dataaddr if (byte(*dataaddr)==1) do "? IRQ"
```

## Conditional Breakpoints

Conditional breakpoints provide a fast and easy way to isolate a specific condition or state within the system (or application) you are debugging. By setting a breakpoint on an instruction or memory address and supplying a conditional expression, Visual SoftICE will only trigger if the breakpoint evaluates to non-zero (TRUE). Because the Visual SoftICE expression evaluator handles complex expressions easily, conditional expressions take you right to the problem or situation you want to debug with ease.

Most Visual SoftICE breakpoint commands accept conditional expressions using the following syntax:

```
breakpoint-command [breakpoint options] [IF conditional  
expression][DO "commands"]
```

The IF keyword, when present, is followed by any expression that you want to be evaluated when the breakpoint is triggered. The breakpoint will be ignored if the conditional expression is FALSE (zero). When the conditional expression is TRUE (non-zero), Visual SoftICE stops and displays the reason for the break, which includes the conditional expression.

Most of the following x86 examples contain system-specific values that vary depending on the exact version of Windows you are running.

- ◆ Watch CSRSS HWND objects (type 1) being created:

```
bpw winsrv!HMAllocObject IF (esp+c == 1)
```

- ◆ Watch CSRSS thread info objects (type 6) being destroyed:

```
bpw winsrv!HFreeObject+0x25 IF (byte(esi+8) == 6)
```

- ◆ Watch process object-handle-tables being created:

```
bpw ntoskrnl!ExAllocatePoolWithTag IF (esp+c == 'Obtb')
```

- ◆ Watch a heap block (230CD8) get freed:

```
bpw ntdll!RtlFreeHeap IF (esp+c == 230CD8)
```

## *Conditional Breakpoint Count Functions*

Visual SoftICE supports the ability to monitor and control breakpoints based on the number of times a particular breakpoint has or has not been triggered. You can use the following count functions in conditional expressions:

- ◆ BPCOUNT
- ◆ BPMISS
- ◆ BPTOTAL
- ◆ BPINDEX

### **BPCOUNT**

The value for the BPCOUNT function is the current number of times that the breakpoint has been evaluated as TRUE.

Use this function to control the point at which a triggered breakpoint causes a popup to occur. Each time the breakpoint is triggered, the conditional expression associated with the breakpoint is evaluated. If the condition evaluates to TRUE, the breakpoint instance count (BPCOUNT) increments by one. If the conditional evaluates to FALSE, the breakpoint miss instance count (BPMISS) increments by one.

The fifth time the breakpoint triggers, the BPCOUNT equals five, so the conditional expression evaluates to TRUE and Visual SoftICE stops.

```
BPX myaddr IF (bpcount==5)
```

Use BPCOUNT only on the righthand side of compound conditional expressions for BPCOUNT to increment correctly:

```
BPX myaddr if (eax==1) && (bpcount==5)
```

Due to the early-out algorithm employed by the expression evaluator, the `BPCount==5` expression will not be evaluated unless `EAX==1` (The C language works the same way). Therefore, by the time `BPCount==5` gets evaluated, the expression is TRUE. BPCOUNT will be incremented and if it equals five, the full expression evaluates to TRUE and Visual SoftICE stops. If `BPCount != 5`, the expression fails, BPMISS is incremented and Visual SoftICE will not stop (although BPCOUNT is now 1 greater).

Once the full expression returns TRUE, Visual SoftICE stops, and all instance counts (BPCOUNT and BPMISS) are reset to 0.

**Note:** Do NOT use BPCOUNT before the conditional expression, otherwise BPCOUNT will not increment correctly:

```
BPX myaddr if (bpcount==5) && (eax==1)
```

## BPMISS

The value for the BPMISS expression function is the current number of times that the breakpoint was evaluated as FALSE.

The expression function is similar to the BPCOUNT function. Use it to specify that Visual SoftICE stop in situations where the breakpoint is continually evaluating to FALSE. The value of BPMISS will always be one less than you expect, because it is not updated until the conditional expression is evaluated. You can use the (`>=`) operator to correct this delayed update condition.

```
BPX myaddr if (eax==43) || (bpmiss>=5)
```

Due to the early-out algorithm employed by the expression evaluator, if the expression `eax==43` is ever TRUE, the conditional evaluates to TRUE and Visual SoftICE stops. Otherwise, BPMISS is updated each time the conditional evaluates to FALSE. After five consecutive failures, the expression evaluates to TRUE and Visual SoftICE stops.

## BPTOTAL

The value for the BPTOTAL expression function is the total number of times that the breakpoint was triggered.

Use this expression function to control the point at which a triggered breakpoint causes a stop to occur. The value of this expression is the total number of times the breakpoint was triggered (refer to the Hits field in the output of the BSTAT command) over its lifetime. This value is never cleared.

The first 50 times this breakpoint is triggered, the condition evaluates to FALSE and Visual SoftICE will not pop up. Every time after 50, the condition evaluates to TRUE, and Visual SoftICE stops on this and every subsequent trap.

```
BPX myaddr if (bptotal > 50)
```

You can use BPTOTAL to implement functionality identical to that of BPCOUNT. Use the modulo “%” operator as follows:

```
if (!(bptotal%COUNT))
```

The COUNT is the frequency with which you want the breakpoint to trigger. If COUNT is four, Visual SoftICE stops every fourth time the breakpoint triggers.

## BPINDEX

Use the BPINDEX expression function to obtain the breakpoint index to use with breakpoint actions.

This expression function returns the index of the breakpoint that caused Visual SoftICE to stop. This index is the same index used by the BL, BC, BD, BE, BPE, BPT, and BSTAT commands. You can use this value as a parameter to any command that is being executed as an action.

The following example of a breakpoint action causes the BSTAT command to be executed with the breakpoint that caused the action to be executed as its parameter:

```
BPX myaddr do "bstat bpindex"
```

This example shows a breakpoint that uses an action to create another breakpoint:

```
BPX myaddr do "t;bpx @esp if(tid==_tid) do \"bc bpindex\";g"
```

BPINDEX is intended to be used with breakpoint actions, and causes an error if it is used within a conditional expression. Its use outside of actions is allowed, but the result is unspecified and you should not rely on it.

## **Using Local Variables in Conditional Expressions**

Visual SoftICE lets you use local variable names in conditional expressions as long as the type of breakpoint is an execution breakpoint (BPX or BPM X). Visual SoftICE does not recognize local symbols in conditional expressions for other breakpoint types, such as BPIO or BPMD RW, because they require an execution scope. This type of breakpoint is not tied to a specific section of executing code, so local variables have no meaning.

When using local variables in conditional expressions, functions typically have a prologue where local variables are created and an epilogue where they are destroyed. You can access local variables after the prologue code completes execution and before the epilogue code begins execution. Function parameters are also temporarily inaccessible using symbol names during prologue and epilogue execution, because of adjustments to the stack frame.

To avoid these restrictions, set a breakpoint on either the first or last source code line within the function body. We will use the following “foobar function” to explain this concept.

### **Foobar Function**

```
1:DWORD foobar ( DWORD foo )
2:{  
3:DWORDfooTmp=0;  
4:  
5:if(foo)
6:{  
7:fooTmp=foo*2;
8:}else{
9:fooTmp=1;
10:}  
11:  
12:return fooTmp;
13:}
```

Source code lines 1 and 2 are outside the function body. These lines execute the prologue code. If you use a local variable at this point, you receive the following symbol error:

```
>BPX foobar if(foo==1)
error: Undefined Symbol (foo)
```

Set the conditional on the source code line 3, where the local variable fooTmp is declared.

Source code line 13 marks the end of the function body. It also begins epilogue code execution; thus, local variables and parameters are out of scope.

Although it is possible to use local variables as the input to a breakpoint command, such as BPMD RW, you should avoid doing this. Local variables are relative to the stack, so their absolute address changes each time the function scope where the variable is declared executes. When the original function scope exits, the address tied to the breakpoint no longer refers to the value of the local variable.

## **Referencing the Stack in Conditional Breakpoints**

If you create your symbol file with full symbol information, you can access function parameters and local variables through their symbolic names, as described in “Using Local Variables in Conditional Expressions” on page 203. If, however, you are debugging without full symbol information, you need to reference function parameters and local variables on the stack.

**Note:** The following section is specific to x86 32-bit flat application or system code.

Function parameters are passed on the stack, so you need to de-reference these parameters through the ESP or EBP registers. Which one you use depends on the function’s prologue and where you set the actual breakpoint in relation to that prologue.

Most 32-bit functions have a prologue of the following form:

```
PUSH    EBP  
MOV     EBP,ESP  
SUB    ESP,size (locals)
```

Which sets up a stack frame as follows:

- ◆ Use either the ESP or EBP register to address parameters. Using the EBP register is not valid until the PUSH EBP and MOV EBP, ESP instructions are executed. Also note that once space for local variables is created (SUB ESP, size) the position of the parameters relative to ESP needs to be adjusted by the size of the local variables and any saved registers.
- ◆ Typically you set a breakpoint on the function address, for example:

```
BPX IsWindow
```

<p>Stack Top</p> <p>Current EBP →</p> <p>Stack Bottom</p> <p>Current ESP →</p>	PARAM n	ESP+(n*4), or EBP+(n*4)+4	Pushed by caller
	PARAM #2	ESP+8, or EBP+C	
	PARAM #1	ESP+4, or EBP+8	
	RET EIP	= Stack pointer on entry	
	SAVE EBP	= Base pointer (PUSH EBP, MOV EBP,ESP)	Call prologue
	LOCALS+size-1		
	LOCALS+0	= Stack pointer after prologue (SUB ESP, size (locals))	
	SAVE EBX	optional save of 'C' registers	
Registers saved by compiler	SAVE ESI		Registers saved by compiler
	SAVE EDI	= Stack pointer after registers are saved	

When this breakpoint is triggered, the prologue has not been executed, and parameters can easily be accessed through the ESP register. At this point, use of EBP is not valid.

**Note:** This assumes a stack-based calling convention with arguments pushed right-to-left.

To be sure that de-referencing the stack in a conditional expression operates as you would expect, use the following guidelines.

- ◆ If you set a breakpoint at the exact function address, for example, BPX IsWindow, use `ESP+(param# * 4)` to address parameters, where param# is 1...n.
- ◆ If you set a breakpoint inside a function body (after the full prologue has been executed), use `EBP+(param# * 4)+4` to address parameters, where param# is 1...n. Be sure that the routine does not use the EBP register for a purpose other than a stack-frame.
- ◆ Functions that are assembly-language based or are optimized for frame-pointer omission may require that you use the ESP register, because EBP may not be set up correctly.

**Note:** For x86, once the space for local variables is allocated on the stack, the local variables can be addressed using a negative offset from EBP. The first local variable is at `EBP-4`. Simple data types are typically Dword sized, so their offset can be calculated in a manner similar to function parameters. For example, with two pointer local variables, one will be at `EBP-4` and the other will be at `EBP-8`.

## **Performance**

Conditional breakpoints have some overhead associated with run-time evaluation. Under most circumstances you see little or no effect on performance when using conditional expressions. In situations where you set a conditional breakpoint on a highly accessed data variable or code sequence, you may notice slower system performance. This is due to the fact that every time the breakpoint is triggered, the conditional expression is evaluated. If a routine is executed hundreds of times per second (such as ExAllocatePool), the fact that any type of breakpoint with or without a conditional is trapped and evaluated with this frequency results in some performance degradation.

## **Duplicate Breakpoints**

Once a breakpoint is set on an address, you cannot set another breakpoint on the same address. With conditional expressions, however, you can create a compound expression using the logical operators (`&&`) or (`||`) to test more than one condition at the same address.

## **Elapsed Time**

Visual SoftICE supports using the target processor time stamp counter. Every time Visual SoftICE stops due to a breakpoint, the elapsed time displays since the last time Visual SoftICE stopped. The time displays after the break reason in seconds, milliseconds, or microseconds, and this can be seen in the command page, or breakpoint page history log.

Most processor time stamp counters are highly accurate, but you must keep the following issues in mind:

- ◆ There is overhead involved in stopping Visual SoftICE, which may impact the results very slightly.
- ◆ If a hardware interrupt occurs before the breakpoint goes off, all the interrupt processing time is included.
- ◆ Certain processors will vary the clock rate dynamically, for power management reasons.

## Breakpoint Statistics

Visual SoftICE collects statistical information about each breakpoint, including the following:

- ◆ Total number of hits, breaks, misses, and errors
- ◆ Current hits and misses

Use the BSTAT command to display this information. Refer to the *Visual SoftICE Command Reference* for more information on the BSTAT command.

## Referring to Breakpoints in Expressions

You can combine the prefix “BP” with the breakpoint index to use as a symbol in an expression. This works for all BPX and BPM breakpoints. Visual SoftICE uses the actual address of the breakpoint.

To disassemble code at the address of the breakpoint with index 0, use the command:

```
U BP0
```

## Manipulating Breakpoints

The Breakpoint page allows you to enable, disable, or delete a breakpoint by right-clicking on that breakpoint and selecting the desired operation (**Enable**, **Disable**, or **Delete**) from the pop-up menu. For any breakpoint that supports disabling, you can also double-click the breakpoint entry icon to toggle it between enabled and disabled states. Likewise, highlighting the breakpoint entry icon and pressing **Delete** will remove that breakpoint.

Visual SoftICE also provides a variety of commands for manipulating breakpoints such as listing, modifying, deleting, enabling, and disabling breakpoints. Breakpoints are identified by breakpoint index numbers, which are uniquely assigned. [Table 7-1](#) describes the breakpoint manipulation commands.

**Table 7-1.** Visual SoftICE Breakpoint Manipulation Commands

Command	Description
BD	Disable a breakpoint.
BE	Enable a breakpoint.
BL	List current breakpoints.
BC	Clear (remove) a breakpoint.

**Note:** Refer to the *Visual SoftICE Command Reference* for more information on each of these commands.

## Using Embedded Breakpoints

It may be helpful for you to embed a breakpoint in your program source rather than setting a breakpoint with Visual SoftICE. To embed a breakpoint in your program, do the following:

- ◆ Place an INT 1 or INT 3 instruction at the desired point in the program source.
- ◆ To enable Visual SoftICE to pop up on such embedded breakpoints, use one of the following commands:
  - ◊ **I1HERE ON** for INT 1 breakpoints
  - ◊ **I3HERE ON** for INT 3 breakpoints

# Chapter 8

## Using Expressions



- ◆ Expression Values
- ◆ Supported Operators
- ◆ Forming Expressions
- ◆ Expression Evaluator Type System

### Expression Values

The Visual SoftICE expression evaluator determines the values of expressions used with Visual SoftICE commands and conditional breakpoints. It provides full operator precedence; support for standard C language arithmetic, bit-wise, and logical operators; predefined macros, functions, and casts for data type conversion; and access to common Visual SoftICE and operating system values.

The Visual SoftICE expression evaluator parses and evaluates expressions similarly to the way a C or C++ language compiler translates expressions. If you are comfortable with either language, you are already familiar with the grammar and syntax of Visual SoftICE expressions.

There are no limitations on the complexity of an expression. You can combine multiple operators, operands, and expressions to create compound expressions for conditional breakpoints or expression evaluation.

This example uses a compound expression to trigger a breakpoint if the first parameter (`ESP+4`) passed to the `IsWindow( )` API function is an `HWND` with the value of `0x10022` or `0x1001E`. If either of the two expressions is `TRUE`, the conditional expression is `TRUE`, and the breakpoint triggers:

```
BPX IsWindow if (esp+4 == 10022) || (esp+4 == 1001E)
```

## Supported Operators

The Visual SoftICE expression evaluator supports the following operators sorted by type:

**Note:** Use the SET EE\_IMPL\_DEREF command to control the expression evaluator's behavior regarding dereferencing. If you have set EE\_IMPL\_DEREF to *on*, and the expression evaluator encounters an expression containing a symbol that is a pointer, it will use the value it points to for evaluation. If you have set EE\_IMPL\_DEREF to *off*, and the expression evaluator encounters an expression containing a symbol that is a pointer, it will use the address of the pointer for evaluation.

**Table 8-1.** Visual SoftICE Indirection Operators

Indirection Operators	Example
*	*eax (gets the Dword value pointed to by eax)
& <i>symbol-name</i>	&Foo (gets the address of the symbol Foo)
[ ] (array subscript)	Foo[ 2 ] (gets the second element of the array Foo)

**Table 8-2.** Visual SoftICE Math Operators

Math Operators	Example
+	eax + 1
-	ebp - 4
*	ebx * 4
/	Symbol / 2
% (modulo)	eax % 3
<< (logical shift left)	b1 << 1 (result is bl shifted left by 1)
>> (logical shift right)	eax >> 2 (result is eax shifted right by 2)

**Table 8-3.** Visual SoftICE Bitwise Operators

Bitwise Operators	Example
& (bitwise AND)	eax & F7
(bitwise OR)	Symbol   4
^ (bitwise XOR)	ebx ^ 0xFF

**Table 8-4.** Visual SoftICE Logical Operators

Logical Operators	Example
! (logical NOT)	<code>! eax</code>
&& (logical AND)	<code>eax &amp;&amp; ebx</code>
(logical OR)	<code>eax    ebx</code>
== (compare equality)	<code>Symbol == 4</code>
!= (compare inequality)	<code>Symbol != al</code>
<	<code>eax &lt; 7</code>
>	<code>bx &gt; cx</code>
<=	<code>ebx &lt;= Symbol</code>
>=	<code>Symbol &gt;= Symbol</code>

## Operator Precedence

Operator precedence within the Visual SoftICE expression evaluator is equivalent to the C language operator precedence. Operator precedence plays a crucial part in evaluating expressions, so the order in which you input expression operators can have a dramatic result on the final result of the expression. To override the default operator precedence to produce a desired result, use parentheses to force the order of evaluation.

## Forming Expressions

**Tip:** Use the ? or EVAL (evaluate expression) command to display the result of any expression.

The Visual SoftICE expression evaluator accepts a variety of operands, such as symbols, register names, user-defined symbols, and numbers, that you can combine with any Visual SoftICE operator. Visual SoftICE places an emphasis on providing flexibility of expression, so input is as natural as possible.

## Numbers

The Visual SoftICE expression evaluator accepts the following numeric inputs.

**Table 8-5.** Visual SoftICE Expression Inputs

Input	Description
Hexadecimal	<p>The default radix for input and output is controlled by the SET RADIX command. The valid character set for hexadecimal numbers is [0-9, A-F]. Hexadecimal input can be optionally preceded by the standard C language radix identifier: 0x. Examples of valid hexadecimal numbers include:</p> <p>FF, ABC, 0x123, 0xFFFF0000</p> <p>The symbolic form of a valid hexadecimal number could conflict with a symbol name. For example, ABC. Use the 0x form to ensure that the number is not misinterpreted as a symbol name.</p> <p>Prefixing 0x to input forces it to be evaluated as hexadecimal input regardless of the default radix.</p> <p>You can also use the HEX(n) and DEC(n) functions to force interpretation to a given radix.</p>

## Registers

Visual SoftICE supports multiple names for the target register sets. You control the target register set name using the SET REGNAME command.

A register name or alias might conflict with a symbol. To force input to be evaluated as a register name, use the REG(n) function.

## Symbols

Symbol names are the symbolic representation of an address or value. They are defined in symbol tables, export tables, or via Visual SoftICE's NAME command, during debugging.

Symbol names in Visual SoftICE differ from symbols defined in C or C++ programs. All compilers add some form of decoration to the names defined in a program, and this decoration often includes characters which are not valid in C/C++ symbol names. Visual SoftICE therefore accepts a wider range of characters in symbol names than a compiler. **Table 8-6** shows the characters which may be found in a legal symbol name. Symbols must begin with one of the characters marked valid as first symbol characters in the table.

**Table 8-6.** "Legal" Symbol Characters

Characters	
A..Z and a..z	Yes
0..9	No
dollar sign (\$)	Yes
underscore (_)	Yes
exclamation point (!)	No
scope operator (:)	Yes

The scope operator (:) is allowed in symbols. However, note that the "operator" is in this context simply part of the symbol name, and is not functioning as a true operator. Any number of scope operators are allowed in a symbol name, so namespaces and nested classes will function properly.

Each symbol file loaded into Visual SoftICE is placed in a separate table, and only one symbol table can be "active" at a time. (Refer to the TABLE command in the *Visual SoftICE Command Reference* for more information on changing the active table.)

To specify a symbol from an inactive symbol table in an expression, you may precede the symbol with the table name, followed by an exclamation point, followed by the symbol name. For example:

```
table-name!symbol-name
```

Symbols that are defined by the NAME command are always active, because Visual SoftICE treats these symbol sources as global.

## **Built-in Casts and Functions**

Visual SoftICE predefines a number of casts and functions for use in expressions. They can be used within expressions to modify values or translate data types.

Use casts (or functions that do not take arguments) just like symbols from a symbol table. Functions that accept arguments operate on user-specified values, looking and behaving like C language functions and have the following form:

```
FUNC (arg-list)
```

The following casts are supported by Visual SoftICE:

Table 8-7. Visual SoftICE Casts

Name	Description	Example
Byte	Get low-order byte	? (Byte) 0x1234 = 0x34
UChar	Convert to unsigned Char	? (UChar) 0x12345678 = 0x78
Bool	Get 1 bit value	? (Bool) 0x1234 = 0x1
Word	Get low-order word	? (Word) 0x12345678 = 0x5678
Dword	Get low-order dword	? (Dword) 0xFF = 0x000000FF
Qword	Convert to quad word	? (Qword) 0x12345678 = 0x12345678
UlongLong	Convert to unsigned double long	? (UlongLong) 0x12345678 = 0x12345678
Short	Convert to short (INT16)	? (Short) 0x12345678 = 0x5678
UShort	Convert to unsigned short	? (UShort) 0x12345678 = 0x5678
Long	Convert byte or word to signed long	? (Long) 0xFF = 0x00000000000000FF ? (Long) 0xFFFF = 0x000000000000FFFF
ULong	Convert to unsigned long	? (ULong) 0xFF = 0x00000000000000FF

Visual SoftICE also has a few useful variables:

**Table 8-8.** Visual SoftICE Predefined Values

Name	Description	Example
IRQL	Windows OS IRQ Level	? IRQL = unsigned-char
Process	KPEB (Kernel Process Environment Block) of the Active OS process	? process
Thread	KTEB (Kernel Thread Environment Block) of the Active OS thread	? thread
PID	Active process Id	? pid == Test32Pid
TID	Active thread Id	? tid == Test32MainTid
BPCount	Breakpoint instance count. For these BP functions, refer to "Conditional Breakpoint Count Functions" on page 200	bp <bp params> IF bpcount == 0x10
BPTotal	Breakpoint total count	bp <bp params> IF bptotal > 0x10
BPMiss	Breakpoint instance miss count	bp <bp params> IF bpmiss == 0x20
BPIIndex	Current Breakpoint Index #	bp <bp params> DO "bd bpindex"

Visual SoftICE also has a few predefined macro-like functions:

**Table 8-9.** Visual SoftICE Predefined Macro-like Functions

Name	Description	Example
HiByte	Get high-order byte	? HiByte(0x1234) = 0x12
HiWord	Get high-order word	? HiWord(0x12345678) = 0x1234

## Expression Evaluator Type System

The Visual SoftICE expression evaluator uses a very basic type system that categorizes all expression values into one of the following types:

Table 8-10. Visual SoftICE Expression Types

Type	Example
Literal	1, 0x80000000, 'ABCD'
Register	EAX, DS, ESP
Symbol-type	PoolHitTag, IsWindow
Address-type	FS:18, &Symbol, WIN32k!CreateCompatiblePublicDC

In most cases, you can ignore the distinction between types as it is only important to Visual SoftICE. In the cases of symbol-type and address-type, however, there are important semantics or restrictions.

### Symbol Type

The symbol-type is used for symbol names that are in export or symbol tables. In general, the type represents the linear address of a symbol within a code or data segment. The symbol type also represents the contents of memory at that linear address. This is similar to the use of a variable in a C program, but because Visual SoftICE is a debugger and not a compiler, there are a few semantic differences. Visual SoftICE determines whether you mean *contents-of* or *address-of* based on the context of how you use the symbol/variable in an expression. In general, the way Visual SoftICE treats a symbol seems completely natural, not unlike that of the C compiler; but, in cases where you are not sure how Visual SoftICE interprets the symbol, you can explicitly state:

address-of (&Symbol) or contents-of (\*Symbol)

**Note:** Refer also to the SET EE\_IMPL\_DEREF command.

When symbol-types are used in expressions, Visual SoftICE will, in most cases, present the result of the expression in the correct type. For example, given an array of integers declared like this:

```
int TinyArray[ ] = { 1, 2, 3, 4 };
```

The expression:

```
?TinyArray[ 1 ]
```

will cause Visual SoftICE to display the second element of the array, which will be of type **int**.

Alternately, if you have a pointer-to-char expression declared like this:

```
char *str = "Twas Brillig"
```

the expression

```
*str
```

will result in the following display:

```
<char> = 0x54, 'T', 84
```

## **Address Type**

Visual SoftICE treats a symbol as an address-type if you use it in an expression where an address-type is legal and it makes sense to use an address. Otherwise, Visual SoftICE automatically indirections the symbol, taking the contents of the memory the symbol represents. You can also control this behavior with the SET EE\_IMPL\_DEREF command.

There are many operations that are illegal or do not make sense for address-types such as multiplication and division, so a majority of the operators used with the symbol-type act like a C compiler and automatically take the contents-of at the address for the symbol.

## **Evaluating Symbols**

When data type information is available, using the ? (evaluate expression) command with a symbol yields the contents of the symbol rather than the address of the symbol. For example, MyVariable is an integer variable containing the value 5, so you get the following:

```
? MyVariable  
int=0x5, "\0\0\0\x05"
```

To get the address of MyVariable, use the following:

```
? &MyVariable
```

If you use a symbol in conjunction with a command other than '?,' be sure to add the address of the '&' operator where needed. For example, the data display command (D) takes an address as a parameter, so to display the contents of a symbol, you should add the '&' operator:

```
dd &MyVariable
```

## **Pointer Arithmetic with Symbols**

When Visual SoftICE performs arithmetic on a symbol whose type is an address, it will perform C-style pointer arithmetic by scaling the second operand by the size of the first. So, given this declaration:

```
long Numbers[] = { 1, 2, 3, 4 };  
long *ptr = Numbers;
```

The Visual SoftICE command

```
? ptr + 1
```

will be equivalent to the same expression in C. Thus, the offset (1) will be scaled by the size of the type pointed to by ptr; in this case, 4 bytes. This causes Visual SoftICE to display the second element of the Numbers array.

## **Array Symbols In Expressions**

Visual SoftICE's array operator allows you to evaluate and display individual members of arrays. It has a couple of limitations, however. First of all, Visual SoftICE does not allow multi-dimensional array expressions. Entering ? mychars[1][1], for example, will produce an error.

Secondly, unlike C and C++, Visual SoftICE does not treat pointers and arrays as equivalent. Using an array operator on a pointer type will therefore produce unpredictable results.

# Chapter 9

## Exploring Windows



- ◆ Overview
- ◆ Inside the Windows Kernel
- ◆ Win32 Subsystem

### Overview

Without qualification, the Windows NT operating system family (Windows NT, Windows 2000, and Windows XP) represents an incredible feat of software engineering and system design. It is hard to imagine a design of such complexity reaching all of its goals, including three of the most difficult: portability, reliability, and extensibility, without compromising either interfaces or implementation. Yet, somehow the system engineers at Microsoft who design and develop the Windows operating system family have managed to keep each and every component of these systems smoothly interlocked, not unlike the precision gears of a finely-made watch. If you are going to write Windows applications, you should explore what lies beneath your application code: the operating system. The knowledge you gain from the time you invest to go beneath your application and into the depths of the system, will benefit both you and the application or driver that you are creating.

This chapter provides a quick overview of the more pertinent and interesting aspects of the basic Windows Operating System. By combining this information with available reference material and a little exploration using a debugger, you should be able to gain a basic understanding of how the components of Windows fit together.

For the purposes of this chapter, the use of the term “Windows” refers to all of the operating systems, released by Microsoft, based on the Windows NT kernel. This includes: Windows 2000, Windows XP, Windows Server 2003, and early pre-releases of Windows Code-named Longhorn.

## ***Resources for Advanced Debugging***

Microsoft provides several resources for advanced Windows debugging including: checked build, the Windows NT DDK, symbol files, driver verifier, and kernel debugger extensions.

### **Checked Build**

If you are not currently using the checked build (that is, the debug version) of Windows, you are missing a lot of valuable information and debugging support that the operating system provides. The checked build contains a wealth of information that is absent from the free build (retail version). This includes basic debug messages, special flags used by the kernel components that allow you to trace the system’s operation, and relatively strict sanity checking of most system API calls. The size and layout of system data structures as well as the implementation of system APIs in the checked build are nearly identical to that of the free build. This allows you to learn and explore using the more verbose checked build, but still feel completely comfortable if you end up debugging under the free build.

It is also possible to use individual components from the checked build on a free build installation. This is often helpful when trying to pin down a crash or other problem that is happening inside an OS component. Using checked build components in this way is as simple as copying the checked build module onto the target system.

All in all, if you want to write more robust applications and drivers, use the checked build.

## Windows DDK

The Windows DDK contains header files, sample code, online help, and special tools that let you query various kernel components. The most obvious and useful resource is NTDDK.H. Although there is quite a bit of information missing from this header file, enough pertinent information is available to make it worth studying. Besides the basic data structures needed for device driver development, system data structures are described (some completely, others briefly, many not at all). There are also many API prototypes and type enumerations that are useful for both exploration and development. There are also useful comments about the system design, as well as restrictions and limitations.

Most of the other header files in the DDK are specific to the more esoteric aspects of the system, but WDM.H, NTDEF.H, BUGCODES.H, and NTSTATUS.H are generally useful.

The Windows DDK includes a few utilities that are of general interest. For example, POOLMON.EXE allows you to monitor system pool usage, and OBJDIR.EXE provides information on the Object Manager hierarchy and information about a specific object within the hierarchy. Visual SoftICE provides similar functionality with the OBJDIR, DEVICE, and DRIVER commands. The utility DRIVERS.EXE, like the Visual SoftICE DRIVER command, lists all drivers within the system, including basic information about the driver. Some versions of the Windows DDK include a significantly more powerful version of the standard PSTAT.EXE utility. PSTAT is a Win32 console application that provides summary information on processes and threads. Included with the Win32 SDK and the Visual C++ compiler, are two utilities worth noting: PVIEW and SPY++. Both provide information on processes and threads, and SPY++ provides HWND and CLASS information.

The Windows DDK also includes help files and reference manuals for device driver development, as well as sample code. The sample code is most useful, because it provides you with the information necessary for creating actual Windows device drivers. Simply find something in your area of interest, build that sample, and step through it with Visual SoftICE.

## Symbol Files

Debug files come in one of two formats. Depending on which version of the operating system you are using, you may either have .DBG debug files or .PDB debug files. For operating systems prior to Win2k all OS symbols were released in the .DBG format. Prior to Win2k-sp3 most OS files were released in .PDB form with the exception of ntoskrnl and a few other key system files. For all operating systems after Win2k-sp3, the debug format has been exclusively .PDB files.

Microsoft provides a separate debug file for every distributed executable file with both the checked and free builds of the Windows operating system. This includes the systems components that make up the kernel executive, device drivers, Win32 system DLLs, sub-system processes, control panel applets, and even accessories and games. The .DBG files contain basic debug information similar to the PUBLIC definitions of a .MAP file. Every API and global variable, exported or otherwise, has a basic definition (for example, name, section and offset). The .PDB-format symbol files include most information on most structures; the older .DBG files do not.

Information on locals is not provided in Microsoft's public symbol files, but having access to a public definition for each API makes debugging through system calls a lot easier.

To examine the symbols in the symbol files, issue the SYM command. To get known types, issue the TYPES command. To get a breakout of a structure, issue the TYPES *-v structname* command. Within Visual SoftICE you can also cast a block of memory to a structure type by casting an address to a type, for example:

? (\_KTEB)address

or

WD address structname

Microsoft has introduced a technology called "Symbol Server" that makes it very easy to get the matching symbols for a given binary. Symbol Server was introduced with the release of Windows XP because of the large scale use of Windows Update. The basic mechanism behind symbol server is that it maintains a collection of debug files that are stored on a server and are uniquely identifiable to a given binary through time/date stamps, GUIDS, files size, and age. Microsoft also supplies the server creation software so that you can setup your own local symbol server for your own binaries. Visual SoftICE integrates directly with symbol servers.

Regardless of your specific area of interest, symbols for the following key system components are important to have available.

Table 9-1. Key System Component Symbols

Component	Description
NTOSKRNL.EXE	The Windows Kernel. (Most of the operating system resides here.)
HAL.DLL	The Hardware Abstraction Layer. Important primitives for NTOSKRNL.
NTDLL.DLL	Basic implementation of the Win32 API, and functionality traditionally attributed to KERNEL. Also the interface between USER and SYSTEM mode. Essentially replaces KERNEL32.DLL.
CSRSS.EXE	The Win32 subsystem server process. Most subsystem calls are routed through this process.
WIN32K.SYS	A system device driver that minimizes inter-process communication between applications and CSRSS. Provides kernel mode equivalents for many of the Win32 APIs.
USER32.DLL	Basic implementation of USER functionality. Mostly stubs to WIN32K.SYS (via LPC to CSRSS).
KERNEL32.DLL.	Some basic implementation of traditional KERNEL functionality, but mostly stubs to NTDLL.DLL.

## Driver Verifier

Microsoft has started adding large numbers of runtime validation checks to the operating system. The types of checks that it provides varies based upon the OS, but in general items such as pool corruption, IRQL violation errors, and low resource simulation. Each release of Windows NT adds additional options and checks. By default these checks are not enabled and need to be enabled by running the `Verifier.exe` utility. From this utility you choose which drivers to analyze and what items to validate. When Driver Verifier finds an exception case, (for example, overrunning a buffer) it generates a blue screen with a stop code of (usually) `0xc4` or `0xc9` and its bugcheck parameters provide additional information. If you have Visual SoftICE loaded you can debug the case that the verifier flagged. BoundsChecker, part of the DriverStudio product, also supplies most of the same features and functionality as driver verifier but has a full user mode UI, and will not cause a crash.

## Resources

The following resources provide extensive information for developing drivers and applications for Windows:

- ◆ *Microsoft Developers Network (MSDN)*  
MSDN, published quarterly on CD-ROM, contains a wealth of information and articles on all aspects of programming Microsoft operating systems. This is one of the only places where you can find practical information on writing Windows device drivers.
- ◆ *Inside Microsoft Windows 2000* - David A. Solomon, Mark E. Russinovich, Microsoft Press  
*Inside Microsoft Windows 2000* provides a high-level view of the design for the Windows 2000 operating system. Each major sub-system is thoroughly discussed, and many block diagrams illuminate internal data structures, policies, and algorithms. Currently, this is the most definitive work on Windows 2000 operating system internals. You will gain the most benefit from the information in this book if you use Visual SoftICE to explore the actual implementation of the system design, for when you step into OS code with Visual SoftICE, many of the higher-level relationships become clear.
- ◆ *Advanced Windows* - Jeffrey Richter, Microsoft Press  
*Advanced Windows* is an excellent resource for the systems programmer developing Win32 applications and system code. Richter presents extensive discussions of processes, threads, memory management, and synchronization objects. Relevant sample code and utilities are also provided.
- ◆ *Programming the Windows Driver Model* - Walter Oney, Microsoft Press  
*Programming the Windows Driver Model* is an excellent resource and the definitive resource for the device driver programmer.
- ◆ *Undocumented Windows 2000 Secrets* - Sven B. Schreiber, Addison Wesley  
*Undocumented Windows 2000 Secrets* focuses on undocumented interfaces and APIs, and is a good introduction to exploratory debugging on Windows.

## Inside the Windows Kernel

To gain a basic understanding of Windows, look at the platform from many different perspectives. A general knowledge of how Windows works at different levels enables you to understand the constraints and assumptions involved in designing other aspects of the operating system.

This section explains the most critical component of the operating system, the Windows Kernel. It describes how Windows configures the core operating system data structures, such as the IDT and TSS, and how to use corresponding Visual SoftICE commands to illustrate the Windows configuration of the CPU. It also examines a general map of the Windows system memory area, describing important system data structures and examining the critical role they play within the operating system.

A majority of the information in this section is based on the implementation details of the following two modules:

- ◆ **Hardware Abstraction Layer (HAL.DLL)**

HAL is the Windows hardware abstraction layer. Its purpose is to isolate as many hardware platform dependencies as possible into one module. This makes the Windows kernel code highly portable.

Various parts of the kernel use platform dependent code, but only for performance considerations.

The primary responsibility of the HAL is to deal with very low-level hardware control such as Interrupt controller programming, hardware I/O, and multiprocessor inter-communication. Many of the HAL routines are dedicated to dealing with specific bus types (PCI, EISA, ISA) and bus adapter cards. HAL also controls basic fault handling and interrupt dispatch.

◆ The Kernel (NTOSKRNL.EXE)

The vast majority of the Windows operating system resides in the Windows Kernel, or Kernel Executive. This is the kernel-level functionality that all other system components, such as the Win32 subsystem, are built upon. The Kernel Executive Services cover a broad range of functionality, including:

- ◊ Memory Management
  - ◊ Object Management
  - ◊ Process and Thread creation and manipulation
  - ◊ Process and Thread scheduling
  - ◊ Local Procedure Call (LPC) facilities
  - ◊ Security Management
  - ◊ Exception handling
  - ◊ VDM hardware emulation
  - ◊ Synchronization primitives, such as Semaphores and Mutants
  - ◊ Run Time Library
  - ◊ File System
  - ◊ Power Management
  - ◊ Multi Processor Synchronization
- ◆ I/O subsystems

## Managing the Intel Architecture

One of the fundamental requirements of starting a protected-mode operating system is the setup of CPU architecture, policies, and address space that the operating system will use. System initialization is coordinated between NTLDR, NTDETECT, NTOSKRNL, and HAL. Use the following Visual SoftICE commands to obtain a general idea of how Windows uses the Intel architecture to provide a secure and robust environment.

Table 9-2. Visual SoftICE Architecture Commands

Command	Description
IT/IDT	Display information on the Interrupt Descriptor Table
GDT	Display information on the Global Descriptor Table
MSR	Displays information on the Model Specific Registers

**Note:** The *Visual SoftICE Command Reference* provides detailed information about using each command.

## IT/IDT (Interrupt Descriptor Table)

Windows creates an IDT for 255 interrupt vectors and maps it into the system linear address space. The first 48 interrupt vectors are generally used by the kernel to trap exceptions, but certain vectors provide operating system services or other special features. Use the Visual SoftICE IT/IDT command to view the Windows Interrupt Descriptor Table.

Table 9-3. Interrupt Descriptor Table

Interrupt #	Purpose
2	NMI. A Task gate is installed here so the OS has a clean set of registers, page-tables, and level 0 stack. This enables the operating system to continue processing long enough to throw a Blue Screen.
8	Double Fault. A Task gate is installed here so the OS has a clean set of registers, page-tables, and level 0 stack. This enables the operating system to continue processing long enough to throw a Blue Screen.
2A	Service to get the current tick count.
2B,2C	Direct thread switch services (older versions of Windows NT).
2D	Debug service.
2E	Execute System Service. Prior to Windows XP, Windows used INT 2E to transition from user to system mode. Since Windows XP, this mechanism has been replaced on newer processors with the faster SYSENTER/SYSEXIT instructions, but the old mechanism is still in place. For more information, refer to the NTCALL command in the <i>Visual SoftICE Command Reference</i> .
30-37	Primary Interrupt Controller (IRQ0-IRQ7) on older PIC-based machines 30 - HAL clock interrupt (IRQ0) on older PIC-based machines.
38-3F	Secondary Interrupt Controller (IRQ8-IRQ15) on older PIC-based machines.

On older machines using the 8259 PIC for controlling interrupts, interrupt vectors 0x30 - 0x3F are mapped by the primary and secondary interrupt controllers, so hardware interrupts for IRQ0 through IRQ15 are vectored through these IDT entries. Most machines produced today use interrupt controllers based on Intel's Advanced Programmable Interrupt Controller (APIC) specification. Such systems are not limited to 15 hardware interrupts. While the IDT itself is the same on these systems, the mapping of hardware interrupts to interrupt numbers in the IDT is not. To determine which interrupt number in the IDT the OS has assigned to a given hardware interrupt, you can use the Visual SoftICE IT/IDT command, which will read this information out of the APIC. The vector column of the IDT output will tell you which IDT entry to look at.

In many cases, these hardware interrupt vectors are not hooked, so the system assigns default stub routines for each one. As devices require the use of these hardware interrupts, the device driver requests to be connected. When the interrupt is no longer needed, the device driver requests to be disconnected.

The default stubs are named KiUnexpectedInterrupt#, where # represents the unexpected interrupt. To determine which interrupt vector is assigned to a particular stub, add 0x30 to the UnexpectedInterrupt#. For example, KiUnexpectedInterrupt2 is actually vectored through IDT vector 32 (0x30 + 2).

Drivers may install and uninstall interrupt handlers as necessary, using IoConnectInterrupt and IoDisconnectInterrupt. These routines create special thunk objects, allocated from the Non-Pageable Pool, which contain data and code to manage simultaneous use of the same interrupt handler by one or more drivers.

## GDT (Global Descriptor Table)

Windows is a flat memory architecture. Thus while it still needs to use selectors, it uses them minimally. Most Win32 applications and drivers are completely unaware that selectors even exist.

The following image shows screen output from the Visual SoftICE GDT command, which shows the selectors in the Global Descriptor Table.

```
VSI>gdt
Global Descriptor Table - Base Address: 0x8003f000, Limit: 3ff
Count: 24
```

Selector	Type	Address	Limit	DPL	Granularity	Present	
0x8	Code: Execute/Readable (accessed)		0x0	0xffffffff	0x0	Page	P
0x10	Data: Read-Write (accessed)		0x0	0xffffffff	0x0	Page	P
0x1b	Code: Execute/Readable (accessed)		0x0	0xffffffff	0x3	Page	P
0x23	Data: Read-Write (accessed)		0x0	0xffffffff	0x3	Page	P
0x28	32bit TSS (busy)		0x80042000	0x20ab	0x0	Byte	P
0x30	Data: Read-Write (accessed)		0xffffdff000	0x1fff	0x0	Page	P
0x3b	Data: Read-Write (accessed)		0x0	0xffff	0x3	Byte	P
0x43	Data: Read-Write		0x400	0xffff	0x3	Byte	P
0x50	32bit TSS (available)		0x80558700	0x68	0x0	Byte	P
0x58	32bit TSS (available)		0x80558768	0x68	0x0	Byte	P
0x60	Data: Read-Write (accessed)		0x22f30	0xffff	0x0	Byte	P
0x68	Data: Read-Write		0xb8000	0x3fff	0x0	Byte	P
0x70	Data: Read-Write		0xfffff7000	0x3ff	0x0	Byte	P
0x78	Code: Execute/Readable		0x80400000	0xffff	0x0	Byte	P
0x80	Data: Read-Write		0x80400000	0xffff	0x0	Byte	P
0x88	Data: Read-Write		0x0	0x0	0x0	Byte	P
0xa0	32bit TSS (available)		0x867ba098	0x68	0x0	Byte	P
0xe0	Code: Conforming, Execute/Readable (accessed)		0xf788e000	0xffff	0x0	Byte	P
0xe8	Data: Read-Write		0x0	0xffff	0x0	Byte	P
0xf0	Code: Execute-Only		0x804d972c	0x3b95b	0x0	Byte	P
0xf8	Data: Read-Write		0x0	0xffff	0x0	Byte	P
0x100	Data: Read-Write (accessed)		0xf789e000	0xffff	0x0	Byte	P
0x108	Data: Read-Write (accessed)		0xf789e000	0xffff	0x0	Byte	P
0x110	Data: Read-Write (accessed)		0xf789e000	0xffff	0x0	Byte	P

Note that the first four selectors address the entire 4GB linear address range. These are flat selectors that Win32 applications and drivers use. The first two selectors have a DPL of zero and are used by device drivers and system components to map system code, data, and stacks. The selectors 1B and 23 are for Win32 applications and map user level code, data, and stacks. These selectors are constant values and the Windows NT system code makes frequent references to them using their literal values.

The selector value 30h addresses the Kernel Processor Control Region and is usually mapped at a base address of 0xFFDFF000. When executing system code, this selector is stored in the FS segment register. Among its many other purposes, the Processor Control Region maintains the current kernel mode exception frame at offset 0.

Similarly, the selector value 3Bh is a user-mode selector that maps the current user thread environment block (UTEB). This selector value is stored in the FS segment register when executing user level code and has the current user-mode exception frame at offset 0. The base address of this selector varies depending on which user-mode thread is running. When a thread switch occurs, the base address of this GDT selector entry is updated to reflect the current UTEB.

## **Windows System Memory Map (on 32-bit Intel Architecture)**

Windows reserves the upper 2GB of the linear address space for system use. The address range 0x80000000 - 0xFFFFFFFF maps system components such as device drivers, system tables, system memory pools, and system data structures such as threads and processes. (It is also possible to change this behavior by adding the /3gb switch to your boot.ini and have 3 gigs available to user mode and 1GB available to the kernel.) While you cannot create an exact map of the Windows system memory space, you can categorize areas that are set aside for specific usage. The following System Memory Map diagram gives you a rough idea of where operating system information is located. Remember that a majority of these system areas could be mapped anywhere within the system address space, but are generally in the address ranges shown.

- ◆ **System Code area**

Boot drivers and the NTOSKRNL and HAL components are loaded in the System Code address space. Non-boot drivers are loaded in the NonPaged system address space near the top of the linear address space. You can use the Visual SoftICE IMAGE and IMAGEMAP commands to examine the base address and extents of boot drivers loaded in this memory area. This is also where the TSS, IDT, and GDT system data structures are mapped.

- ◆ **System View area**

Under Windows, the System View address space maps the global tables for GDI and USER objects. You can use the Visual SoftICE OBJTAB command to view information about the USER object table.

◆ System Tables Area

This region of linear memory maps process page tables and related data structures. This is one of the few areas of system memory that is not truly global, in that each process has unique page tables. When Windows executes a process context switch, the physical address of the process Page Directory is extracted from the kernel process environment block (KPEB) and loaded into the CR3 register. This causes the process page tables to be mapped in this memory area. Although the linear addresses remain the same, the physical memory used to back this area contains process-specific values. In Visual SoftICE terminology, the Page Directory is essentially an Address Context. When you use the Visual SoftICE ADDR command to change to a specific process context, you are *loading the Page Directory information for this process*.

To manage the mapping of linear memory to physical memory, Windows reserves a 4MB region of the system linear address space for Page Tables. This 4MB region represents the entire range of memory necessary to fully define a Page Directory and complete set of page tables. The need for a 4MB region can be calculated given that there is one Page Directory structure which contains entries for 1024 Page Tables. To map a 4GB linear address space, each Page Table must map a 4MB region of linear address space (4GB /1024). Each Page Table is a multiple of the CPU page size (which is 4KB under Windows), so multiplying 1024 by 4096 (the page size) yields the expected 4MB value. Thus an operating system that uses paging and a 4KB page size requires 4MB of memory to map the entire address space.

**Note:** 64-bit versions of Windows and different architecture versions use differing page directory structures and page sizes.

The diagram on the next page shows a system memory map for Windows.

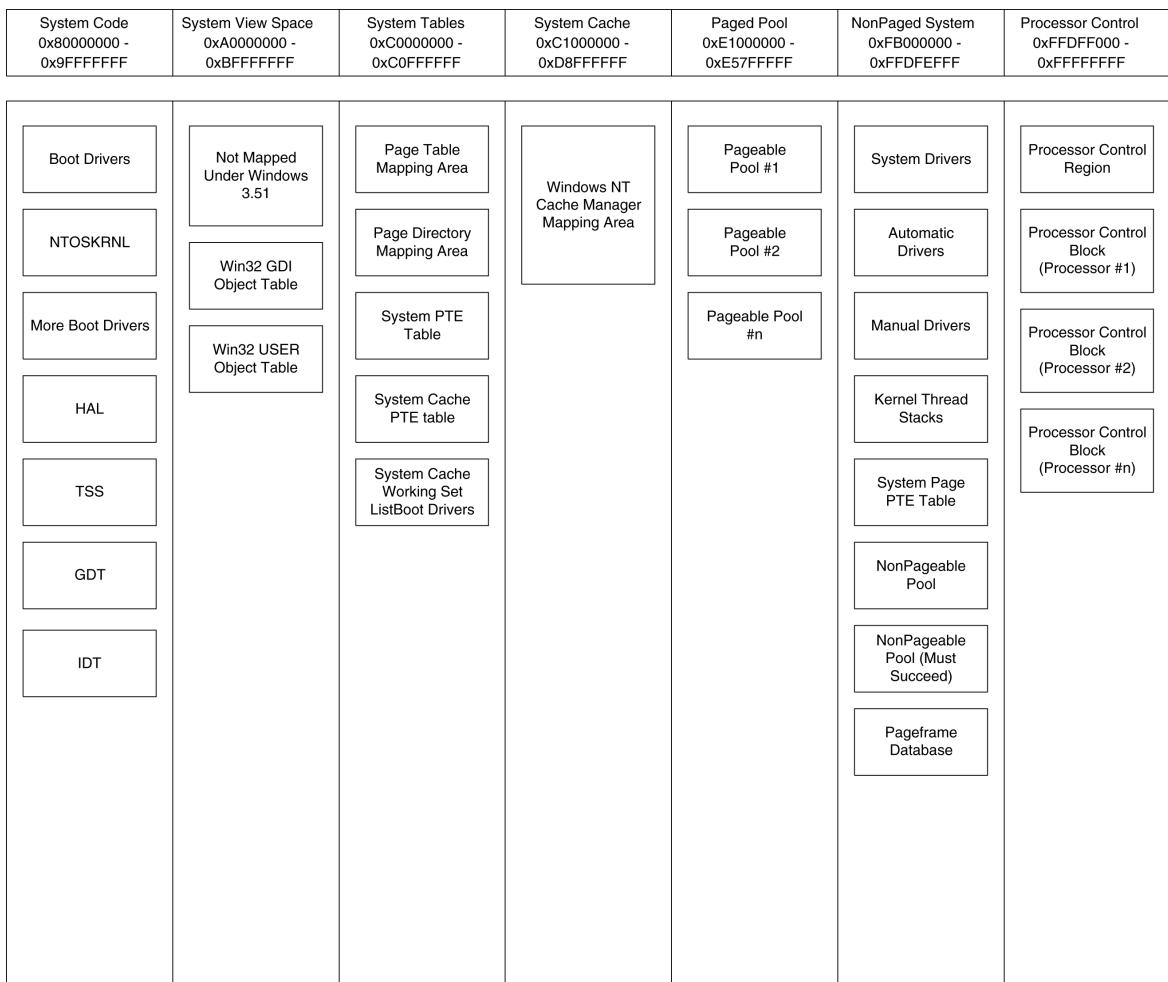


Figure 9-1. Windows System Memory Map

In this design, the Page Directory is actually performing two functions. In addition to being the Page Directory, representing 4GB, it also serves as a page table, representing 4MB in the address range of 0xC0000000 - 0xC03FFFFF. The Page Directory maps the 4MB region where the process page tables are mapped (0xC0000000-0xC03FFFFF), so the Page Directory entry that maps this area must point to itself. If you use the Visual SoftICE PAGE command, the physical address of the Page Directory displayed at the top of the command output matches the physical address for the entry that maps the 0xC0000000 - 0xC03FFFFF memory range. If you use the Visual SoftICE ADDR command to obtain the CR3 (the CR3 register contains the physical address of the Page Directory) value for the current process and supply this value as input to the Visual SoftICE PHYS command, all the linear addresses that are mapped to the physical address of the Page Directory are displayed. One of the addresses is 0xC0300000.

- ◇ Use the PAGE command in Visual SoftICE to display the top level page table info (include type, address, and page size for the platform):

```
VSI>page  
Page Size: 0x1000 bytes
```

Entry	Address	Type
PDE	0x39000	(Global)

- ◇ Then use the PAGE command passing a virtual address of interest to show the page table layout and data for this virtual address:

```
VSI>page 0xf770d36a
```

Entry	Physical	Data	Physical	Page(PPN)	Attributes
PDE	0x0	0x1031963	0x1031000 (0x1031)	P A RW S	
-PTE	0x69f236a	0x69f2121	0x69f2000 (0x69f2)	P A R S G	

- ◇ You can validate that this mapping with the PHYS command, which displays all virtual addresses that map to a physical address.

```
VSI>phys 0x69f236a  
0xf770d36a
```

## System Page Table Entries and Abstractions

The acronym, PTE, which appears in various places on the system map, stands for Page Table Entry. Each PTE describes one page of memory, including its physical address and attributes. Because Windows also runs on non-Intel platforms, and because the operating system may need to extend the types of page-level protection beyond what any particular CPU may provide, Windows virtualizes the CPU PTE with a software model of varying depth, which allows additional attributes and access control.

By overloading the meaning of an attribute bit within this abstraction, the operating system can gain control on a page fault, and examine the extended attributes of the corresponding entry to determine why the operating system requested that the fault occur. Throughout NTOSKRNL, manipulations are performed on the entry abstraction, and translated to the actual CPU PTE type. Note that the operating system also compares the entry to its corresponding CPU PTE to ensure their consistency. This effectively prevents an application or device driver from directly manipulating the page table entries.

- ◆ **Paged Pool Area:** The Paged Pool system memory area is where ntoskrnl!ExAllocatePool and its related functions allocate memory that can be paged to disk. This is in direct contrast to the Non-Paged pool area. Non-Paged pool allocations are never paged to disk and are designed for routines such as Interrupt Handlers that need high performance or need a guarantee that a piece of information is always available for use.

Windows makes extensive use of the Paged pools, as this is where most operating system objects are created. Note that the starting address and the size and number of paged pools is determined dynamically during system initialization.

Although there is one Paged Pool area, there are multiple paged pools. The number is determined during system initialization. Paged pool allocations occur with relatively high frequency and those accesses must be thread safe, so having one data structure which must be owned exclusively by one thread during memory allocation or deallocation creates a bottleneck. To avoid potential traffic jams and reduced system performance, multiple pool descriptors are created, each with its own private data structures, including an executive spinlock for thread synchronization. Thus, the more paged pools created, the more threads that can perform paged pool allocations simultaneously, increasing the throughput of the system. An important design note, in case you plan on using similar techniques in your driver or application, is that the overhead for a Paged Pool (or Non-Paged Pool) descriptor is very minimal. Thus it's practical for four or five of them to exist. However, determine that an actual bottleneck exists before creating elaborate schemes to solve a non-existent problem.

- ◆ **Non-Paged System Area:** This linear region is intended for system components and data structures that need to be present in memory at all times. This includes non-boot drivers, kernel mode thread stacks, two Non-Paged memory pools, and the Page Frame Database. Although it is contradictory to say that items in the Non-Paged System area can become not present; the truth is that they can be. Specifically, kernel thread stacks and process address spaces can be made not present, and often are.

The Non-Paged pool is similar to the Paged Pool with the exception that objects created in the Non-Paged pool are not discarded from memory for any reason. The Non-Paged pool is used to allocate key system data structures such as kernel process and thread environment blocks. There is a second Non-Paged pool used for memory allocations that *must succeed*. At system initialization, NTOSKRNL reserves a small amount of physical memory for critical allocations, and saves this memory for use by the must succeed pool. The size of an allocation from the must succeed pool must be less than one page. If the must succeed allocation cannot be satisfied, or the requested allocation size is larger than a page, the system throws a *Blue Screen*.

- ◆ **Processor Control Region:** At the high end of the system memory area is the Processor Control Region. Here, Windows maintains Processor Control Block (PRCB) data structures for each processor within the system and a global data structure, the Processor Control Region that reflects the current state of the system. The Processor Control Region (PCR) contains key pieces of information about the current state of the system, such as the currently running kernel thread; the current interrupt request level (IRQL); the current exception frame; base addresses of the IDT, TSS, and GDT; and kernel thread stack pointers. Small portions of the PCR and PRCB data structures are documented in NTDDK.H.

In many cases, device driver writers need to know the current IRQL at which they are executing. Although you could look inside the PCR data structure at offset 0x24, it is simpler to use the Visual SoftICE intrinsic function, *IRQL*, as follows:

```
? IRQL
<uchar> = 0x2, 2
```

The most common piece of data accessed from the PRCB is the current kernel thread pointer. 32-bit x86 code that accesses the current thread is usually of the form:

```
mov reg, FS:[124].
```

Remember that while executing in system mode, the FS register is set to a GDT selector whose base address points to the beginning of the PCR.

For more extensive information on the current thread use the following commands:

VSI>thread thread

TID	Krnl TEB	StackBtm	StkTop	StackPtr	User TEB	Process(Id)
0071	FF0889E0	FC42A000	FC430000	FC42FE5C	7FFDE000	WINWORD(6A)

The current process is not stored as part of the PCR or PRCB. Windows references the current process through the current thread. Code such as the following obtains the current process pointer:

```
mov    eax,    FS:[124]      ; get the current thread (KTEB)
mov    esi,    [eax+40h]     ; get the threads process pointer (KPEB)
```

# Win32 Subsystem

## Inside CSRSS

The Win32 subsystem server process CSRSS implements the Win32 API. The Win32 API provides many different types of service, including functionality traditionally attributed to the original Windows components KERNEL, USER, and GDI. Although these standard modules exist in the form of 32-bit DLLs under Windows NT, most of the core functionality is actually implemented in WINSRV.DLL and WIN32K.SYS within the CSRSS process. Calls that are traditionally associated with one of the standard Windows components are typically implemented as stubs that call other modules, for example, NTDLL.DLL, or use inter-process communication to CSRSS for servicing.

Most USER and GDI API calls have their functionality implemented in USER32 and GDI32 modules that are loaded into your application's address space. This allows the most common services to execute as simple function calls. The WIN32K.SYS module allows USER and GDI services to execute more efficiently through a simple transition from user to system mode. Depending on which processor and OS you are using, this can occur through a SYSENTER instruction or through an int 2e. Having WIN32K.SYS as a device driver that provides application services allows Windows to maintain a high level of encapsulation and robustness, while providing a much more efficient pseudo client-server service architecture.

Although CSRSS executes as a separate process, it still has a big impact on the address space of every Win32 application. If you use the Visual SoftICE HEAP command on your process, you will notice at least two heaps that your application did not specifically create, but were created on its behalf. The first is the default process heap that was created during process initialization. The second is a heap specifically created by CSRSS. There may be other heaps in your application address space that were not created by your process. These heaps are generally located very high in the user-mode address space and appear if you use the Visual SoftICE ADDRESSMAP command, but do not appear in the output of the HEAP command. The reason for this is quite simple: for each user-mode process, a list of process heaps is maintained and the Visual SoftICE HEAP command uses this list to enumerate the heaps for a process. If the heap was not created by or on behalf of your application, it does not appear in the process heap list. The Visual SoftICE ADDRESSMAP command traverses the user-mode address space for your application identifying regions of memory that are mapped.

Heaps that exist in the process address space, but that are not enumerated in the process heap list, were mapped into the process address space by another process. In most cases, this mapping is done by CSRSS. During subsystem initialization, CSRSS creates a heap at a well-known base address. When new processes are created, this heap is mapped into their address spaces at the same well-known base address. Theoretically, mapping the heap of one process at the same base address of another process allows both processes to use that heap. In practice, there are issues that might prevent this from working under all circumstances – synchronization being one such issue. Note that under newer versions of Windows, more than one heap may be mapped into the process address space, and those heaps may be mapped at different base addresses in different processes. Also, new versions of the operating system use heaps that are created in the system address space, and these heaps are sometimes mapped into the user address space. Windows allows the creation of heaps within the system address space using APIs exported from NTOSKRNL. These APIs are similar to the same APIs exported from the user-mode module, NTDLL.DLL.

## *USER and GDI Objects*

The protected Win32 subsystem process, CSRSS, provides a majority of the traditional USER functionality. APIs and data structures provided by the WINSRV.DLL and WIN32K.SYS modules manage window classes and window data structures, as well as many other USER data types.

The following USER object types exist. Object type IDs are listed in parentheses.

---

FREE (0)	Object Entry is unused/invalid.
HWND (1)	Window Objects.
MENU (2)	Windows MENU object.
ICON/CURSOR (3)	Windows ICON or CURSOR object.
DEFERWINDOWPOS (4)	Object returned by the BeginDeferWindowPosition API.
HOOK (5)	Windows Hook thunk.
THREADINFO (6)	CSRSS Client Thread Instance Data.
CLIPBOARD FORMAT (7)	Registered Clipboard Formats.
CPD (8)	Call Procedure Data thunk.

---

---

<b>ACCELERATOR (9)</b>	Accelerator Table Object.
<b>WINDOW STATION (0xD)</b>	
<b>KEYBOARD LAYOUT (0xE)</b>	Object to describe a keyboard layout.
<b>DDEOBJECT (0xA)</b>	DDE Objects such as strings.

---

Rather than maintaining per-process data structures for USER and GDI object types, CSRSS maintains a master handle table for all processes. The USER and GDI objects are segregated into two different tables that have the same basic structure and semantics. WINSRV provides distinct Handle Manager APIs for managing the two different tables. You can identify the handle manager API names by the HM prefix in front of the API name, and the GDI specific routines by the “g” appended to this prefix. The routine HAllocObject creates USER object types, while HmgAlloc is a GDI object type API that creates GDI object types.

The management of USER and GDI handles is relatively straightforward, and its design is a good example of how to implement basic management of abstract object types. Specifically, this API uses a simple, but robust, technique for creating unique handles and managing reference counts. The design also provides for handle opaqueness which prevents applications, including USER32 and CSRSS, from directly manipulating the objects outside the handle manager. Preventing clients, including itself, from directly manipulating the object data allows the handle manager to ensure that reference counts and synchronization issues are managed correctly.

The master object tables maintained by the Handle Manager are growable arrays of fixed size entries. The following table lists the fields for an object table. Only columns with **bold** field headers are part of the entry. The columns with *italicized* headers are for illustration only.

<i>Entry</i>	<b>Object Pointer (DWORD)</b>	<b>Owner (DWORD)</b>	<b>Type (BYTE)</b>	<b>Flags (BYTE)</b>	<b>Instance Count (WORD)</b>	<i>Handle Value</i>
0	NULL	NULL	FREE (0)	00	0001	00010000
1	HEAP *	HEAP *	DESKTOP (0C)	00	0001	00010001
2	HEAP *	HEAP *	HWND (04)	01	0003	00030002

The Object Pointer field points to the actual object data. This pointer is generally from one of the CSRSS heaps or the Paged Pool. The type field is the enumeration for the object type. The Instance Count field creates unique handles. The Flags field is used by the Handle Manager to note special conditions, such as when a thread locks an object for exclusive use.

## How Handle Values Are Created

Initially, all object table Instance counts are set to 1. When a new Object Entry is allocated, the Instance Count is combined with the table index to create a unique handle value. When references are made to an object, the table entry portion of the handle is extracted and used to index into the table. As part of the handle validation, the instance count is extracted from the table entry and compared to the handle being validated. If the instance count does not match the table entry instance count, the handle is bogus. The following example illustrates these concepts:

To create an object handle from an object table entry:

```
ObjectHandle = TableEntryIndex + (InstanceCount << 16);
```

To validate an object handle:

```
ObjectTable[LOWORD(handle)].InstanceCount ==  
HIWORD(handle);
```

When an object is destroyed, all fields are reinitialized to zero and the current Instance Count for that entry is incremented by one. Thus, when the object table entry is reused, it generates a different handle value for the new object.

**Note:** The actual object type is not part of the object handle value. This means that given an object handle, an application cannot directly determine its type. It is necessary to dereference the object table entry to obtain the object type.

This technique for creating unique handle values is simple and efficient, and makes validation trivial. Imagine the case where a process creates a window and obtains a handle to that window. During subsequent program execution, the process destroys the window but retains the handle value. If the process uses the handle after the window is destroyed, the handle value is invalid and the type it points to has an object type of FREE. This condition is caught, and the program is not be able to use the handle successfully. In the meantime, if another process creates a new object, it is likely that the entry originally for the now destroyed window will be reused. If the original program uses the invalid window handle, the handle instance counts no longer match, and the validation fails.

Object tables are not process specific, so USER and GDI object handles values are not unique to a specific process. HWND handles are unique across the entire Win32 subsystem. One process never has an HWND handle value that is duplicated in any other process.

## USER Object Table

Use the Visual SoftICE OBJTAB command to display all the object entries within the USER object table. The OBJTAB command is relatively flexible, allowing a handle or table entry index to be specified. It also supports the display of objects by type using abbreviations for the object type names. To see a list of object type names that the OBJTAB command can use, specify the -H option on the OBJTAB command line.

The Object field can reference the object specific data for an object table entry. All objects have a generic header that is maintained by the object manager, which includes the object handle value and a thread reference count. Most object types also contain a pointer to a desktop object and/or a pointer to its owner.

## Monitoring USER Object Creation

If you do a considerable amount of Win32 application development, the HAllocObject API is a convenient place to monitor creation of object types such as windows. Use the Visual SoftICE MACRO command to create a breakpoint template that can trap creation of specific object types as follows:

```
MACRO obx = "bpw winsrv!HAllocObject if (esp->c == %1)"
```

The HAllocObject API is implemented in WINSRV.DLL and the object type being created is the third parameter, which translates to Dword ptr esp [ 0Ch ]. The syntax “esp->c” dereferences the requested object type, and is equivalent to \*(esp+c). The “%1” portion of the conditional expression is a place holder for argument replacement. When you execute the OBX macro, the argument provided is inserted into the macro stream at the “%1”:

```
OBX 1 -> bpx winsrv!HAllocObject if (esp->c == 1)
```

When this breakpoint is instantiated, it traps all calls to HAllocObject that creates window object types.

## 32-bit x86 Windows Process Address Space

The address space for a user-mode process is mapped into the lower 2GB of linear memory at addresses 0x00000000 - 0x7FFFFFFF. The upper 2GB of linear memory is reserved for the operating system kernel and device drivers.

In general, each Win32 application's process address space has the following regions of linear memory mapped for the corresponding purpose.

Table 9-4. Process Address Space

Linear Address Range	Purpose
0x00000000 - 0x0000FFFF	Protected region. Useful for detecting NULL pointer writes.
0x00010000	Default load address for Win32 processes.
0x70000000 - 0x78000000	Typical range for Win32 subsystem DLLs to be loaded.
0x7FFB0000 - 0x7FFD3FFF	ANSI and OEM code pages. Unicode translation table(s).
0x7FFDE000 - 0x7FFDEFFF	Primary user-mode thread environment block.
0x7FFDF000 - 0x7FFDFFFF	User-mode process environment block (UPEB).
0x7FFE0000 - 0x7FFE0FFF	Message queue region.
0x7FFF0000 - 0x7FFFFFFF	Protected region.

Under Windows, the lowest and highest 64KB regions in the user-mode address space are reserved and are never mapped to physical memory. The 64KB at the bottom of the linear address space is designed to help catch writes through NULL pointers.

The default load address for processes under Windows is 0x10000. Processes often change their load address to a different base address. Use the linker or the REBASE utility to set the default load address of a DLL or EXE.

The linear range at 0x70000000 is an approximation of the area where Win32 subsystem modules load. Use the Visual SoftICE IMAGE, IMAGEMAP, or ADDRESSMAP commands to obtain information on modules loaded in this range.

The user process environment block is always mapped at 0x7FFDF000, while the process's primary user-mode thread environment block is one page below that at 0x7FFDE000. As a process creates other worker threads, they are mapped on page boundaries at the current, highest unused linear address.

The following use of the Visual SoftICE THREAD command shows how each subsequent thread is placed one page below the previous thread:

```
VSI>thread winword

TID      Krnl TEB      StackBtm   StkTop      StackPtr    User TEB      Process( Id )
006B    FFA7FDA0    FEAD7000   FEADB000   FEADAE64  7FFDE000  WINWORD( 83 )
007C    FF0A0AE0    FEC2A000   FEC2D000   FEC2CE18  7FFDD000  WINWORD( 83 )
009C    FF04E4E0    FC8F9000   FC8FC000   FC8FBE18  7FFDC000  WINWORD( 83 )
```

To find out more about the user-mode address space of a process, use the Visual SoftICE ADDRESSMAP command. The ADDRESSMAP command provides a high-level view of the linear regions that were reserved and/or committed. From its output you see the process heaps, modules, and memory-mapped files, as well as the thread stacks and thread environment blocks.

## Heap API

### Heap Architecture

Every user-mode application directly or indirectly uses the Heap API routines, which are exported from KERNEL32 and NTDLL. Heaps are designed to manage large areas of linear memory and sub-allocate smaller memory blocks from within this region. The core implementation of the Heap API routine is contained within NTDLL, but some of the application interfaces such as HeapCreate and HeapValidate are exported from KERNEL32. For some API routines, such as HeapFree, there is no code implementation within KERNEL32, so they are fixed by the loader to point at the actual implementation within NTDLL.

**Note:** The technique of fixing an export in one module to the export of another module is called 'Snapping.'

Although the Heap API routines used by applications are relatively straightforward and designed for ease of use, the implementation and data structures underneath are quite sophisticated. The management of heap memory has come quite a long way from the standard C run-time library routines malloc() and free().

Specifically, the Heap API handles allocations of large, non-contiguous regions of linear memory, which are used for sub-allocation and to optimize coalescing of adjacent blocks of free memory. The Heap API also performs fast look-ups of best-fit block sizes to satisfy allocation requests, provides thread-safe synchronization, and supplies extensive heap information and debugging support.

The primary heap data structure is large, at approximately 1400 bytes, for a free build and twice that for a checked build. This does not include the size of other data structures that help manage linear address regions. A vast majority of this overhead is attributed to 128 doubly-linked list nodes that manage free block chains. Small blocks, less than 1KB in size, are stored with other blocks of the same size in doubly linked lists. This makes finding a best-fit block very fast. Blocks larger than 1KB are stored in one sorted, doubly-linked list. This is an obvious example of a time versus space trade-off, which could be important to the performance of your application.

To understand the design and implementation of the Heap API, it is important to realize that a Win32 heap is not necessarily composed of one section of contiguous linear memory. For growable heaps, it might be necessary to allocate many linear regions, using VirtualAlloc, which will generally be non-contiguous. Special data structures track all the linear address regions that comprise the heap. These data structures are called Heap Segments. Another important aspect of the Heap API design is the use of the two-stage process of reserving and committing virtual memory that is provided by the VirtualAlloc and related APIs. Managing which memory is reserved and which memory is committed requires special data structures known as Uncommitted Range Tables, or UCRs for short.

The Ntdll!RtlCreateHeap() API implements heap creation and initialization. This routine allocates the initial virtual region where the heap resides and builds the appropriate data structures within the heap. The heap data structure and Heap Segment #1 reside within the initial page of the virtual memory that is initially allocated for the heap. Heap Segment #1 resides just beyond the heap header. Heap Segment #1 is initialized to manage the initial virtual memory allocated for the heap. Any committed memory beyond Heap Segment #1 is immediately available for allocation through HeapAlloc(). If any memory within Heap Segment #1 is reserved, a UCR table entry is used to track the uncommitted range.

**Note:** Kernel32!HeapAlloc() is 'Snapped' to Ntdll!RtlAllocateHeap.

Besides the 128 free lists mentioned above, the heap header data structure contains 8 UCR table entries, which should be sufficient for small heaps, although as many UCRs as are necessary can be created. It also contains a table for sixteen (16) Heap Segment pointers. A heap can never have more than sixteen segments, as no provision is made for allocating extra segments entries. If the heap requires thread synchronization, the heap header appends a critical section data structure to the end of the fixed size portion of the heap header preceding Heap Segment #1.

The diagram on the next page is a high-level illustration of how a typical heap is constructed, and how the most important pieces relate to each other.

The left side of the diagram represents a region of virtual memory that is allocated for the heap. The heap header appears at the beginning of the allocated memory and is followed by Heap Segment #1. The first entry within the heap's segment table points to this data structure. Committed memory immediately follows Heap Segment #1. This memory is initially marked as a free block. When an allocation request is made, assuming this block of memory is large enough, a portion is used to satisfy the allocation and the remainder continues to be marked as a free block. Beyond the committed region is an area of memory that is reserved for future use. When an allocation request requires more memory than is currently committed, a portion of this area is committed to satisfy the request.

Heap Segment #1 tracks the virtual memory region initially allocated for the heap. The starting address for the heap segment equals to the base address of the heap and the end range points to the end of the allocated memory. A portion of the heap in the diagram is in a reserved state, that is, it has not been committed, so the heap segment uses an available UCR entry to track the area. When memory must be committed to satisfy an allocation request, all UCR entries maintained by a particular segment are examined to determine if the size of the uncommitted range is large enough to satisfy the allocation. To increase performance, the heap segment tracks the largest available UCR range and the total number of uncommitted pages within the virtual memory region of the heap segment.

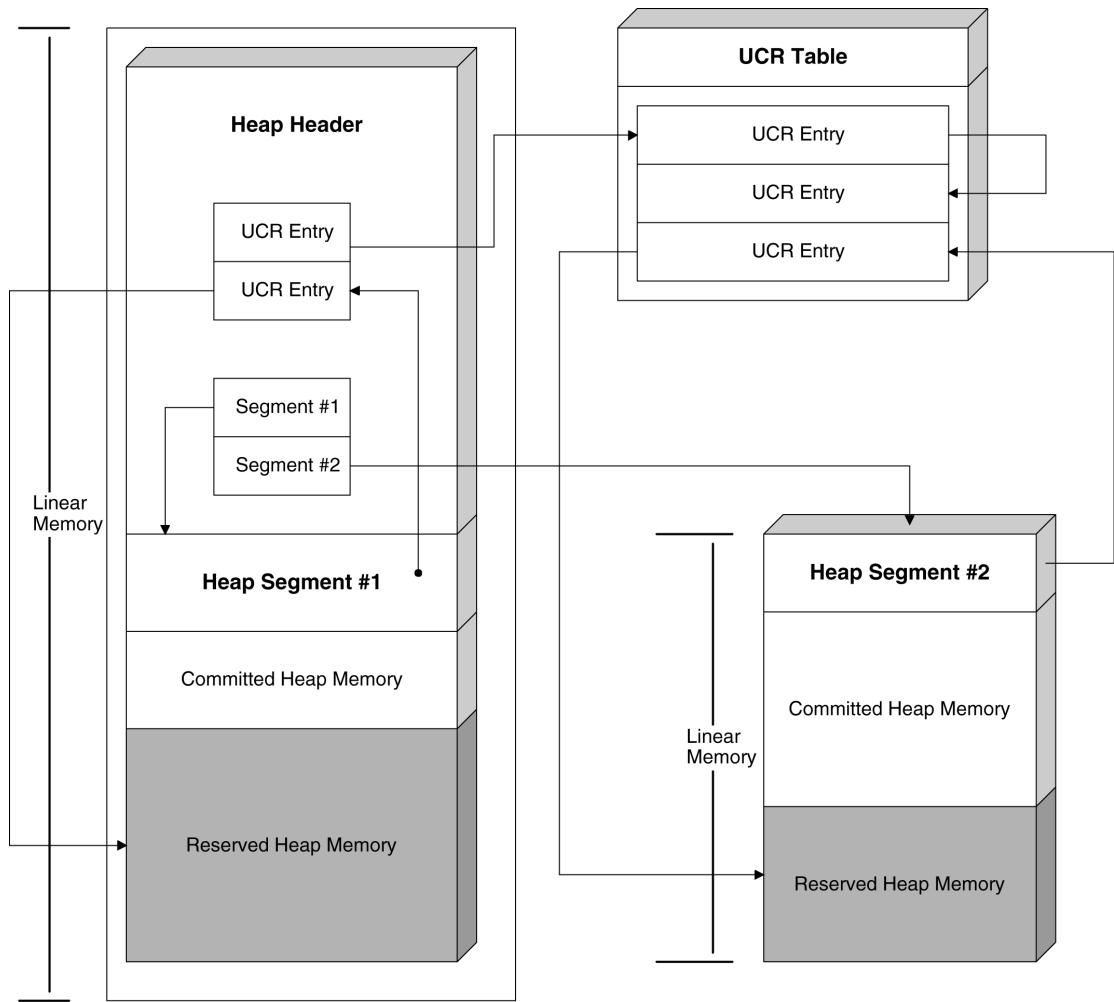


Figure 9-2. Typical Heap Construction

On the right side of the diagram, a second area of virtual memory was allocated and is managed by Heap Segment #2. Additional heap segments are created when an allocation request exceeds the size of the largest uncommitted range within the existing segment. This is only true if the size of the requested allocation is less than the heap's VMThreshold. When the requested allocation size exceeds the VMThreshold, the heap block is directly allocated through VirtualAlloc and a new heap segment is not created.

As mentioned previously, a small number of UCR entries are provided within the heap header. For illustration purposes, this diagram shows a UCR TABLE entry that was allocated specifically to increase the number of UCR entries that are available. The need to create an extra UCR table is generally rare, and is usually a sign that a large number of segments were created or that the heap segments are fragmented.

Fragmentation of virtual memory can occur when the Heap API begins decommitting memory during the coalescing of free blocks. Decommitting memory is the term used to describe reverting memory from a committed state to a reserved or uncommitted state. When a free block spans more than one physical page, that page becomes a candidate for being decommitted. If certain decommit threshold values are satisfied, the Heap manager begins decommitting free pages. When those pages are not contiguous with an existing uncommitted range, a new UCR entry must be used to track the range.

The following examples use the Visual SoftICE HEAP command to examine the default heap for the Explorer process.

- 1 Use the -S option of the HEAP32 command to display segment information for the default heap:

Base	Id	Cmmt/Psnt/Rsvd	Segments	Flags	Process
00140000	01	001C/0018/00E4	1	00000002	Explorer
01		00140000-00240000	001C/0018/00E4	E4000	

:heap32 -s 140000

Diagram annotations:

- A line points from the "Segments" column to the value "1".
- A line points from the "Flags" column to the value "00000002".
- A line points from the "Process" column to the value "Explorer".
- A line points from the "Largest" label to the memory range "00140000-00240000".
- A line points from the "Heap segment count" label to the "Segments" column.
- A line points from the "Heap segment memory range" label to the memory range "00140000-00240000".

- 2** Use the -X option of the HEAP command to display extended information about the default heap:

```
:heap32 -x 140000
```

Extended Heap Summary for heap 00140000 in Explorer					
Heap Base:	140000	Heap Id:	1	Process:	Explorer
Total Free:	6238	Alignment:	8	Log Mask:	10000
Seg Reserve:	100000	Seg Commit:	2000		
Committed:	112k	Present:	96k	Reserved:	912k
Flags:	GROWABLE				
DeCommit:	1000	Total DeC:	10000	VM Alloc:	7F000
Default size of a heap segment		Default size for commits		VM threshold	

- 3** Use the -B option of the HEAP command to display the base addresses of heap blocks within the default heap:

```
:heap32 -b 140000
```

Base	Type	Size	Seg#	Flags
00140000	HEAP	580	01	
00140580	SEGMENT	38	01	
001405B8	ALLOC	30	01	

In the above output, you can see how the heap header is followed by Heap Segment #1 and that the first allocated block is just beyond the Heap Segment data structure.

## Managing Heap Blocks

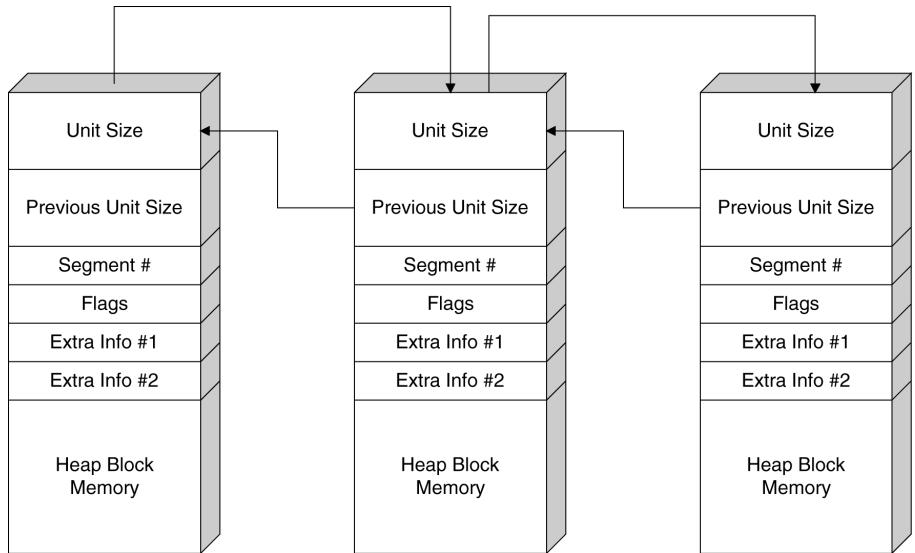
As discussed in the preceding section, the Heap API uses the Win32 Virtual Memory API routines to allocate large regions of the linear address space and uses heap segments to manage committed and uncommitted ranges. The actual sub-allocation engine that manages the allocation and deallocation of the memory blocks used by your application is built on top of this functionality. To track allocated and free blocks, the Heap API creates a header for each block.

The diagram on the next page illustrates how the heap manager tracks blocks of *contiguous* memory. The heap manager also tracks non-contiguous free blocks in doubly-linked lists, but the node pointers for the next and previous links are not stored in the block header. Instead, the heap manager uses the first two Dwords within the heap block memory area.

As shown in Figure 12-3, each block stores its unit size as well as the unit size of the previous block. The unit size represents the number of heap units occupied by the heap block. The previous unit size is the number of heap units occupied by the previous heap block. Using these two values, the heap manager is able to walk contiguous heap blocks.

Heap units represent the base granularity of allocations made from a heap. The size of an allocation request is rounded upwards as necessary, so that it is an even multiple of this granularity. Rather than using a granularity of 1 byte, the heap manager uses a granularity of 8 bytes. This means that all allocations are an even multiple of 8 bytes, and that allocation sizes can be converted to units by round up and dividing by 8. For example, if a process requests an allocation of 32 bytes, the number of units is  $32 / 8 = 4$ . If the allocation request was 34 bytes, the allocation size is rounded upward to an even multiple of 8. In this example, the 34 bytes requested would be rounded to an allocation of 40 bytes, or 5 units. The process requesting the allocation is unaware of any rounding to satisfy unit granularity and proceeds as if the allocation request of 34 bytes was actually 34 bytes.

By using a unit size of 8, the types of allocation made by most applications can be recorded using one word value with the restriction that the maximum size of a heap block, in units, is the largest unsigned short or 0xFFFF. This makes the theoretical maximum size of a heap block in bytes,  $0xFFFF * 8$ , or 524,280 bytes. (This limitation is documented in the Win32 HeapAlloc API documentation.) Does that mean that a program cannot allocate a heap block greater than 512k? Well, yes and no. A heap block larger than 512k cannot be allocated, but there is nothing to prevent the Heap API from using VirtualAlloc to allocate a region of linear memory to satisfy the request. This is exactly what the heap manager does if the size of the requested allocation exceeds the heaps VMThreshold. The value of VMThreshold is stored in the heap header and by default is 520,192 bytes (or 0xFE000 units). When the heap manager allocates a large heap block using VirtualAlloc, the resulting structure is referred to as a Virtually Allocated Block (VAB).



**Figure 9-3.** Tracking Blocks of Contiguous Memory

The heap manager walks contiguous heap blocks by converting the current heap block's unit size into bytes and adding that to the heap block's base address. The address of the previous heap block is calculated in a similar manner, converting the unit size of the previous block to bytes and subtracting it from the heap block's base address. The heap manager walks contiguous heap blocks during coalescing free blocks, sub-allocating a smaller block from a larger free block, and when validating a heap or heap entry.

Unit sizes are important for free block list management as the array of 128 doubly-linked lists inside the heap header track free blocks by unit size. Free blocks that have a unit size in the range from 1 to 127 are stored in the free list at the corresponding array index. Thus, all free blocks of unit size 32 are stored in `Heap->FreeLists[32]`. Because it is not possible to have a heap block that is 0 units, the free list at array index zero stores all heap blocks that are larger than 127 units; these entries are sorted by size in ascending order. Because a majority of allocations made by a process are less than 128 units (1024 bytes or 1K), this is a fast way to find an exact or best fit block to satisfy an allocation. Blocks of 128 units or greater are allocated much less frequently, so the overhead of doing a linear search of one free list does not have a large impact on the overall performance of most applications.

The flags field within the heap block header denotes special attributes of the block. One bit is used to mark a block as allocated versus free.

Another is used if it is a VAB. Another is used to mark the last block within a committed region. The last block within a committed region is referred to as a sentinel block, and indicates that no more contiguous blocks follow. Using this flag is much faster than determining if a heap block address is valid by walking the heap segment's UCR chain. Another flag is used to mark a block for free or busy-tail checking. When a process is debugged, the heap manager marks the block in certain ways. Thus, when an allocated block is released or a free block is reallocated, the heap manager can determine if the heap block was overwritten in any way.

The extra info fields of the heap block header have different usage depending on whether the block is allocated or free. In an allocated block, the first field records the number of extra bytes that were allocated to satisfy granularity or alignment requirements. The second field is a pseudo-tag. Heap tags and pseudo tags are beyond the scope of this discussion.

For a free block, the extra info fields hold byte and bit-mask values that access a free-list-in-use bit-field maintained within the heap header. This bit-field provides quicker lookups when a small block needs to be allocated. Each bit within the bit-field represents one of the 127 small block free lists, and if the corresponding bit is set, that free list contains one or more free entries. A zero bit means that a free entry of that size is not available and a larger block will need to be sub-allocated from. The first extra info field holds the byte index into the bit-field array. The heap block memory array is also different depending on the allocated state of the free block. For allocated blocks, this is the actual memory used by your application. For free blocks, the first two Dwords (1 unit) are used as next and previous pointers that link free blocks together in a doubly-linked list. If the process that allocated the heap block is being debugged, an allocated heap block also contains a busy-tail signature at the end of the block. Free blocks are marked with a special tag that can detect if a stray pointer writes into the heap memory area, or the process continues to use the block after it was deallocated.

The following diagram shows the basic architecture of an allocated heap block.

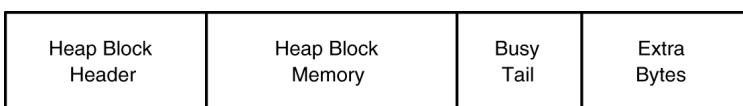


Figure 9-4. Basic Architecture of an Allocated Heap Block

The portion labeled *Extra Bytes* is memory that was needed to satisfy the heap unit size or heap alignment requirements. This memory area should not be used by the allocating process, but the heap manager does not directly protect this area from being overwritten. The busy-tail signature appears just beyond the end of the memory allocated for use by the process. If an application writes beyond the size of the area requested, this signature is destroyed and the heap manager signals the debugger with a debug message and an INT 3. It is possible for a process to write into the extra bytes area without disturbing the busy-tail signature. In this case, the overwrite is not caught. The Heap API provides an option for initializing heap memory to zero upon allocation. If this option is not specified when debugging, the heap manager fills the allocated memory block with a special signature. You can use this signature to determine if the memory block was properly initialized in your code.

The following diagram shows the basic architecture of a free heap block.

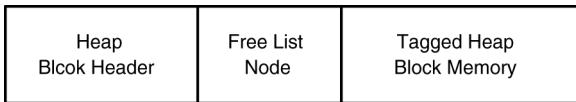


Figure 9-5. Basic Architecture of a Free Heap Block

When a block is deallocated and the process is being debugged, the heap manager writes a special signature into the heap memory area. When the block is allocated at some point in the future, the heap manager checks that the tag bytes are intact. If any of the bytes was changed, the heap manager outputs a debug message and executes an INT 3 instruction. This is a good thing if the debugger you are using traps the INT 3, but most debuggers ignore this debug-break because it was not set by the debugger. As an aside, having the Free List Node pointers at the beginning of the memory block is somewhat flawed, because a program that continues to use a free block is more likely to overwrite data at the beginning of the block than data at the end. Because these pointers are crucial to navigating the heap, an invalid pointer eventually causes an exception. When this exception occurs, it can be quite difficult to track this overwrite back to the original free block.

The following two examples show how to use the Visual SoftICE HEAP command to aid in monitoring and debugging Win32 heap issues.

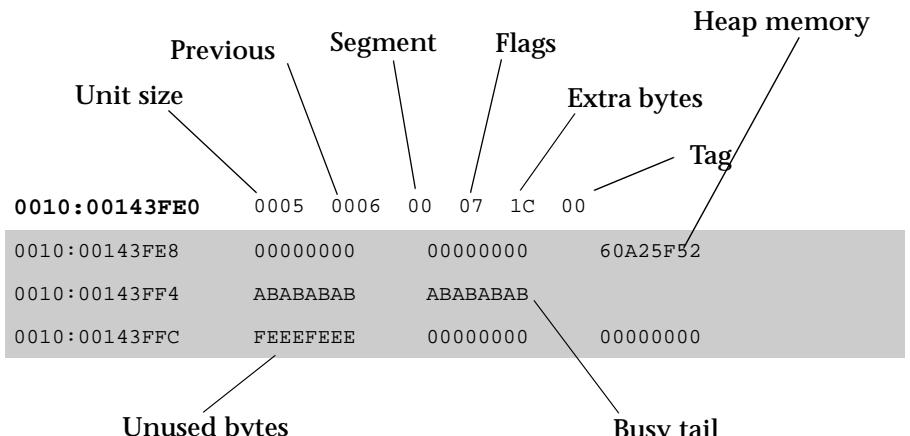
The first example uses the HEAP command to walk all the entries for the heap based at 0x140000. The -B option of the HEAP command causes the base address and size information to display as the heap manager would view the information. Without the -B option, the HEAP command shows base addresses and sizes as viewed by the application that allocated the memory. The output is abbreviated for clarity and the two heap blocks that appear in bold type are used to examine the heap block header in the second example.

```
:HEAP32 -b 140000

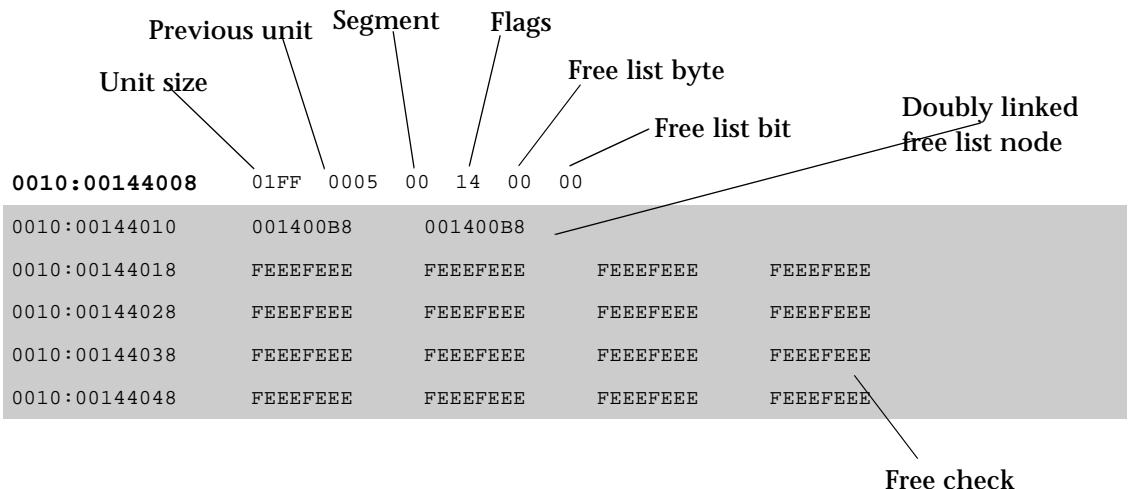
Base      Type       Size    Seg#   Flags
00140000  HEAP       580     01
00140580  SEGMENT    38      01      TAGGED | BUSYTAIL
001405B8  ALLOC      40      01
...
00143FE0  ALLOC      28      01      TAGGED | BUSYTAIL
00144008  FREE       FF8     01      FREECHECK | SENTINEL
```

To examine the contents of an allocated heap block and a free block, the second example dumps memory at the base address of the heap block at 0x143FE0. Enough memory is dumped to show the subsequent block, which is a free block at address 0x144008.

- 1 The heap block header fields from the memory dump at address 0x143FE0 are identified with call-outs. This heap block is 5 units in size (40 bytes) and 0x1C bytes of that size is overhead for the heap block header (1 unit), busy-tail (1 unit), unit alignment (1 Dword), and an extra unit left over from a previous allocation.



The heap block immediately following this is a free block that begins at address 0x144008. This block is 0x1FF units and the size of the previous block is 5 units. For free blocks 1KB or larger (80+ units), the Free List byte position and bit-mask values are not used and are zero. The flag for this heap block indicates that it is a sentinel (bit 4, or 0x10).



Immediately following the heap header is the location where the heap manager has placed a doubly-linked list node for tracking free blocks. The pointer values for the next and previous fields of the node are both 0x1400B8. After the free list node, the heap manager tagged all the blocks memory with a special signature that is validated the next time the block is allocated, coalesced with another block, or a heap validation is performed.

# Appendix A

## Troubleshooting Visual SoftICE



### Troubleshooting

If you encounter the following problems, try the corresponding solutions. If you encounter further difficulties, contact the Technical Support Center.

Problem	Solution
In Visual SoftICE, I want to connect to the target machine but I do not see it in the list of network target machines.	<p>There may be several reasons for this. First, you may not have installed the target portion of Visual SoftICE on that target machine.</p>
	<p>If you are confident that you have properly installed the Visual SoftICE target components on the target machine, then go to the configuration tool and make certain that you have properly configured a connection.</p>
	<p>If you have a properly configured connection on that target machine, and it is active, then you will need to do more extensive troubleshooting. Contact the Technical Support Center with all applicable details.</p>
I would like to set a breakpoint on a known symbol, but when set, it does not appear to be on the correct address. This can even crash the target.	<p>If breakpoints that are set on symbolic addresses fail or are located at the wrong place in memory, it is generally due to a symbol problem (having the wrong symbols loaded, or not having sufficient data available to fixup an address in an image). Certain images are optimized after compilation and linking, and the PDB may provide only part of the location information, where the image (PE) provides fixups and relocation data. Therefore, it is always important to have copies of the image you are debugging on the master.</p>
	<p>For a complete discussion of this issue and how to fix it, including examples, refer <i>Troubleshooting Visual SoftICE</i> in the online help.</p>
I want to set a break point on something in my source, but when I open the source code and attempt to set the break point, it fails to set.	<p>This is because the source is not mapped where Visual SoftICE can determine the location of the image file to break on.</p>
	<p>To fix this, you can either pre-load symbols into Visual SoftICE using the ADDSYM command, or you can set the breakpoint on the loading of that module (using the BLOAD command) and have them found automatically. Once the module is in memory, you can set breakpoints on other locations.</p>
I am noticing that the values in the Command page do not appear to be correct. For example, I entered the CPU command and Visual SoftICE returned the value for my processor speed as 585.	<p>Most values returned in the Command page are displayed in hexadecimal format by default. Try converting the value you receive to decimal, and reinterpret the results. The 585 processor speed in your example would actually be 1400 decimal.</p>
	<p>Use the SET RADIX command to configure the displayed and input radix to whatever format you are expecting to see.</p>

Problem	Solution
After installing Visual SoftICE on my target machine, it will no longer boot. It reaches a specific part of the boot sequence and then goes no further.	Try connecting with the master and check the status of the target to see if it has reached a breakpoint. A major cause of this behavior is the “stop on boot” setting. Once the target is up and running, use DSConfig to check the “stop on boot” setting for the target.
I want to connect to the target machine, but do not want to wait for it to finish booting. How do I accomplish this?	Use the WCONNECT command.
In the Command page, I receive far too much information regarding the steps and other trivialities I do not care about. How do I stop or decrease these messages?	Use the SET command by itself to review your settings. You probably have a more detailed message level reporting configured. Use the SET MSGLEVEL command to decrease or turn off message reporting in the Command page.
I'm trying to stack-walk on an IA64/x64 machine, and I keep getting a message that <b>Unwind Information is Unavailable</b> .	<p>This is a very common issue when trying to stack-walk on an IA64/x64 machine. This results from one of three problems: Either no pdb file is loaded for the specific image, the unwind data section of the image is paged out, or you set your EXEPATH and need to reload the table.</p> <p>First, make certain that the correct pdb file is loaded for the image. If the error persists, try copying the image file into a folder, and then set your executable path (EXEPATH) to include that folder. If you set your EXEPATH and still cannot get unwind information, reload the table.</p>
After installing the Visual SoftICE network debugging, I lost my Ethernet connection.	<p>This may not be a problem. It may be just a misunderstanding of the Visual SoftICE network debugging. Visual SoftICE uses a dedicated network that becomes the connection for debugging only. So if your connection to the debugger is working, but connection to the internet is not, then this is normal.</p> <p>We suggest that the target machine have two network cards: one PCI network card for normal network operation, and another PCI network card for the Visual SoftICE debugger. The only other alternative is to use the VNIC driver in addition to the installed connection.</p>



# Appendix B

## Kernel Debugger Extensions



- ◆ Debugger Extension Overview and VSI Support
- ◆ Controlling Debugger Extension DLLs
- ◆ Using Debugger Extension Commands

### Debugger Extension Overview and VSI Support

There are two different types of KD extension DLLs:

- ◆ “Old style” WinDbg Extensions: Extensions that call routines in `wdbgexts.h`
- ◆ “New style” DbgEng extensions: Extensions that call routines in `dbgeng.h` and `wdbgexts.h`

Visual SoftICE supports both KD extension DLLs.

### Controlling Debugger Extension DLLs

There are several commands for controlling the debugger extension DLLs:

- ◆ **!Module.load** (Load Extension DLL) loads a new DLL.  
Example: `!kdextx86.load`
- ◆ **!Module.unload** (Unload Extension DLL) unloads a DLL.  
Example: `!kdextx86.unload`
- ◆ **kdlist** (List Debugger Extensions) displays all loaded debugger extension modules in their default search order.

You can also load an extension DLL by using the full **!module.extension** syntax the first time you issue a command from that module, but you have to manually unload a KD extension DLL that has been loaded. Even disconnecting from a target will not unload KD extension DLLs automatically. Refer to the “[Using Debugger Extension Commands](#)” section later in this appendix for details.

The extension DLLs that you are using must match the operating system of the target computer. The extension DLLs that ship with WinDbg are each placed in a different subdirectory of the installation directory according to the OS version and target mode (Release version or Debug version). For example, nt4fre, w2kchk, and winxp. You must make sure you use the right version, and you must set the right KD extension search path (using either the path definitions in the **Settings** dialog, or via the **SET KDEXTPATH** command at the command line).

You should be as specific as possible when setting the KD extension search path, such as using `C:\Program Files\Debugging Tools for Windows\winxp` instead of `C:\Program Files\Debugging Tools for windows`, otherwise Visual SoftICE will look for the first name-matched KD Extension DLL and load it.

**Note:** If you write your own debugger extensions, you should place them in a new directory and add that directory to the debugger extension path.

## Using Debugger Extension Commands

The use of debugger extension commands is very similar to the use of other commands. A debugger extension command is an entry point in a KD Extensions DLL called by the debugger.

You invoke debugger extensions via the following syntax:

`![module.]extension [arguments]`

---

**Required:** The module name should not be followed with the .dll file name extension.

---

If the module has not already been loaded, it will be loaded into the debugger using a call to `LoadLibrary(module)`. After the debugger has loaded the extension library, it calls the `GetProcAddress` function to locate the extension name in the extension module. The extension name is case-sensitive and must be entered exactly as it appears in the extension module's .def file. If the extension address is found, the extension is called.

If the module name is not specified, the debugger will search the loaded extension modules for this export. The latest loaded DLL is searched first.

When an extension module is unloaded, it is removed from the search order. When an extension module is loaded, it is added to the beginning of the search order.

You can use the **klist** (List Debugger Extensions) command to display a list of all loaded extension modules in their current search order.

If you attempt to execute an extension command that is not in any of the loaded extension modules, you will get an `SI_E_NAME_NOT_FOUND` error message.



# Glossary



## Datum

Artificial data elements contrived to make the user's life easier. For example, BPID, and \_TID (current thread ID). Datums are time and target context sensitive.

## Interrupt Descriptor Table (IDT)

Table pointed to by the IDTR register, which defines the interrupt/exception handlers. Use the IDT command to display the table.

## MAP file

Human-readable file containing debug data, including global symbols and usually line number information.

## MMX

Multimedia extensions to the Intel Pentium and Pentium-Pro processors.

## Object

Represents any hardware or software resource that needs to be shared as an object. Also, the term section is sometimes called an object. Refer to [Section](#).

## One-Shot Breakpoint

Breakpoint that only goes off once. It is cleared after the first time it goes off or the next time Visual SoftICE pops up for any reason.

## Ordinal Form

When a symbol table is not relocated, it is said to be in its ordinal form; in this state, the selectors are section numbers or segment numbers (for 16 bit).

## **Relocate**

Adjust program addresses to account for the program's actual load address.

## **Section**

In the PE file format, a chunk of code or data sharing various attributes. Each section has a name and an ordinal number.

## **Sticky Breakpoint**

Breakpoint that remains until you remove it. It remains even through unloading and reloading of your program.

## **SYM File**

File containing debug data, including global symbols and usually line number information. The SYM file is usually derived from a MAP file.

## **Symbol Table**

Visual SoftICE-internal representation of the debugging information, for example, symbols and line numbers associated with a specific module.

## **Virtual Breakpoint**

Breakpoint that can be set on a symbol or a source line that is not yet loaded in memory.

# Index



## Numerics

1394 [41](#), [46](#)

## A

### About

- the context bar [62](#)
- the status bar [61](#)

### Accessing Images [180](#)

### Active Table [181](#)

### ADDR command [231](#)

### Address

- space [242](#)
- type [216](#)

### Audience [ix](#)

### Automatic

- changes in mode state [88](#)
- loading [183](#)
- output redirection [78](#)
- unloading [183](#)

### Available Targets [51](#)

## B

### BC command [208](#)

### BD command [208](#)

### BE command [208](#)

### Bitwise operators [210](#)

### BL command [208](#)

### Breakpoint

- action, setting [199](#)
- options [193](#)
- page [92](#)
- statistics [207](#)

## Breakpoints [56](#)

- duplicate [206](#)
- embedded [208](#)
- expressions [207](#)
- INT 1 and INT 3 [208](#)
- manipulating [207](#)
- statistics [207](#)
- types [192](#)
- using [191](#)
- virtual [198](#)

### Browsing Available Targets [51](#)

### BSTAT command [207](#)

### Built-in functions [213](#)

## C

### CardBus [43](#), [44](#)

### Character constants [212](#)

### Checked build [220](#)

### Choosing Your Version [1](#)

### Command

- comments [76](#)
- syntax output redirection [76](#)

### Command Page [97](#)

- features [76](#)

### Commands

- BC [208](#)
- BD [208](#)
- BE [208](#)
- BL [208](#)

### BSTAT [207](#)

### Conditional Breakpoints [199](#)

- referencing the stack [204](#)

### Conditional breakpoints

- count functions [200](#)

### Connecting to Targets [53](#)

### Context Bar Controls [62](#)

### Context Menus [63](#)

Copy 79  
CSRSS 237  
Current Context 57  
Cursor Types 60  
Customer Assistance xiii  
Customizing  
    the status bar 71  
    toolbars 72  
Cut 79

## D

Data Item Drag and Drop 80  
Debug Message Page 103  
Debug Toolbar 63  
Debugging  
    features 5  
    resources 220  
Dedicated PCI Ethernet 41, 43  
    NIC 43  
DEVICE command 221  
Disassembly Page 121  
Drag and Drop 79  
DRIVER command 221  
Dual Machine Debugging 3  
Duplicate Breakpoints 206

## E

Eaddr function 216  
EBP register 204  
Elapsed Time 206  
Embedded breakpoints 208  
ESP register 204  
Event Page 127  
Execution Breakpoints 193  
Export 186  
Expression  
    evaluator 209  
        built-in functions 213  
        expression values 216  
        forming expressions 211  
        numbers 212  
        operators 210  
        registers 212  
        symbols 212  
    types 216  
    values 209  
        address-type 216  
        literal-type 216  
        register-type 216  
        symbol-type 216

Expressions  
    breakpoints 207  
    forming 211

## F

Filtering Messages 104  
Firewire 41, 46  
Fixed Address Breakpoints 198  
Fonts & Colors Settings 73  
Forming expressions 211  
Functions  
    built-in 213  
    expression evaluator 215

## G

GDI objects 238  
GDT command 228  
General Global Settings 65  
Global  
    settings 65  
Global Descriptor Table 226, 228

## H

Handle values 240  
Heap  
    API 243  
    architecture 243  
    blocks 248  
HEAP32 command 237, 247

## I

I/O Breakpoints 195  
IDT command 227  
Image Search 186  
Image-Relative Breakpoints 198  
Indirection operators 210  
INT 1 instruction breakpoints 208  
INT 3 instruction breakpoints 208  
Integrated MS Symbol Server Access 184  
Intel architecture 226  
Interrupt  
    breakpoints 195  
    Descriptor Table 226, 227

## K

KD Extension  
    DLLs 259  
    Overview 259  
KD Extension Paths 186

KD Extensions 259

## Kernel

Windows 225

Kernel Debugger Extensions 259

Keyboard Definitions 50

Keyboard Settings 69

## L

Live Mode On Connection 88

Loading Tables 181

## Local

copies of images 180

variables in conditional expressions 203

Locals Page 132

Logical operators 211

## M

MACRO command 241

Manipulating breakpoints 207

Matching Tables and Images 182

Math operators 210

## Memory

breakpoints 194

map of system memory 230

Page 135

MS Symbol Server

access 184

paths 186

## N

NonPaged System area 235

NTCALL command 227

NTOSKRNL.EXE 226

## O

OBJDIR command 221

OBJTAB command 230, 241

OHCI 41, 45

On Demand Symbol Handling 6

## Operators

bitwise 210

expression evaluator 210

indirection 210

logical 211

math 210

precedence 211

Organization x

Output Redirection 76, 78

## P

Pad Level Page Drag and Drop 79

Pads 85

Page Access 54

PAGE command 233

Paged Pool System area 234

Pages 85

Paste 79

Path Settings 186

Paths 67

Persistent Loading Tables 184

Per-Workspace Settings 66

notes 187

PHYS command 233

Precedence operators 211

Preface ix

Preferences 64

Pre-Loading Tables 184

Print 81

preview 81

Process address space 242

Process List Page 145

Processor Control Region 236

PTE 234

## R

Referencing the Stack 204

Registers Page 151

## S

Saving Contents 80

Script

execution 82

search path 186

Scripts 68

Serial 41

Setting

breakpoint actions 199

general paths 186

Visual SoftICE paths 186

Single Machine Debugging 2

SoftICE 1, 2

Source

Page 156

status bar 62

symbol state icons 59

search 186

Stack

frame 204

Page 167

referencing in conditional breakpoints 204

Status Bar  
examples 61  
settings 71  
notes 187

Supported Symbols 180

Symbol  
location 179  
management 179  
table state 59

Symbol Files 222

Symbol Handling, On Demand 6

Symbol Management 179

Symbols 212  
in Visual SoftICE 185  
type 216

System

Code area 230  
memory map 230  
Tables System area 231  
View System area 230

System Page Table Entries 234

## T

Table Loading Controls 183

Target

state 57  
transport overview 41

Target Browser Utility 52

Target Menu 52

Targets 51, 53

Text Scratch Page 171

THREAD command 243

Toolbar

settings 71  
notes 187

Troubleshooting 255

  Visual SoftICE 255

Typographical Conventions xii

## U

UHCI 41, 45

Understanding Breakpoint Contexts 198

Universal PCI Ethernet 41, 44

Unloading Tables 181

USB

  host controller 41, 45  
  NIC 41, 45

Useful Documentation xii

USER

object creation 241  
Object Table 241  
objects 238

User

  interface overview 49  
  symbol search path 186

Using Exports 185

## V

Virtual breakpoints 198

Virtual NIC Driver 41

Visual SoftICE 1, 3

  features 5  
  global settings 66  
  icons 54  
  overview 5  
  symbol table status 188  
  UI overview 49  
  user interface 85  
    overview 49

Visual SoftICE User Interface

  Overview 49

VSI

  symbol table status 188  
  system activity messages 187

## W

Watch Page 174

Win32 subsystem 237

Window Message Breakpoints 197

Windows

  components 237  
  DDK 221  
  exploring 219  
  kernel 225  
  references 224  
  system memory map 230  
Workspace Save and Load 74

Workspace Uses 51

Workspaces 50

Writing Scripts for VSI 83